



# DCS216 Operating Systems

## Lecture 05 Processes (1)

**Mar 11<sup>th</sup>, 2024**

**Instructor: Xiaoxi Zhang**  
**Sun Yat-sen University**



## ■ Content

- Process Concept
  - The Process
  - Process States
  - Process Control Block (PCB)
  - Threads
- Operations on Processes
  - Process Creation
  - Process Termination
- Unix and Linux Examples
- Process Scheduling
- Context Switch



## ■ Process Concept

### ■ What is a Process?

- (Informally) A process is a **program** in **execution**, or a **running program**.
- Program is a **passive** entity, a lifeless thing: it just sits there on the disk.
- Process is an **active** entity: it consumes CPU and Memory.
- Program becomes process when it is **loaded** into memory



## ■ Process Concept

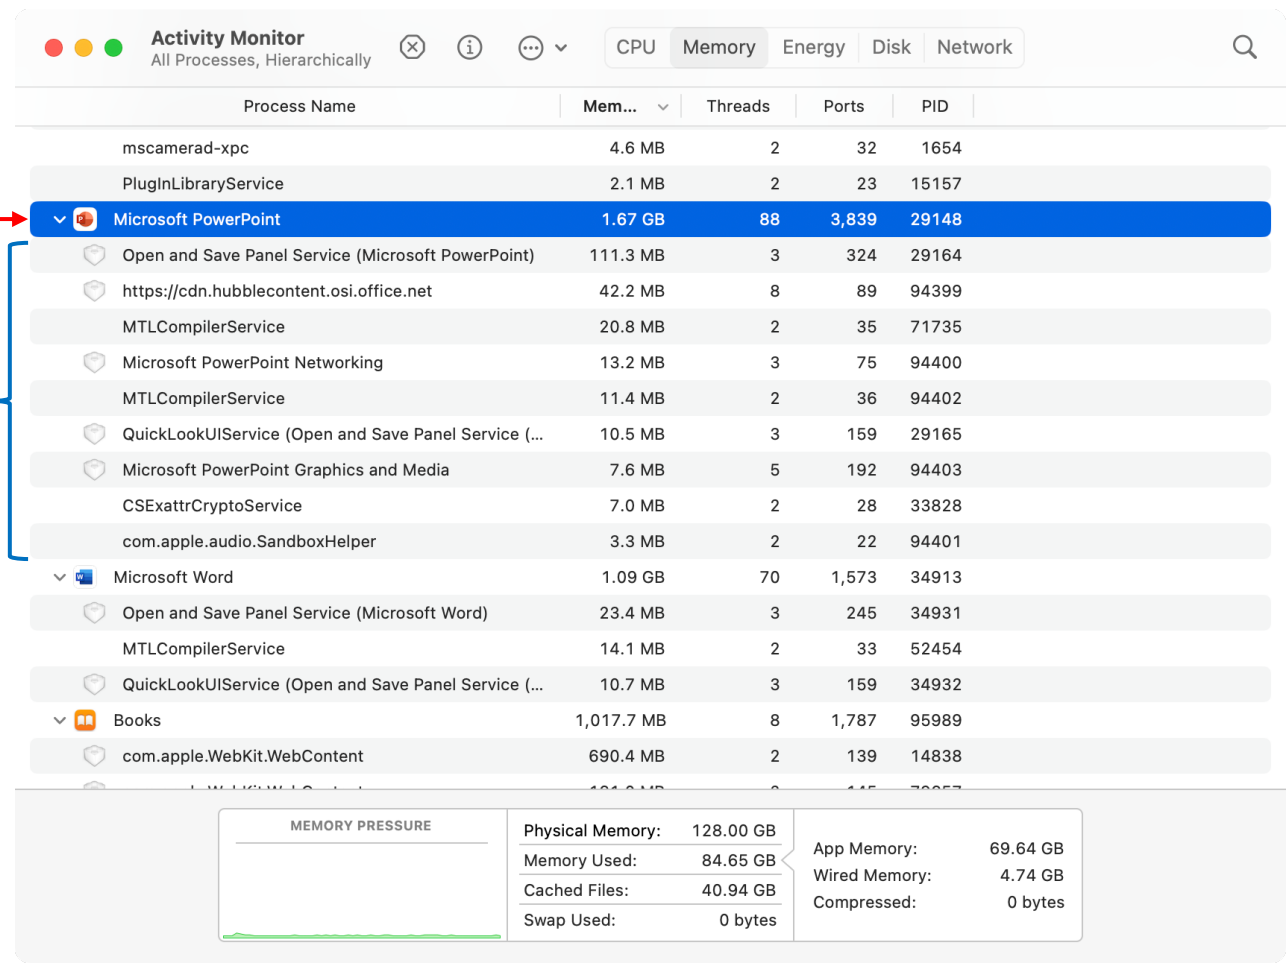
- What is a Process?
  - (Informally) A process is a **program** in **execution**, or a **running program**.
  - Program is a **passive** entity, a lifeless thing: it just sits there on the disk.
  - Process is an **active** entity: it consumes CPU and Memory.
  - Program becomes process when it is **loaded** into memory
- Definition: **execution environment with restricted rights**
  - One or more **threads** executing in a single **address space**
  - Owns file descriptors, network connections
- Instance of a running program
  - When you run an executable file, it runs in its own process
  - **Application**: one or more **processes** working together
- Processes are protected from each other
- OS protected from processes
- In modern OSes, anything outside of kernel runs in a process

## ■ Process Concept

- Instance of a running program
  - When you run an executable file, it runs in its own process
  - **Application**: one or more **processes** working together

Application  
PowerPoint

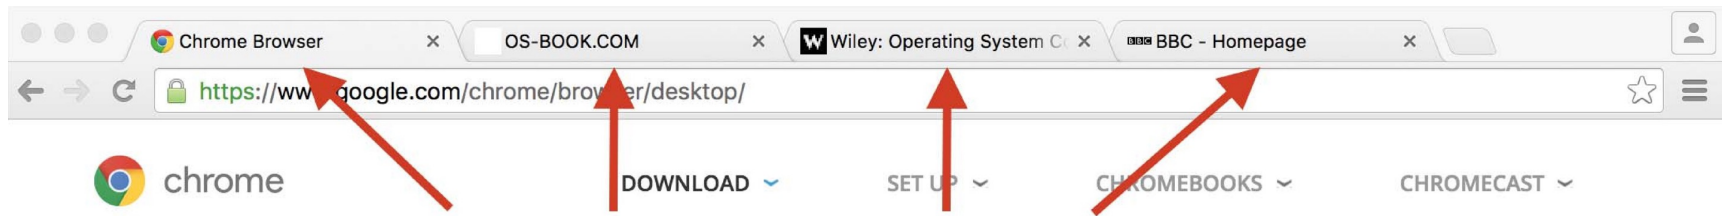
Processes  
that constitute  
PowerPoint





## ■ Multiprocess Application Example – Google Chrome Browser

- Many web browsers ran as single process
  - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is **multiprocess** with 3 different types of processes:
  - **Browser** process – manages user interface, disk and network I/O
  - **Renderer** process – renders (渲染) webpages, deals with HTML, JavaScript. A new renderer created for each website opened
    - Runs in **sandbox** (沙盒) restricting disk and network I/O, minimize security risks
  - **Plug-in** process – for each type of plug-in (插件)



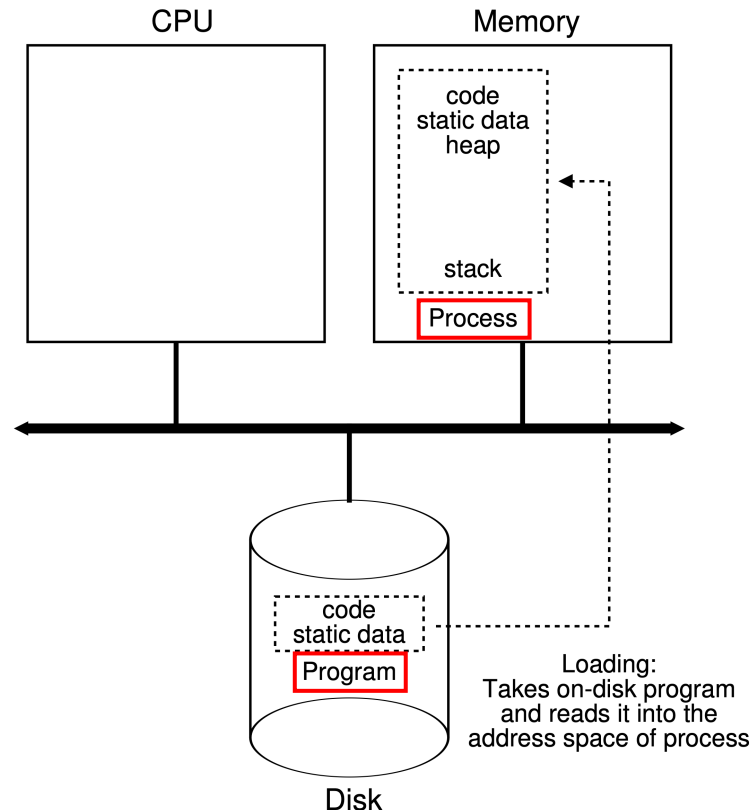
Each tab represents a separate process.



## ■ Process Concept

### ■ What is a Process?

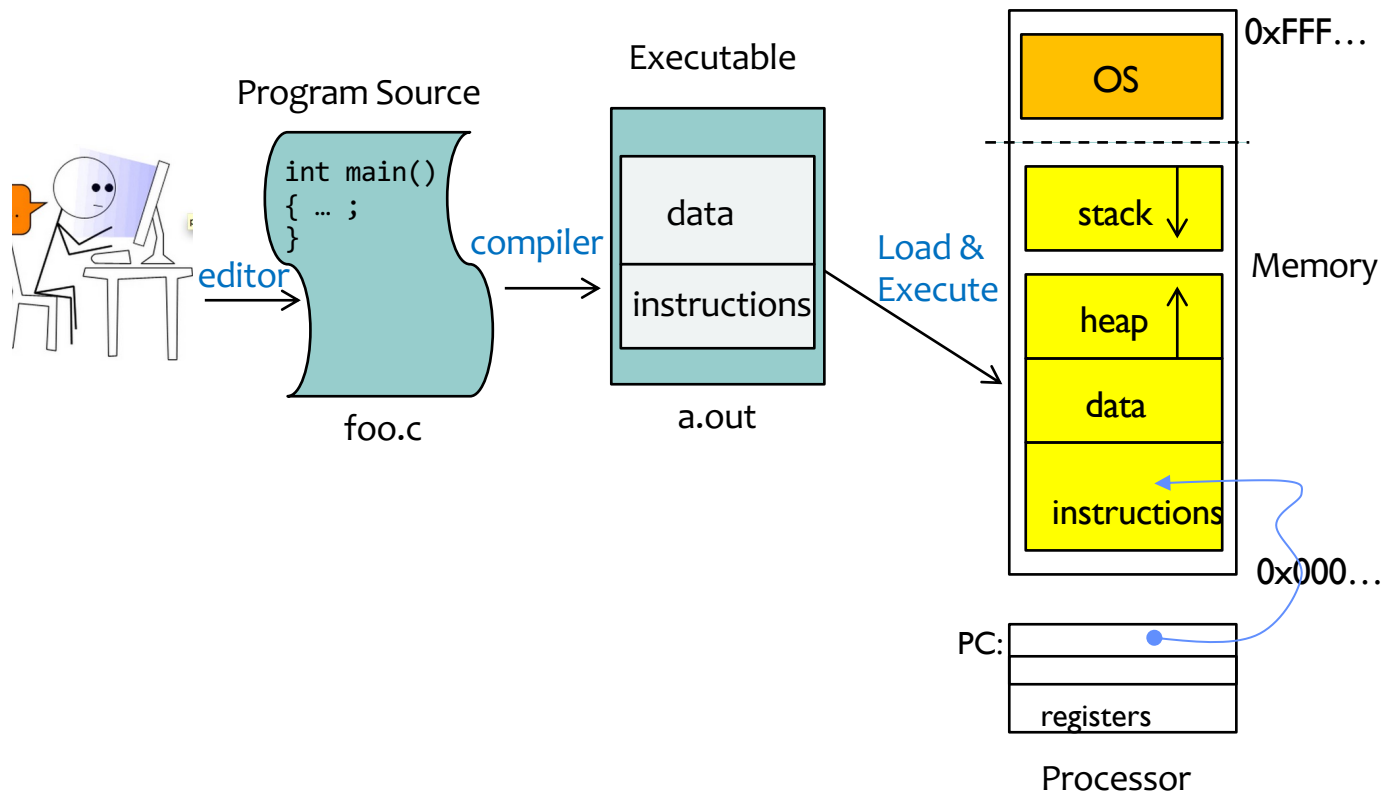
- (Informally) A process is a **program** in **execution**, or a **running program**.
- Program is a **passive** entity, a lifeless thing: it just sits there on the disk.
- Process is an **active** entity: it consumes CPU and Memory.
- Program becomes process when it is **loaded** into memory





## ■ Running of a program

- Load instruction and data segments of executable file into memory
- Create stack and heap
- “Transfer control to program”
- Provide services to program
- While protecting OS and program

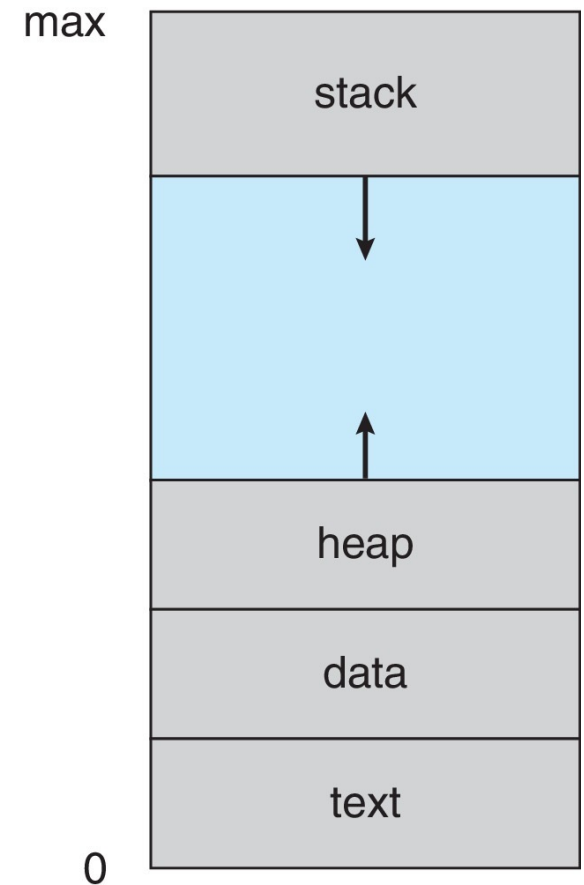






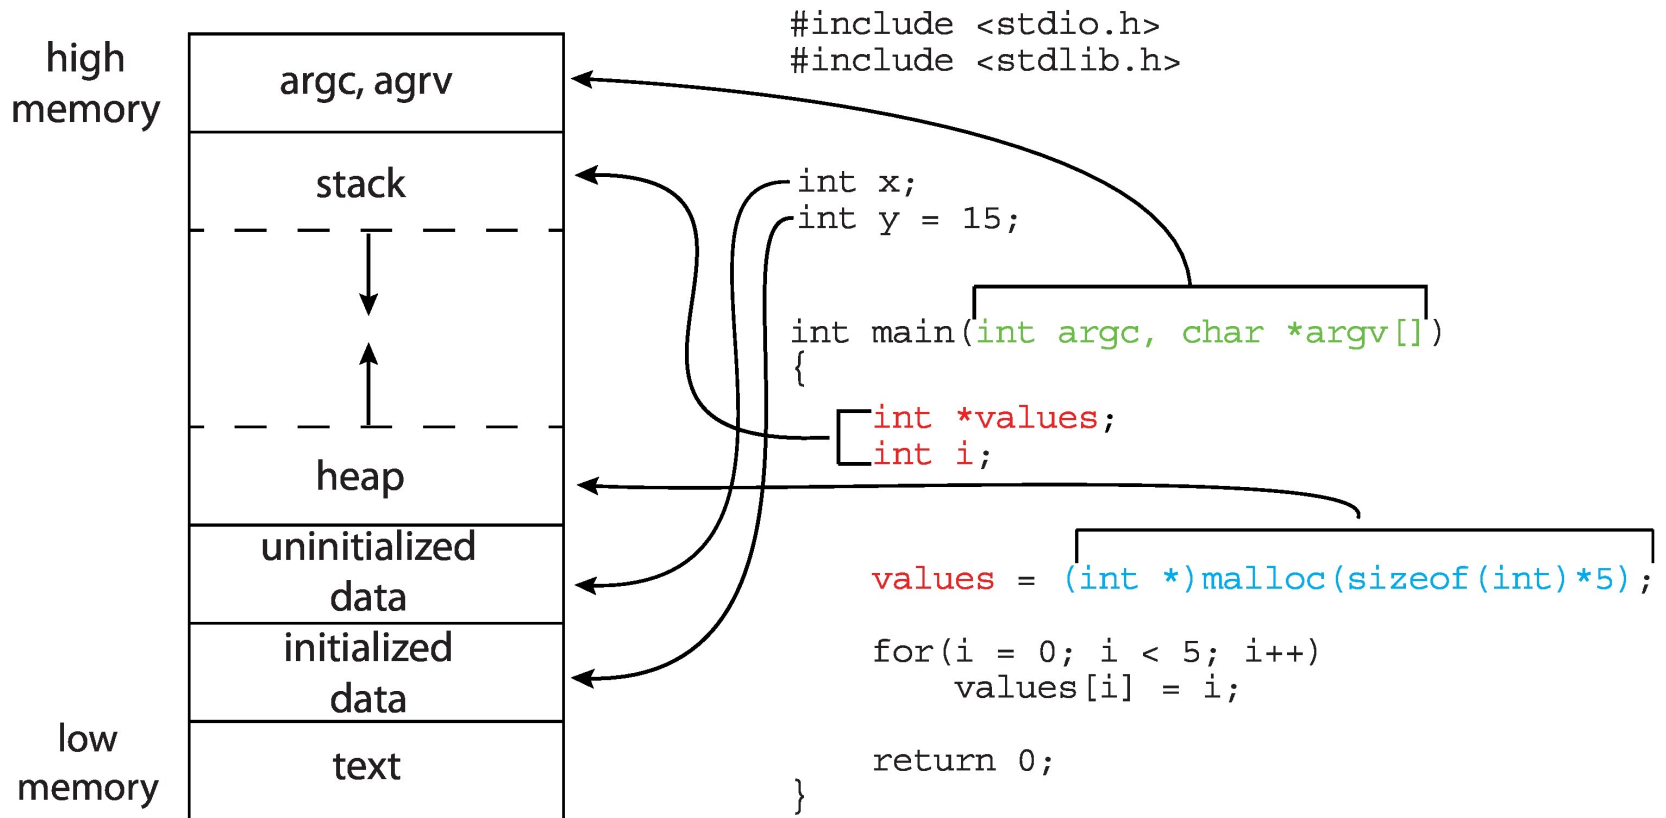
## ■ Process Consists of

- **Text** Section: the program code
- **Data** Section: global variables
- **Stack** Section: temporary data storage when invoking functions, e.g., *function parameters*, *return addresses*, *local variables*
- **Heap** Section: dynamically allocated memory during runtime via `malloc()`.
- Current Activity:
  - **Program Counter** (Instruction Pointer)
  - **Processor Registers**





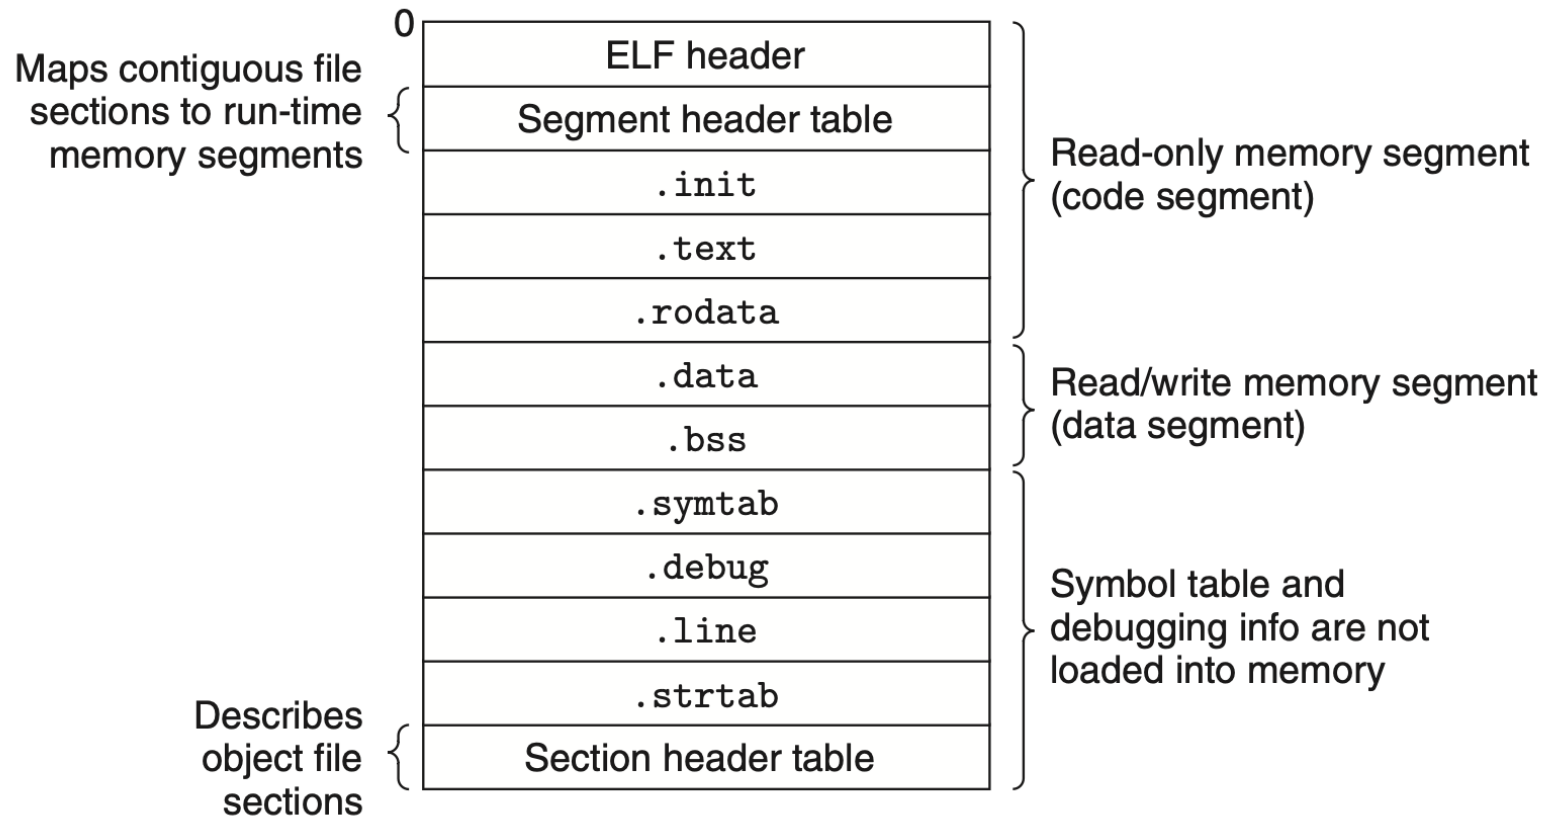
## ■ Process Memory Layout





## ■ Typical ELF Executable Object File on Linux

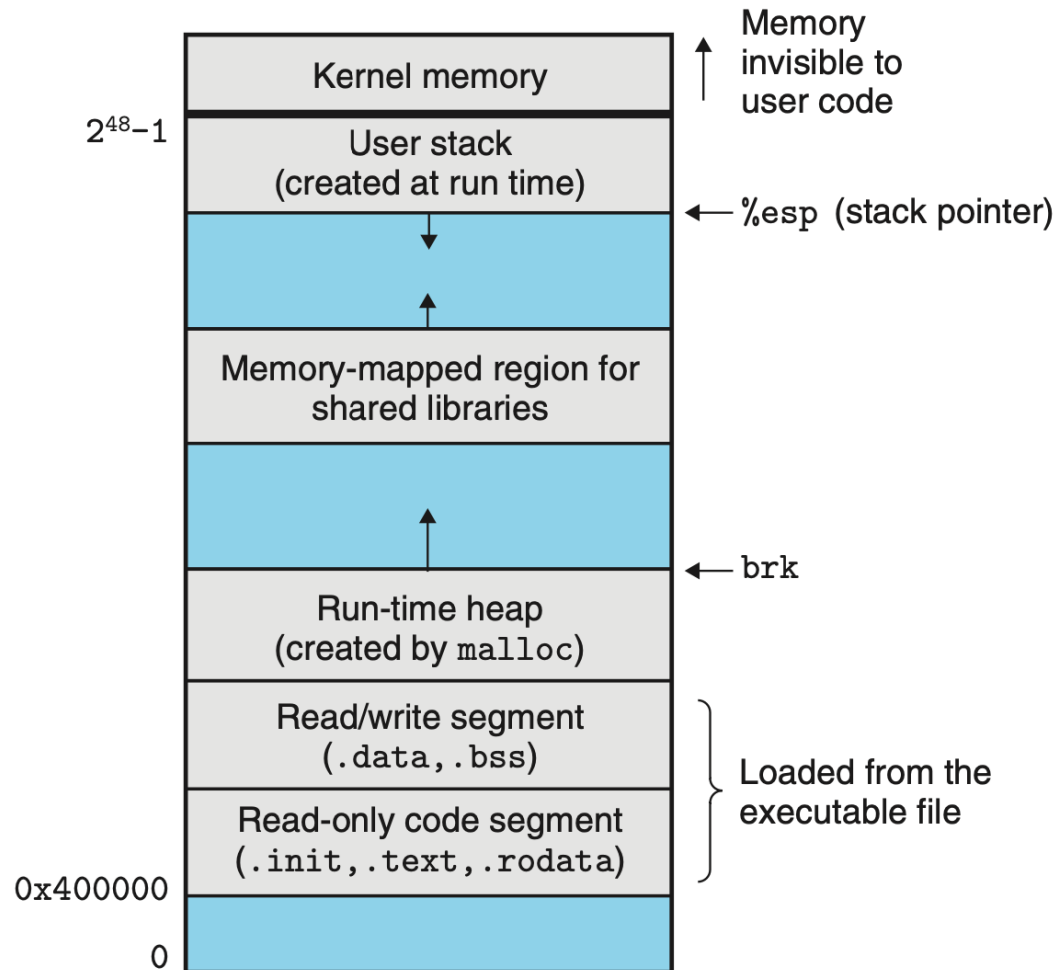
- Exercise: Use `objdump` to inspect various sections of an executable file (a program)





## ■ Process Memory Layout on Linux

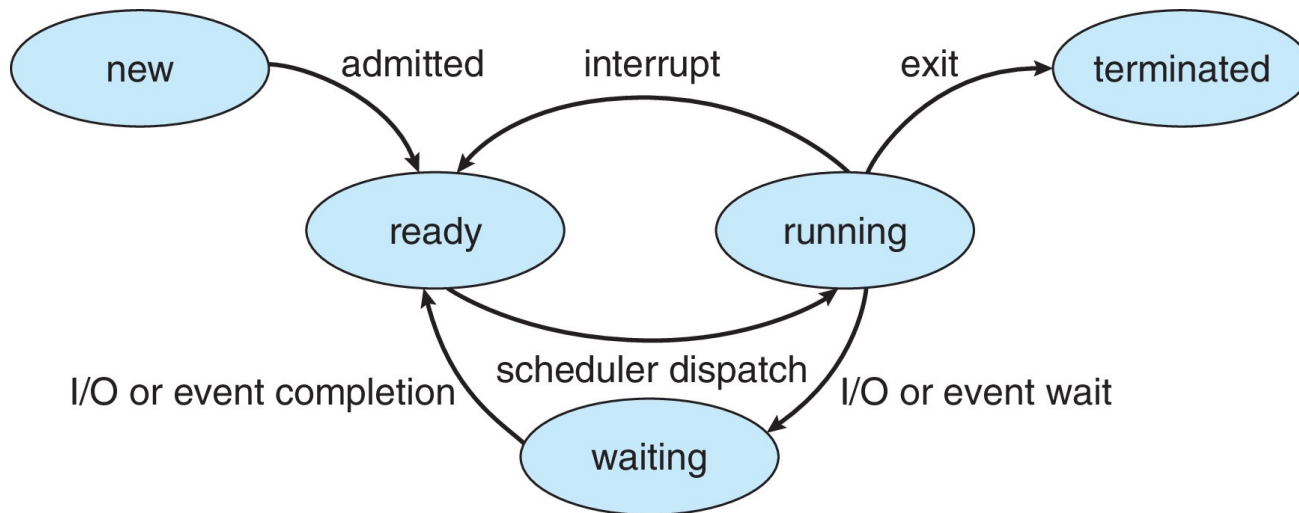
- Exercise: Use `gdb` to inspect memory layout of a process





## ■ Process States

- As a process executes, it changes state
  - **New:** The process is being created
  - **Running:** Instructions are being executed
  - **Waiting (Blocked):** The process is waiting for some event to occur
  - **Ready:** The process is waiting to be assigned to a processor
  - **Terminated:** The process has finished execution





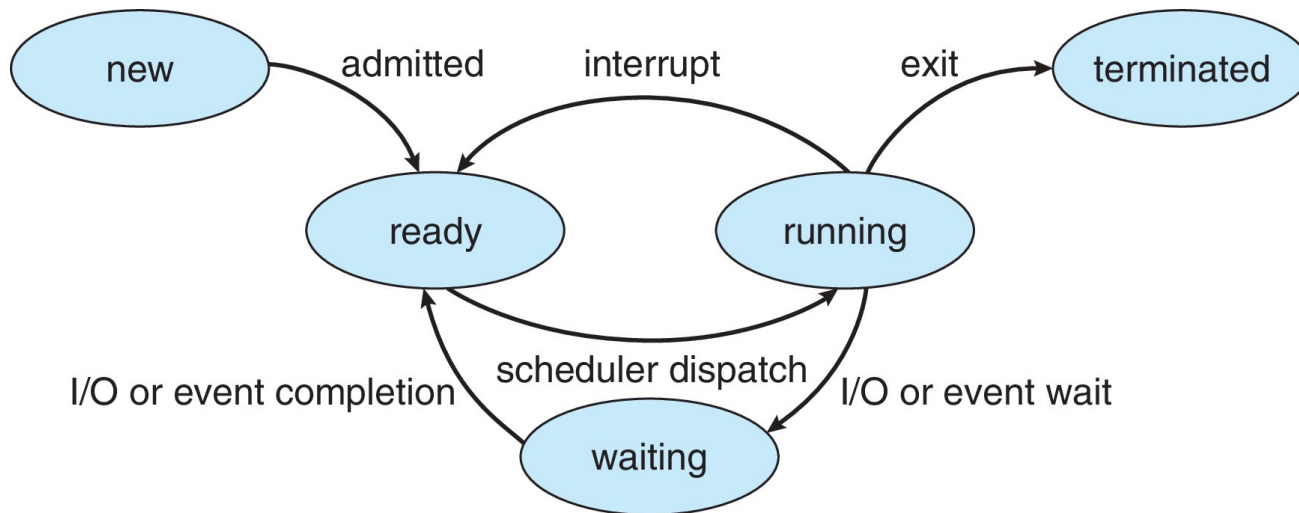
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
```

Using gcc to generate executable file **ps\_state**. The **program ps\_state** now sits somewhere on the disk. There is no **process** of the **program ps\_state** running at the moment.





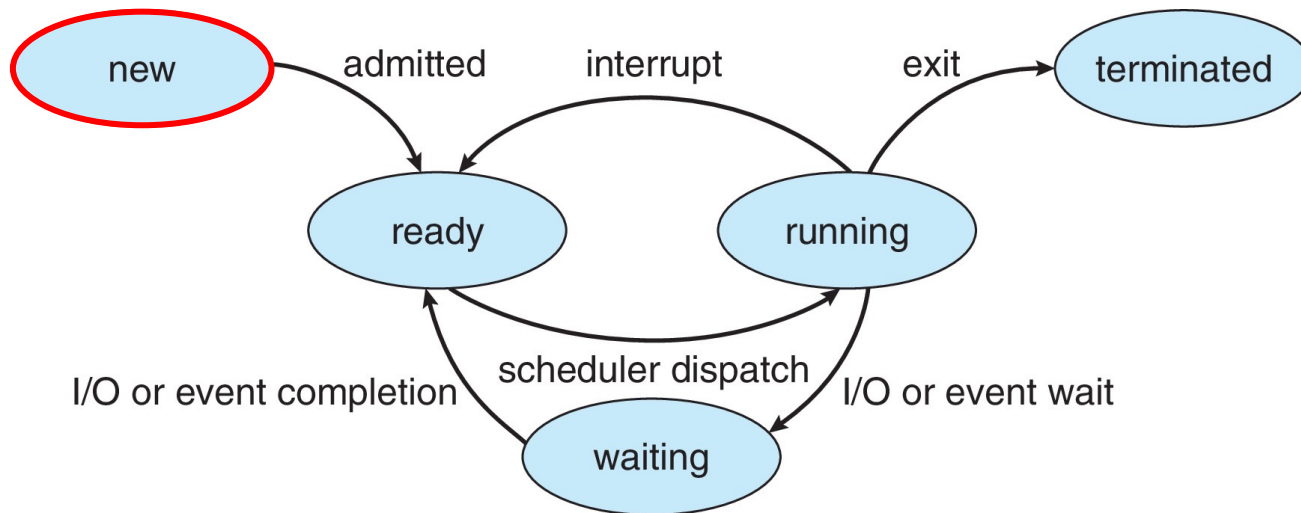
## Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
```

Execute the program ps\_state.  
The process is admitted and being created by the OS.





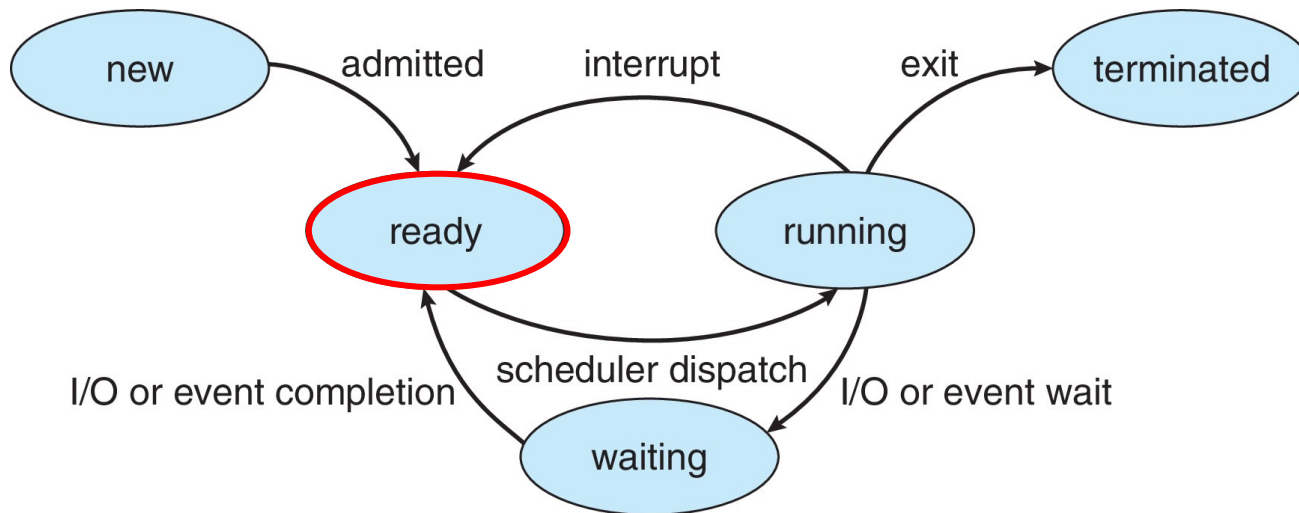
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
```

The process has been created and ready to run.







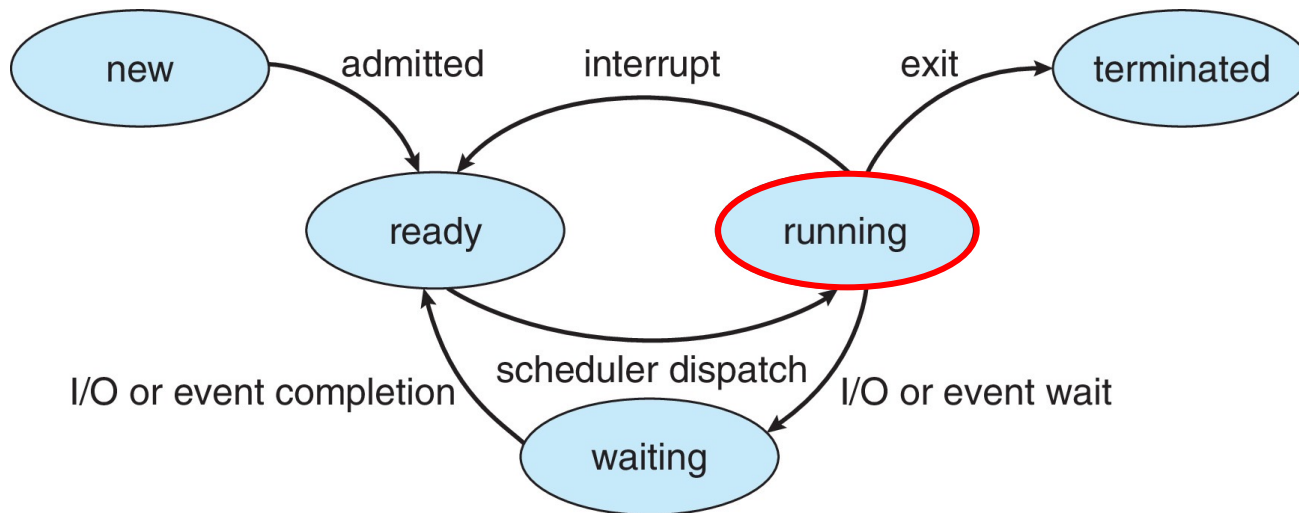
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
```

The scheduler has chosen to run this process on one of the CPU(s).





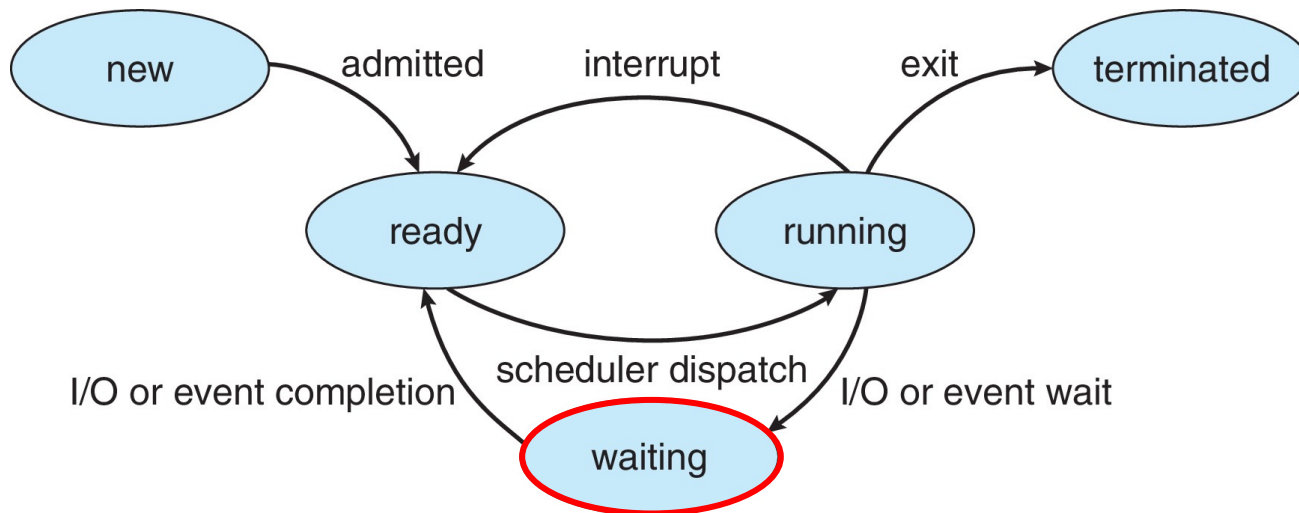
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
```

The process has requested I/O.  
The scheduler moves this process from running to waiting list.  
The process is waiting for I/O to complete.





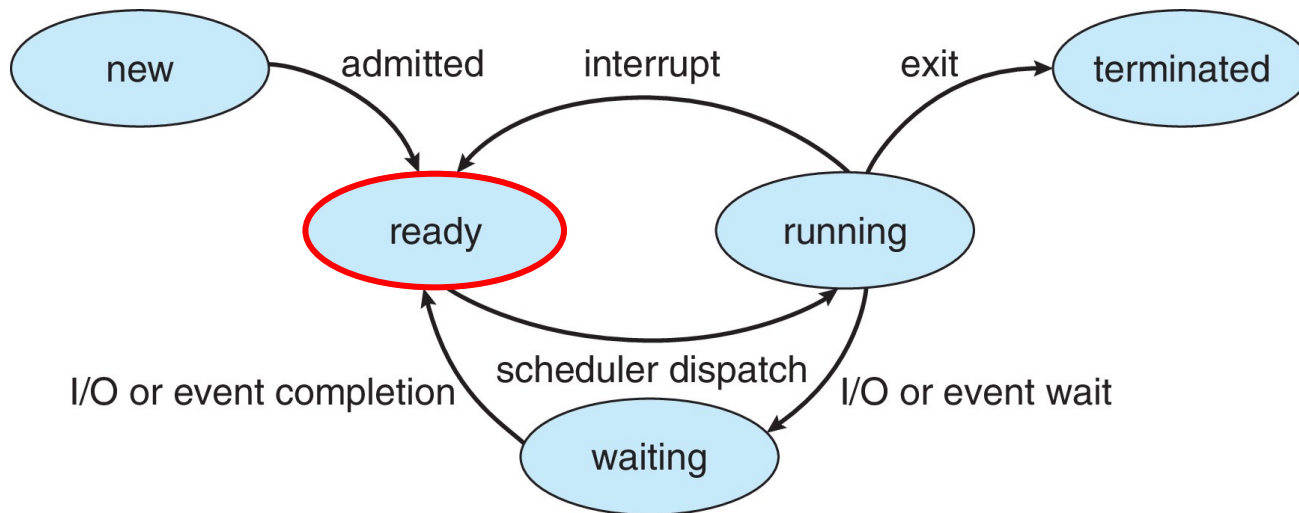
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
```

I/O completed, sends an interrupt to the CPU. The CPU scheduler moves this process to ready list. But this process is not running yet.





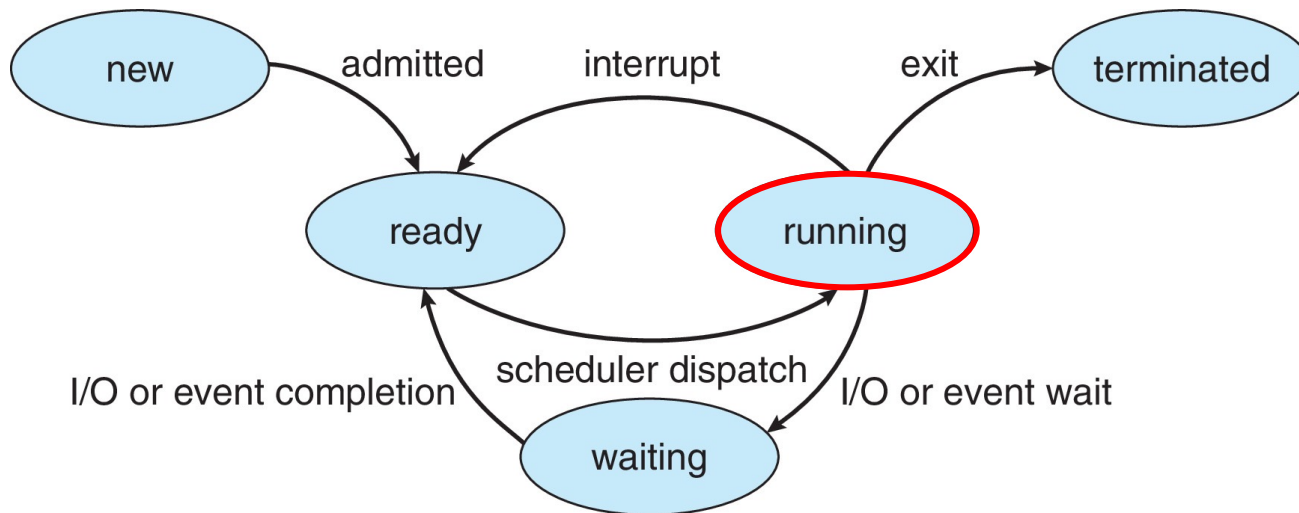
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
```

This process is scheduled to run, performing normal function logic.





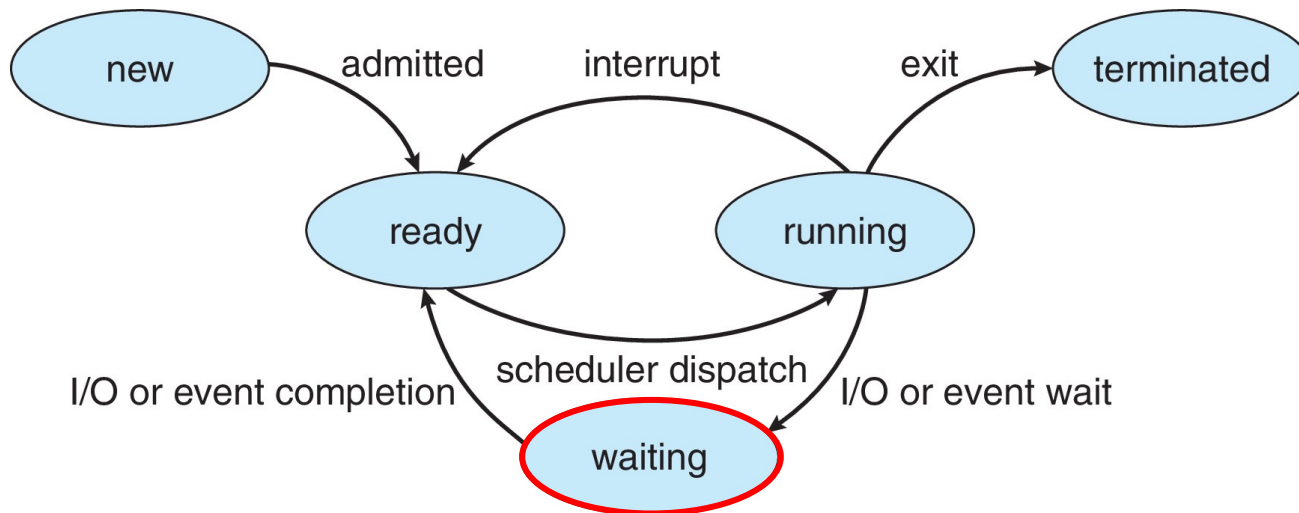
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
```

The process has requested I/O again.  
The scheduler moves this process from running to waiting list.  
The process is waiting for I/O to complete.





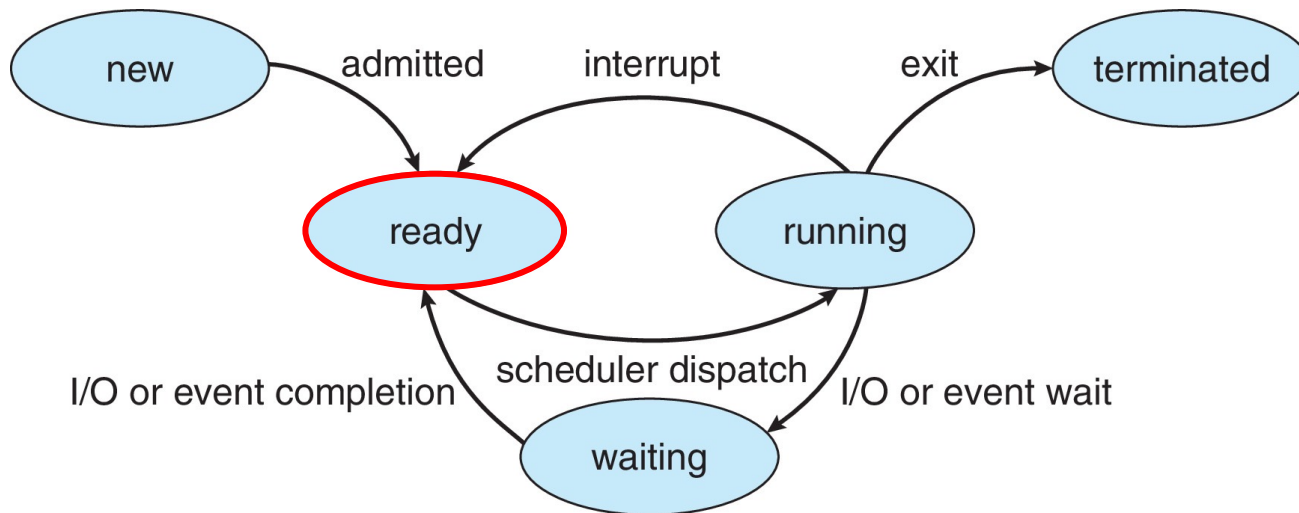
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
n*n: 9
```

I/O completed, sends an interrupt to the CPU. The CPU scheduler moves this process to ready list. But this process is not running yet.





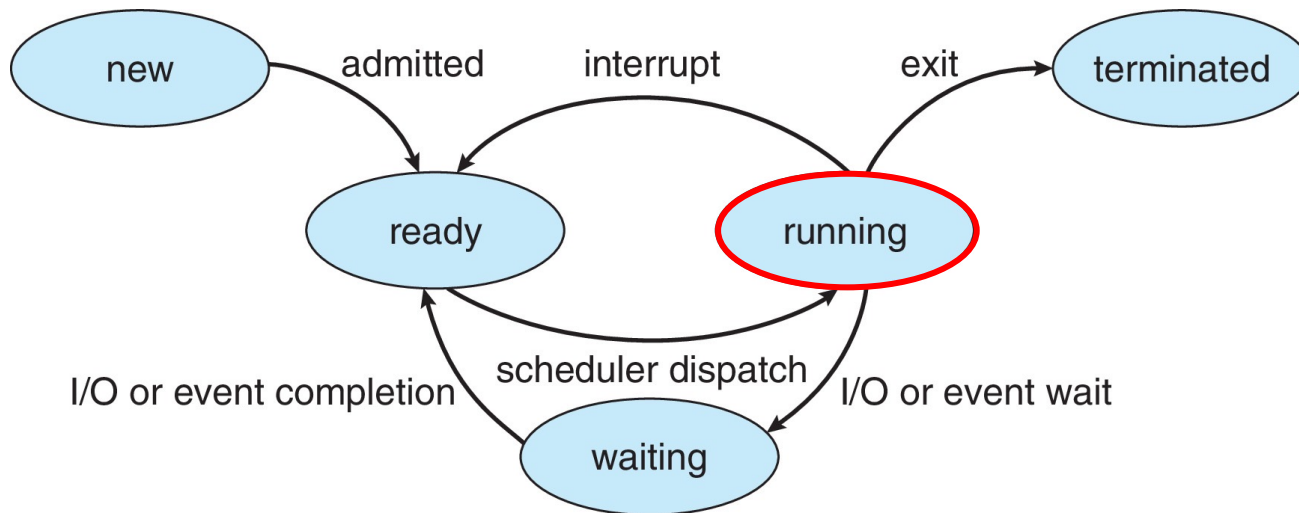
## Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
n*n: 9
```

This process is scheduled to run, performing normal function logic.





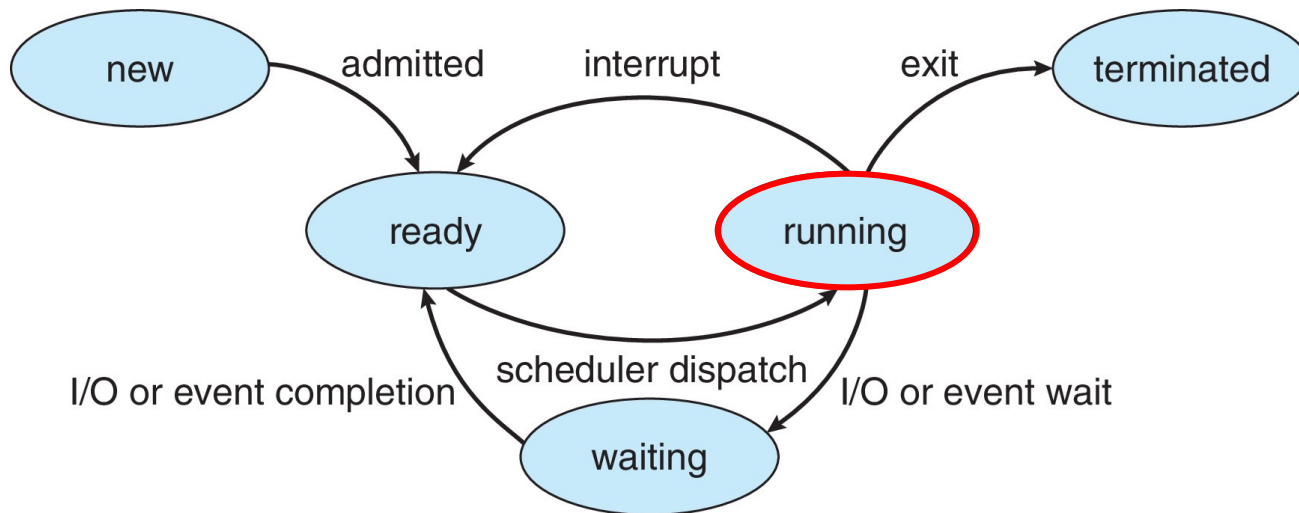
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
n*n: 9
```

This process has reached the end of program execution, e.g., invoking `exit()` system call to terminate.







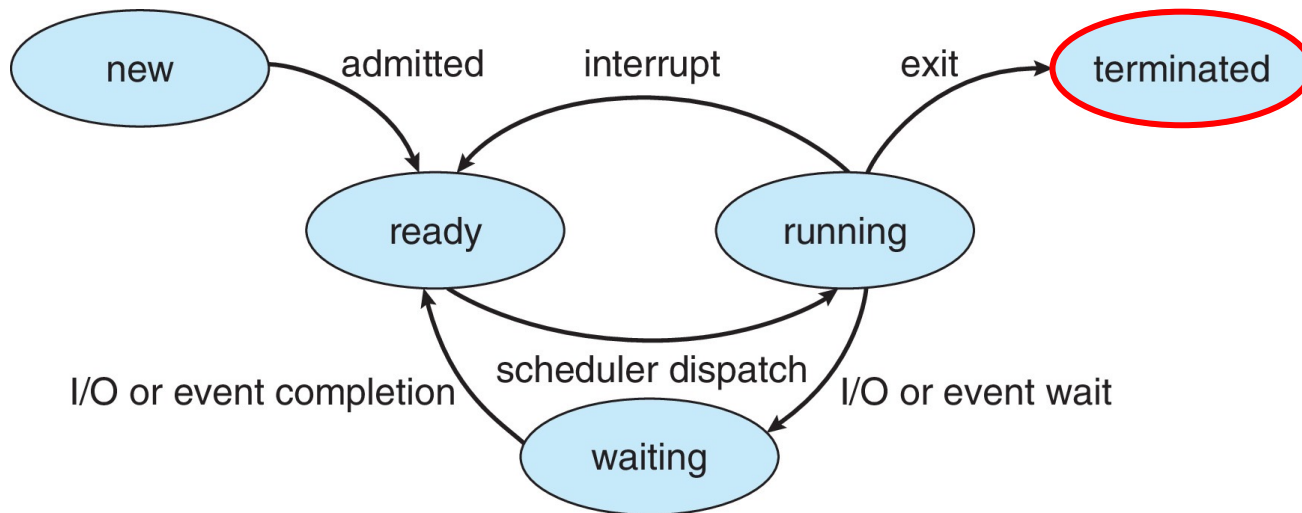
## ■ Process States Example

```
#include <stdio.h>

int main() {
    int n;
    scanf("%d", &n);
    n = n * n;
    printf("n*n: %d\n", n);
    return 0;
}
```

```
$ gcc -o ps_states ps_states.c
$ ./ps_states
3
n*n: 9
```

This process is now terminated by the OS.

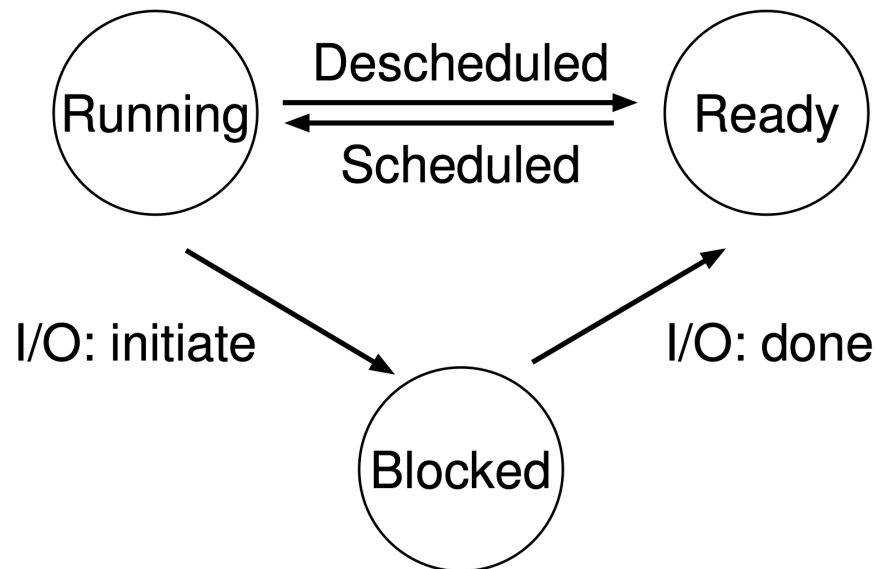




## ■ Process States

### ■ A Simplified View:

- **Running:** Instructions are being executed
- **Waiting (Blocked):** The process is waiting for some event to occur
- **Ready:** The process is waiting to be assigned to a processor





## ■ Process Control Block (PCB)

(Information associated with each process):

- **Process State:** running, waiting, etc.
- **Program Counter:** location of instruction to execute next
- **CPU Registers:** contents of all process-centric registers
- **CPU Scheduling Info:** priorities, scheduling queues
- **Memory Management Info:** memory allocated to process
- **Accounting Info:** CPU used, clock time elapsed, time limits
- **I/O Info:** I/O devices allocated to process, list of open files

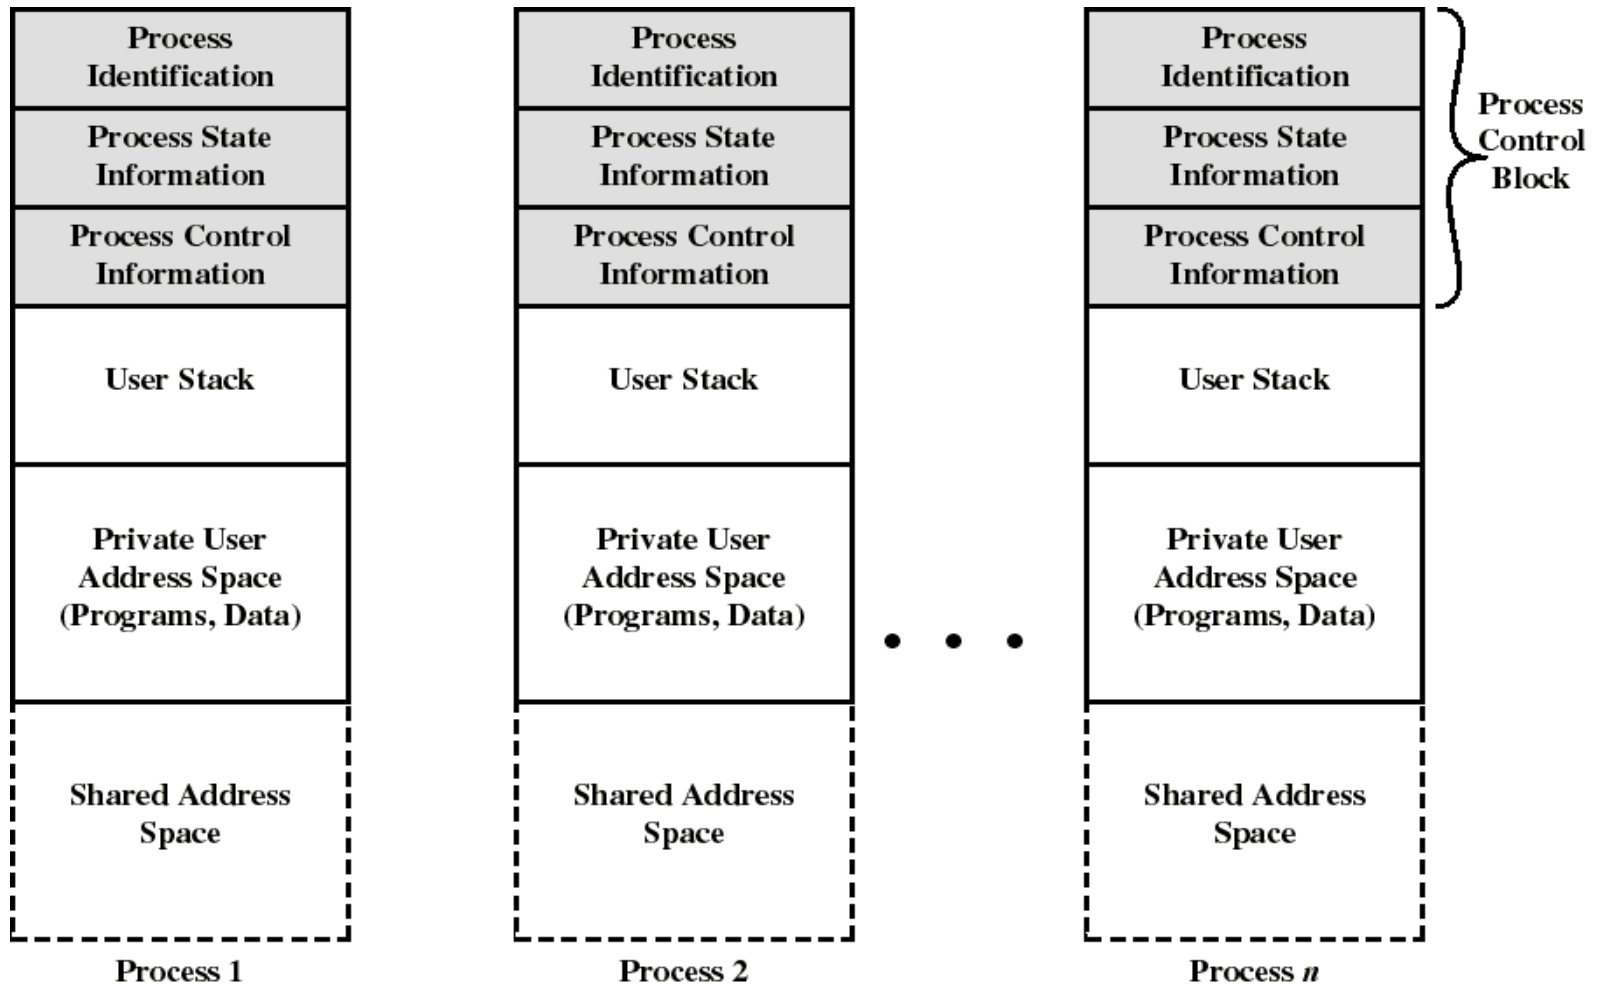
process state
process number
program counter
registers
memory limits
list of open files
...



## ■ Process Control Block (PCB)

### ■ Where exactly is the PCB?

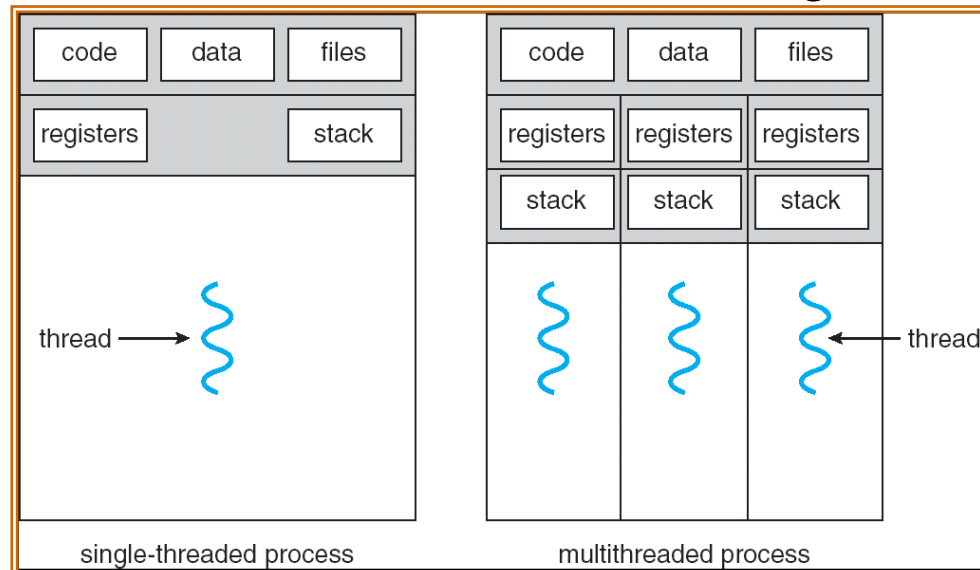
- In the **kernel**'s memory space, part of the kernel's data structures.





## Threads

- The process model discussed so far has implied that a process is a running program that performs a single **thread** of execution
- Modern OSes have extended the process concept to allow a process to have multiple **threads** of execution
  - thus enabling a process to perform more than one task at a time
  - multiple threads of control → **threads**
  - PCB extended to contain multiple program counters
- *More on Threads in later chapters when dealing with concurrency*





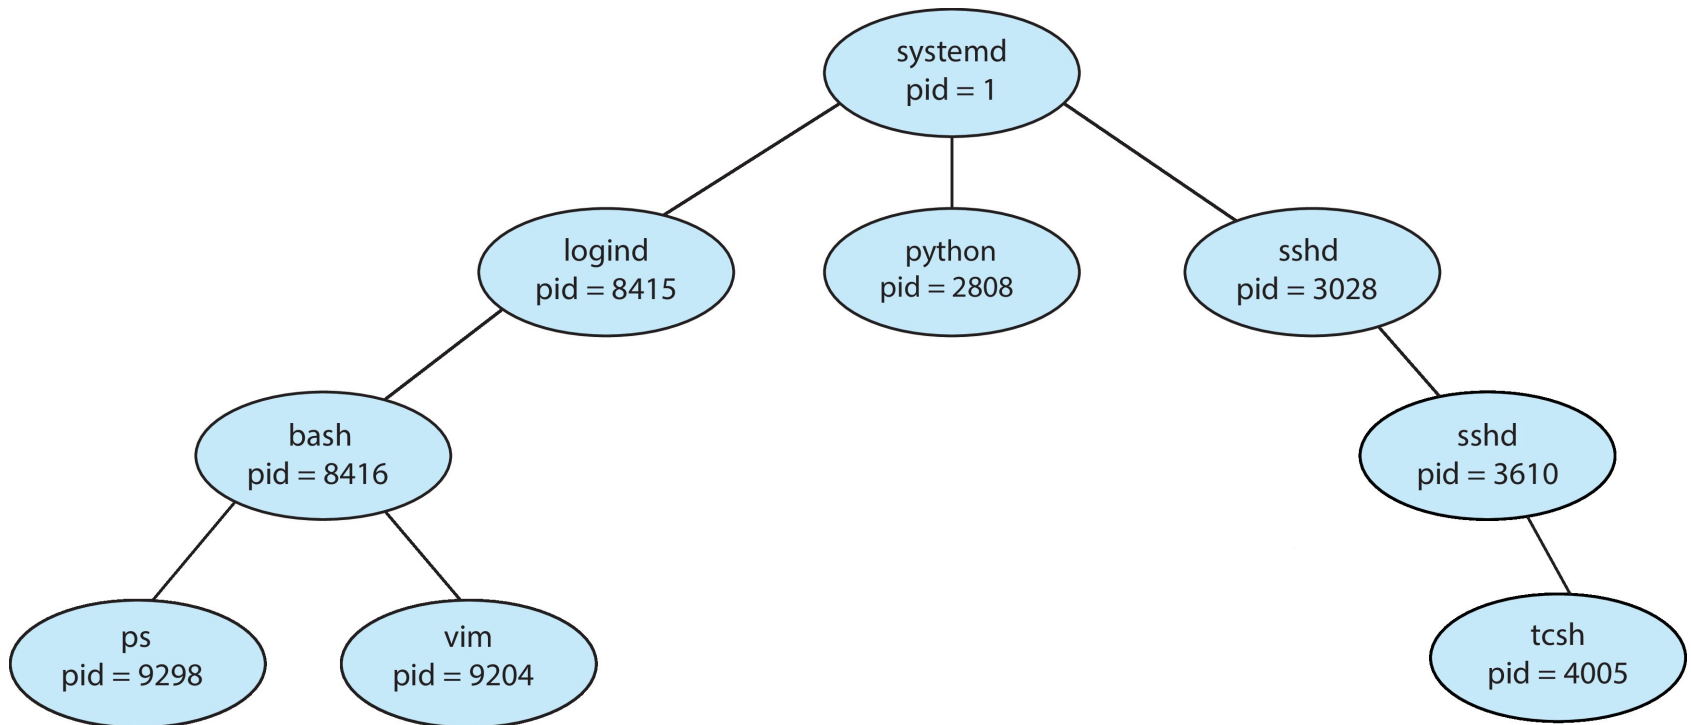
## ■ Operations on Processes

- Processes in most systems can execute concurrently
- OS can spawn (衍生) new processes (Of course!)
- Processes themselves should also be able to spawn new processes
  - via system calls
- OS must provide mechanisms for:
  - Process Creation
  - Process Termination



## ■ Process Creation

- **Parent** process create **children** processes, which in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via **process identifier (pid)**





## ■ Process Creation

- **Parent** process create **children** processes, which in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via **process identifier (pid)**

```
$ pstree
systemd--ModemManager--2*[{ModemManager}]
      |
      |--agetty
      |--containerd--12*[{containerd}]
      |--cron
      |--dockerd--13*[{dockerd}]
      |--rsyslogd--3*[{rsyslogd}]
      |--sshd--sshd--sshd--bash--tmux: client
      |--systemd--(sd-pam)
              |--dbus-daemon
              |--pulseaudio--2*[{pulseaudio}]
      |--systemd-journal
      |--systemd-logind
      |--systemd-resolve
      |--systemd-timesyn--{systemd-timesyn}
      |--tmux: server--bash--emacs
              |--6*[bash]
              |--bash--vim
              |--bash--pstree
      |--upowerd--2*[{upowerd}]
```





## ■ Process Creation

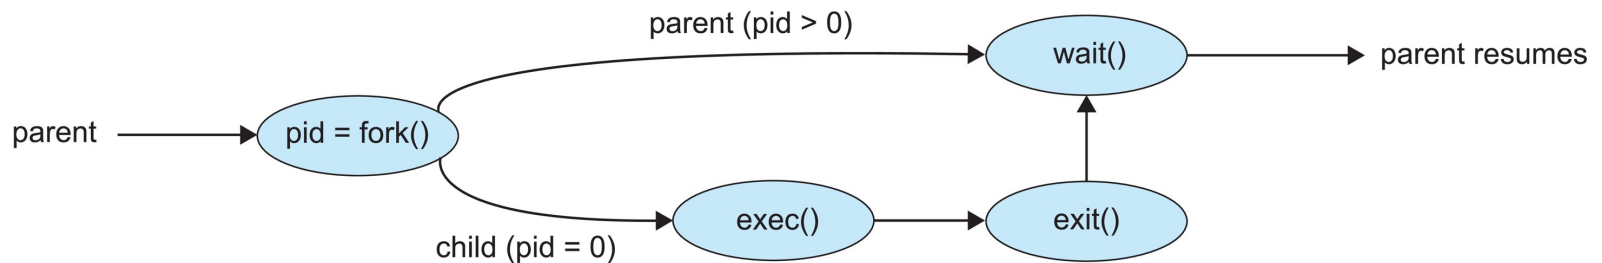
- **Parent** process create **children** processes, which in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via **process identifier (pid)**
- Possible resource sharing:
  - Parent and children processes **share all** resources
  - Children processes **share subset** of parent's resources
  - Parent and children processes **share no** resources
- Possible Execution:
  - Parent and children processes execute **concurrently**
  - Parent **waits** until children processes terminate
- Address space:
  - Child duplicate of parent
  - Child has a program loaded into it (different address space)



## ■ Process Creation

### ■ UNIX API

- `fork()` system call creates new process
- `exec()` system call used after `fork()` to replace the process' memory space with new program
- Parent process calls `wait()` waiting for the child to terminate





## ■ UNIX fork()

- `pid_t fork(void)`; causes creation of a new process.

```
/* p1_fork.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", getpid());

    pid_t rc = fork();

    if (rc < 0) {          /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("Hi, I am child (pid: %d)\n",
            getpid());
    } else {              /* Parent (old process) */
        printf("Hi, I am parent of %d (pid: %d)\n",
            rc, getpid());
    }

    return 0;
}
```

```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am parent of 96648 (pid: 96647)
Hi, I am child (pid: 96648)
```



## ■ UNIX fork()

- `pid_t fork(void)`; causes creation of a new process.

```
/* p1_fork.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", getpid());

    pid_t rc = fork();

    if (rc < 0) {          /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("Hi, I am child (pid: %d)\n",
               getpid());
    } else {              /* Parent (old process) */
        printf("Hi, I am parent of %d (pid: %d)\n",
               rc, getpid());
    }

    return 0;
}
```

```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am parent of 96648 (pid: 96647)
Hi, I am child (pid: 96648)
```

Parent and child run concurrently. Either parent or child could get executed first. This could also happen:



```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am child (pid: 96648)
Hi, I am parent of 96648 (pid: 96647)
```



## ■ UNIX wait()

- `pid_t wait(int *stat_loc);` suspends execution of current process until its child process has completed.

```
/* p2_fork_wait.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", getpid());

    pid_t rc = fork();

    if (rc < 0) { /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("Hi, I am child (pid: %d)\n",
            getpid());
    } else { /* Parent (old process) */
        int rc_wait = wait(NULL);
        printf("Hi, I am parent of %d (pid: %d)\n",
            rc_wait, getpid());
    }

    return 0;
}
```

```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am parent of 96648 (pid: 96647)
Hi, I am child (pid: 96648)
```

```
$ gcc -o p2_fork_wait p2_fork_wait.c
$ ./p2_fork_wait
Hello world (pid: 96874)
Hi, I am child (pid: 96875)
Hi, I am parent of 96875 (pid: 96874)
```



## ■ UNIX wait()

- `pid_t wait(int *stat_loc);` suspends execution of current process until its child process has completed.

```
/* p2_fork_wait.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", getpid());

    pid_t rc = fork();

    if (rc < 0) { /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("Hi, I am child (pid: %d)\n",
            getpid());
    } else { /* Parent (old process) */
        int rc_wait = wait(NULL);
        printf("Hi, I am parent of %d (pid: %d)\n",
            rc_wait, getpid());
    }

    return 0;
}
```

```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am parent of 96648 (pid: 96647)
Hi, I am child (pid: 96648)
```

```
$ gcc -o p2_fork_wait p2_fork_wait.c
$ ./p2_fork_wait
Hello world (pid: 96874)
Hi, I am child (pid: 96875)
Hi, I am parent of 96875 (pid: 96874)
```



`wait()` system call **blocks** the parent process until the child process finishes execution, thus making sure **child finishes before parent**.



## ■ UNIX exec()

- `int exec(...);` transforms the calling process into a new process

```
/* p3_fork_wait_exec.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", getpid());

    pid_t rc = fork();

    if (rc < 0) { /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("Hi, I am child (pid: %d)\n",
               getpid());
        execl("/bin/date", "date", NULL);
        printf("This line shouldn't print.\n");
    } else { /* Parent (old process) */
        int rc_wait = wait(NULL);
        printf("Hi, I am parent of %d (pid: %d)\n",
               rc_wait, getpid());
    }
    return 0;
}
```

```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am parent of 96648 (pid: 96647)
Hi, I am child (pid: 96648)
```

```
$ gcc -o p2_fork_wait p2_fork_wait.c
$ ./p2_fork_wait
Hello world (pid: 96874)
Hi, I am child (pid: 96875)
Hi, I am parent of 96875 (pid: 96874)
```

```
$ gcc -o p3_fork_wait_exec p3_fork_wait_exec.c
$ ./p3_fork_wait_exec
Hello world (pid: 97477)
Hi, I am child (pid: 97478)
Sun Mar 10 06:32:20 CST 2024
Hi, I am parent of 97478 (pid: 97477)
```



## ■ UNIX exec()

- `int exec(...);` transforms the calling process into a new process

```
/* p3_fork_wait_exec.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", getpid());

    pid_t rc = fork();

    if (rc < 0) { /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("Hi, I am child (pid: %d)\n",
            getpid());
        execl("/bin/date", "date", NULL);
        printf("This line shouldn't print.\n");
    } else { /* Parent (old process) */
        int rc_wait = wait(NULL);
        printf("Hi, I am parent of %d (pid: %d)\n",
            rc_wait, getpid());
    }

    return 0;
}
```

```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am parent of 96648 (pid: 96647)
Hi, I am child (pid: 96648)
```

```
$ gcc -o p2_fork_wait p2_fork_wait.c
$ ./p2_fork_wait
Hello world (pid: 96874)
Hi, I am child (pid: 96875)
Hi, I am parent of 96875 (pid: 96874)
```

```
$ gcc -o p3_fork_wait_exec p3_fork_wait_exec.c
$ ./p3_fork_wait_exec
Hello world (pid: 97477)
Hi, I am child (pid: 97478)
Sun Mar 10 06:32:20 CST 2024
Hi, I am parent of 97478 (pid: 97477)
```

child now runs a different program





## ■ UNIX exec()

- `int exec(...);` transforms the calling process into a new process

```
/* p3_fork_wait_exec.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", getpid());

    pid_t rc = fork();

    if (rc < 0) { /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("Hi, I am child (pid: %d)\n",
            getpid());
        execl("/bin/date", "date", NULL);
        printf("This line shouldn't print.\n");
    } else { /* Parent (old process) */
        int rc_wait = wait(NULL);
        printf("Hi, I am parent of %d (pid: %d)\n",
            rc_wait, getpid());
    }
    return 0;
}
```

```
$ gcc -o p1_fork p1_fork.c
$ ./p1_fork
Hello world (pid: 96647)
Hi, I am parent of 96648 (pid: 96647)
Hi, I am child (pid: 96648)
```

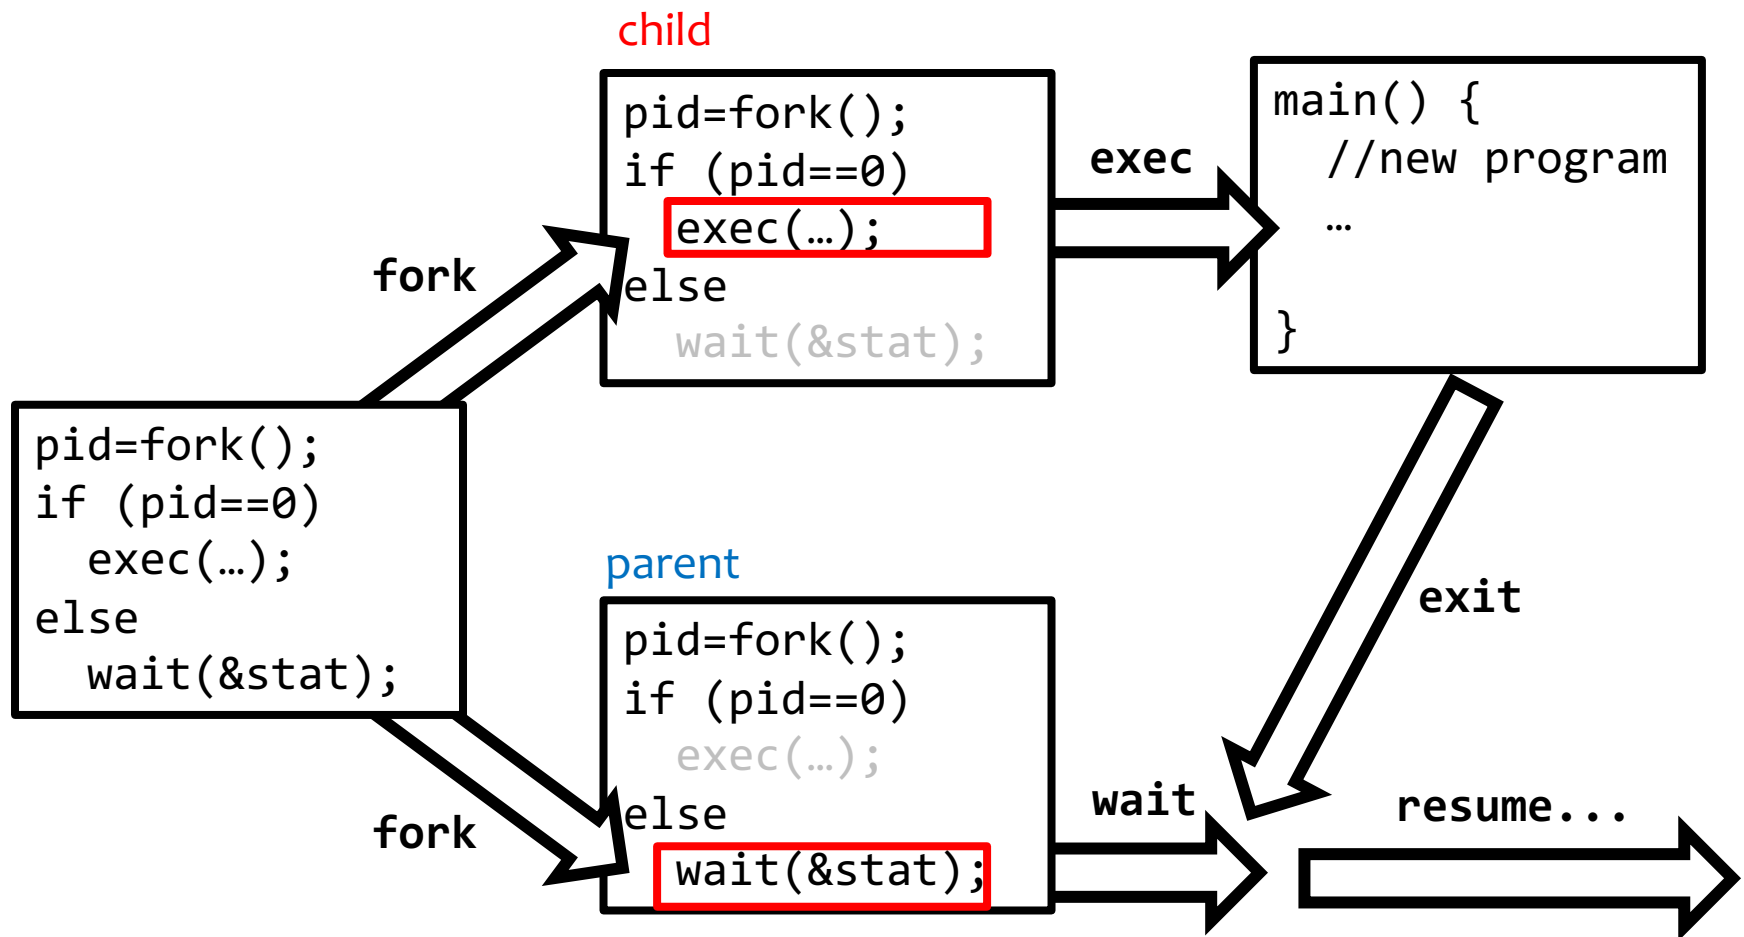
```
$ gcc -o p2_fork_wait p2_fork_wait.c
$ ./p2_fork_wait
Hello world (pid: 96874)
Hi, I am child (pid: 96875)
Hi, I am parent of 96875 (pid: 96874)
```

```
$ gcc -o p3_fork_wait_exec p3_fork_wait_exec.c
$ ./p3_fork_wait_exec
Hello world (pid: 97477)
Hi, I am child (pid: 97478)
Sun Mar 10 06:32:20 CST 2024
Hi, I am parent of 97478 (pid: 97477)
```

program code overwritten, old program instructions discarded.



## ■ UNIX fork(), exec(), wait(), exit()





## ■ UNIX `exec()`

- `int exec(...)`; transforms the calling process into a new process
- `exec()` is actually a family of functions (``man 3 exec``):
  - `execl("/bin/ls", "ls", NULL);`
  - `execvp("ls", "ls", NULL);`
  - `execle("/bin/ls", "ls", NULL, NULL);`
  - `execv("/bin/ls", param_list);`
  - `execvp("ls", param_list);`
  - `execvpe("ls", param_list, envp_list);`
- `execve()` is the actual system call (``man 2 execve``) for UNIX process creation. The family of `exec()` functions are C Library Functions that call `execve()`.
- But anyway, we still call `exec()` a **system call** for historical reasons and by convention.



## ■ Process Creation

### ■ UNIX API

- `fork()` system call creates new process
- `exec()` system call used after `fork()` to replace the process' memory space with new program
- Parent process calls `wait()` waiting for the child to terminate

### ■ Windows API

- A single `CreateProcess()` system call
  - Complex (ugly) interface



## ■ Process Creation

### ■ UNIX API

- `fork()` system call creates new process
- `exec()` system call used after `fork()` to replace the process' memory space with new program
- Parent process calls `wait()` waiting for the child to terminate

### ■ Windows API

- A single `CreateProcess()` system call
  - Complex (ugly) interface (**10 parameters**)

```
// Start the child process.
if (!CreateProcess( NULL,    // No module name (use command line)
    argv[1],          // Command line
    NULL,              // Process handle not inheritable
    NULL,              // Thread handle not inheritable
    FALSE,             // Set handle inheritance to FALSE
    0,                 // No creation flags
    NULL,              // Use parent's environment block
    NULL,              // Use parent's starting directory
    &si,                // Pointer to STARTUPINFO structure
    &pi )               // Pointer to PROCESS_INFORMATION structure
)
```



## ■ Process Termination

- Process executes last statement and then asks the OS to delete using the `exit()` system call.
  - returns status data from child to parent (via `wait()`)
  - process' resources are deallocated by OS
- Parent may terminate the execution of children processes
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - The parent is exiting
    - Some OSes do not allow a child to continue if its parent terminates
    - Cascading Termination – all children terminated



## ■ Process Termination

- Some OSes do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
  - **Cascading Termination.** All children, grandchildren, etc., are terminated
  - The termination is initiated by the Operating System.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process
  - `pid = wait(&status);`
- If no parent waiting (did not invoke `wait()`), then the terminated child process is a **zombie**
- If parent terminated without invoking `wait()`, then the (**still running**) child process is an **orphan**



## ■ Zombie (僵尸进程)

- A zombie is a process that has **completed execution** but still has an entry in the process table. It often occurs when the parent process **didn't call `wait()`** to read the exit status of the child process.

```
/* p4_zombie.c */
int main(int argc, char *argv[]) {
    pid_t rc = fork();

    if (rc < 0) { /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("    Hi, I am child (pid: %d)\n",
               getpid());
        printf("    Child process exiting...\n");
        exit(0);
    } else { /* Parent (old process) */
        printf("Hi, I am parent of %d (pid: %d)\n",
               rc, getpid());
        printf("Parent sleeping...\n");
        sleep(60);
        exit(0);
    }
    return 0;
}
```

```
$ gcc -o p4_zombie p4_zombie.c
$ ./p4_zombie
Hi, I am parent of 4120 (pid: 4118)
Parent sleeping...
    Hi, I am child (pid: 4120)
    Child process exiting...
```

```
$ ps
  PID TTY          TIME CMD
 3855 ttys010      0:00.38 -zsh
 4118 ttys011      0:00.00 ./p4_zombie
 4120 ttys011      0:00.00 <defunct>
```

Child process is a zombie





## ■ Zombie (僵尸进程)

- The `fork()` system call creates a new process (child)
- The child process **exits** immediately after printing its message
- The parent process, instead of calling `wait()`, goes to sleep for 60s.
- During the parent's sleep, the child becomes a **zombie**

```
/* p4_zombie.c */
int main(int argc, char *argv[]) {
    pid_t rc = fork();

    if (rc < 0) { /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("    Hi, I am child (pid: %d)\n",
               getpid());
        printf("    Child process exiting...\n");
        exit(0);
    } else { /* Parent (old process) */
        printf("Hi, I am parent of %d (pid: %d)\n",
               rc, getpid());
        printf("Parent sleeping...\n");
        sleep(60);
        exit(0);
    }
    return 0;
}
```

```
$ gcc -o p4_zombie p4_zombie.c
$ ./p4_zombie
Hi, I am parent of 4120 (pid: 4118)
Parent sleeping...
    Hi, I am child (pid: 4120)
    Child process exiting...
```

```
$ ps
  PID TTY          TIME CMD
 3855 ttys010      0:00.38 -zsh
 4118 ttys011      0:00.00 ./p4_zombie
 4120 ttys011      0:00.00 <defunct>
```

Child process is a zombie



## ■ Orphan (孤儿进程)

- An orphan process occurs when a child process is still running after its parent process has finished. In most OSes, when a parent process terminates, any running child processes are adopted by a special process called `init` or `systemd`, with the pid of 1.

```
/* p5_orphan.c */
int main(int argc, char *argv[]) {
    pid_t rc = fork();

    if (rc < 0) {          /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("    Hi, I am child (pid: %d, ppid: %d)\n", getpid(), getppid());
        sleep(10);
        printf("    Hi, I am child (pid: %d, ppid: %d)\n", getpid(), getppid());
    } else {               /* Parent (old process) */
        printf("Hi, I am parent of %d (pid: %d)\n", rc, getpid());
        sleep(1);
        printf("Parent exiting...\n");
        exit(0);
    }
    return 0;
}
```

```
$ gcc -o p5_orphan p5_orphan.c
$ ./p5_orphan
Hi, I am parent of 4600 (pid: 4599)
    Hi, I am child (pid: 4600, ppid: 4599)
Parent exiting...
    Hi, I am child (pid: 4600, ppid: 1)
```

After parent process terminates, the child process is adopted (领养) by `systemd` (pid:1)



## ■ Orphan (孤儿进程)

- The `fork()` system call creates a new process (child)
- The parent process prints its message, `sleep(1)`, and then exits.
- The child process prints its message and sleeps for **10s**. During child's sleep, the parent process already terminates, leaving the child an **orphan**.

```
/* p5_orphan.c */
int main(int argc, char *argv[]) {
    pid_t rc = fork();

    if (rc < 0) {          /* Fork failed */
        fprintf(stderr, "Fork Failed\n");
        exit(1);
    } else if (rc == 0) { /* Child (new process) */
        printf("    Hi, I am child (pid: %d, ppid: %d)\n", getpid(), getppid());
        sleep(10);
        printf("    Hi, I am child (pid: %d, ppid: %d)\n", getpid(), getppid());
    } else {              /* Parent (old process) */
        printf("Hi, I am parent of %d (pid: %d)\n", rc, getpid());
        sleep(1);
        printf("Parent exiting...\n");
        exit(0);
    }
    return 0;
}
```

```
$ gcc -o p5_orphan p5_orphan.c
$ ./p5_orphan
Hi, I am parent of 4600 (pid: 4599)
    Hi, I am child (pid: 4600, ppid: 4599)
Parent exiting...
    Hi, I am child (pid: 4600, ppid: 1)
```

After parent process terminates, the child process is adopted (领养) by systemd (pid:1)



**Thank you!**