



DCS216 Operating Systems

Lecture 08 Inter-process Communication (2)

Mar 20th, 2024

Instructor: Xiaoxi Zhang
Sun Yat-sen University



■ Content

- Overview
- Shared-Memory Systems
- Message-Passing Systems
- Pipes
- Communication in Client-Server Systems
 - Sockets
 - Remote Procedure Calls (RPCs)



■ Content

- Overview
- Shared-Memory Systems
- Message-Passing Systems
 - Message-Passing Design Principles
 - System V Message Queues
 - POSIX Message Queues
- Pipes
 - Unnamed Pipes
 - Named Pipes
- Communication in Client-Server Systems
 - Sockets
 - Remote Procedure Calls (RPCs)

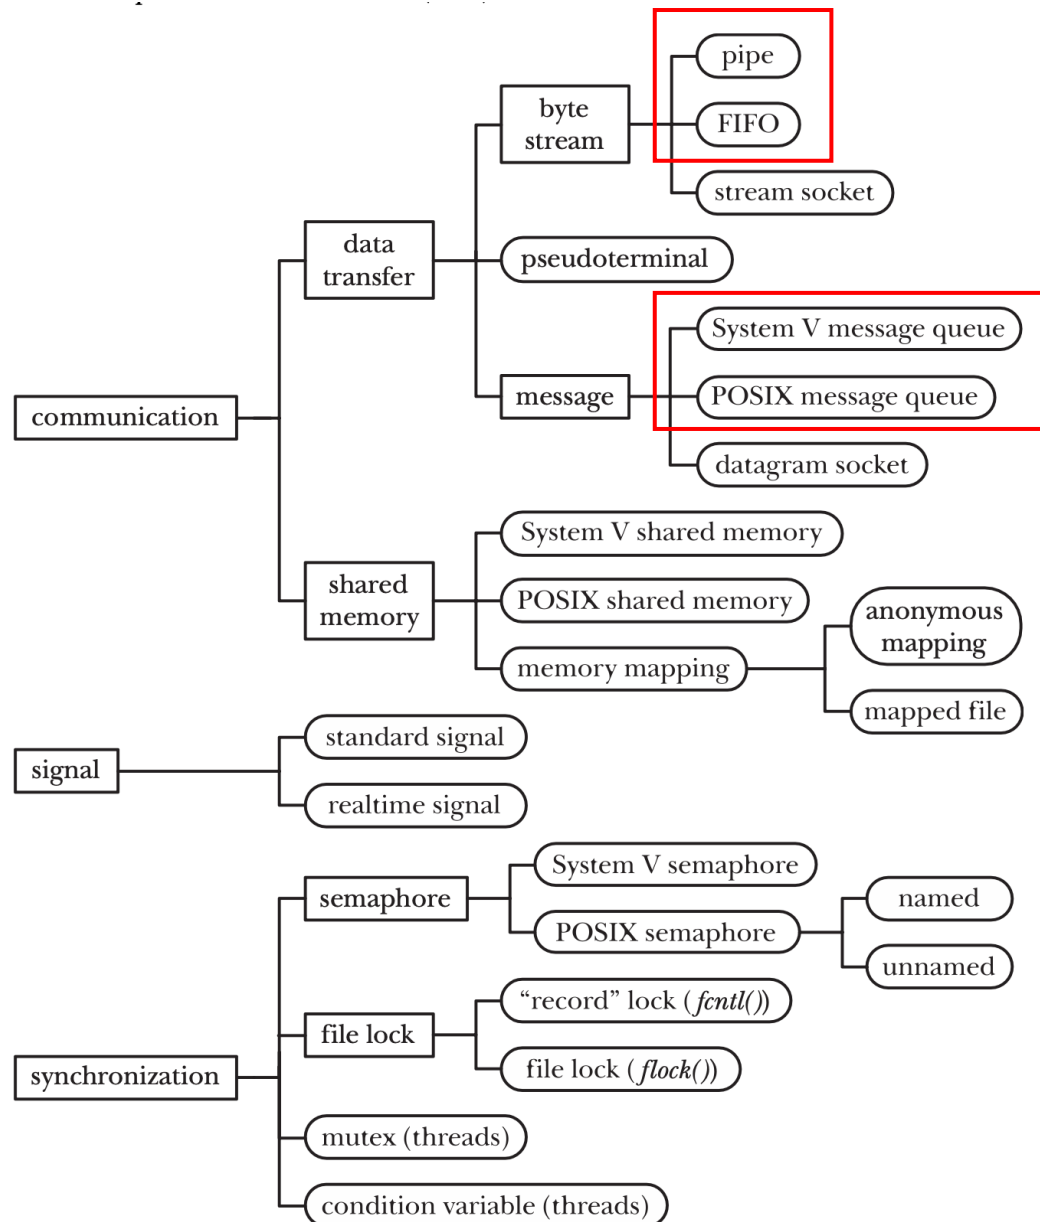


Figure 43-1: A taxonomy of UNIX IPC facilities



■ Message-Passing Systems

- Message passing provides a mechanism to allow processes to communicate and synchronize their actions **without sharing the same address space**, i.e. shared variables. (无需依靠共享变量来同步)



■ Message-Passing Systems

- Message passing provides a mechanism to allow processes to communicate and synchronize their actions **without sharing the same address space**, i.e. shared variables. (无需依靠共享变量来同步)
 - Communication takes place by means of **messages** exchanged between cooperating processes
 - Useful for exchanging **small amounts of data**
 - Typically implemented using **system calls**, and thus require more time-consuming task of kernel intervention
 - **Easier** to implement in a distributed system **than shared memory**
 - Message queues reside in **kernel space** with **race condition or resource confliction handled by the kernel**.




■ Message-Passing Systems

- Message size can be either **fixed** or **variable**
 - **Variable-sized** messages require a more **complex** system-level implementation
 - but the programming task is much **simpler**
 - **Fixed-sized** messages is **easier** to implement on the system-level
 - but it makes the task of programming more **difficult**
- Common kind of **tradeoff** seen throughout operating system design
 - Unlike **algorithm design**, which typically focuses on finding an efficient way to solve a specific problem (mathematical)
 - **System design** involves creating an architecture for software systems, which often require **balancing** various goals and constraints (e.g., scalability, reliability, maintainability, cost and performance)



■ Message-Passing Systems

- Message passing provides (at least) two basic operations:
 - **send**(message) or **send**(dest, message)
 - **receive**(message) or **receive**(src, message)






■ Message Format

■ Header

- Message Type
- Message Length
- Destination ID & Source ID
- Control Information
 - What to do if run out of buffer space?
 - Sequence numbers
 - Priority



■ Body

- Message Content

■ Queueing discipline: usually FIFO, but can also include priorities

■ Example: a message format containing only message type and message content:

```
struct message {  
    long mtype;      // message type  
    char mtext[80];  // message text  
};
```



■ Message-Passing Systems

- If processes P and Q wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via **send/receive**
- Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communication processes?
 - What is the capacity of a link?
 - Is the message size fixed or variable?
 - Is a link unidirectional or bidirectional?



■ Message-Passing Systems

- Logical Implementation of Communication Link
 - **Direct** or **Indirect** Communication
 - **Synchronous** or **Asynchronous**
 - **Automatic** or **Explicit** buffering



■ Direct Communication

- Processes must name each other explicitly:
 - **send**(P, message) – send a message to process P
 - **receive**(Q, message) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bidirectional



■ Indirect Communication

- Messages are directed and received from mailboxes (ports)
 - Each mailbox has a unique ID
 - Processes can communicate only if they share a mailbox (port)
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bidirectional



■ Indirect Communication

- Messages are directed and received from mailboxes (ports)
 - Each mailbox has a unique ID
 - Processes can communicate only if they share a mailbox (port)
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bidirectional
- Operations
 - Create a new mailbox (port)
 - Send and receive messages through mailbox
 - Delete a mailbox
- Primitives:
 - **send**(A, message) – send a message to mailbox A
 - **receive**(A, message) – receive a message from mailbox A



■ Indirect Communication

■ Primitives:

- **send**(A, message) – send a message to mailbox A
- **receive**(A, message) – receive a message from mailbox A

■ Mailbox sharing

- P₁, P₂, and P₃ share mailbox A
- P₁ sends;
- P₂ and P₃ receive
- Who gets the message?

■ Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



■ Synchronization

- Message passing may be either blocking or non-blocking
 - **Blocking** (阻塞) is considered **synchronous** (同步)
 - Blocking send – the sender is blocked until the message is received
 - Blocking receive – the receiver is blocked until message is available
 - **Non-blocking** (非阻塞) is considered **asynchronous** (异步)
 - Non-blocking send – the sender sends the message and continue
 - Non-blocking receive – the receiver receives:
 - A valid message, or
 - Null message
 - Different combinations possible
 - If both send and receive are blocking, we have a **rendezvous** (汇聚点).



■ Buffering

- Queue of messages attached to the link.
- Implemented in one of three ways:
 - **Zero Capacity** – no messages are queued on a link
 - Sender must **wait** for receiver (rendezvous)
 - **Bounded Capacity** – finite length of N messages
 - Sender must **wait** if link is full
 - **Unbounded Capacity** – infinite length
 - Sender **never waits**



■ IPC Models

■ Shared Memory

- Direct Sharing (**System V Standard**) **System** Calls:
 - `shmget()`, `shmat()`, `shmdt()`, `shmctl()`
- Indirect Sharing (**POSIX Standard**) **Library** Calls:
 - `shm_open()`, `shm_unlink()`, `ftruncate()`, `mmap()`

■ Message Passing

■ Pipes

- Unnamed Pipe: `pipe()`
- Named Pipe: `mkfifo()`

■ Message Queues

- System V: `msgget()`, `msgsnd()`, `msgrcv()`, `msgctl()`
- POSIX: `mq_open()`, `mq_close()`, `mq_send()`, `mq_receive()`

- Sockets: `socket()`, `bind()`, `listen()`, `accept()`, `connect()`,
`send()`, `recv()`

- Signals: `signal()`, `sigaction()`



■ IPC Models

■ Shared Memory

- Direct Sharing (System V Standard) System Calls:
 - `shmget()`, `shmat()`, `shmdt()`, `shmctl()`
- Indirect Sharing (POSIX Standard) Library Calls:
 - `shm_open()`, `shm_unlink()`, `ftruncate()`, `mmap()`

■ Message Passing

■ Pipes

- Unnamed Pipe: `pipe()`
- Named Pipe: `mkfifo()`

■ Message Queues

- System V: `msgget()`, `msgsnd()`, `msgrcv()`, `msgctl()`
- POSIX: `mq_open()`, `mq_close()`, `mq_send()`, `mq_receive()`

- Sockets: `socket()`, `bind()`, `listen()`, `accept()`, `connect()`,
`send()`, `recv()`

- Signals: `signal()`, `sigaction()`



■ System V Message Queues API

- **Get** a System V message queue identifier

- `int msgget(key_t key, int msgflg);`

- **Send** data into a message queue

- `int msgsnd(int msqid, const void *msgp, size_t msgsz, int msgflg);`

- **Receive** data from a message queue

- `ssize_t msgrcv(int msqid, void *msgp, size_t msgsz, long msgtyp, int msgflg);`

- **Perform** various operations on a message queue

- `int msgctl(int msqid, int cmd, struct msqid_ds *buf);`
 - E.g., ``msgctl(msgid, IPC_RMID, NULL)`` **removes** the message queue ``msgid``

- **CLI Utility:** Show Message Queue Info

- ``ipcs -q``



■ System V Message Queues Example

<pre>/* sender.c */ #include "message.h" int main() { int msgid; struct message msg; // Create a message queue msgid = msgget(KEY, 0666 IPC_CREAT); // Prepare the message msg.mtype = 1; // Set the message type strcpy(msg.mtext, "Hello Message..."); // Send the message msgsnd(msgid, &msg, sizeof(msg.mtext), 0); printf("Sent message: %s\n", msg.mtext); return 0; }</pre>	<pre>/* receiver.c */ #include "message.h" int main() { int msgid; struct message msg; // Access the message queue msgid = msgget(KEY, 0666); // Receive the message msgrcv(msgid, &msg, sizeof(msg.mtext), 0, 0); printf("Recv message: %s\n", msg.mtext); // Optionally, remove the message queue msgctl(msgid, IPC_RMID, NULL); return 0; }</pre>
<pre>/* message.h */ #include <stdio.h> #include <stdlib.h> #include <string.h> #include <sys/ipc.h> #include <sys/msg.h> #define KEY 0x42 struct message { long mtype; // message type char mtext[80]; // message text };</pre>	



■ System V Message Queues Example

```
/* sender.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Create a message queue
    msgid = msgget(KEY, 0666 | IPC_CREAT);
    // Prepare the message
    msg.mtype = 1; // Set the message type
    strcpy(msg.mtext, "Hello Message...");
    // Send the message
    msgsnd(msgid, &msg, sizeof(msg.mtext), 0);
    printf("Sent message: %s\n", msg.mtext);
    return 0;
}
```

```
/* message.h */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#define KEY 0x42
struct message {
    long mtype; // message type
    char mtext[80]; // message text
};
```

```
/* receiver.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Access the message queue
    msgid = msgget(KEY, 0666);
    // Receive the message
    msgrcv(msgid, &msg, sizeof(msg.mtext), 0,
           0);
    printf("Recv message: %s\n", msg.mtext);
    // Optionally, remove the message queue
    msgctl(msgid, IPC_RMID, NULL);
    return 0;
}
```

```
$ ./sender
Sent message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
0x00000042   1        zxx      666      80            1
$ ./receiver
Recv message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
```



■ System V Message Queues Example

```
/* sender.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Create a message queue
    msgid = msgget(KEY, 0666 | IPC_CREAT);
    // Prepare the message
    msg.mtype = 1; // Set the message type
    strcpy(msg.mtext, "Hello Message...");
    // Send the message
    msgsnd(msgid, &msg, sizeof(msg.mtext), 0);
    printf("Sent message: %s\n", msg.mtext);
    return 0;
}
```

```
/* message.h */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#define KEY 0x42
struct message {
    long mtype; // message type
    char mtext[80]; // message text
};
```

```
/* receiver.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Access the message queue
    msgid = msgget(KEY, 0666);
    // Receive the message
    msgrcv(msgid, &msg, sizeof(msg.mtext), 0,
           0);
    printf("Recv message: %s\n", msg.mtext);
    // Optionally, remove the message queue
    msgctl(msgid, IPC_RMID, NULL);
    return 0;
}
```

```
$ ./sender
Sent message: Hello Message...
$ ipcs -q
----- Message Queues -----
key      msqid    owner    perms    used-bytes    messages
0x00000042 1        zxx      666      80            1
$ ./receiver
Recv message: Hello Message...
$ ipcs -q
----- Message Queues -----
key      msqid    owner    perms    used-bytes    messages
```



■ System V Message Queues Example

```
/* sender.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Create a message queue
    msgid = msgget(KEY, 0666 | IPC_CREAT);
    // Prepare the message
    msg.mtype = 1; // Set the message type
    strcpy(msg.mtext, "Hello Message...");
    // Send the message
    msgsnd(msgid, &msg, sizeof(msg.mtext), 0);
    printf("Sent message: %s\n", msg.mtext);
    return 0;
}
```

```
/* message.h */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#define KEY 0x42
struct message {
    long mtype; // message type
    char mtext[80]; // message text
};
```

```
/* receiver.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Access the message queue
    msgid = msgget(KEY, 0666);
    // Receive the message
    msgrcv(msgid, &msg, sizeof(msg.mtext), 0,
           0);
    printf("Recv message: %s\n", msg.mtext);
    // Optionally, remove the message queue
    msgctl(msgid, IPC_RMID, NULL);
    return 0;
}
```

```
$ ./sender
Sent message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
0x00000042   1        zxx      666      80            1
$ ./receiver
Recv message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
```




■ System V Message Queues Example

```
/* sender.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Create a message queue
    msgid = msgget(KEY, 0666 | IPC_CREAT);
    // Prepare the message
    msg.mtype = 1; // Set the message type
    strcpy(msg.mtext, "Hello Message...");
    // Send the message
    msgsnd(msgid, &msg, sizeof(msg.mtext), 0);
    printf("Sent message: %s\n", msg.mtext);
    return 0;
}
```

```
/* message.h */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#define KEY 0x42
struct message {
    long mtype; // message type
    char mtext[80]; // message text
};
```

```
/* receiver.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Access the message queue
    msgid = msgget(KEY, 0666);
    // Receive the message
    msgrcv(msgid, &msg, sizeof(msg.mtext), 0, 0);
    printf("Recv message: %s\n", msg.mtext);
    // Optionally, remove the message queue
    msgctl(msgid, IPC_RMID, NULL);
    return 0;
}
```

```
$ ./sender
Sent message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
0x00000042    1        zxx      666      80            1

$ ./receiver
Recv message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
```



■ System V Message Queues Example

```
/* sender.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Create a message queue
    msgid = msgget(KEY, 0666 | IPC_CREAT);
    // Prepare the message
    msg.mtype = 1; // Set the message type
    strcpy(msg.mtext, "Hello Message...");
    // Send the message
    msgsnd(msgid, &msg, sizeof(msg.mtext), 0);
    printf("Sent message: %s\n", msg.mtext);
    return 0;
}
```

```
/* message.h */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#define KEY 0x42
struct message {
    long mtype; // message type
    char mtext[80]; // message text
};
```

Sender and receiver
don't need to run
concurrently. Msg
stays in the kernel
buffer...

```
/* receiver.c */
#include "message.h"
int main() {
    int msgid;
    struct message msg;
    // Access the message queue
    msgid = msgget(KEY, 0666);
    // Receive the message
    msgrcv(msgid, &msg, sizeof(msg.mtext), 0,
           0);
    printf("Recv message: %s\n", msg.mtext);
    // Optionally, remove the message queue
    msgctl(msgid, IPC_RMID, NULL);
    return 0;
}
```

```
$ ./sender
Sent message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
0x00000042   1        zxx      666      80            1
$ ./receiver
Recv message: Hello Message...
$ ipcs -q
----- Message Queues -----
key          msqid    owner    perms    used-bytes    messages
```




■ POSIX Message Queues API

- Similar to File I/O API:
 - `open()`, `close()`, `read()`, `write()`
- Open a message queue descriptor
 - `mqd_t mq_open(const char *name, int oflag);`
- Close a message queue descriptor
 - `int mq_close(mqd_t mqdes);`
- Send a message to a message queue
 - `int mq_send(mqd_t mqdes, const char *msg_ptr, size_t msg_len, unsigned int msg_prio);`
- Receive a message from a message queue
 - `ssize_t mq_receive(mqd_t mqdes, char *msg_ptr, size_t msg_len, unsigned int *msg_prio);`
- See ``man mq_overview``

■ Pipes

- Pipes are the **oldest** form of IPC and provided by **all** UNIX systems.
 - Historically, pipes have been **half duplex** (i.e., data flows in only one direction). Some systems provide full-duplex pipes, but not recommended.
 - Pipes can be used only between parent and child processes.
 - Despite these limitations, half-duplex pipes are still the most commonly used form of IPC
- A pipe is one form of Message Passing IPC.



Pipes


- A pipe is created by calling the `pipe()` function.
 - `int pipe(int fd[2]);`
- Two file descriptors are returned through the `fd` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#define MAXLINE 80

int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    write(fd[1], "Hello, world!", 14);
    len = read(fd[0], buf, MAXLINE);
    printf("buf(%d): %s\n", len, buf);
    return 0;
}
```

```
$ gcc -o pipe1 pipe1.c
$ ./pipe1
buf(14): Hello, world!
```



■ Pipes


- Two file descriptors are returned through the `fd` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing
- A pipe **in a single process** is next to useless.

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#define MAXLINE 80

int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    write(fd[1], "Hello, world!", 14);
    len = read(fd[0], buf, MAXLINE);
    printf("buf(%d): %s\n", len, buf);
    return 0;
}
```

```
$ gcc -o pipe1 pipe1.c
$ ./pipe1
buf(14): Hello, world!
```




■ Pipes

- Two file descriptors are returned through the `fd` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing
- A pipe **in a single process** is next to useless.
- followed by `fork()` to create IPC channel between parent & child

```
int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    pid_t rc = fork();
    if (rc == 0) {          /* Child */
        close(fd[0]);
        write(fd[1], "Hello, world!", 14);
    } else {               /* Parent */
        close(fd[1]);
        len = read(fd[0], buf, MAXLINE);
        printf("buf(%d): %s\n", len, buf);
    }
    return 0;
}
```




Pipes

- Two file descriptors are returned through the `fd` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing
- A pipe **in a single process** is next to useless.
- followed by `fork()` to create IPC channel between parent & child

```
int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    pid_t rc = fork();
    if (rc == 0) {          /* Child */
        close(fd[0]);
        write(fd[1], "Hello, world!", 14);
    } else {                /* Parent */
        close(fd[1]);
        len = read(fd[0], buf, MAXLINE);
        printf("buf(%d): %s\n", len, buf);
    }
    return 0;
}
```

```
$ gcc -o pipe2 pipe2.c
$ ./pipe2
buf(14): Hello, world!
```




■ Pipes

- Two file descriptors are returned through the ``fd`` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing
- A pipe **in a single process** is next to useless.
- followed by ``fork()`` to create IPC channel between parent & child

```
int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    pid_t rc = fork();
    if (rc == 0) {          /* Child */
        close(fd[0]);
        write(fd[1], "Hello, world!", 14);
    } else {               /* Parent */
        close(fd[1]);
        len = read(fd[0], buf, MAXLINE);
        printf("buf(%d): %s\n", len, buf);
    }
    return 0;
}
```

```
$ gcc -o pipe2 pipe2.c
$ ./pipe2
buf(14): Hello, world!
```




■ Pipes

- Two file descriptors are returned through the ``fd`` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing
- A pipe **in a single process** is next to useless.
- followed by ``fork()`` to create IPC channel between parent & child

```
int main() {  
    int len;  
    int fd[2];  
    char buf[MAXLINE];  
  
    pipe(fd);  
    pid_t rc = fork();  
    if (rc == 0) {          /* Child */  
        close(fd[0]);  
        write(fd[1], "Hello, world!", 14);  
    } else {                /* Parent */  
        close(fd[1]);  
        len = read(fd[0], buf, MAXLINE);  
        printf("buf(%d): %s\n", len, buf);  
    }  
    return 0;  
}
```

```
$ gcc -o pipe2 pipe2.c  
$ ./pipe2  
buf(14): Hello, world!
```




Pipes

- Two file descriptors are returned through the `fd` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing
- A pipe **in a single process** is next to useless.
- followed by `fork()` to create IPC channel between parent & child

```
int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    pid_t rc = fork();
    if (rc == 0) { /* Child */
        close(fd[0]);
        write(fd[1], "Hello, world!", 14);
    } else { /* Parent */
        close(fd[1]);
        len = read(fd[0], buf, MAXLINE);
        printf("buf(%d): %s\n", len, buf);
    }
    return 0;
}
```

```
$ gcc -o pipe2 pipe2.c
$ ./pipe2
buf(14): Hello, world!
```




Pipes

- Two file descriptors are returned through the `fd` argument:
 - `fd[0]` is open for reading
 - `fd[1]` is open for writing
- A pipe **in a single process** is next to useless.
- followed by `fork()` to create IPC channel between parent & child

```
int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    pid_t rc = fork();
    if (rc == 0) { /* Child */
        // close(fd[0]);
        write(fd[1], "Hello, world!", 14);
    } else { /* Parent */
        // close(fd[1]);
        len = read(fd[0], buf, MAXLINE);
        printf("buf(%d): %s\n", len, buf);
    }
    return 0;
}
```

It might seem OK to leave `fd[0]` in **child** and `fd[1]` in **parent** alone without closing them. But it's **recommended** to **close unused file descriptors** to ensure proper pipe **termination**.




■ Pipes

- A pipe is created by calling the `pipe()` function.
- followed by `fork()` to create IPC channel between parent & child
- also called **Unnamed Pipes/Ordinary Pipes**: cannot be accessed from outside the process that created it.

```
int main() {
    int len;
    int fd[2];
    char buf[MAXLINE];

    pipe(fd);
    pid_t rc = fork();
    if (rc == 0) { /* Child */
        close(fd[0]);
        write(fd[1], "Hello, world!", 14);
    } else { /* Parent */
        close(fd[1]);
        len = read(fd[0], buf, MAXLINE);
        printf("buf(%d): %s\n", len, buf);
    }
    return 0;
}
```

```
$ gcc -o pipe2 pipe2.c
$ ./pipe2
buf(14): Hello, world!
```




Pipes

- A pipe is created by calling the `pipe()` function.
- followed by `fork()` to create IPC channel between parent & child
- also called **Unnamed Pipes/Ordinary Pipes**: cannot be accessed from outside the process that created it.

```
int main() {  
    int len;  
    int fd[2];  
    char buf[MAXLINE];  
  
    pipe(fd);  
    pid_t rc = fork();  
    if (rc == 0) { /* Child */  
        close(fd[0]);  
        write(fd[1], "Hello, world!", 14);  
    } else { /* Parent */  
        close(fd[1]);  
        len = read(fd[0], buf, MAXLINE);  
        printf("buf(%d): %s\n", len, buf);  
    }  
    return 0;  
}
```

Both `write()` and `read()` are **blocking** function calls, but eventually they will have a **rendezvous**(汇聚点). Hence no need to use `sleep()` or `wait()`

```
$ gcc -o pipe2 pipe2.c  
$ ./pipe2  
buf(14): Hello, world!
```



■ Pipes

- Consider the following program. How to make sure child prints first?

```
int main(int argc, char **argv) {  
    int rc = fork();  
    if (rc < 0) {  
        fprintf(stderr, "fork failed.\n");  
        exit(1);  
    } else if (rc == 0) {  
        printf("hello\n");  
    } else {  
        printf("goodbye\n");  
    }  
    return 0;  
}
```

■ Pipes


- Consider the following program. How to make sure child prints first?
 - Can you do this **without** calling ``wait()`` in parent?

```
int main(int argc, char **argv) {  
    int rc = fork();  
    if (rc < 0) {  
        fprintf(stderr, "fork failed.\n");  
        exit(1);  
    } else if (rc == 0) {  
        printf("hello\n");  
    } else {  
        wait(NULL);  
        printf("goodbye\n");  
    }  
    return 0;  
}
```


■ Pipes

- Consider the following program. How to make sure child prints first?
 - Use ``pipe()`. Make use of `read()`. and `write()`. rendezvous!`

```
int main(int argc, char **argv) {
    int fd[2];
    int rc = fork();
    if (rc < 0) {
        fprintf(stderr, "fork failed.\n");
        exit(1);
    } else if (rc == 0) {
        close(fd[0]);
        printf("hello\n");
        write(fd[1], "k", 1);
        close(fd[1]);
    } else {
        close(fd[1]);
        char buf[2];
        read(fd[0], buf, 1);
        printf("goodbye\n");
        close(fd[0]);
    }
    return 0;
}
```




Pipes

- Consider the following program. How to make sure child prints first?
 - Use `pipe()`. Make use of `read()` and `write()` **rendezvous!**

```
int main(int argc, char **argv) {
    int fd[2];
    int rc = fork();
    if (rc < 0) {
        fprintf(stderr, "fork failed.\n");
        exit(1);
    } else if (rc == 0) {
        close(fd[0]);
        printf("hello\n");
        write(fd[1], "k", 1);
        close(fd[1]);
    } else {
        close(fd[1]);
        char buf[2];
        read(fd[0], buf, 1);
        printf("goodbye\n");
        close(fd[0]);
    }
    return 0;
}
```

Blocking calls



■ Pipes

- Example (Shell Commands): **List** processes of current user and **sort** them by their PID using pipes (``cmd1 | cmd2``)


```
$ ps
  PID TTY          TIME CMD
1464792 pts/245    00:00:00 ps
1464794 pts/245    00:00:00 sort
 693845 pts/245    00:00:00 bash
1464793 pts/245    00:00:00 tee
1464791 pts/245    00:00:00 wc
```



Pipes

- Example (Shell Commands): **List** processes of current user and **sort** them by their PID using pipes (``cmd1 | cmd2``)



```
$ ps
  PID TTY          TIME CMD
1464792 pts/245    00:00:00 ps
1464794 pts/245    00:00:00 sort
 693845 pts/245    00:00:00 bash
1464793 pts/245    00:00:00 tee
1464791 pts/245    00:00:00 wc
$ ps | sort -k1n
  PID TTY          TIME CMD
 693845 pts/245    00:00:00 bash
1464791 pts/245    00:00:00 wc
1464792 pts/245    00:00:00 ps
1464793 pts/245    00:00:00 tee
1464794 pts/245    00:00:00 sort
```



Pipes

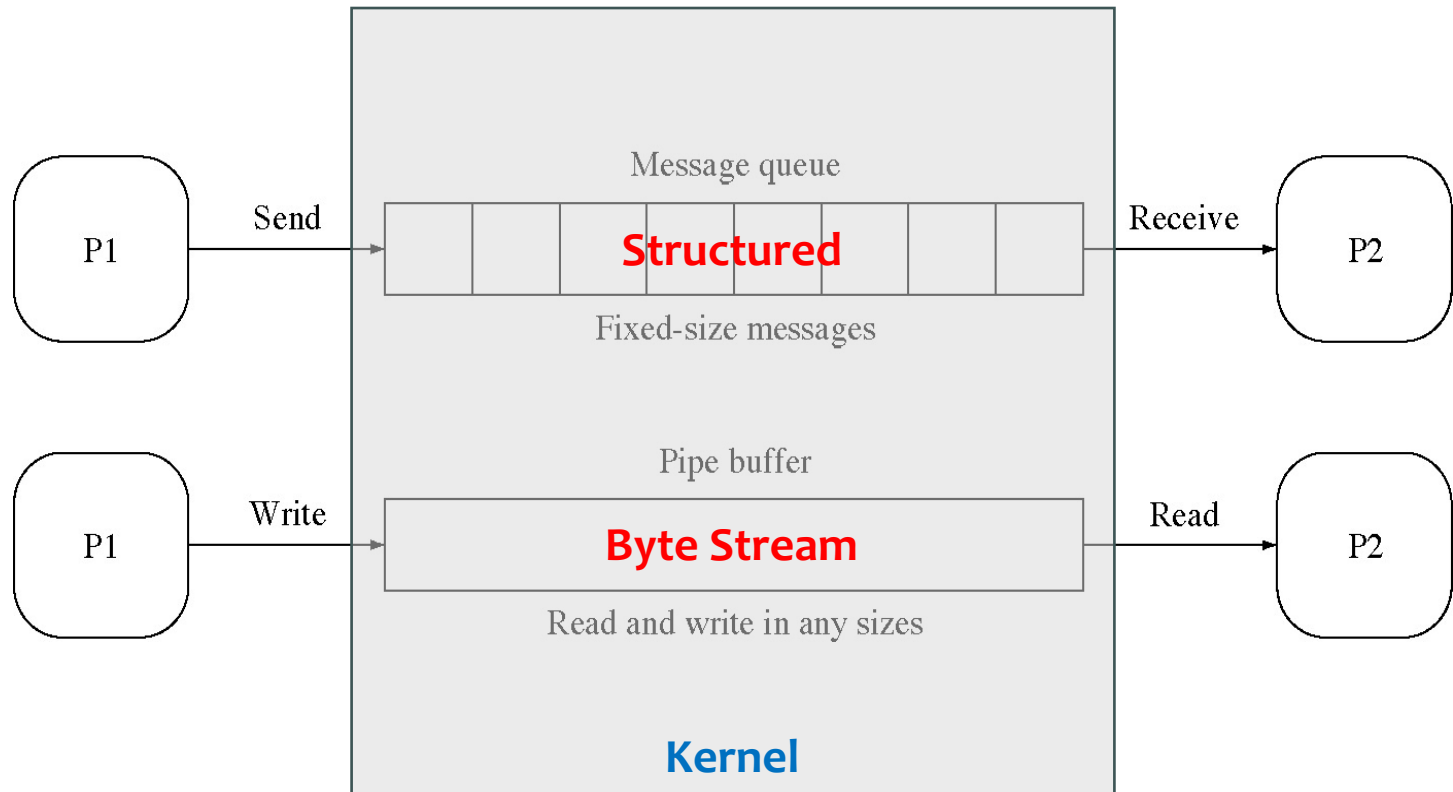
- Example (Shell Commands): **List** processes of current user and **sort** them by their PID using pipes (``cmd1 | cmd2``)

```
$ ps
  PID TTY          TIME CMD
1464792 pts/245    00:00:00 ps
1464794 pts/245    00:00:00 sort
 693845 pts/245    00:00:00 bash
1464793 pts/245    00:00:00 tee
1464791 pts/245    00:00:00 wc
$ ps | sort -k1n
  PID TTY          TIME CMD
 693845 pts/245    00:00:00 bash
1464791 pts/245    00:00:00 wc
1464792 pts/245    00:00:00 ps
1464793 pts/245    00:00:00 tee
1464794 pts/245    00:00:00 sort
$ ps | wc -l
6
```



■ Pipes

■ Message Queues vs Pipes



■ Pipes

- Acts as a conduit (管道) allowing two processes to communicate
- **Unnamed Pipes/Ordinary Pipes:** cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
 - Windows calls these *Anonymous Pipes*.
- Limitations:
 - Ordinary pipes exist only while processes are communicating.
 - If either process terminates, the pipe no longer exists.
 - Ordinary pipes are unidirectional and half-duplex.
 - Ordinary pipes require parent-child relationship.
 - Synchronous and **blocking** operations: reads and writes to pipes are generally blocking, e.g., a process that reads from an empty pipe will be blocked until there is data to read.

■ Pipes


- Acts as a conduit (管道) allowing two processes to communicate
- **Unnamed Pipes/Ordinary Pipes:** cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
 - Windows calls these *Anonymous Pipes*.
- **Named Pipes(命名管道):** provides a much more **powerful** communication tool compared with Ordinary Pipes. Named Pipes are also called **FIFOs** in UNIX systems.
- Named pipes extend the concept of Ordinary Pipes by providing a named, persistent endpoint in the file system that can be accessed by multiple, potentially unrelated processes.

■ Named Pipes

- Named pipes are referred to as **FIFOs** in UNIX systems.
- Once created, they appear as typical **files** in the file system.
- A FIFO is created with the ``mkfifo()`` system call
 - and manipulated with ordinary ``open()``, ``read()``, ``write()``, ``close()`` system calls
- A FIFO continues to exist until it is explicitly deleted from the file system.
 - Named pipes are **persistent** (cont. to exist after processes terminate).
 - Named pipes are **bidirectional**
 - although bidirectional, only **half-duplex** transmission is permitted
 - ..to achieve **full-duplex**, two **FIFOs** are typically used.
 - Named pipes **don't require parent-child relationship**
 - allows several readers and writers
 - Named pipes can be configured as **non-blocking** (defaults to blocking)
 - Named pipes work for processes on the same machine
 - ..for communication over networks, **sockets** must be used

Named Pipes

- Example (Shell Commands): Suppose we want to process the output of ``ps`` **twice** using ``sort`` and ``wc`` without using a temp file.
 - Since ``ps`` output changes all the time, if we execute ``ps | sort -k1n`` and ``ps | wc -l`` separately, we are not dealing with the same ``ps`` output.



■ Named Pipes

- Example (Shell Commands): Suppose we want to process the output of ``ps`` **twice** using ``sort`` and ``wc`` without using a temp file.

```
$ mkfifo fifo1
```

FIFO

Named Pipes

- Example (Shell Commands): Suppose we want to process the output of `ps` twice using `sort` and `wc` without using a temp file.

```
$ mkfifo fifo1
$ ls -l fifo1
prw-rw-r-- 1 zxx zxx 0 Mar 19 08:39 fifo1
$ file fifo1
fifo1: fifo (named pipe)
```

'p' in the file permission bits indicates Pipe.


FIFO

■ Named Pipes

- Example (Shell Commands): Suppose we want to process the output of ``ps`` twice using ``sort`` and ``wc`` without using a temp file.

```
$ mkfifo fifo1  
$ wc -l < fifo1 &  
[1] 1464791
```

< **redirects** the output of `fifo1` to ``wc`` (as its stdin)
& **indicates** that ``wc`` runs in the background, so that we can run another program ``sort`` in the foreground




Named Pipes

- Example (Shell Commands): Suppose we want to process the output of `ps` twice using `sort` and `wc` without using a temp file.

```
$ mkfifo fifo1  
$ wc -l < fifo1 &  
[1] 1464791
```

`wc` runs in the background with job id [1]. We can use `jobs` to list background processes.


< redirects the output of `fifo1` to `wc` (as its stdin)
& indicates that `wc` runs in the background, so that we can run another program `sort` in the foreground



Named Pipes

- Example (Shell Commands): Suppose we want to process the output of `ps` twice using `sort` and `wc` without using a temp file.


```
$ mkfifo fifo1
$ wc -l < fifo1 &
[1] 1464791
$ ps | tee fifo1 | sort -k1n
    PID TTY          TIME CMD
  693845 pts/245    00:00:00 bash
 1464791 pts/245    00:00:00 wc
 1464792 pts/245    00:00:00 ps
 1464793 pts/245    00:00:00 tee
 1464794 pts/245    00:00:00 sort
      6      24      180
```



Named Pipes

- Example (Shell Commands): Suppose we want to process the output of `ps` twice using `sort` and `wc` without using a temp file.

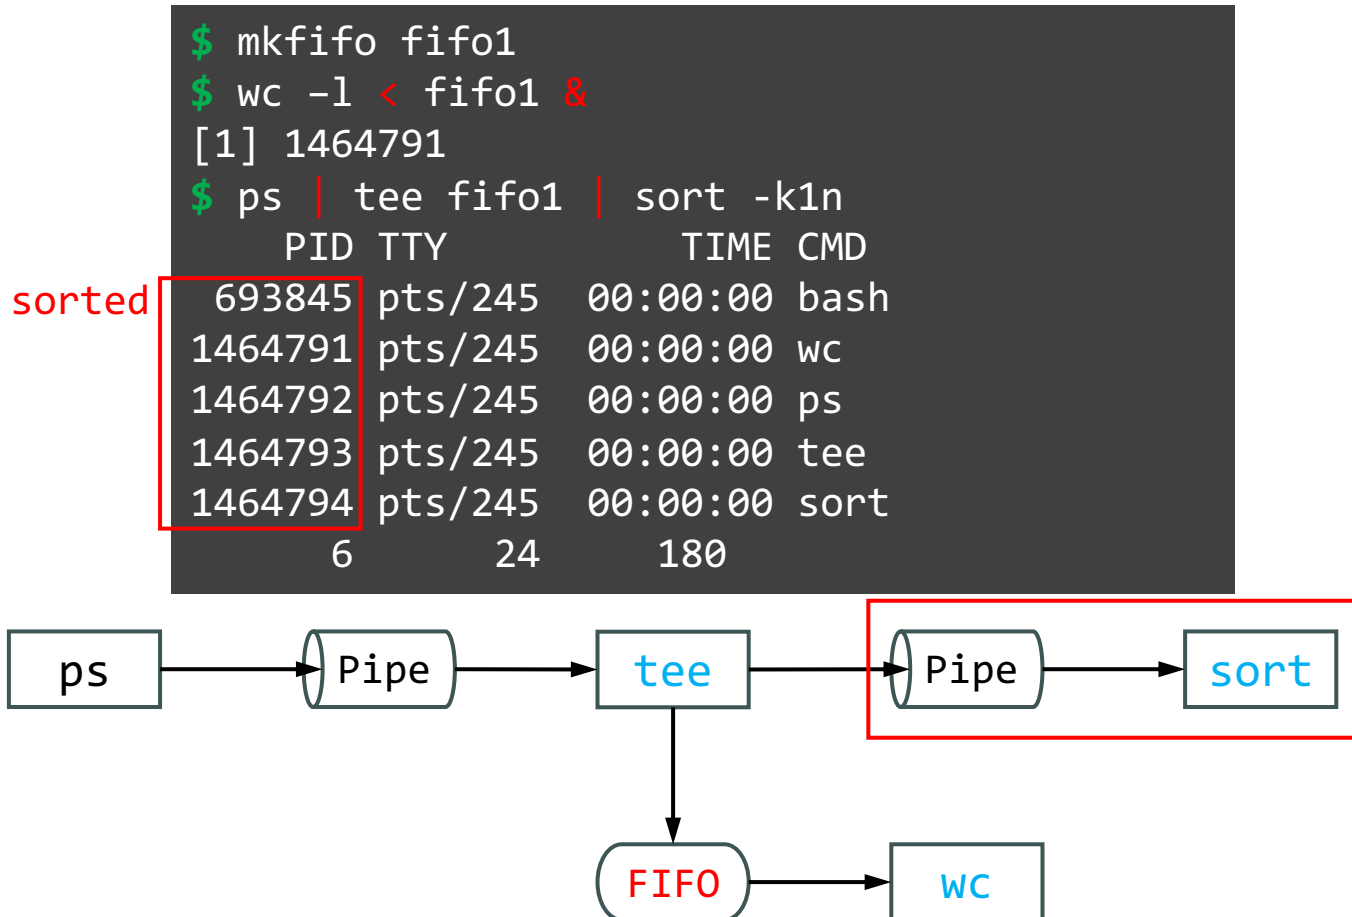
```
$ mkfifo fifo1
$ wc -l < fifo1 &
[1] 1464791
$ ps | tee fifo1 | sort -k1n
      PID TTY          TIME CMD
  693845 pts/245    00:00:00 bash
 1464791 pts/245    00:00:00 wc
 1464792 pts/245    00:00:00 ps
 1464793 pts/245    00:00:00 tee
 1464794 pts/245    00:00:00 sort
      6      24      180
```



The program `tee file` duplicates **input** to **output** and **file**, hence the shape **T**

Named Pipes

- Example (Shell Commands): Suppose we want to process the output of `ps` twice using `sort` and `wc` without using a temp file.




Named Pipes

- Example (Shell Commands): Suppose we want to process the output of `ps` twice using `sort` and `wc` without using a temp file.

```
$ mkfifo fifo1
$ wc -l < fifo1 &
[1] 1464791
$ ps | tee fifo1 | sort -k1n
      PID TTY          TIME CMD
  693845 pts/245    00:00:00 bash
 1464791 pts/245    00:00:00 wc
 1464792 pts/245    00:00:00 ps
 1464793 pts/245    00:00:00 tee
 1464794 pts/245    00:00:00 sort
```

`wc -l` output





Thank you!