



# DCS216 Operating Systems

## Lecture 10 Threads (1)

**Mar 27<sup>th</sup>, 2024**

**Instructor: Xiaoxi Zhang**  
**Sun Yat-sen University**

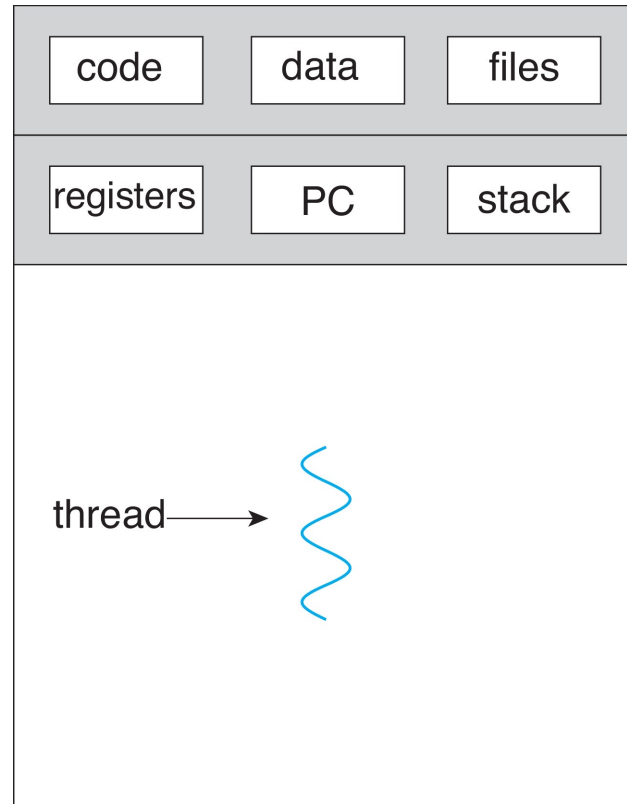


## ■ Content

- Overview
- Multicore Programming
- User Threads and Kernel Threads
- Multithreading Models
- Thread Libraries
  - POSIX Pthreads
- Implicit Threading
  - Thread Pools
  - OpenMP
- Threading Issues
- Examples
  - Linux `clone()`
  - Windows Threads

## ■ Single-threaded Process

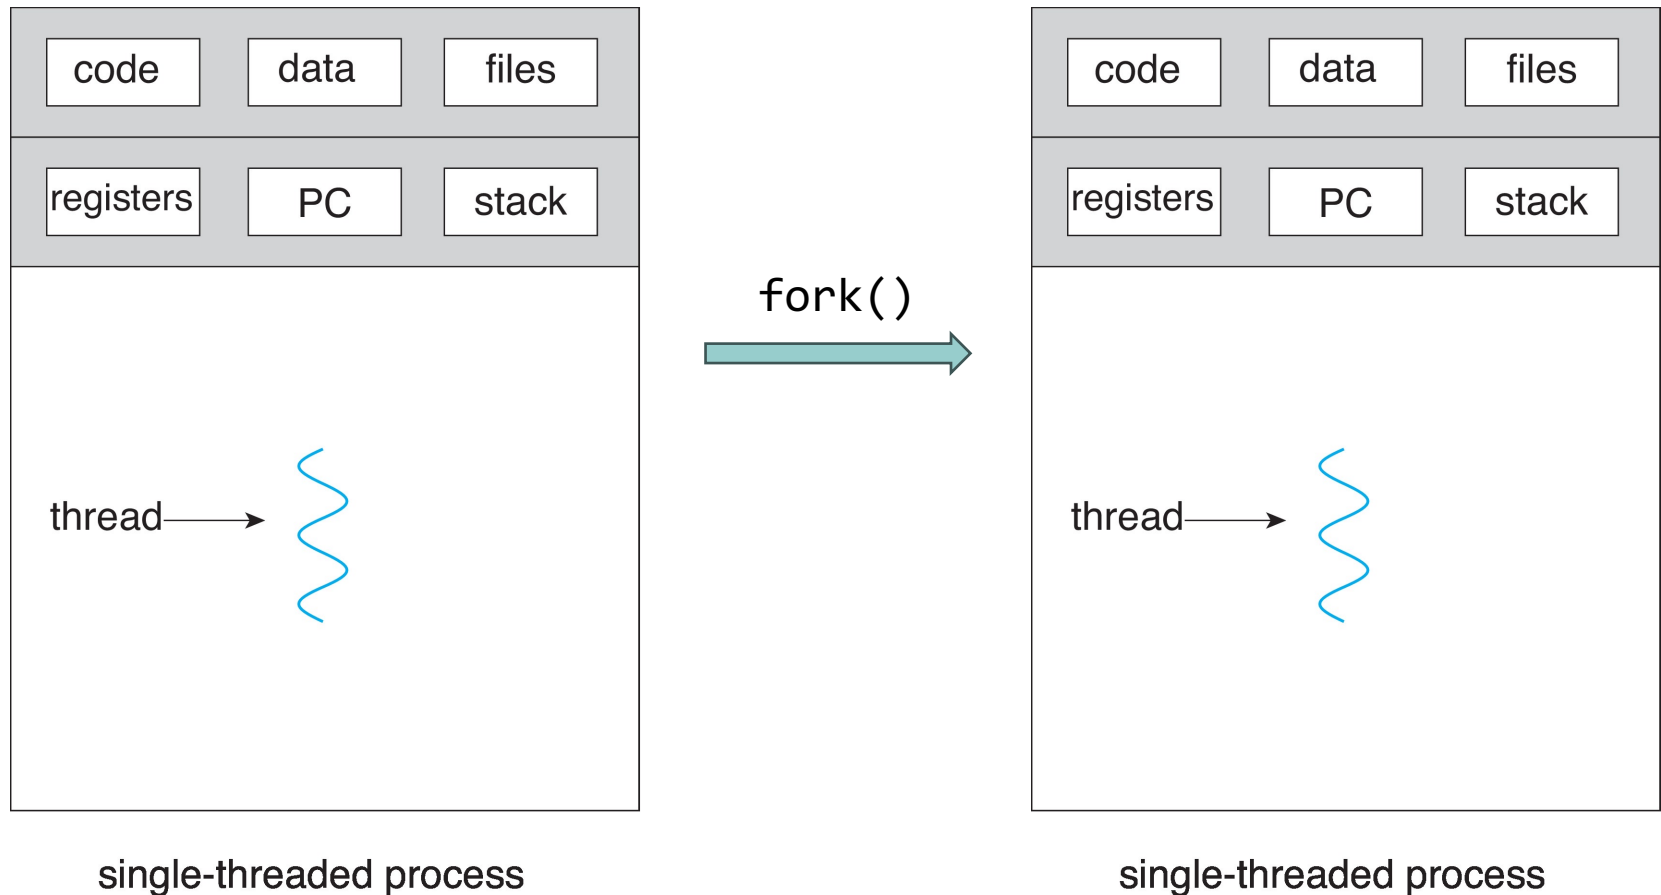
- So far, we have only considered a process to have a single unique sequence of execution context, aka., single-threaded process.



single-threaded process

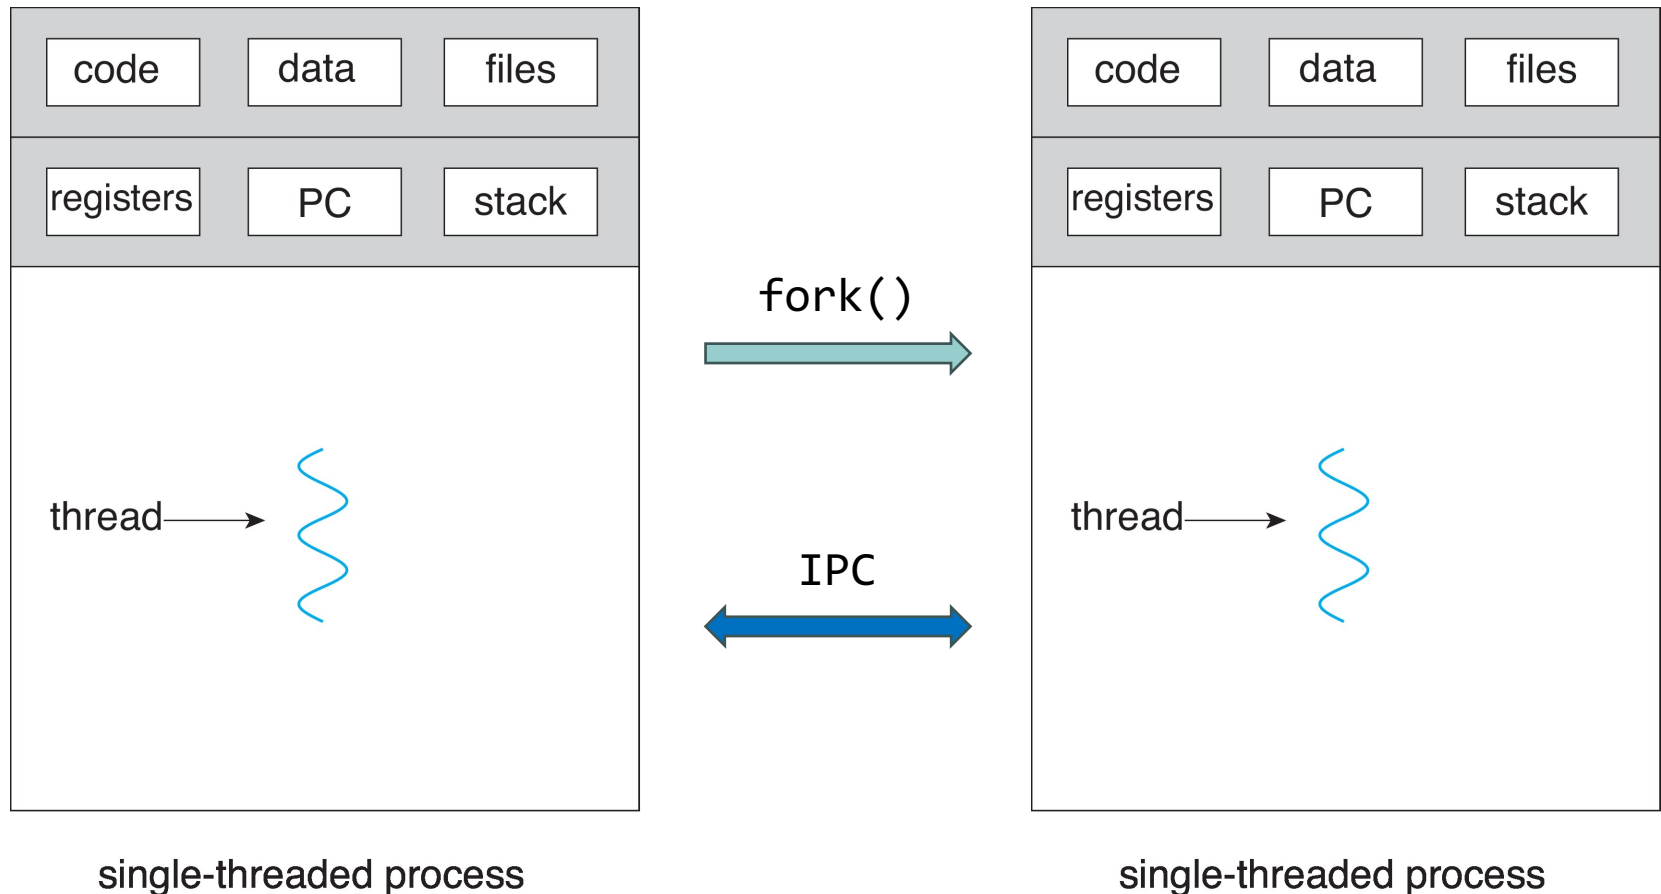
## ■ Single-threaded Process

- In cases where we have to perform a separate task, we `fork()` and `exec()` another process



## ■ Single-threaded Process

- In cases where we have to perform a separate task, we `fork()` and `exec()` another process, possibly communicating via IPC.





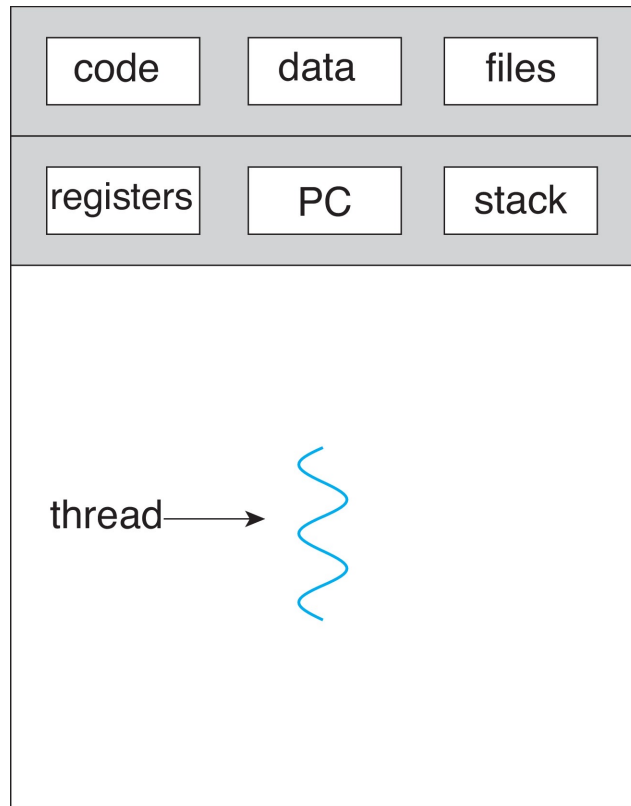
## ■ Single-threaded Process

- In cases where we have to perform a separate task, we `fork()` and `exec()` another process, possibly communicating via IPC.
  - Creating and managing processes is generally considered **expensive**
    - E.g., `fork()` system call
  - Context switching **between processes** is pure overhead, and **expensive**
    - Store PCB(old) and Load PCB(new)
  - IPC between processes is **expensive**
    - Setup Shared Memory (system calls)
    - Message Passing (system calls for every `read()` and `write()`)
    - Explicit coordination and management required from programmers
- Is there a better (cheaper) way to simply spawn a new task?
  - Multi-threaded Process




## ■ Single-threaded and Multi-threaded Processes

- An oversimplified View: **Process**  $\sim$  **Address Space** + **Threads**
- Each thread represents single unique execution context
  - Program Counter, Registers, Execution Flags, Stack



single-threaded process



multithreaded process



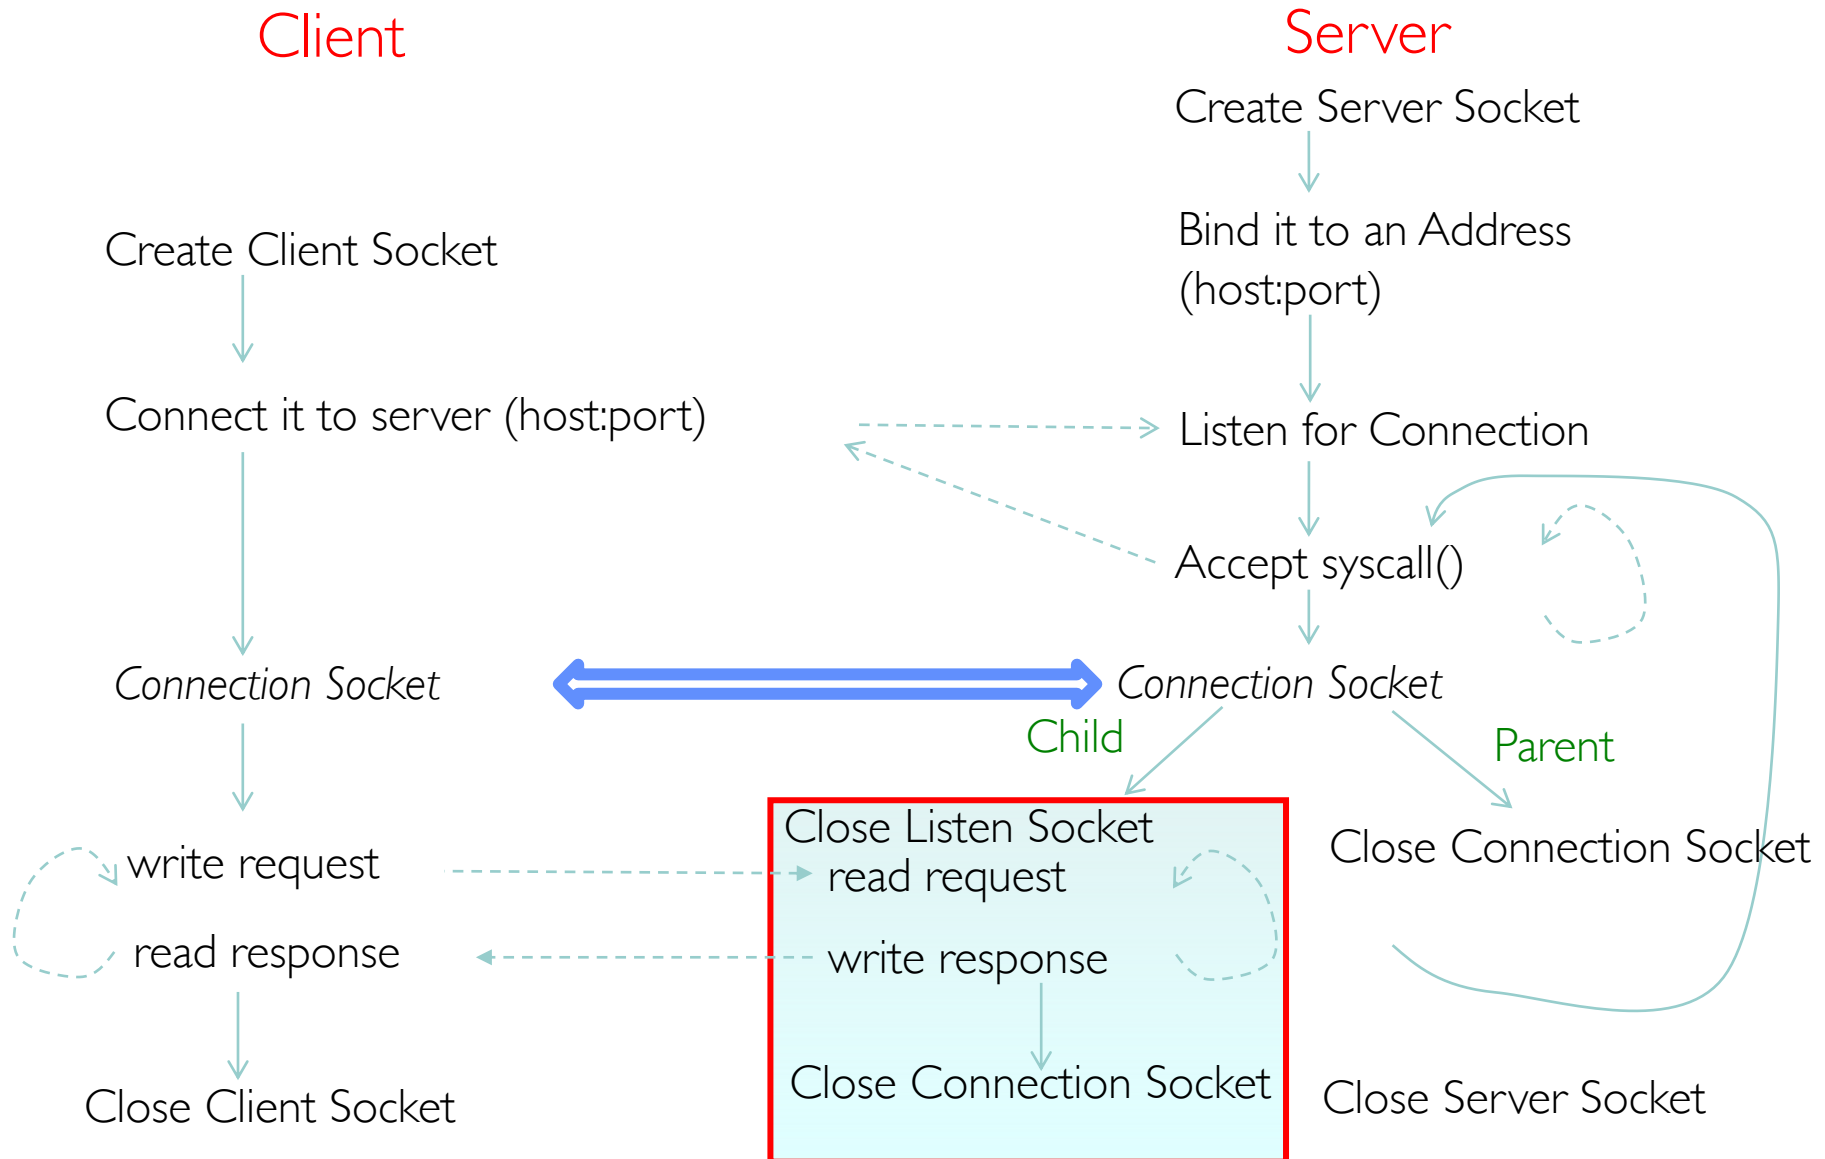
## ■ Single-threaded and Multi-threaded Processes

- An oversimplified View: **Process**  $\sim$  **Address Space** + **Threads**
- Each thread represents single unique execution context
  - Program Counter, Registers, Execution Flags, Stack
- Use threads to spawn a new task (vs. Processes)
  - Creating and managing threads is less expensive than processes
    - No need to allocate a separate memory space or duplicate resources since they are shared among threads within a process
  - Context switching between threads is less expensive
    - Store **TCB**(old) and Load **TCB**(new)
      - **TCB** == **T**hread **C**ontrol **B**lock contains way less info than **PCB**
  - Communication between threads is cheaper than IPC
    - Threads share the same address space → direct communication
- Okay! But there is no such thing as a free lunch. What's the catch?
  - Well, with Threads, you sacrificed **Protection** provided by processes.





## ■ Sockets with Protection and Concurrency

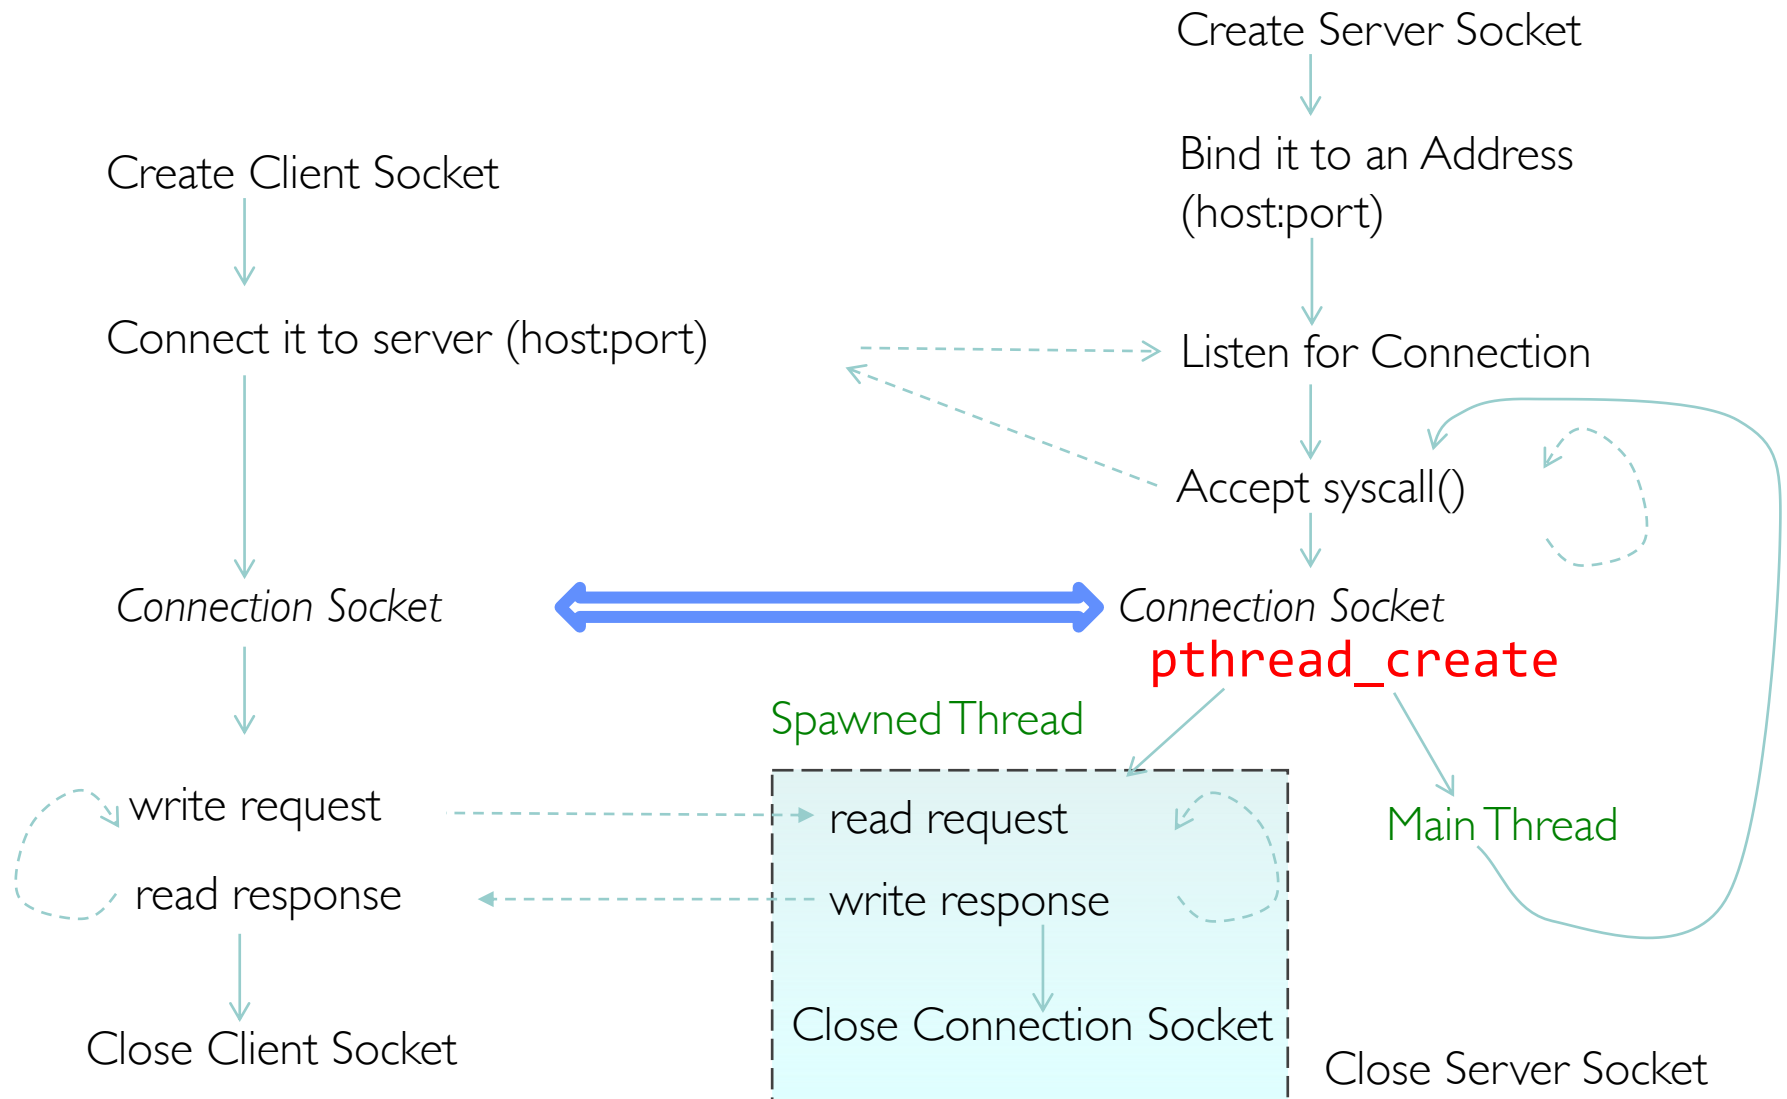




## ■ Sockets with Concurrency, **without** Protection

Client

Server





## ■ Motivation for Threads

- Most modern applications are multithreaded
  - Threads run within application
- Multiple tasks with the application can be implemented by several threads, e.g., a **word processor**, such as Microsoft Word.app:
  - Update display
  - Responding to keystrokes from the user
  - Spell checking
  - ...
- Process creation is heavy-weight.
- Thread creation is light-weight.
- Can simplify code, improve efficiency
- Kernels are generally multithreaded




## ■ Motivation for Threads

- Operating Systems must handle **multiple things at once** (MTAO)
  - Processes, interrupts, background system maintenance
- Networked servers must handle MTAO
  - Multiple connections handled simultaneously
- Parallel programs must handle MTAO
  - To achieve better performance
- Programs with user interface often must handle MTAO
  - To achieve user responsiveness while doing computation
- Network and disk bound programs must handle MTAO
  - To hide network/disk latency
  - Sequence steps in access or communication

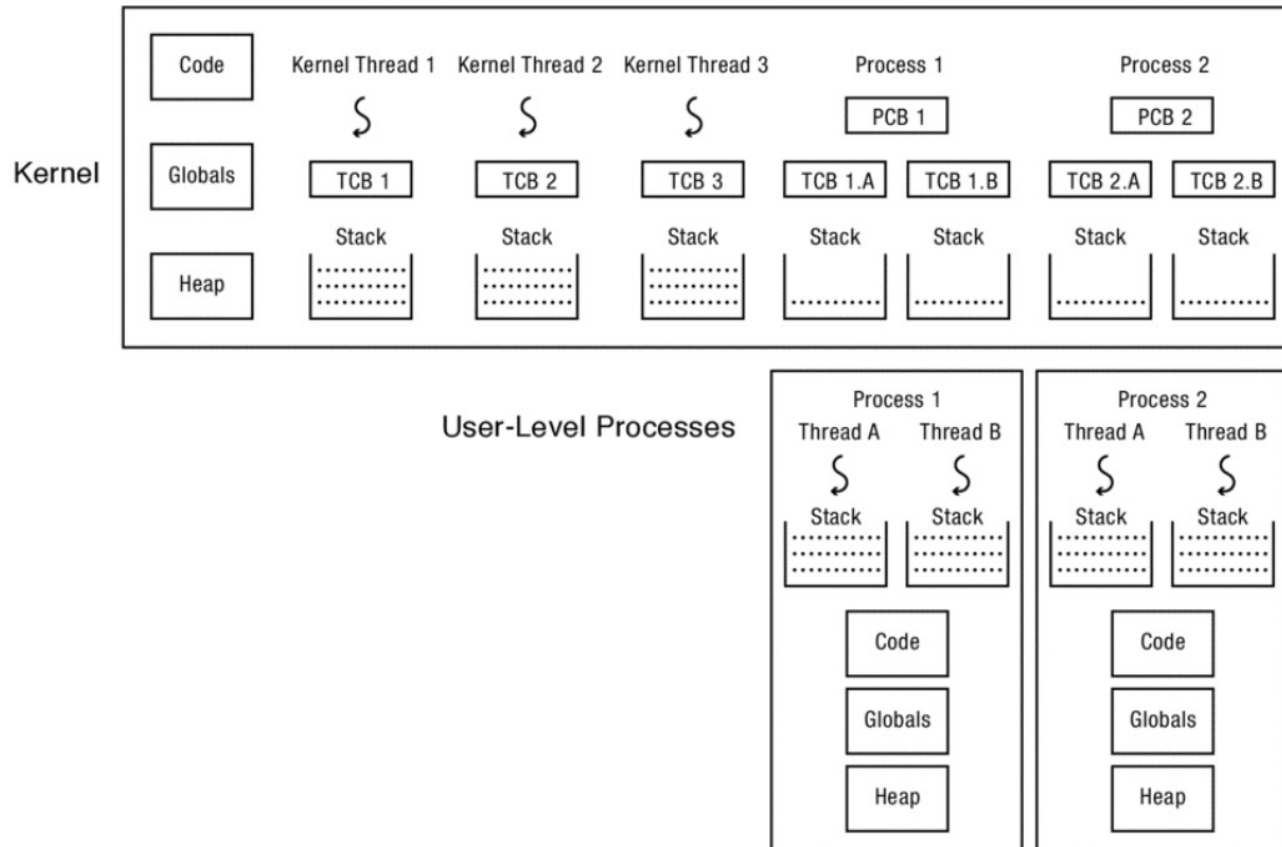
## ■ Memory Layout with Two Threads

- Within the same Address Space of a Process
- Two sets of CPU registers
- Two sets of Stacks



## ■ Memory Layout with Two Threads

- Within the same Address Space of a Process
- Two sets of CPU registers
- Two sets of Stacks





## ■ Benefits of Threads

### ■ Responsiveness (响应能力)

- may allow continued execution if part of the process is blocked, especially important for user interfaces

### ■ Resource Sharing (资源共享)

- threads share resources of process, easier than IPC such as shared memory or message passing
- Inter-thread communication and synchronization is faster than IPC

### ■ Economy (经济性/低系统开销)

- cheaper than process creation (no need to create new address space)
- thread switching lower overhead than process context switching


### ■ Scalability (可扩展性/可伸缩性)

- process can take advantage of multicore architectures, i.e., multiple threads running in parallel on different cores.

## ■ Benefits of Threads

### ■ Example 1: Word Processor

- One thread displays UI and texts
- A second thread reads user inputs from keyboard (*usually blocks*)
- A third thread perform disk I/O (*usually blocks*)

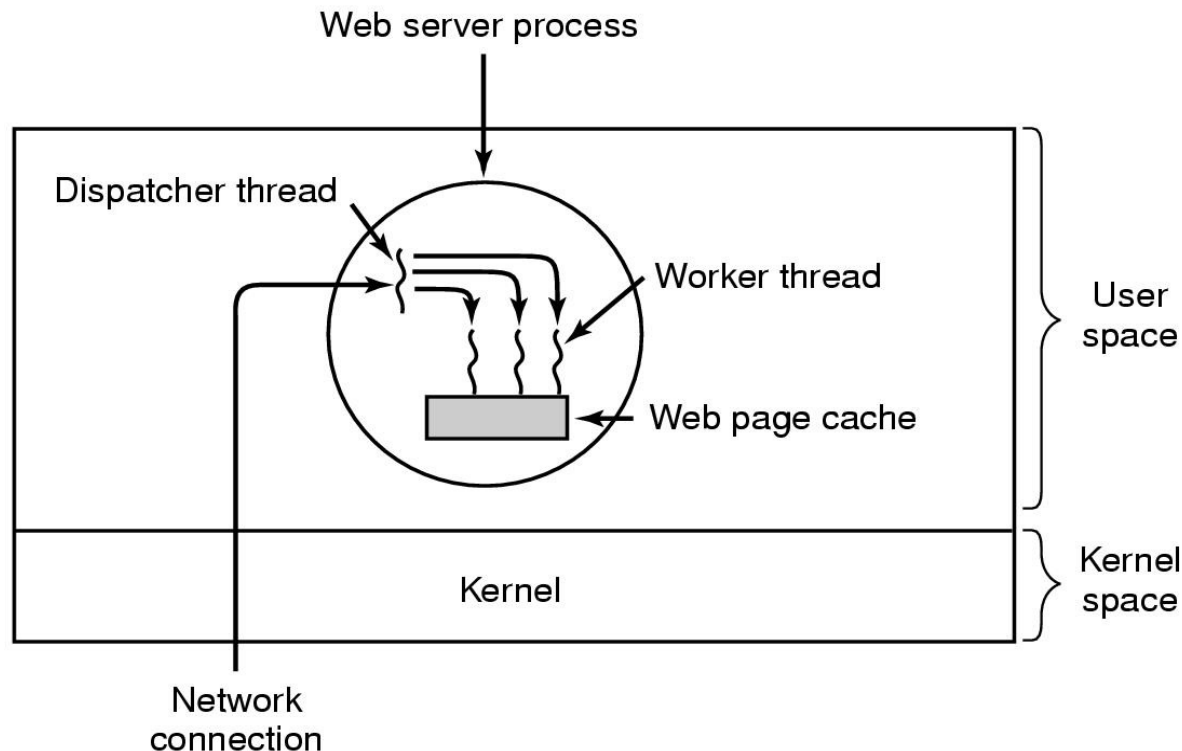




## ■ Benefits of Threads

### ■ Example 2: A File/Web server

- The server needs to handle requests over a short period. Hence it is more efficient to create (and destroy) a single thread for each request.
- On a SMP machine, multiple threads can possibly be executed simultaneously on multiple CPU cores.



## ■ Benefits of Threads

### ■ Example 2: A File/Web server

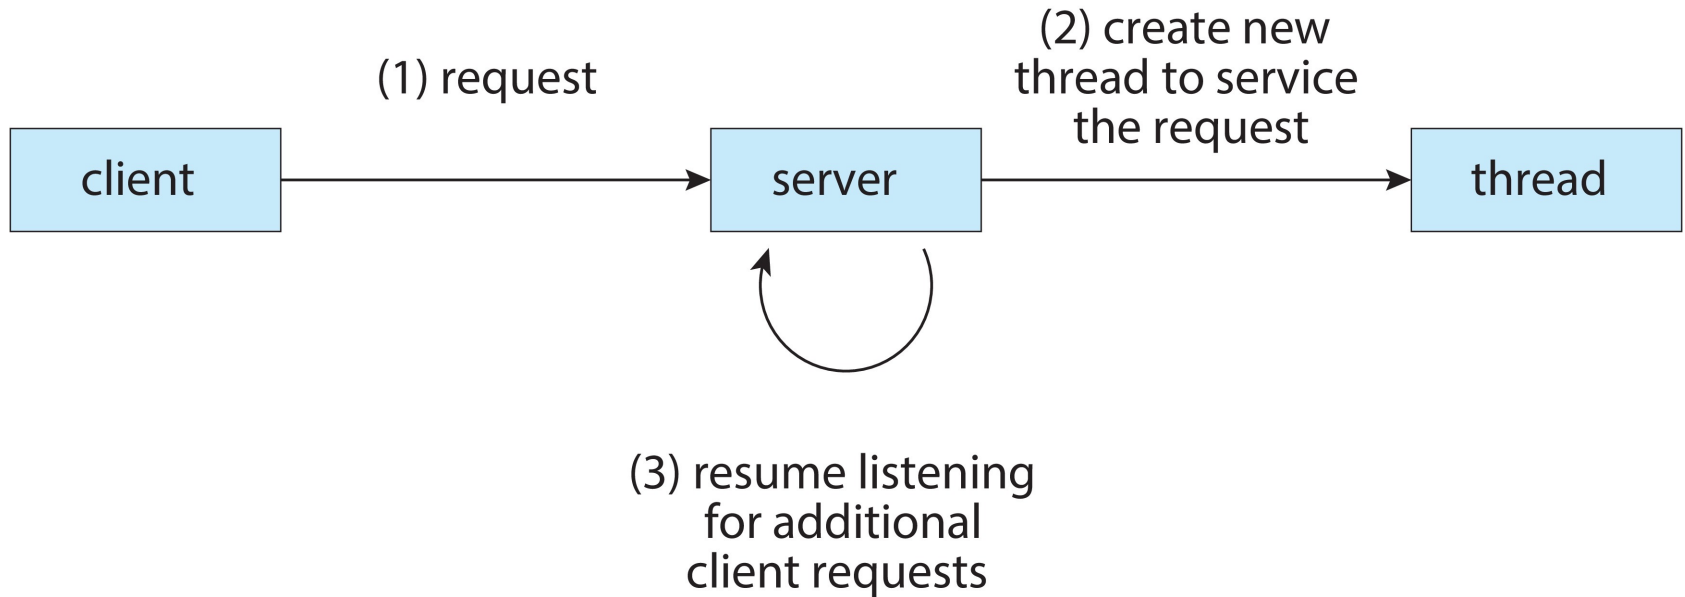
#### ■ Dispatcher thread (main thread):

```
while (true) {  
    get_next_request(&buf);  
    handoff_work(&buf);  
}
```

#### ■ Worker thread

```
while (true) {  
    wait_for_work(&buf);  
    look_for_page_in_cache(&buf, &page);  
    if (page_not_in_cache(&page))  
        read_page_from_disk(&buf, &page);  
    return_page(&page);  
}
```

## ■ Multithreaded Server Architecture



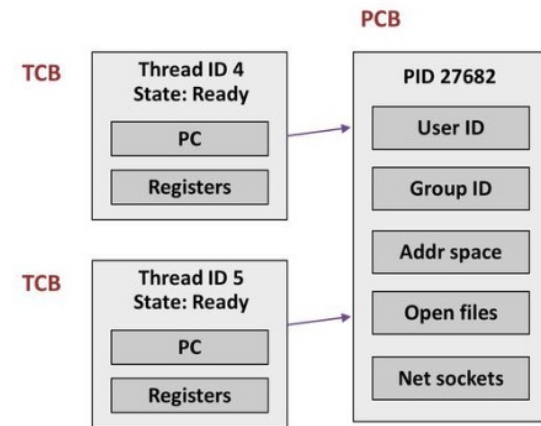
## ■ Benefits of Threads

### ■ Two important implications

- Threaded applications often run **faster** than non-threaded applications
  - context switching between kernel and user-space are avoided

```
// Simplified pseudocode for process context switching
void context_switch_process(Process *current, Process *next) {
    // Save current process state to its PCB
    save_state_to_pcb(&current->pcb);
    // Switch memory address space to the next process
    switch_memory_address_space(next->memory_address_space);
    // Restore the state of the next process from its PCB
    restore_state_from_pcb(&next->pcb);
    // Update system scheduler and process pointers
    current_process = next;
}
```

```
// Simplified pseudocode for thread context switching within the same process
void context_switch_thread(Thread *current, Thread *next) {
    // Save current thread state to its TCB
    save_state_to_tcb(&current->tcb);
    // Restore the state of the next thread from its TCB
    restore_state_from_tcb(&next->tcb);
    // Update system scheduler and thread pointers
    current_thread = next;
}
```





## ■ Benefits of Threads

- Two important implications
  - Threaded applications often run **faster** than non-threaded applications
    - context switching between kernel and user-space are avoided
  - Threaded applications are **harder** to develop
    - although simple, clean designs can help
- Additionally, the assumption is that the development environment provides a **Threads Library** for developers to use.
  - Most modern OSes do



## ■ Multicore Systems

- Multicore architectures place multiple cores on a single chip.
  - Each core appears as a separate processor to the OS
  - Whether the cores appear across CPU chips or within CPU chips, we call these systems **multicore** or **multiprocessor** systems.
  - Almost all systems are multicore nowadays.
    - E.g., your phone has **>6 cores**; our code-server has **128 cores**!

0[0.0]	8[0.0]	16[0.0]	24[0.0]	32[0.0]	40[0.0]	48[0.0]	56[0.0]	64[0.0]	72[0.0]	80[0.0]	88[0.0]	96[0.0]	104[0.0]	112[0.0]	120[0.0]
1[0.0]	9[0.0]	17[0.0]	25[0.0]	33[0.0]	41[0.0]	49[0.0]	57[0.0]	65[0.0]	73[0.0]	81[0.0]	89[0.0]	97[0.0]	105[0.0]	113[0.0]	121[0.0]
2[0.0]	10[0.0]	18[0.0]	26[0.0]	34[0.0]	42[0.0]	50[0.0]	58[0.0]	66[0.0]	74[0.0]	82[0.0]	90[0.0]	98[0.0]	106[0.0]	114[0.0]	122[0.0]
3[0.0]	11[0.0]	19[0.0]	27[0.0]	35[0.0]	43[0.0]	51[0.0]	59[0.0]	67[0.0]	75[0.0]	83[0.0]	91[0.0]	99[0.0]	107[0.0]	115[0.0]	123[0.0]
4[0.0]	12[0.0]	20[0.0]	28[0.0]	36[0.0]	44[0.0]	52[0.0]	60[0.0]	68[0.0]	76[0.0]	84[0.0]	92[0.0]	100[0.0]	108[0.0]	116[0.0]	124[0.0]
5[0.0]	13[0.0]	21[0.0]	29[0.0]	37[0.0]	45[0.0]	53[0.0]	61[0.0]	69[0.0]	77[0.0]	85[0.0]	93[0.0]	101[0.0]	109[0.0]	117[0.0]	125[0.0]
6[0.0]	14[0.0]	22[0.0]	30[0.0]	38[0.0]	46[0.0]	54[0.0]	62[0.0]	70[0.0]	78[0.0]	86[0.0]	94[0.0]	102[0.0]	110[0.0]	118[0.0]	126[0.0]
7[0.0]	15[0.0]	23[0.0]	31[0.0]	39[0.0]	47[0.0]	55[0.0]	63[0.0]	71[0.0]	79[0.0]	87[0.0]	95[0.0]	103[0.0]	111[0.0]	119[0.0]	127[0.0]
Mem[     ]								4.49G/126G							
Swp[ ]								0K/8.00G							
								Tasks: 173; 1 running							
								Load average: 6.11 2.02 0.71							
								Uptime: 00:00:40							
PIDΔUSER															
PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command						
1	root	20	0	162M	10240	7168	S	0.0	0.0	0:01.49	init				
1430	root	19	-1	621M	128M	128M	S	0.0	0.1	0:00.31	systemd-journald				
1471	root	RT	0	282M	25600	8192	S	0.0	0.0	0:00.02	multipathd -d -s				
1474	root	20	0	26500	6144	4096	S	0.0	0.0	0:00.14	systemd-udev				
1797	systemd-n	20	0	16128	8192	7168	S	0.0	0.0	0:00.03	systemd-networkd				
1806	_rpc	20	0	8108	3072	3072	S	0.0	0.0	0:00.00	rpcbind -f -w				
1807	systemd-r	20	0	25540	12288	8192	S	0.0	0.0	0:00.06	systemd-resolved				
1808	systemd-t	20	0	89364	6144	6144	S	0.0	0.0	0:00.04	systemd-timesyncd				
1859	clk	20	0	706M	49580	32768	S	0.0	0.0	0:00.72	node /usr/lib/code-server				
3379	clk	20	0	642M	56744	35840	S	0.0	0.0	0:00.56	node /usr/lib/code-server/out/node/entry				


- Multithreaded programming provides a mechanism for more **efficient use** of these multiple computing cores and improved **concurrency**.



## ■ Multicore Systems


### ■ Concurrent execution on single-core system:

- execution of threads are **interleaved** over time



### ■ Parallelism on a multi-core system:

- some threads can run **in parallel**.





## ■ Concurrency is not Parallelism

- **Concurrency** is about **handling** multiple things at once (**MTAO**)
- **Parallelism** is about **doing** multiple things **simultaneously**

- **It is possible to have concurrency without parallelism.**

Example: Two threads on a single-core system

- ... execute concurrently ...
  - ... but not in parallel
- 
- Each thread handles or manages a separate thing or task
  - But those tasks are not necessarily executing simultaneously





- Challenges in programming for **Multicore** systems:
  - Identifying Tasks
  - Balance
  - Data Splitting
  - Data Dependency
  - Testing and debugging



## ■ Challenges in programming for **Multicore** systems:

### ■ Identifying Tasks


- This involves examining applications to find areas that can be divided into separate, concurrent tasks. Ideally, tasks are independent of one another and thus can run in parallel on individual cores.

### ■ Balance

### ■ Data Splitting

### ■ Data Dependency

### ■ Testing and debugging





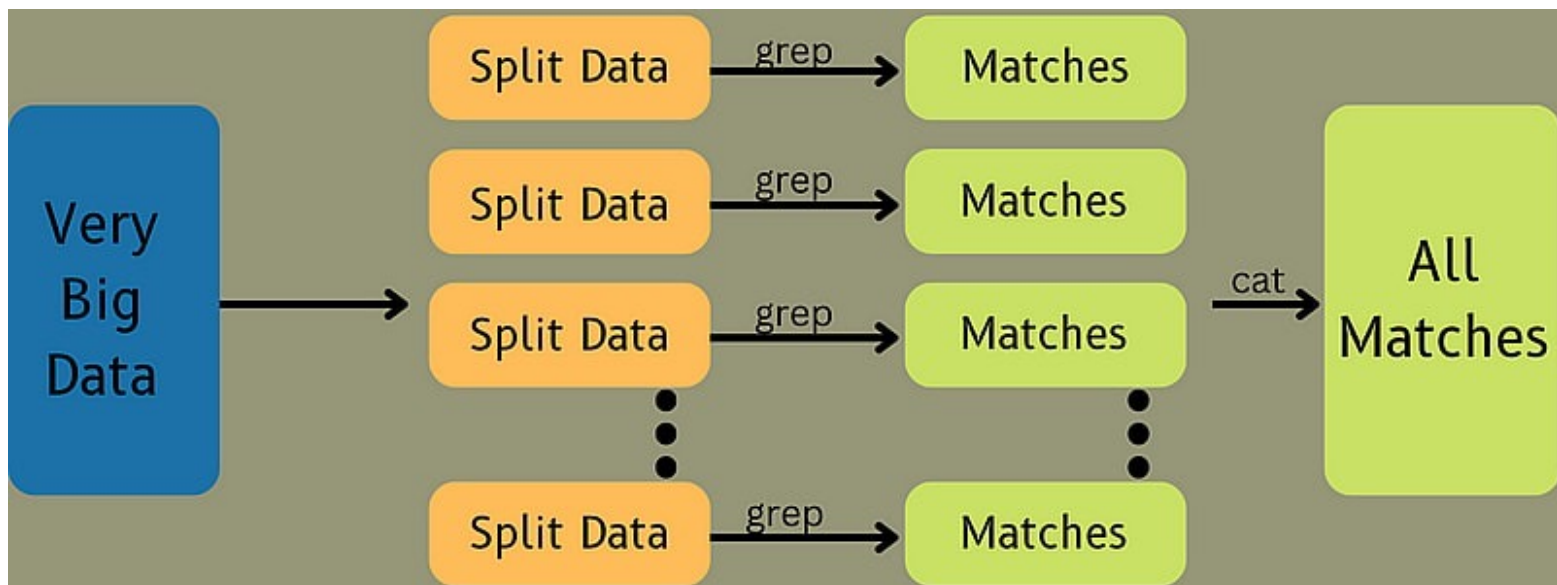
## ■ Challenges in programming for **Multicore** systems:

- Identifying Tasks
- **Balance**
  - While identifying tasks that can run in parallel, programmers must also ensure that the tasks perform equal work of equal value. In some instances, a certain task may not contribute as much value to the overall process as other tasks. Using a separate execution core to run that task may not be worth the cost.
- Data Splitting
- Data Dependency
- Testing and debugging



## ■ Challenges in programming for **Multicore** systems:

- Identifying Tasks
- Balance
- **Data Splitting**
  - Just as applications are divided into separate tasks, the data accessed and manipulated by the tasks must be divided to run on separate cores.
- Data Dependency
- Testing and debugging

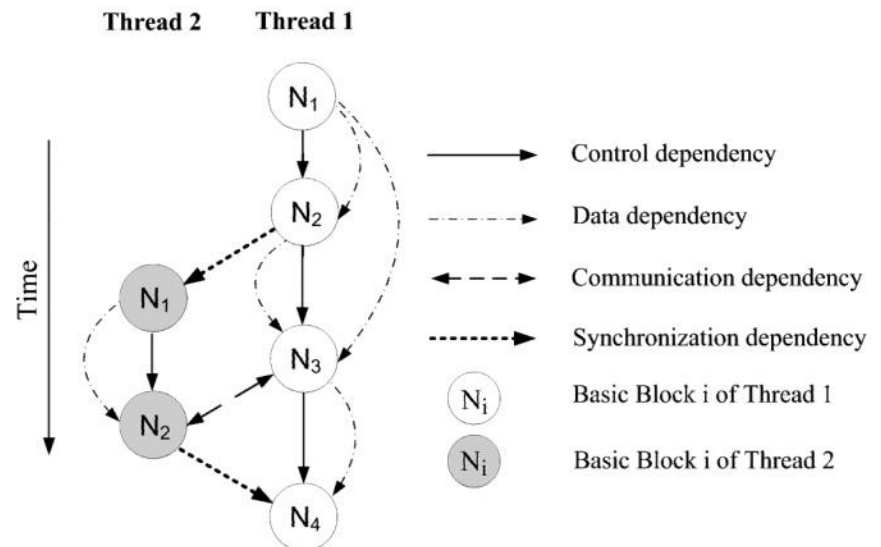


## ■ Challenges in programming for Multicore systems:

- Identifying Tasks
- Balance
- Data Splitting
- **Data Dependency**

- The data accessed by the tasks must be examined for dependencies between two or more tasks. When **one task depends on data from another**, programmers must ensure that the execution of the tasks is synchronized to accommodate the data **dependency**.

- Testing and debugging





## ■ Challenges in programming for **Multicore** systems:

- Identifying Tasks
- Balance
- Data Splitting
- Data Dependency
- **Testing and debugging**
  - When running in parallel on multiple cores, **many different execution paths are possible**. Testing and debugging such concurrent programs is inherently more difficult than on single-threaded applications.


```
(gdb) b 16 thread 2
Breakpoint 3 at 0x555555556477: file mat_mult.cpp, line 16.
(gdb) n
[Thread 0x7ffff624e700 (LWP 8629) exited]
15         for (int i = 0; i < mat1[row].size(); ++i) {
(gdb) n
[Thread 0x7ffff7250700 (LWP 8623) exited]
[Thread 0x7ffff6a4f700 (LWP 8624) exited]

Thread 2 "mat_mult" hit Breakpoint 3, dotProduct (
    mat1=std::vector of length 2, capacity 2 = {...},
    mat2=std::vector of length 2, capacity 2 = {...},
    result=std::vector of length 2, capacity 2 = {...}, row=0, col=0)
    at mat_mult.cpp:16
16         result[row][col] += mat1[row][i] + mat2[j][col];
(gdb) █
```

## Types of Parallelism


### Data Parallelism

- focuses on distributing subsets of the **same** data across multiple computing cores and performing the **same** operation on each core



### Task Parallelism

- involves distributing not data but **tasks** (threads) across multiple computing cores. Each thread is performing a unique operation.
- Different threads may be operating on the same data, or different data







## ■ Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial (non-parallel) and parallel components
- **S** is the serial portion; **N** processing cores
- $speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$
- For example, if an application is 75% parallel / 25% serial, moving from 1 to 2 cores results in a speedup of  $\frac{1}{0.25 + \frac{0.75}{2}} = 1.6$
- As N approaches infinity, speedup approaches  $\frac{1}{S}$ 
  - For example, if 50% is serial, then the maximum speedup is 2.0 times.
- Serial portion of an application has disproportionate effect on performance gained by adding additional cores.






## ■ Amdahl's Law



$$\text{speedup} = \frac{1}{0.25 + \frac{0.75}{2}} = 1.6$$



$$\text{speedup} = \frac{1}{0.25 + \frac{0.75}{4}} = 2.286$$




$$\text{speedup} = \frac{1}{0.25 + \frac{0.75}{\infty}} = 4$$



## Amdahl's Law

- Serial portion of an application has disproportionate effect on performance gained by adding additional cores.





## ■ User Threads

- Management done by user-level threads library
- The kernel is unaware of the existence of User Level Threads
  - When one thread makes a **blocking** syscall, the whole process will **block**
- Thread switching does not require kernel mode privileges
- Scheduling is application specific


## ■ Kernel Threads

- Supported and managed directly by the kernel
- Kernel maintains context information for the process and threads
- Switching between threads requires kernel intervention



## ■ Multithreading Models

- Support for threads may be provided either at the user level, for **user level threads** (ULTs), or by the kernel, for **kernel level threads** (KLT).
  - **User threads** are supported above the kernel and are managed without kernel support
  - **Kernel threads** are supported and managed directly by the OS




- Relationship between user threads and kernel threads:
  - **Many-to-One** model
  - **One-to-One** model (*most common*)
  - **Many-to-Many** model



## ■ Many-to-One Model


- Many user-level threads mapped to a single kernel-level thread
  - The kernel is unaware of user-level threads
- One thread blocking causes all threads to block
- Multiple threads may not run in parallel on multicore system because only one may be in the kernel at a time
- *Few systems currently use this model*
- Examples:
  - Solaris Green Threads
  - GNU Portable Threads





## ■ One-to-One Model


- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process often restricted due to overhead
- Examples
  - Linux
  - Windows





## ■ Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the OS to create a sufficient number of kernel threads
- More flexible
- Much more difficult to implement than other models
- Examples
  - Windows with the *ThreadFiber* package
  - Otherwise not very common





## ■ Thread Libraries

- **Thread libraries** provide programmers with API for creating and managing threads
- Two primary ways of implementing threads:
  - Library entirely in user space
    - no system calls
  - Kernel-level library supported by the OS
    - code and data structures for the library exist in kernel space
    - Invoking an API function results in a system call to the kernel





## ■ Thread Libraries

- Thread libraries provide programmers with API for creating and managing threads
- Two primary ways of implementing threads:
  - Library entirely in user space
    - no system calls
  - Kernel-level library supported by the OS
    - code and data structures for the library exist in kernel space
    - Invoking an API function results in a system call to the kernel
- Three main thread libraries are in use today:
  - **POSIX Pthreads**
  - Windows Thread Library
  - Java Thread API



## ■ Pthreads

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- May be provided either as user-level or kernel-level
- **Specification**, not **implementation**
- API specifies behavior of thread library, implementation is up to development of the library
  - Linux uses Native POSIX Thread Library (NPTL) as their pthread impl
  - FreeBSD primarily uses `libthr` as its pthread library impl
- Common in UNIX operating system (Linux, macOS)



## ■ Pthreads API

### ■ Thread Creation:

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine) (void *), void *arg);  
  
/* Returns 0 on success, or a positive error number on error */
```

### ■ Thread Completion:


```
int pthread_join(pthread_t thread, void **retval);  
  
/* Returns 0 on success, or a positive error number on error */
```

### ■ Others:

```
void pthread_exit(void *retval);      /* Terminates calling thread */  
  
int pthread_kill(pthread_t thread, int sig); /* Send sig to thread */  
  
pthread_t pthread_self(void);         /* Return current thread ID */  
  
int pthread_detach(pthread_t thread); /* Detach thread, no need to join */
```

## ■ Pthreads: Fork-Join Pattern

- Main thread creates (forks) multiple threads by passing them **args** to work on
- ... and then **joins** with them, collecting results.



## ■ Pthreads Example

```
/* t0.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }
    printf("sum: %ld\n", sum);
    printf("Thread finished its work.\n");
    return NULL;
}

int main(int argc, char **argv) {
    pthread_t tid;

    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, NULL);
    printf("main(): Thread returned.\n");

    return 0;
}
```

```
$ ./t0
main(): creating new thread...
Thread started working...
sum: 55
Thread finished its work.
main(): Thread returned.
```

## ■ Pthreads Example

```
/* t0.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }
    printf("sum: %ld\n", sum);
    printf("Thread finished its work.\n");
    return NULL;
}

int main(int argc, char **argv) {
    pthread_t tid;

    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, NULL);
    printf("main(): Thread returned.\n");

    return 0;
}
```

```
$ ./t0
main(): creating new thread...
Thread started working...
sum: 55
Thread finished its work.
main(): Thread returned.
```

Why is **void** pointer used as argument?

**Answer:** It allows us to pass in **any** type of argument. We have to explicitly cast the **(void \*)arg** into the desired type. (in the case of `t0.c`, the **long** type).

## ■ Pthreads Example

```
/* t0.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    struct data *obj = (struct data *)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }
    printf("sum: %ld\n", sum);
    printf("Thread finished its work.\n");
    return NULL;
}

int main(int argc, char **argv) {
    pthread_t tid;
    struct data obj;
    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)&obj);
    pthread_join(tid, NULL);
    printf("main(): Thread returned.\n");

    return 0;
}
```

```
$ ./t0
main(): creating new thread...
Thread started working...
sum: 55
Thread finished its work.
main(): Thread returned.
```

Why is **void** pointer used as argument?

**Answer:** It allows us to pass in **any** type of argument. We have to explicitly cast the **(void \*)arg** into the desired type. (in the case of `t0.c`, the **long** type).

When passing a compound data type, such as **struct**, the pointer to **struct** must be used.

## ■ Pthreads Example

```
/* t0.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }
    printf("sum: %ld\n", sum);
    printf("Thread finished its work.\n");
    return NULL;
}

int main(int argc, char **argv) {
    pthread_t tid;

    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, NULL);
    printf("main(): Thread returned.\n");

    return 0;
}
```

```
$ ./t0
main(): creating new thread...
Thread started working...
sum: 55
Thread finished its work.
main(): Thread returned.
```

Recall that threads in the same process share many resources, most notably the **global variables**.



## ■ Pthreads Example

```
/* t0.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }
    printf("sum: %ld\n", sum);
    printf("Thread finished its work.\n");
    return NULL;
}

int main(int argc, char **argv) {
    pthread_t tid;

    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, NULL);
    printf("main(): Thread returned.\n");

    return 0;
}
```

```
$ ./t0
main(): creating new thread...
Thread started working...
sum: 55
Thread finished its work.
main(): Thread returned.
```

Recall that threads in the same process share many resources, most notably the **global variables**.

## ■ Pthreads Example

```
/* t1.c */
#include <pthread.h>
#include <stdio.h>
long sum = 0;
void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }
    printf("sum: %ld\n", sum);
    printf("Thread finished its work.\n");
    return NULL;
}
int main(int argc, char **argv) {
    pthread_t tid;

    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, NULL);
    printf("main(): Thread returned.\n");
    printf("sum: %ld\n", sum);
    return 0;
}
```

```
$ ./t1
main(): creating new thread...
Thread started working...
Thread finished its work.
main(): Thread returned.
sum: 55
```

Recall that threads in the same process share many resources, most notably the **global variables**.

## ■ Pthreads Example

```
/* t1.c */
#include <pthread.h>
#include <stdio.h>
long sum = 0;
void *threadfun(void *arg) {
    printf("Thread started working...\n");

    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }

    printf("Thread finished its work.\n");
    return NULL;
}
int main(int argc, char **argv) {
    pthread_t tid;

    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, NULL);
    printf("main(): Thread returned.\n");
    printf("sum: %ld\n", sum);
    return 0;
}
```

```
$ ./t1
main(): creating new thread...
Thread started working...
Thread finished its work.
main(): Thread returned.
sum: 55
```

Recall that threads in the same process share many resources, most notably the **global variables**.

## ■ Pthreads Example

```
/* t2.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }

    printf("Thread finished its work.\n");
    return (void *)sum;
}

int main(int argc, char **argv) {
    pthread_t tid;
    void *res;
    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, &res);
    printf("main(): Thread returned.\n");
    printf("sum: %ld\n", (long)res);
    return 0;
}
```

```
$ ./t2
main(): creating new thread...
Thread started working...
Thread finished its work.
main(): Thread returned.
sum: 55
```

Recall that threads in the same process share many resources, most notably the **global variables**.

Communication between threads can also be achieved by passing the return value to the main thread via

```
pthread_join(pthread_t thread,
              void **retval);
```

## ■ Pthreads Example

```
/* t2.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }

    printf("Thread finished its work.\n");
    return (void *)sum;
}

int main(int argc, char **argv) {
    pthread_t tid;
    void *res;
    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, &res);
    printf("main(): Thread returned.\n");
    printf("sum: %ld\n", (long)res);
    return 0;
}
```

```
$ ./t2
main(): creating new thread...
Thread started working...
Thread finished its work.
main(): Thread returned.
sum: 55
```

Recall that threads in the same process share many resources, most notably the **global variables**.

Communication between threads can also be achieved by passing the return value to the main thread via

```
pthread_join(pthread_t thread,
              void **retval);
```

Note that `retval`'s type is `void **`, a pointer to a generic pointer. **Explicit cast** is required to interpret `retval`.

## ■ Pthreads Example

```
/* t2.c */
#include <pthread.h>
#include <stdio.h>

void *threadfun(void *arg) {
    printf("Thread started working...\n");
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }

    printf("Thread finished its work.\n");
    pthread_exit((void *)sum);
}

int main(int argc, char **argv) {
    pthread_t tid;
    void *res;
    printf("main(): creating new thread...\n");
    pthread_create(&tid, NULL,
                  threadfun, (void *)10);
    pthread_join(tid, &res);
    printf("main(): Thread returned.\n");
    printf("sum: %ld\n", (long)res);
    return 0;
}
```

```
$ ./t2
main(): creating new thread...
Thread started working...
Thread finished its work.
main(): Thread returned.
sum: 55
```

When exiting the thread, we can also  
**explicitly** call

```
void pthread_exit(void *retval);
```

## ■ Pthreads Example

```
/* t3.c */
#include <pthread.h>
#include <stdio.h>
void *threadfun(void *arg) {
    printf("Thread %ld started working...\n",
        pthread_self());
    long sum = 0;
    long n = (long)arg;
    for (int i = 1; i <= n; i++) {
        sum = sum + i;
    }
    printf("sum: %ld\n", sum);
    printf("Thread finished its work.\n");
    return NULL;
}
int main(int argc, char **argv) {
    pthread_t tid[10];
    printf("main(): creating new thread...\n");
    for (int t = 0; t < 10; t++)
        pthread_create(&tid[t], NULL,
            threadfun, (void *)(t*10));
    for (int t = 0; t < 10; t++)
        pthread_join(tid[t], NULL);
    printf("main(): Thread returned.\n");
    return 0;
}
```

```
$ ./t3
main(): creating new thread...
Thread 138983810004544 started working...
sum: 0
Thread finished its work.
Thread 138983801611840 started working...
sum: 55
Thread finished its work.
Thread 138983793219136 started working...
sum: 210
Thread finished its work.
Thread 138983784826432 started working...
sum: 465
Thread finished its work.
Thread 138983776433728 started working...
sum: 820
Thread finished its work.
...
```

We often use an array of `pthread_t` to store the tid of multiple threads.



**Thank you!**