



# DCS216 Operating Systems

## Lecture 14 Synchronization (3)

**Apr 15<sup>th</sup>, 2024**

**Instructor: Xiaoxi Zhang**  
**Sun Yat-sen University**



## ■ Content

- Monitors
  - Condition Variables
- Synchronization within the Kernel
- POSIX Synchronization
- Alternative Approaches



## ■ Problems with Semaphores

- Although semaphores provide a convenient and effective mechanism for synchronization, using them incorrectly can result in errors that are difficult to detect.
- Incorrect use of semaphore operations:
- **Examples 1:** Suppose that a program interchanges the order in which the `wait()` and `signal()` operations on the semaphore `mutex` are executed, resulting in the following execution (**mutual exclusion** no longer holds):

```
signal(mutex);  
...  
// critical section  
...  
wait(mutex);
```



## ■ Problems with Semaphores

- Although semaphores provide a convenient and effective mechanism for synchronization, using them incorrectly can result in errors that are difficult to detect.
- Incorrect use of semaphore operations:
- **Examples 2:** Suppose that a program replaces `signal()` with `wait()` : (thread will permanently block on 2<sup>nd</sup> call to `wait()`)

```
wait(mutex);  
...  
// critical section  
...  
wait(mutex);
```



## ■ Problems with Semaphores

- Although semaphores provide a convenient and effective mechanism for synchronization, using them incorrectly can result in errors that are difficult to detect.
- Incorrect use of semaphore operations:
- **Examples 3:** the same API (`sem_wait`, `sem_post`) for different purposes leads to confusion

```
void* producer(void* arg) {  
    int next_produced;  
    for (int i = 0; i < 10; i++) {  
        next_produced = i;  
        sem_wait(&empty); // Decrement empty count  
        sem_wait(&mutex);  
        // Add the item to the buffer  
        buffer[in] = next_produced;  
        in = (in + 1) % BUFFER_SIZE;  
        printf("--> Produced %d. ", next_produced);  
        printBuffer();  
  
        sem_post(&mutex);  
        sem_post(&full);  
    }  
}
```



## ■ Problems with Semaphores

- Although semaphores provide a convenient and effective mechanism for synchronization, using them incorrectly can result in errors that are difficult to detect.
- Incorrect use of semaphore operations:
- **Examples 3:** the same API (`sem_wait`, `sem_post`) for different purposes leads to confusion

Programmers can get confused and **interchange** their orders by mistake without noticing, leading to deadlocks or other bugs.

```
void* producer(void* arg) {  
    int next_produced;  
    for (int i = 0; i < 10; i++) {  
        next_produced = i;  
        {  
            sem_wait(&mutex);  
            sem_wait(&empty); // Decrement empty count  
            // Add the item to the buffer  
            buffer[in] = next_produced;  
            in = (in + 1) % BUFFER_SIZE;  
            printf("--> Produced %d. ", next_produced);  
            printBuffer();  
  
            sem_post(&mutex);  
            sem_post(&full);  
        }  
    }  
}
```



## ■ Problems with Semaphores

- Although semaphores provide a convenient and effective mechanism for synchronization, using them incorrectly can result in errors that are difficult to detect.
- Incorrect use of semaphore operations
- **Solution:** (a cleaner separation of functions)
  - Avoid using semaphores for *both* **Mutual Exclusion** and **Scheduling Constraints**
  - Use **Locks** for **Mutual Exclusion**
  - Use **Condition Variables** for **Scheduling Constraints**



## ■ Monitors (管程)

- To deal with such errors caused by incorrect use of semaphores, one strategy is to incorporate simple synchronization tools as high-level constructs.





## ■ Monitors (管程)

- To deal with such errors caused by incorrect use of semaphores, one strategy is to incorporate simple synchronization tools as high-level constructs.
- **Monitor:** A high-level abstraction that provides a convenient and effective mechanism for synchronization
- A Monitor type is an **Abstract Data Type (ADT)** that includes a set of **programmer-defined operations** that are provided with mutual exclusion within the monitor.

```
Monitor monitor_name {  
    /* shared variable declarations */  
  
    function P1 (...) { ... }  
    function P2 (...) { ... }  
    function P3 (...) { ... }  
  
    initialization_code (...) {...}  
}
```




## ■ Monitors (管程)

- **Sharing:** A monitor is shared by concurrent processes/threads
- **Encapsulation & Security:** The representation of a monitor cannot be used directly by the processes/threads
  - Only the functions defined within a monitor can access the shared variables declared within the monitor
- **Mutual Exclusion:** The monitor construct ensures that only one process/thread at a time is active within the monitor.
  - Programmers does not need to code this synchronization explicitly
  - Shared data are protected by placing them within the monitor
  - The monitor locks shared data upon process entry

## Monitors

### Schematic View of a Monitor





## ■ Condition Variables

- With only Mutex Locks and Semaphores, the **monitor** is not sufficiently powerful for modeling some synchronization schemes.
- Sometimes, we wish to check whether a **condition** is true before we continue its execution.

```
Monitor bounded_buffer {  
    /* shared variables */  
    Condition not_full;  
    Condition not_empty;  
  
    void Producer() {  
        wait(not_full);  
        ...  
        signal(not_empty);  
    }  
    void Consumer() {  
        wait(not_empty);  
        ...  
        signal(not_full);  
    }  
}
```

## ■ Condition Variables

- With only Mutex Locks and Semaphores, the **monitor** is not sufficiently powerful for modeling some synchronization schemes.
- Sometimes, we wish to check whether a **condition** is true before we continue its execution.

```
Monitor bounded_buffer {  
    /* shared variables */  
    Condition not_full;  
    Condition not_empty;  
  
    void Producer() {  
        wait(not_full);  
        ...  
        signal(not_empty);  
    }  
    void Consumer() {  
        wait(not_empty);  
        ...  
        signal(not_full);  
    }  
}
```

Producer should wait when the buffer is full. It should continue only if (`not_full == true`)



## ■ Condition Variables

- With only Mutex Locks and Semaphores, the **monitor** is not sufficiently powerful for modeling some synchronization schemes.
- Sometimes, we wish to check whether a **condition** is true before we continue its execution.

```
Monitor bounded_buffer {  
    /* shared variables */  
    Condition not_full;  
    Condition not_empty;  
  
    void Producer() {  
        wait(not_full);  
        ...  
        signal(not_empty);  
    }  
    void Consumer() {  
        wait(not_empty);  
        ...  
        signal(not_full);  
    }  
}
```

After the Producer creates an item, it should notify (signal) the Consumer that the buffer is now `not_empty`.

## ■ Condition Variables

- With only Mutex Locks and Semaphores, the **monitor** is not sufficiently powerful for modeling some synchronization schemes.
- Sometimes, we wish to check whether a **condition** is true before we continue its execution.

```
Monitor bounded_buffer {  
    /* shared variables */  
    Condition not_full;  
    Condition not_empty;  
  
    void Producer() {  
        wait(not_full);  
        ...  
        signal(not_empty);  
    }  
    void Consumer() {  
        wait(not_empty);  
        ...  
        signal(not_full);  
    }  
}
```

After the Producer creates an item, it should notify (signal) the Consumer that the buffer is now `not_empty`, potentially unblocking Consumer.



## ■ Condition Variables

- condition `x, y`;
- Two operations are allowed on a condition variable:
  - `x.wait()` : the calling thread that invokes `wait()` is suspended
    - ...until `x.signal()`.
  - `x.signal()` : resumes (wakes up) one of the threads (if any)
    - ...that invoked `x.wait()`
    - If no `x.wait()` on the condition variable `x`, then it has no effect on condition variable `x`.
    - In contrast to `semaphore.signal()`, which always affects the **state** of the **semaphore**.
    - In other words, condition variables are **stateless**, while semaphores are **stateful**.



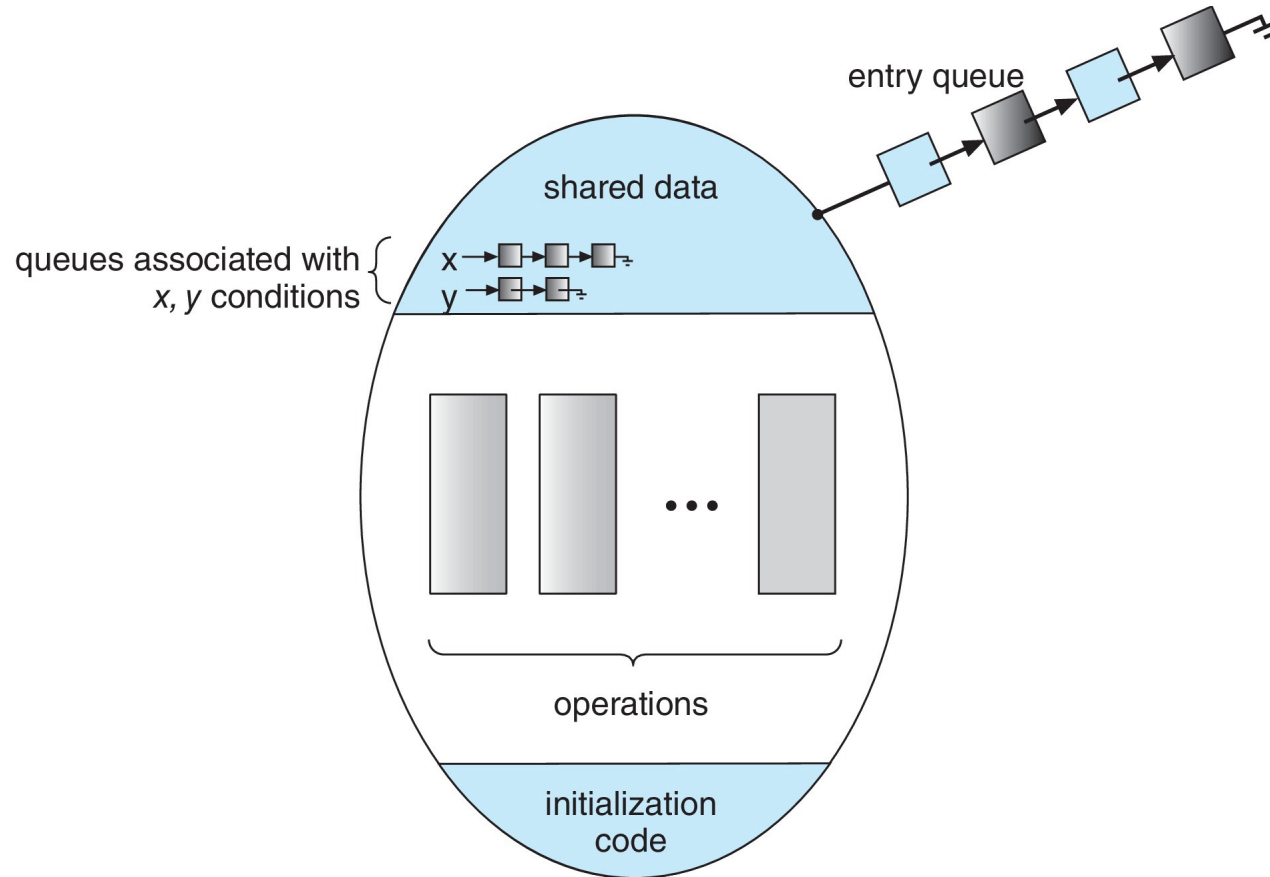


## ■ Condition Variables

- condition `x, y`;
- Two operations are allowed on a condition variable:
  - `x.wait()` : the calling thread that invokes `wait()` is suspended
    - ...until `x.signal()`.
  - `x.signal()` : resumes (wakes up) one of the threads (if any)
    - ...that invoked `x.wait()`
    - If no `x.wait()` on the condition variable `x`, then it has no effect on condition variable `x`.
    - In contrast to `semaphore.signal()`, which always affects the **state** of the **semaphore**.
    - In other words, condition variables are **stateless**, while semaphores are **stateful**.
  - `x.broadcast()` : resumes all threads that invoked `x.wait()`.
    - Often seen in other implementations, such as `pthread_cond_broadcast()` and Java Threads API

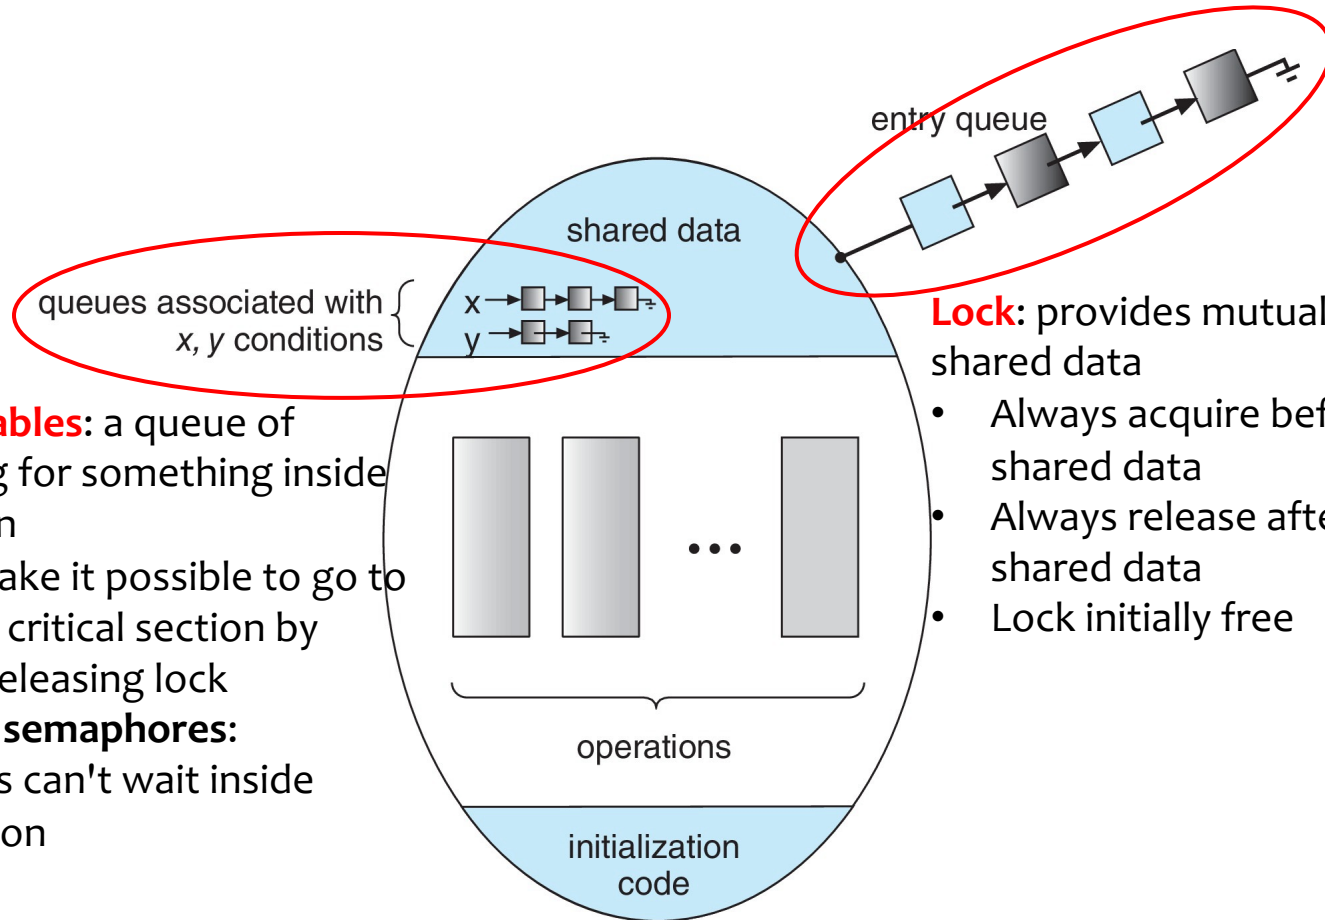
## Monitors

### Schematic View of a Monitor with Condition Variables



## Monitors

### Schematic View of a Monitor with Condition Variables



**Lock:** provides mutual exclusion to shared data

- Always acquire before accessing shared data
- Always release after finishing with shared data
- Lock initially free

**Condition Variables:** a queue of threads waiting for something inside a critical section

- **Key idea:** make it possible to go to sleep inside critical section by atomically releasing lock
- **Contrast to semaphores:** Semaphores can't wait inside critical section

**Monitor:** a **lock** with zero or more **condition variables** for managing concurrent access to shared data.



## ■ POSIX API for Condition Variables

### ■ Data type:

- `pthread_cond_t cond;`

### ■ Basic operations:

- `pthread_cond_wait(&cond, &mutex);`

- **Atomically** put current thread to **block** on the condition variable `cond` and release the *mutex lock* `mutex`

- `pthread_cond_signal(&cond);`

- Wake up one of the threads (if any) that blocks on `cond`.

- `pthread_cond_broadcast(&cond);`

- Wake up all threads (if any) that block on `cond`.



## ■ POSIX Condition Variables Example

```
/* parent_wait.c */
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

void *child(void *arg) {
    printf("Child\n");
    // How to indicate we are done?
    return NULL;
}

int main() {
    pthread_t p;
    printf("Parent BEGIN\n");
    pthread_create(&p, NULL, child, NULL);
    // How to wait for child?
    printf("Parent END\n");
    sleep(1);
    return 0;
}
```

```
$ ./parent_wait
Parent BEGIN
Parent END
Child
```

## ■ POSIX Condition Variables Example

```
/* parent_wait2.c */
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

int done = 0;

void *child(void *arg) {
    printf("Child\n");
    // How to indicate we are done?
    done = 1;
    return NULL;
}

int main() {
    pthread_t p;
    printf("Parent BEGIN\n");
    pthread_create(&p, NULL, child, NULL);
    pthread_join(p, NULL);
    // How to wait for child?
    while (done == 0)
        ; // Spin
    printf("Parent END\n");
    sleep(1);
    return 0;
}
```

```
$ ./parent_wait2
Parent BEGIN
Child
Parent END
```

1. The most natural approach is to use a shared global variable **done** as a flag to indicate whether child thread is done.
2. Before child process terminate, set **done = 1**
3. In the parent thread, keep checking the value of **done** before proceeding.

**Disadvantage:** Hugely inefficient as the parent spins and wastes CPU cycles.

## ■ POSIX Condition Variables Example

```
/* parent_wait3.c */
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

void *child(void *arg) {
    printf("Child\n");
    // How to indicate we are done?
    pthread_exit(NULL);
}

int main() {
    pthread_t p;
    printf("Parent BEGIN\n");
    pthread_create(&p, NULL, child, NULL);
    // How to wait for child?
    pthread_join(p, NULL);
    printf("Parent END\n");
    sleep(1);
    return 0;
}
```

```
$ ./parent_wait3
Parent BEGIN
Child
Parent END
```

Remember `pthread_join()` and `pthread_exit()`?

Obviously, they solve this problem perfectly.

But, how are they implemented under the hood?



## ■ POSIX Condition Variables Example

```
/* parent_wait4.c */
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

void *child(void *arg) {
    printf("Child\n");
    // How to indicate we are done?
    my_thread_exit();
}

int main() {
    pthread_t p;
    printf("Parent BEGIN\n");
    pthread_create(&p, NULL, child, NULL);
    // How to wait for child?
    my_thread_join();
    printf("Parent END\n");
    sleep(1);
    return 0;
}
```

Here, we demonstrate a simple implementation of `thread_exit()` and `thread_join()` using Monitors, which is basically {**lock + condition variables**}.



## ■ POSIX Condition Variables Example

```
/* parent_wait4.c */
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

void *child(void *arg) {
    printf("Child\n");
    // How to indicate we are done?
    my_thread_exit();
}

int main() {
    pthread_t p;
    printf("Parent BEGIN\n");
    pthread_create(&p, NULL, child, NULL);
    // How to wait for child?
    my_thread_join();
    printf("Parent END\n");
    sleep(1);
    return 0;
}
```

```
$ ./parent_wait4
Parent BEGIN
Child
Parent END
```

```
int done = 0;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;

void my_thread_exit() {
    pthread_mutex_lock(&mutex);
    done = 1;
    pthread_cond_signal(&cond);
    pthread_mutex_unlock(&mutex);
}

void my_thread_join() {
    pthread_mutex_lock(&mutex);
    while (done == 0)
        pthread_cond_wait(&cond, &mutex);
    pthread_mutex_unlock(&mutex);
}
```

## ■ POSIX Condition Variables Example

```
/* parent_wait4.c */
#include <stdio.h>
#include <unistd.h>
#include <pthread.h>

void *child(void *arg) {
    printf("Child\n");
    // How to indicate we are done?
    my_thread_exit();
}

int main() {
    pthread_t p;
    printf("Parent BEGIN\n");
    pthread_create(&p, NULL, child, NULL);
    // How to wait for child?
    my_thread_join();
    printf("Parent END\n");
    sleep(1);
    return 0;
}
```

```
$ ./parent_wait4
Parent BEGIN
Child
Parent END
```

```
int done = 0;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;

void my_thread_exit() {
    pthread_mutex_lock(&mutex);
    done = 1;
    pthread_cond_signal(&cond);
    pthread_mutex_unlock(&mutex);
}

void my_thread_join() {
    pthread_mutex_lock(&mutex);
    while (done == 0)
        pthread_cond_wait(&cond, &mutex);
    pthread_mutex_unlock(&mutex);
}
```



Question: Does it spin?

## ■ POSIX Condition Variables Example

```
int done = 0;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;

void my_thread_exit() {
    pthread_mutex_lock(&mutex);
    done = 1;
    pthread_cond_signal(&cond);
    pthread_mutex_unlock(&mutex);
}

void my_thread_join() {
    pthread_mutex_lock(&mutex);
    while (done == 0)
        pthread_cond_wait(&cond, &mutex);
    pthread_mutex_unlock(&mutex);
}
```

```
int done = 0;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;

void my_thread_exit() {
    done = 1;
    pthread_cond_signal(&cond);
}

void my_thread_join() {
    while (done == 0)
        pthread_cond_wait(&cond, &mutex);
}
```

Is the Mutex Lock necessary?



## ■ POSIX Condition Variables Example

```
// Parent Thread           // Child Thread

// call thread_join()
while (done == 0) {
-----
// call thread_exit()
done == 1;

pthread_cond_signal();
// no effect, since no
// thread is waiting
-----
pthread_cond_wait();
// waits forever
}
```

```
int done = 0;
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;

void my_thread_exit() {
    done = 1;
    pthread_cond_signal(&cond);
}

void my_thread_join() {
    while (done == 0)
        pthread_cond_wait(&cond, &mutex);
}
```

In **Monitors**, it is mandatory to **acquire** the lock before any operation on condition variables. After `wait()` or `signal()`, we must **release** the lock.

**Remember, always hold the lock while `signal()` or `wait()`.**

## ■ Mesa vs. Hoare Semantics

- Why do we CV: :**wait**() inside a **while** loop?

```
void my_thread_join() {  
    pthread_mutex_lock(&mutex);  
    while (done == 0)  
        pthread_cond_wait(&cond, &mutex);  
    pthread_mutex_unlock(&mutex);  
}
```

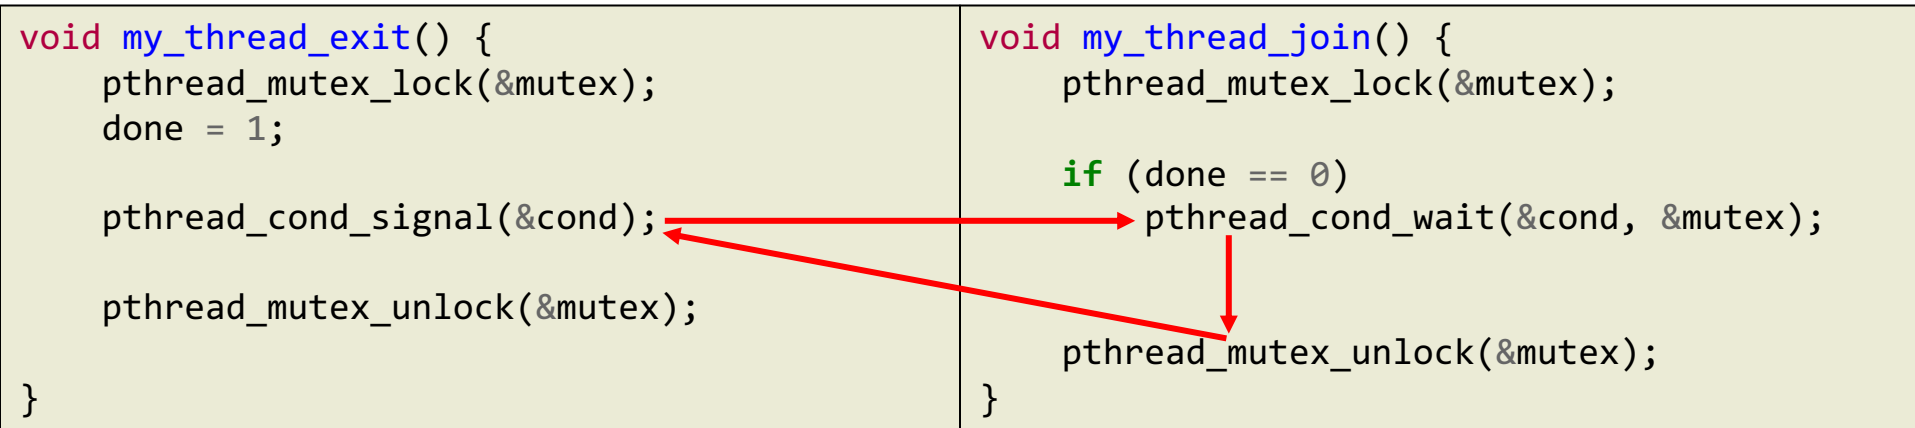
- Why can't we do this (check only once with an **if** statement)?

```
void my_thread_join() {  
    pthread_mutex_lock(&mutex);  
    if (done == 0)  
        pthread_cond_wait(&cond, &mutex);  
    pthread_mutex_unlock(&mutex);  
}
```

- **Answer:** depends on the semantics of Monitor/Condition Variable.
  - **Hoare Semantics:** CV: :**signal**() is immediately (atomically) followed by CV: :**wait**()
  - **Mesa Semantics:** No such guarantee.
    - **Most modern OSes use Mesa Semantics!**

## ■ Hoare Semantics

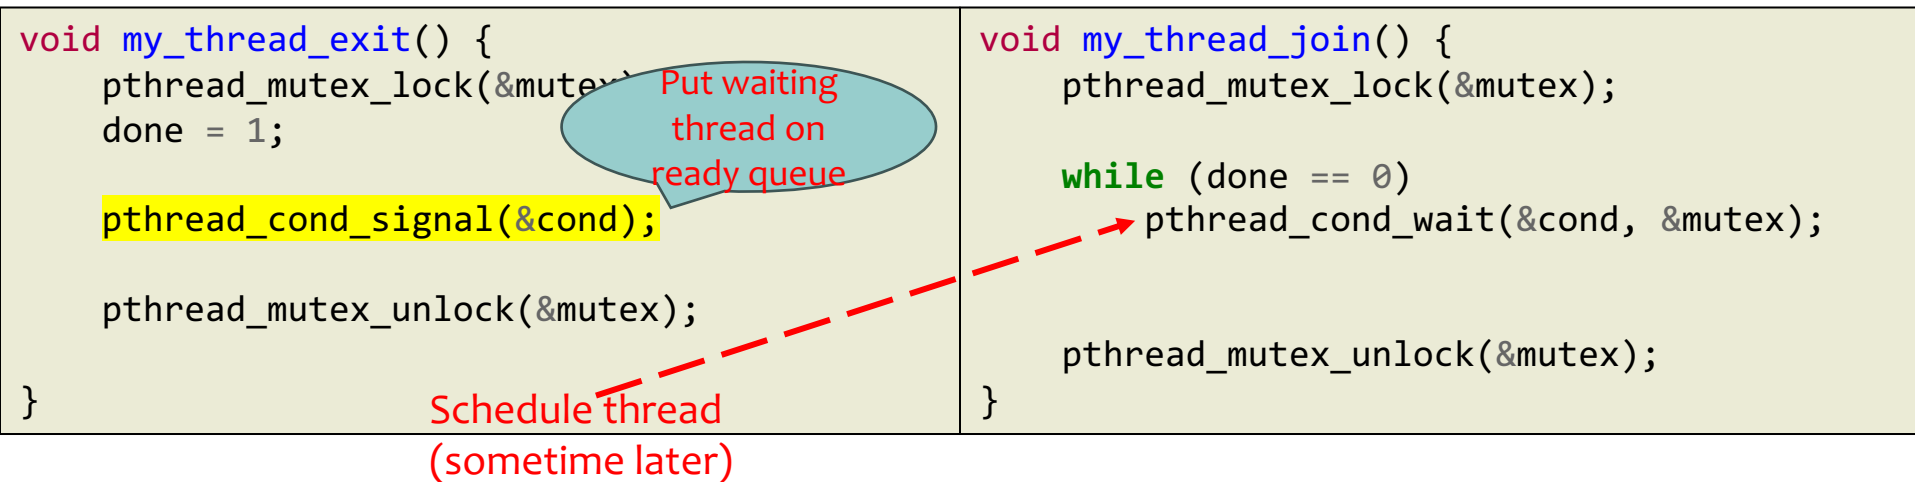
- Signaler gives up lock, CPU to waiter; waiter runs immediately
- Then waiter gives up lock, CPU back to signaler when it exits critical section or if it waits again.



- On first glance, this seems like good semantics
  - Waiter gets to run immediately, condition is still correct!
- Most textbooks talk about Hoare semantics scheduling
  - However, hard to implement, not really necessary
  - Forces a lot of context switching (inefficient)

## ■ Mesa Semantics

- Signaler keeps lock and CPU
- Waiter placed on ready queue with no special priority



- Practically, need to check condition again after `wait()`
  - by the time the waiter gets scheduled, condition may be false again. So, just check again with the while loop
- Most Modern OSes adopt Mesa Semantics
  - Efficient, easier to implement
  - No need for special treatment in the Scheduler.

## ■ Monitors

- Monitors represent the synchronization logic of the program
  - Wait if necessary
  - Signal when change something so any waiting threads can proceed
- Typical structure of monitor-based program (Mesa Semantics):

```
lock
while (need to wait) {
    cond.wait();
unlock
```

```
do something so no need to wait
lock

cond.signal();

unlock
```



## ■ Monitors

### ■ Implementing a Monitor Using Semaphores

- **Monitors:** a **high-level** synchronization construct that allows threads to have both mutual exclusion and the ability to wait(block) for a certain condition to become true. Monitors also automatically handle waking up threads when conditions change.
- **Semaphores:** a **low-level** synchronization primitive that manages access to common resources by using counters. If a semaphore's counter is positive, accessing the resource decrements the counter. If the counter is zero, the thread attempting access is blocked.

## ■ Monitors

### ■ Implementing a Monitor (**Hoare**) Using Semaphores

```
#include <stdio.h>
#include <pthread.h>
#include <semaphore.h>
// Declare semaphores
sem_t mutex, next, sem_done;
int next_count = 0, count_done = 0;

void init_monitor() {
    // Binary semaphore for mutex
    sem_init(&mutex, 0, 1);
    // Binary semaphore for orderly access
    sem_init(&next, 0, 0);
    // Semaphore to signal child completion
    sem_init(&sem_done, 0, 0);
}

void enter_monitor() {
    sem_wait(&mutex);
}

void exit_monitor() {
    if (next_count > 0)
        sem_post(&next);
    else
        sem_post(&mutex);
}
```

```
void wait_child_done() {
    count_done++;
    if (next_count > 0)
        sem_post(&next);
    else
        sem_post(&mutex);
    sem_wait(&sem_done);
    count_done--;
}

void signal_child_done() {
    if (count_done > 0) {
        next_count++;
        sem_post(&sem_done);
        sem_wait(&next);
        next_count--;
    }
}
```

## ■ Monitors

### ■ Implementing a Monitor (**Hoare**) Using Semaphores

```
void *child(void *arg) {
    enter_monitor();
    printf("Child\n");
    signal_child_done();
    exit_monitor();
    return NULL;
}

int main() {
    pthread_t p;

    init_monitor();

    pthread_create(&p, NULL, child, NULL);

    enter_monitor();
    printf("Parent BEGIN\n");
    wait_child_done();
    printf("Parent END\n");
    exit_monitor();

    pthread_join(p, NULL);
    return 0;
}
```

```
void wait_child_done() {
    count_done++;
    if (next_count > 0)
        sem_post(&next);
    else
        sem_post(&mutex);
    sem_wait(&sem_done);
    count_done--;
}

void signal_child_done() {
    if (count_done > 0) {
        next_count++;
        sem_post(&sem_done);
        sem_wait(&next);
        next_count--;
    }
}
```

```
$ ./monitor
Parent BEGIN
Child
Parent END
```




## ■ Kernel Synchronization in Windows

- Uses interrupt masks to protect access to global resources on uniprocessor systems
- Use **spinlocks** on multiprocessor systems
  - Spinlocking-thread will never be preempted
- Also provides **dispatch objects** user-land which may act mutexes, semaphores, events, and timers
  - **Events**
    - An event acts much like a condition variable
  - **Timers** notify one or more thread when time quantum expired
  - Dispatcher objects either **signaled-state** (object available) or **non-signaled state** (thread will block)



## ■ Kernel Synchronization in Windows

- Mutex dispatcher object





## ■ Kernel Synchronization in Linux

- Prior to kernel version 2.6, disables interrupts to implement short critical sections
- Version 2.6 and later: fully preemptive
- Linux provides:
  - Semaphores
  - Atomic integers
  - Spinlocks
  - Reader-writer versions of both
- On single-processor systems, **spinlocks** replaced by enabling and disabling kernel **preemption**

Single Processor	Multiple Processors
Disable kernel preemption	Acquire spin lock
Enable kernel preemption	Release spin lock



## ■ Kernel Synchronization in Linux

- Atomic variables:
  - `atomic_t` is the type for atomic integer

- Consider the variables:

```
atomic_t counter;  
int value;
```

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&amp;counter,5);</code>	<code>counter = 5</code>
<code>atomic_add(10,&amp;counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4,&amp;counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&amp;counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&amp;counter);</code>	<code>value = 12</code>



## ■ POSIX Synchronization

- POSIX Mutex Locks
- POSIX Semaphores
- POSIX Condition Variables





## ■ POSIX Synchronization

### ■ POSIX Mutex Locks

```
#include <pthread.h>

pthread_mutex_t mutex;

/* Create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);

/* Acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* Critical section */

/* Release the mutex lock */
pthread_mutex_unlock(&mutex);
```



## ■ POSIX Synchronization

### ■ POSIX Semaphores

- **Named** Semaphores
- **Unnamed** Semaphores

```
#include <semaphore.h>

sem_t *sem;

/* Named Semaphores */
/* Create the semaphore and initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);

/* Unnamed Semaphores */
/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);

/* Semaphore Wait() / Down() / P() */
sem_wait(sem);

/* Semaphore Signal() / Up() / V() */
sem_post(sem);
```



## ■ POSIX Synchronization

### ■ POSIX Condition Variables

```
#include <pthread.h>

pthread_mutex_t mutex;
pthread_cond_t cond;

/* Initialize mutex lock */
pthread_mutex_init(&mutex, NULL);
/* Initialize condition variable */
pthread_cond_init(&cond, NULL);

/* Typical Monitor Construct: *wait* part */
pthread_mutex_lock(&mutex);
while (a != b)
    pthread_cond_wait(&cond, &mutex);
pthread_mutex_unlock(&mutex);

/* Typical Monitor Construct: *signal* part */
pthread_mutex_lock(&mutex);
a = b;
pthread_cond_signal(&cond);
pthread_mutex_unlock(&mutex);
```



## ■ Alternative Approaches

- Transactional Memory
- OpenMP
- Functional Programming Languages



## ■ Transactional Memory

- Consider a function `update()` that must be called atomically. One option is to use mutex locks:

```
void update ()
{
    acquire();

    /* modify shared data */

    release();
}
```

- However, using synchronization mechanisms such as mutex locks and semaphores involve many potential problems, including deadlock.
- Additionally, as the number of threads increases, traditional locking doesn't scale well, because the level of contention among threads for lock ownership becomes very high.



## ■ Transactional Memory

- A memory transaction is a sequence of read-write operations to memory that are performed **atomically**. A transaction can be completed by adding `atomic{S}` which ensure statements in `S` are executed atomically.

```
void update ()
{
    atomic {
        /* modify shared data */
    }
}
```

- Advantage over locks: the **transactional memory system**, rather than the developer, is responsible for guaranteeing atomicity.
- Transactional memory can be implemented in:
  - Software transactional memory (STM)
  - Hardware transactional memory (HTM)



## ■ OpenMP

- OpenMP is a set of compiler directives and API that support parallel programming.

```
void update(int value) {  
    #pragma omp critical  
    {  
        count += value;  
    }  
}
```

- The code contained within the `#pragma omp critical` directive is treated as a critical section and performed atomically.



## ■ Functional Programming Languages

- Functional programming languages offer a different paradigm than procedural languages in that they do not maintain state.
- Variables are treated as immutable and cannot change state once they have been assigned a value.
- There is increasing interest in functional languages such as Erlang and Scala for their approach in handling data races.





**Thank you!**