



DCS216 Operating Systems

Lecture 15 Synchronization (4)

Apr 17th, 2024

Instructor: Xiaoxi Zhang
Sun Yat-sen University



■ Content

- Liveness
- Starvation vs. Deadlock
- Deadlock in Multithreaded Applications
- System Model & Resource-Allocation Graph
- Deadlock Characterization
- Methods for Handling Deadlocks
 - Deadlock **Prevention**
 - Invalidate one of the necessary conditions
 - Deadlock **Detection**
 - Deadlock Detection Algorithm
 - Recovery from Deadlock
 - Deadlock **Avoidance**
 - Banker's Algorithm

■ Liveness (活性)

- Liveness refers to the properties that a system must satisfy to ensure that processes make progress during their execution life cycle.
 - a guarantee that something good **eventually** happens
- Key aspects of liveness:
 - Progress
 - Freedom from Deadlock
 - Freedom from Starvation
 - Fairness

■ Liveness (活性)

- Liveness refers to the properties that a system must satisfy to ensure that processes make progress during their execution life cycle.
 - a guarantee that something good **eventually** happens
- Key aspects of liveness:
 - **Progress:** ensures that if a process needs to perform an action, it will eventually be able to do so.
 - A failure in this property might result in **deadlock** or **livelock**.
 - **Freedom from Deadlock:** Deadlock occurs when processes are **stuck** waiting indefinitely for resources that are held by each other.
 - **Freedom from Starvation:** This property ensures that every process gets a chance to proceed. Starvation occurs when a process is **perpetually denied** necessary resources, usually because of scheduling or resource allocation policies.
 - **Fairness:** This is often considered as part of liveness, ensuring that all processes are treated in a **fair** manner over time. This means that all processes will eventually be given CPU time and access to resources.




■ Starvation vs Deadlock

■ **Starvation**: thread waits indefinitely

- Example: low-priority thread waiting for resources constantly in use by high-priority threads, also known as **Priority Inversion (优先级反转)**.

■ **Deadlock**: circular waiting for resources

- Thread A owns Res 1, and is waiting for Res 2
- Thread B owns Res 2, and is waiting for Res 1



■ **Deadlock** \Rightarrow **Starvation** but not vice versa

- Starvation can end (but doesn't have to)
- Deadlock can't end without external intervention



■ Single-Lane Bridge Crossing



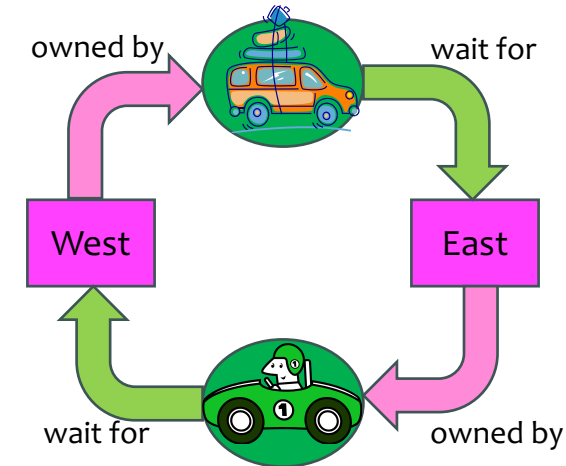

Only one car can
cross at a time

Two lanes converged
into one



■ Single-Lane Bridge Crossing

- Each segment of road can be viewed as a resource
 - Car must own the segment under them
 - Must acquire segment that they are moving into
- For bridge: must acquire both halves
 - Traffic only in one direction at a time



- **Deadlock:** Shown above when two cars from opposite directions meet in the middle
 - Each owns one segment and requests to acquire the other
 - Deadlock resolved if one car backs up (preempt resources and rollback)
 - Several cars may have to be backed up
- **Starvation:** (not deadlock)
 - East-going traffic really fast \Rightarrow no one can go to west



Deadlock with Locks

- This lock pattern exhibits non-deterministic deadlock
 - Sometimes it happens, sometimes it doesn't
- Really hard to debug!

Thread A:

```
acquire(&x);
```

```
acquire(&y);
```

...

```
release(&y);
```

```
release(&x);
```

Thread B:


```
acquire(&y);
```

```
acquire(&x);
```

...

```
release(&x);
```

```
release(&y);
```





Deadlock with Locks (**Unlucky** Case)

- This lock pattern exhibits non-deterministic deadlock
 - Sometimes it happens, sometimes it doesn't
- Really hard to debug!

Thread A:

```
acquire(&x);
```

```
acquire(&y); <stalled>
```

```
<unreachable>
```

```
...
```

```
release(&y);
```

```
release(&x);
```

Thread B:

```
acquire(&y);
```


```
acquire(&x); <stalled>
```

```
<unreachable>
```

```
...
```

```
release(&x);
```

```
release(&y);
```





■ Deadlock with Locks (**Lucky** Case)

- Sometimes, deadlock won't occur with proper scheduling

Thread A:

```
acquire(&x);
```

```
acquire(&y);
```


...

```
release(&y);
```

```
release(&x);
```

Thread B:

```
acquire(&y);
```





■ Deadlock with Locks (**Lucky** Case)

- Sometimes, deadlock won't occur with proper scheduling

Thread A:

```
acquire(&x);
```

```
acquire(&y);
```

```
...
```

```
release(&y);
```

```
release(&x);
```

Thread B:

```
acquire(&y);
```


```
<y acquired>
```

```
acquire(&x);
```

```
...
```

```
release(&x);
```

```
release(&y);
```





Deadlock with Locks

```
/* deadlock.c */
pthread_mutex_t lock1, lock2;

void *threadA(void *arg) {
    pthread_mutex_lock(&lock1);
    usleep(100);
    pthread_mutex_lock(&lock2);

    printf("Thread A working...\n");

    pthread_mutex_unlock(&lock2);
    pthread_mutex_unlock(&lock1);
    pthread_exit(NULL);
}

void *threadB(void *arg) {
    pthread_mutex_lock(&lock2);
    usleep(100);
    pthread_mutex_lock(&lock1);

    printf("Thread B working...\n");

    pthread_mutex_unlock(&lock1);
    pthread_mutex_unlock(&lock2);
    pthread_exit(NULL);
}
```

```
int main() {
    pthread_mutex_init(&lock1, NULL);
    pthread_mutex_init(&lock2, NULL);

    pthread_t p1, p2;
    pthread_create(&p1, NULL, threadA, NULL);
    pthread_create(&p2, NULL, threadB, NULL);

    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    return 0;
}
```

```
$ ./deadlock
<process stuck...>
^C
```

■ Livello

```
/* livelock.c */
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

pthread_mutex_t lock1, lock2;

void* threadA(void* arg) {
    int done = 0;
    while (!done) {
        pthread_mutex_lock(&lock1);
        printf("Thread A acquired Lock 1\n");
        // Try to acquire Lock2
        if (pthread_mutex_trylock(&lock2) == 0) {
            printf("Thread A acquired Lock 2\n");
            // Simulate work
            pthread_mutex_unlock(&lock2);
            printf("Thread A released Lock 2\n");
            done = 1;
        }
        usleep(100);
        pthread_mutex_unlock(&lock1);
        printf("Thread A released Lock 1\n");
    }
    pthread_exit(0);
}
```

```
$ ./livelock
Thread A acquired Lock 1
Thread B acquired Lock 2
Thread A released Lock 1
Thread A acquired Lock 1
Thread B released Lock 2
Thread B acquired Lock 2
...
<continue forever>
^C
```

```
void* threadB(void* arg) {
    int done = 0;
    while (!done) {
        pthread_mutex_lock(&lock2);
        printf("Thread B acquired Lock 2\n");
        // Try to acquire Lock1
        if (pthread_mutex_trylock(&lock1) == 0) {
            printf("Thread B acquired Lock 1\n");
            // Simulate work
            pthread_mutex_unlock(&lock1);
            printf("Thread B released Lock 1\n");
            done = 1;
        }
        usleep(100);
        pthread_mutex_unlock(&lock2);
        printf("Thread B released Lock 2\n");
    }
    pthread_exit(0);
}
```



■ Other Types of Deadlock

- Threads often block waiting for resources
 - Mutex Locks
 - Terminals
 - Printers
 - Memory
 - CD Drives
- Threads often block waiting for other threads
 - Pipes
 - Sockets
- Deadlocks can occur on any of these!



Deadlock with Memory Space

- If there're only **3MB** of free space, we get the same deadlock situation

Thread A:

```
malloc_wait(2MB);  
malloc_wait(2MB);
```

...


```
free(2MB);  
free(2MB);
```

Thread B:

```
malloc_wait(2MB);  
malloc_wait(2MB);
```

...

```
free(2MB);  
free(2MB);
```





Deadlock with Memory Space

- If there're only **3MB** of free space, we get the same deadlock situation

Thread A:

```
malloc_wait(2MB);
```

```
malloc_wait(2MB);
```

```
...
```

```
free(2MB);
```

```
free(2MB);
```

Thread B:


```
malloc_wait(2MB);
```

```
malloc_wait(2MB);
```

```
...
```

```
free(2MB);
```


```
free(2MB);
```





■ Dining-Philosopher Problem


- Five chopsticks for five philosophers
 - Philosopher will grab any one they can.
 - One chopstick at a time.
 - Need two chopsticks to eat





■ Dining-Philosopher Problem

- Five chopsticks for five philosophers
 - What if they all grab the one chopstick on their **right** at the same time?
 - **Deadlock!**






■ Dining-Philosopher Problem

■ Deadlock

- How to fix deadlock?
- How to prevent deadlock?





■ Four (necessary) requirements for Deadlock

- Mutual Exclusion
- Hold and Wait
- No Preemption
- Circular Wait



■ Four (necessary) requirements for Deadlock

■ Mutual Exclusion

- Only one thread at a time can use a resource

■ Hold and Wait

- Thread holding at least one resource is waiting to acquire additional resources held by other threads

■ No Preemption

- Resources are released only voluntarily by the thread holding that resources, after thread is finished with it


■ Circular Wait

- There exists a set $\{T_1, T_2, \dots, T_n\}$ of waiting threads
 - T_1 is waiting for a resource held by T_2
 - T_2 is waiting for a resource held by T_3
 - ...
 - T_n is waiting for a resource held by T_1



■ System Model

- A set of n Threads $\{T_1, T_2, \dots, T_n\}$
- Resource types $\{R_1, R_2, \dots, R_m\}$
 - CPU cycles, Memory space, I/O devices
- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows:
 - Request() / Use() / Release()



■ Resource-Allocation Graph: (V, E)

- V is partitioned into two types:
 - $T = \{T_1, T_2, \dots, T_n\}$ are the set of **threads** in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$ are the set of **resource types** in the system.
- E can be categorized into two types:
 - **Request** Edge – directed edge $T_i \rightarrow R_j$
 - **Assignment** Edge – directed edge $R_j \rightarrow T_i$



Resource-Allocation Graph Example

The sets T , R , and E :


- $T = \{T_1, T_2, T_3\}$
- $R = \{R_1, R_2, R_3, R_4\}$
- $E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3, R_1 \rightarrow T_2, R_2 \rightarrow T_2, R_2 \rightarrow T_1, R_3 \rightarrow T_3\}$

Resource instances:

- 1 instance of resource type R_1
- 2 instances of resource type R_2
- 1 instance of resource type R_3
- 3 instances of resource type R_4

Thread states:

- T_1 is holding an instance of R_2 and is waiting for an instance of R_1
- T_2 is holding an instance of R_1 and an instance of R_2 , and is waiting for an instance of R_3
- T_3 is holding an instance of R_3





Resource-Allocation Graph Example

The sets T , R , and E :


- $T = \{T_1, T_2, T_3\}$
- $R = \{R_1, R_2, R_3, R_4\}$
- $E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3, R_1 \rightarrow T_2, R_2 \rightarrow T_2, R_2 \rightarrow T_1, R_3 \rightarrow T_3\}$

Resource instances:

- 1 instance of resource type R_1
- 2 instances of resource type R_2
- 1 instance of resource type R_3
- 3 instances of resource type R_4

Thread states:

- T_1 is **holding an instance of R_2** and is **waiting for an instance of R_1**
- T_2 is holding an instance of R_1 and an instance of R_2 , and is waiting for an instance of R_3
- T_3 is holding an instance of R_3





Resource-Allocation Graph Example

The sets T , R , and E :


- $T = \{T_1, T_2, T_3\}$
- $R = \{R_1, R_2, R_3, R_4\}$
- $E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3,$
 $R_1 \rightarrow T_2, R_2 \rightarrow T_2, R_2 \rightarrow T_1, R_3 \rightarrow T_3\}$

Resource instances:

- 1 instance of resource type R_1
- 2 instances of resource type R_2
- 1 instance of resource type R_3
- 3 instances of resource type R_4

Thread states:

- T_1 is holding an instance of R_2 and is waiting for an instance of R_1
- T_2 is **holding an instance of R_1 and an instance of R_2** ,
and is **waiting for an instance of R_3**
- T_3 is holding an instance of R_3





Resource-Allocation Graph Example

The sets T , R , and E :


- $T = \{T_1, T_2, T_3\}$
- $R = \{R_1, R_2, R_3, R_4\}$
- $E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3, R_1 \rightarrow T_2, R_2 \rightarrow T_2, R_2 \rightarrow T_1, R_3 \rightarrow T_3\}$

Resource instances:

- 1 instance of resource type R_1
- 2 instances of resource type R_2
- 1 instance of resource type R_3
- 3 instances of resource type R_4


Thread states:

- T_1 is holding an instance of R_2 and is waiting for an instance of R_1
- T_2 is holding an instance of R_1 and an instance of R_2 , and is waiting for an instance of R_3
- T_3 is **holding an instance of R_3**






■ Resource-Allocation Graph Example




Simple Resource
Allocation Graph




Allocation Graph
with **Deadlock**



Resource-Allocation Graph Example




Simple Resource
Allocation Graph




Allocation Graph
with **Deadlock**



Resource-Allocation Graph Example




Simple Resource
Allocation Graph



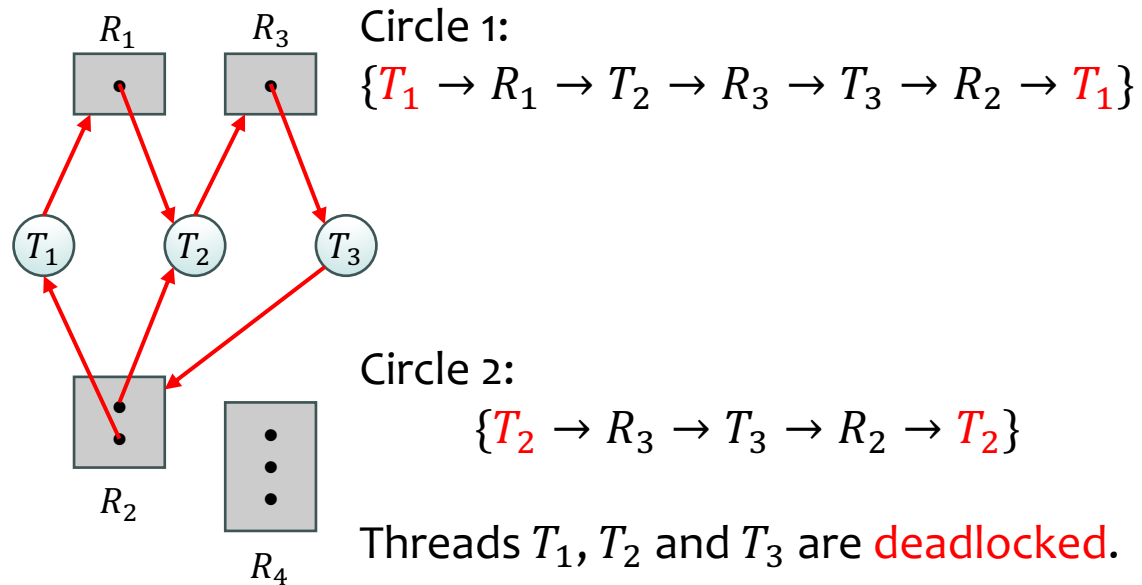
Allocation Graph
with **Deadlock**



Resource-Allocation Graph Example




Simple Resource
Allocation Graph




Allocation Graph
with **Deadlock**




Resource-Allocation Graph Example



Simple Resource
Allocation Graph




Allocation Graph with
Deadlock



Allocation Graph with
Cycle, but No Deadlock



■ Resource-Allocation Graph Example



Circle:


$\{T_1 \rightarrow R_1 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1\}$

...but no **deadlock**...

Allocation Graph with
Cycle, but No **Deadlock**



Resource-Allocation Graph Example



Allocation Graph with
Cycle, but No **Deadlock**

Circle:


$\{T_1 \rightarrow R_1 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1\}$

...but no **deadlock**...

Because thread T_4 may **release** its instance of resource type R_2 ...



Resource-Allocation Graph Example



Allocation Graph with
Cycle, but No **Deadlock**

Circle:

$\{T_1 \rightarrow R_1 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1\}$

...but no **deadlock**...


Because thread T_4 may **release** its instance of resource type R_2 ...

That instance of resource type R_2 can then be allocated to T_3 , breaking the cycle.




Resource-Allocation Graph


- If graph contains no cycles \Rightarrow **no deadlock**
- If graph contains a cycle \Rightarrow
 - if **only one** instance per resource type, then **deadlock**.
 - if **several** instances per resource type, **possibility** of **deadlock**.



Simple Resource
Allocation Graph



Allocation Graph
with **Deadlock**



Allocation Graph with
Cycle, but No **Deadlock**




Deadlock Detection Algorithm

- Data structures: ($[x]$ represents row vector of size m)
 - $[Avail]$: Number of available resources of each type
 - $[Alloc_i]$: Current resources held by thread i ($i \in [1, n]$)
 - $[Request_i]$: Current requests from thread i ($i \in [1, n]$)
- See if tasks can eventually terminate on their own

```
add all  $T[i]$  to UNFINISHED
do {
    done = true
    for each  $T[i]$  in UNFINISHED {
        if ( $[Request_i] \leq [Avail]$ ) {
            remove  $T[i]$  from UNFINISHED
             $[Avail] = [Avail] + [Alloc_i]$ 
            done = false
        }
    }
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>		UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	0	0	
T1	0	1	1	0			
T2	1	0	0	0			
T3	1	0	0	1			
T4	0	1	0	0			

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	0 0	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	0 0	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	0 0	
T1	0	1	1 0			T1
T2	1	0	0 0			T2
T3	1	0	0 1			T3
T4	0	1	0 0			T4

- See if tasks can eventually terminate on their own

```

add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)

```



- Nodes left in UNFINISHED \Rightarrow **deadlocked**




Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	0 0	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

```
add all T[i] to UNFINISHED
do {
  done = true
  for each T[i] in UNFINISHED {
    if ([Requesti] <= [Avail]) {
      remove T[i] from UNFINISHED
      [Avail] = [Avail] + [Alloci]
      done = false
    }
  }
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>		UNFINISHED
	R1	R2	R1	R2	0	0	
T1	0	1	1	0			T1
T2	1	0	0	0			T2
T3	1	0	0	1			T3
T4	0	1	0	0			T4


- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```

do {
    done = true
    for each  $T[i]$  in UNFINISHED {
        if ( $[Request_i] \leq [Avail]$ ) {
            remove  $T[i]$  from UNFINISHED
             $[Avail] = [Avail] + [Alloc_i]$ 
            done = false
        }
    }
} while (done == false)

```



- Nodes left in UNFINISHED \Rightarrow **deadlocked**




Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 0	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

```
add all T[i] to UNFINISHED
do {
  done = true
  for each T[i] in UNFINISHED {
    if ([Requesti] <= [Avail]) {
      remove T[i] from UNFINISHED
      [Avail] = [Avail] + [Alloci]
      done = false
    }
  }
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 0	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 0	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 0	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 1	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>		UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1	1	
T1	0	1	1	0			T1
T2	1	0	0	0			T2
T3	1	0	0	1			T3
T4	0	1	0	0			T4


- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```

do {
    done = true
    for each  $T[i]$  in UNFINISHED {
        if ( $[Request_i] \leq [Avail]$ ) {
            remove  $T[i]$  from UNFINISHED
             $[Avail] = [Avail] + [Alloc_i]$ 
            done = false
        }
    }
} while (done == false)

```



- Nodes left in UNFINISHED \Rightarrow **deadlocked**



Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 1	
T1	0	1	1 0			T1
T2	1	0	0 0			T2
T3	1	0	0 1			T3
T4	0	1	0 0			T4

- See if tasks can eventually terminate on their own

add all $T[i]$ to UNFINISHED

```
do {  
    done = true  
    for each  $T[i]$  in UNFINISHED {  
        if ( $[Request_i] \leq [Avail]$ ) {  
            remove  $T[i]$  from UNFINISHED  
             $[Avail] = [Avail] + [Alloc_i]$   
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>		UNFINISHED
	R1	R2	R1	R2	1	1	
T1	0	1	1	0			T1
T2	1	0	0	0			T2
T3	1	0	0	1			T3
T4	0	1	0	0			T4

- See if tasks can eventually terminate on their own

```

add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)

```



- Nodes left in UNFINISHED \Rightarrow **deadlocked**



Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>		UNFINISHED
	R1	R2	R1	R2	1	2	
T1	0	1	1	0			T1
T2	1	0	0	0			T2
T3	1	0	0	1			T3
T4	0	1	0	0			T4


- See if tasks can eventually terminate on their own

```

add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)

```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 2	
<i>T1</i>	0	1	1	0		<i>T1</i>
<i>T2</i>	1	0	0	0		<i>T2</i>
<i>T3</i>	1	0	0 1			<i>T3</i>
<i>T4</i>	0	1	0	0		<i>T4</i>


- See if tasks can eventually terminate on their own

```

add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Request]i <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloc]i
            done = false
        }
    }
} while (done == false)

```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	1 2	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own


add all $T[i]$ to UNFINISHED

```

do {
    done = true
    for each  $T[i]$  in UNFINISHED {
        if ( $[Request_i] \leq [Avail]$ ) {
            remove  $T[i]$  from UNFINISHED
             $[Avail] = [Avail] + [Alloc_i]$ 
            done = false
        }
    }
} while (done == false)

```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	2 2	UNFINISHED
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4


- See if tasks can eventually terminate on their own

```

add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)

```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm

	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>		UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	2	2	
<i>T1</i>	0	1	1	0			<i>T1</i>
<i>T2</i>	1	0	0	0			<i>T2</i>
<i>T3</i>	1	0	0	1			<i>T3</i>
<i>T4</i>	0	1	0	0			<i>T4</i>


- See if tasks can eventually terminate on their own

add all *T*[*i*] to UNFINISHED

```

do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)

```



- Nodes left in UNFINISHED \Rightarrow **deadlocked**



Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	R1	R2	R1	R2	2 2	
T1	0	1	1	0		T1
T2	1	0	0	0		T2
T3	1	0	0	1		T3
T4	0	1	0	0		T4

- See if tasks can eventually terminate on their own

add all T[i] to UNFINISHED

```
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Deadlock Detection Algorithm


	<u>Alloc</u>		<u>Request</u>		<u>Avail</u>	UNFINISHED
	<i>R1</i>	<i>R2</i>	<i>R1</i>	<i>R2</i>	2 2	
<i>T1</i>	0	1	1	0		<i>T1</i>
<i>T2</i>	1	0	0	0		<i>T2</i>
<i>T3</i>	1	0	0	1		<i>T3</i>
<i>T4</i>	0	1	0	0		<i>T4</i>

- See if tasks can eventually terminate on their own

add all *T*[*i*] to UNFINISHED

```
do {  
    done = true  
    for each T[i] in UNFINISHED {  
        if ([Requesti] <= [Avail]) {  
            remove T[i] from UNFINISHED  
            [Avail] = [Avail] + [Alloci]  
            done = false  
        }  
    }  
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**






Deadlock Detection Algorithm

- Data structures: ($[x]$ represents row vector of size m)
 - $[Avail]$: Number of available resources of each type
 - $[Alloc_i]$: Current resources held by thread i ($i \in [1, n]$)
 - $[Request_i]$: Current requests from thread i ($i \in [1, n]$)
- See if tasks can eventually terminate on their own

```
add all  $T[i]$  to UNFINISHED
do {
    done = true
    for each  $T[i]$  in UNFINISHED {
        if ( $[Request_i] \leq [Avail]$ ) {
            remove  $T[i]$  from UNFINISHED
             $[Avail] = [Avail] + [Alloc_i]$ 
            done = false
        }
    }
} while (done == false)
```



Time Complexity: $O(m \times n^2)$

- Nodes left in UNFINISHED \Rightarrow **deadlocked**



■ Deadlock Detection Algorithm Usage

- When, and how often, to invoke the detection algorithm?
- It depends on:
 - How **often** is a deadlock likely to occur?
 - How **many** threads will be affected by deadlock when it happens?
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked threads "**caused**" the deadlock.



■ Recovery from Deadlock

■ Process and Thread Termination

- **Method #1:** Abort all deadlocked processes
- **Method #2:** Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort? **Factors:**
 - Priority of the process
 - How long the process has computed and how much longer?
 - What resources has the process used?
 - How many more resources does the process need?
 - How many processes will need to be terminated?

■ Resource Preemption

- **Selecting a victim:** minimize cost
- **Rollback:** return to some safe state, restart the process for that state
- **Starvation:**
 - same process always be picked as victim
 - include number of rollback in cost factor



■ Methods for Handling Deadlocks

- **Ignore** the problem of deadlock altogether and pretend that deadlocks never occur in the system.

- **Ostrich Algorithm** (鸵鸟算法)
 - Used by most OSes (e.g., Linux and Windows)
 - Up to the developers to handle deadlocks.



- **Ensure** that the system will never enter a deadlock state:

- Deadlock **Prevention**

- by constraining how requests for resources can be made.
- simple and direct by structurally eliminating deadlocks

- Deadlock **Avoidance**

- requires that the OS be given additional info in advance concerning which resources a thread will request and use during its lifetime.
- dynamic, sophisticated tracking and management of resources

- **Allow** the system to enter a deadlock state, **detect** it, and then **recover**.



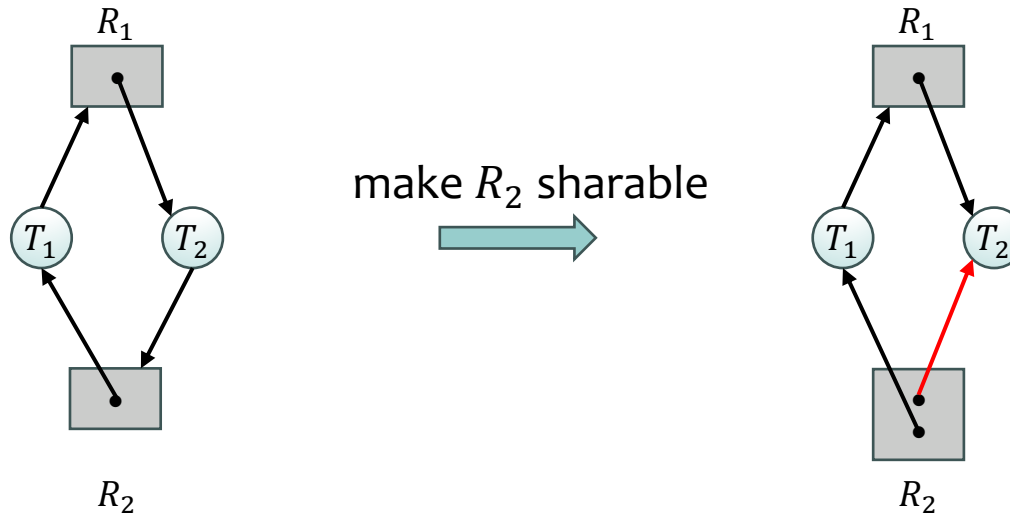
■ Deadlock Prevention

- Invalidate one of the four **necessary** conditions for deadlock:
 - Mutual Exclusion
 - Hold and Wait
 - No Preemption
 - Circular Wait



Deadlock Prevention

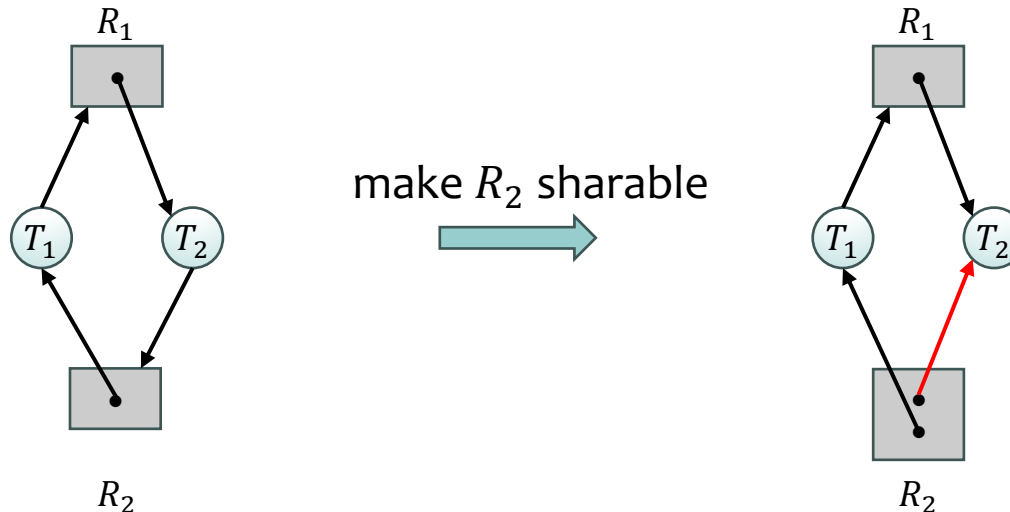
- Invalidate one of the four **necessary** conditions for deadlock:
- **Mutual Exclusion**
 - not required for sharable resources (e.g., read-only files)





Deadlock Prevention

- Invalidate one of the four **necessary** conditions for deadlock:
- **Mutual Exclusion**
 - not required for sharable resources (e.g., read-only files)
 - In general, we cannot prevent deadlocks **simply by denying mutual exclusion**, because some resources (e.g., **mutex locks**) are **intrinsically nonsharable**.



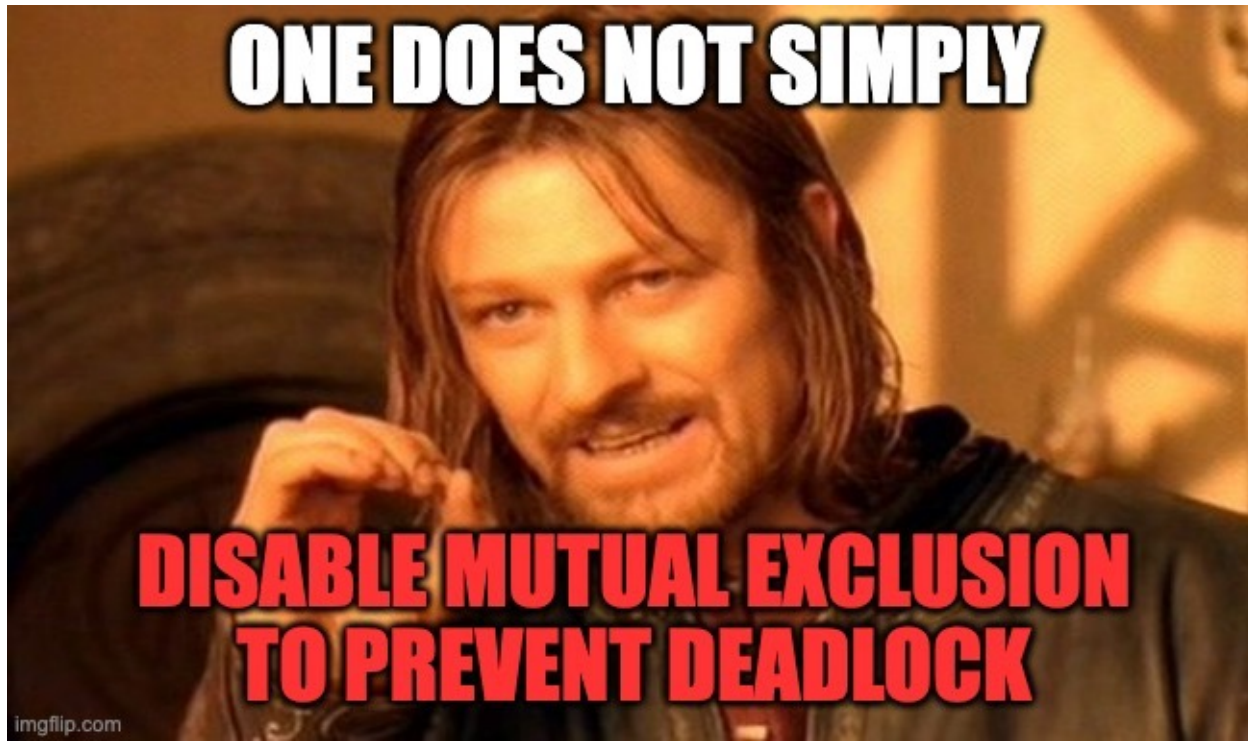


■ Deadlock Prevention

- Invalidate one of the four **necessary** conditions for deadlock:

- **Mutual Exclusion**

- not required for sharable resources (e.g., read-only files)
- In general, we cannot prevent deadlocks **simply by denying mutual exclusion**, because some resources (e.g., **mutex locks**) are **intrinsically nonsharable**.





■ Deadlock Prevention

- Invalidate one of the four **necessary** conditions for deadlock:

- **Hold and Wait**

- A thread must be holding at least one resource and waiting to acquire additional resources that are currently being held by other threads.

- ~~■ **Hold and Wait**~~

- must guarantee that whenever a thread requests a resource, it does not hold any other resources
 - **Method #1:** require threads to request and be allocated all its resources **before** it begins execution
 - **Method #2:** allow a thread to request resources **only** when it has **none** allocated to it.
 - **Disadvantages:**
 - Low resource utilization
 - Starvation is possible



■ Deadlock Prevention

- Invalidate one of the four **necessary** conditions for deadlock:

- **No Preemption**

- Resources cannot be preempted; that is, a resource can be released only voluntarily by the thread holding it, after that thread has completed its task.

- ~~■ **No Preemption**~~

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
 - Preempted resources are added to the list of resources for which the thread is waiting
 - Thread will be restarted only when it can regain its old resources, as well as the new ones that it is requesting



Deadlock Prevention

- Invalidate one of the four **necessary** conditions for deadlock:


■ Circular Wait


- There exists a set $\{T_1, T_2, \dots, T_n\}$ of waiting threads such that

$$T_1 \rightarrow R_? \rightarrow T_2 \rightarrow R_? \rightarrow \dots \rightarrow T_n \rightarrow R_? \rightarrow T_1$$

■ ~~Circular Wait~~ (most common)

- Impose a total **ordering** of all resource types, and require that each thread requests resources in **an order** of enumeration.
- Simply **assign** each resource (i.e., mutex locks) a unique number.
- Resources must be acquired in **ascending order**.
- For example, if $F(\text{mutex1}) = 1$, $F(\text{mutex2}) = 5$, then:

Thread A: 
acquire(&mutex1);
acquire(&mutex2);
...
release(&mutex2);
release(&mutex1);

Thread B: 
acquire(&mutex2);
acquire(&mutex1);
...
release(&mutex1);
release(&mutex2);



■ Deadlock Avoidance

- **Idea:** When a thread requests a resource, the OS checks if it would result in deadlock
 - If not, it grants the resource right away.
 - If so, it waits for other threads to release resources.
- Does it work?
 - No!

Thread A:

`acquire(&x);`

`acquire(&y);`

`...`

`release(&y);`

`release(&x);`

Thread B:

`acquire(&y);`

`acquire(&x);`

`...`

`release(&x);`

`release(&y);`



■ Deadlock Avoidance

- **Idea:** When a thread requests a resource, the OS checks if it would result in deadlock
 - If not, it grants the resource right away.
 - If so, it waits for other threads to release resources.
- Does it work?
 - No!

	<u>Thread A:</u>	<u>Thread B:</u>	
	<code>acquire(&x);</code>	<code>acquire(&y);</code>	
Blocks...	<code>acquire(&y);</code>	<code>acquire(&x);</code>	Wait?
	But it's already too late...
	<code>release(&y);</code>	<code>release(&x);</code>	
	<code>release(&x);</code>	<code>release(&y);</code>	



■ Deadlock Avoidance: Three States

■ Safe State

- System can delay resource acquisition to prevent deadlock


■ Unsafe State

Deadlock avoidance: prevent system from reaching an **unsafe** state

- No deadlock yet...
- But threads can request resources in a pattern that **unavoidably** leads to deadlock

■ Deadlocked State

- There exists a deadlock in the system.
- Also considered "unsafe"






■ Safe State

- When a thread requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence $\langle T_1, T_2, \dots, T_n \rangle$ of **ALL** threads such that for each T_i , the resources that T_i can still request can be satisfied by currently available resources + resources held by all the T_j , with $j < i$.
- In other words:
 - If T_i resource needs are not immediately available, then T_i can wait until all T_j have finished
 - When T_j is finished, T_i can obtain needed resources, execute, return allocated resources, and terminate
 - When T_i terminates, T_{i+1} can obtain its needed resources, and so on...



■ Basic Facts

- If a system is in **safe** state \Rightarrow no deadlocks
- If a system is in **unsafe** state \Rightarrow possibility of deadlocks
- **Deadlock Avoidance** \Rightarrow ensure that a system will never enter an **unsafe** state





Deadlock Avoidance

- **Idea:** When a thread requests a resource, the OS checks if it would result in ~~deadlock~~ **an unsafe state**
 - If not, it grants the resource right away.
 - If so, it waits for other threads to release resources.

Thread A:

```
acquire(&x);  
acquire(&y);  
...  
release(&y);  
release(&x);
```

Thread B:

```
acquire(&y);  
acquire(&x);  
...  
release(&x);  
release(&y);
```

Wait until
Thread A
releases
mutex x



■ Deadlock Avoidance

- Requires that the system has some additional presumed information available
 - Simplest and most useful model requires that each thread declare the maximum number of resources of each type that it may need
 - The deadlock-avoidance algorithm dynamically examines the resource allocation state to ensure that there can never be a circular-wait condition
 - Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes



■ Banker's Algorithm for Avoiding Deadlocks

■ The idea:

- Declare maximum resource needs in advance
- Allow particular threads to proceed if
 - $(\text{available resources} - \text{\#requested}) \geq \text{max remaining that might be needed by any thread}$

■ Banker's algorithm:

- Allocate resources dynamically
- Evaluate each thread, and grant access if some ordering of threads is still deadlock-free afterwards
- **Technique:** pretend each request is granted, then run **deadlock detection algorithm** by substituting:
$$([\text{Request}_i] \leq [\text{Avail}]) \Rightarrow ([\text{Max}_i - \text{Alloc}_i] \leq [\text{Avail}])$$
- **Grant request** if the resulting state is deadlock-free (conservative)




Deadlock Detection Algorithm (Revisit)

- Data structures: ($[x]$ represents row vector of size m)
 - $[Avail]$: Number of available resources of each type
 - $[Alloc_i]$: Current resources held by thread i ($i \in [1, n]$)
 - $[Request_i]$: Current requests from thread i ($i \in [1, n]$)
- See if tasks can eventually terminate on their own

```
add all  $T[i]$  to UNFINISHED
do {
    done = true
    for each  $T[i]$  in UNFINISHED {
        if ( $[Request_i] \leq [Avail]$ ) {
            remove  $T[i]$  from UNFINISHED
             $[Avail] = [Avail] + [Alloc_i]$ 
            done = false
        }
    }
} while (done == false)
```

- Nodes left in UNFINISHED \Rightarrow **deadlocked**





Banker's Algorithm for Avoiding Deadlocks

```
add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)
```

```
add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Maxi - Alloci] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)
```

Banker's algorithm:

- Allocate resources dynamically
- Evaluate each thread, and grant access if some ordering of threads is still deadlock-free afterwards
- Technique:** pretend each request is granted, then run **deadlock detection algorithm** by substituting:
$$([Request_i] \leq [Avail]) \Rightarrow ([Max_i - Alloc_i] \leq [Avail])$$
- Grant request** if the resulting state is deadlock-free (conservative)



Banker's Algorithm for Avoiding Deadlocks

```
add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)
```

```
add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Maxi - Alloci] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)
```

Banker's algorithm:

- Allocate resources dynamically
- Evaluate each thread, and grant access if some ordering of threads is still deadlock-free afterwards
- Technique:** pretend each request is granted, then run **deadlock detection algorithm** by substituting:
$$([Request_i] \leq [Avail]) \Rightarrow ([Max_i - Alloc_i] \leq [Avail])$$
- Grant request** if the resulting state is deadlock-free (conservative)
- Keep system in a "**SAFE**" state: there exists a sequence $\{T_1, T_2, \dots, T_n\}$ with T_1 **requesting all remaining resources**, then T_2 , then T_3 , etc.



Banker's Algorithm for Avoiding Deadlocks

```
add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Requesti] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)
```

```
add all T[i] to UNFINISHED
do {
    done = true
    for each T[i] in UNFINISHED {
        if ([Maxi - Alloci] <= [Avail]) {
            remove T[i] from UNFINISHED
            [Avail] = [Avail] + [Alloci]
            done = false
        }
    }
} while (done == false)
```

Banker's algorithm:

- Allocate resources dynamically
- Evaluate each thread, and grant access if some ordering of threads is still deadlock-free afterwards
- Technique:** pretend each request is granted, then run **deadlock detection algorithm** by substituting:
$$([Request_i] \leq [Avail]) \Rightarrow ([Need_i] \leq [Avail])$$
- Grant request** if the resulting state is deadlock-free (conservative)
- Keep system in a "**SAFE**" state: there exists a sequence $\{T_1, T_2, \dots, T_n\}$ with T_1 **requesting all remaining resources**, then T_2 , then T_3 , etc.



Banker's Algorithm Example

- 5 threads: $\{T_0, T_1, T_2, T_3, T_4\}$
- 3 resource types:
 - A (10 instances), B (5 instances) and C (7 instances)
- Snapshot at current state of the system:

	<u>Alloc</u>			<u>Max</u>			<u>Avail</u>		
	A	B	C	A	B	C	A	B	C
T0	0	1	0	7	5	3	3	3	2
T1	2	0	0	3	2	2			
T2	3	0	2	9	0	2			
T3	2	1	1	2	2	2			
T4	0	0	2	4	3	3			

- Is the system in a **SAFE** state?



Banker's Algorithm Example

- 5 threads: $\{T_0, T_1, T_2, T_3, T_4\}$
- 3 resource types:
 - A (10 instances), B (5 instances) and C (7 instances)
- Snapshot at current state of the system:

	<u>Alloc</u>			<u>Max</u>			<u>Avail</u>		
	A	B	C	A	B	C	A	B	C
T0	0	1	0	7	5	3	3	3	2
T1	2	0	0	3	2	2			
T2	3	0	2	9	0	2			
T3	2	1	1	2	2	2			
T4	0	0	2	4	3	3			

	<u>Need</u>		
	A	B	C
T0	7	4	3
T1	1	2	2
T2	6	0	0
T3	0	1	1
T4	4	3	1

$$\text{Need}_i = [\text{Max}_i - \text{Alloc}_i]$$

- Is the system in a **SAFE** state?
 - Yes. Because there exists a sequence $\langle T_1, T_3, T_4, T_2, T_0 \rangle$ that satisfies the safety requirement.



■ Summary

- Four (necessary) conditions for deadlocks
 - Mutual Exclusion
 - Hold and Wait
 - No Preemption
 - Circular Wait
- Techniques for addressing deadlocks
 - Deadlock **Prevention**:
 - write your code in a way that isn't prone to deadlock.
 - Deadlock **Detection** and **Recovery**:
 - Let it happen, and then figure out how to recover once detected.
 - Deadlock **Avoidance**:
 - Dynamically delay resource requests so deadlock doesn't happen
 - **Banker's Algorithm** provides an algorithmic way to do this
 - Deadlock **Denial**:
 - Ignore the possibility of deadlock (used by most OSes, e.g., Linux)
 - Ostrich Algorithm (鸵鸟算法)



Thank you!