



DCS216 Operating Systems

Lecture 20 Memory (3) Demand Paging

May 15th, 2024

**Instructor: Xiaoxi Zhang
Sun Yat-sen University**




■ Content

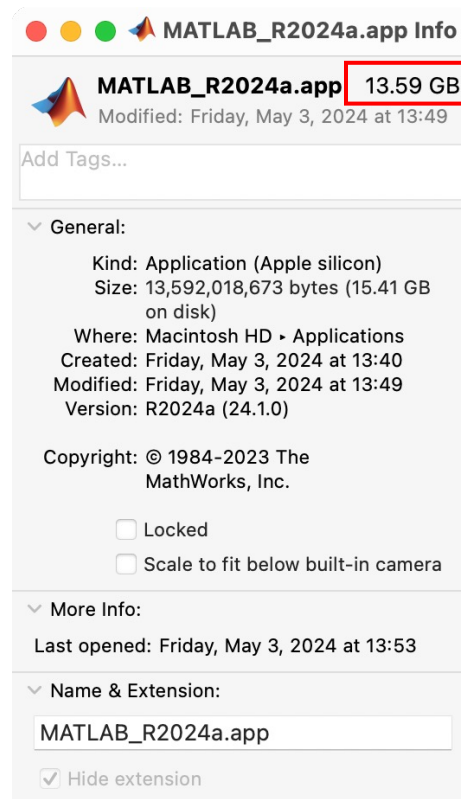
- Background
- Demand Paging
- Copy-on-Write (**COW**)
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- OS Examples
 - Linux
 - Windows
 - Solaris

■ Background

- So far, we require that the **entire** program be loaded into memory **before** it can execute.
- **Question:** How to execute a program that is **larger** (in size) **than** the physical memory available on the system?



MATLAB_R2024a.
app



8 核中央处理器

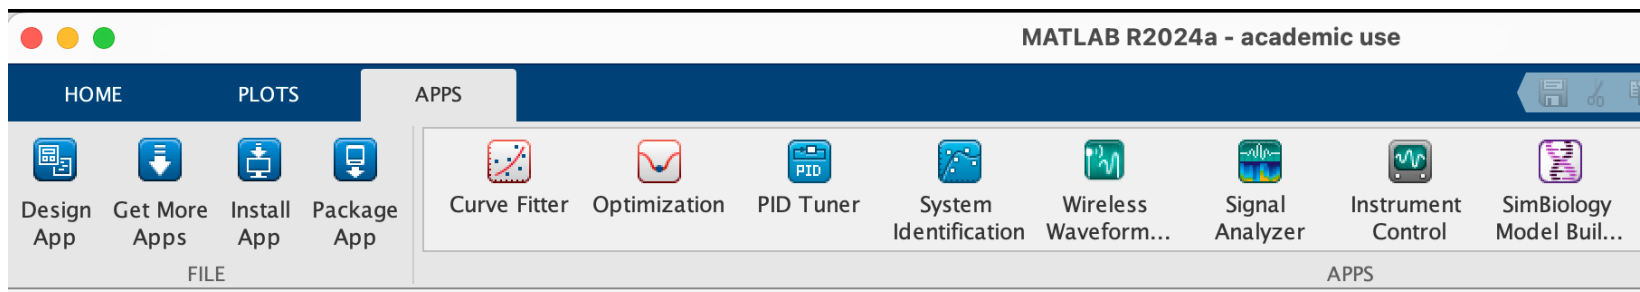
8 核图形处理器

8GB 统一内存

256GB 固态硬盘¹

■ Background

- The requirement that instructions must be in physical memory to be executed seems both **necessary** and **reasonable**.
- However, it limits the size of a program to the size of available physical memory.
- In fact, a closer look at real programs (such as *Matlab*) shows us that, in many cases, the entire program is not needed:
 - Programs often have code to handle unusual error conditions.
 - These errors seldom occur in practice \Rightarrow almost never executed.
 - Certain options and features of a program may be used very rarely.
 - For example, Matlab has many toolboxes installed by default, but rarely used unless explicitly invoked.





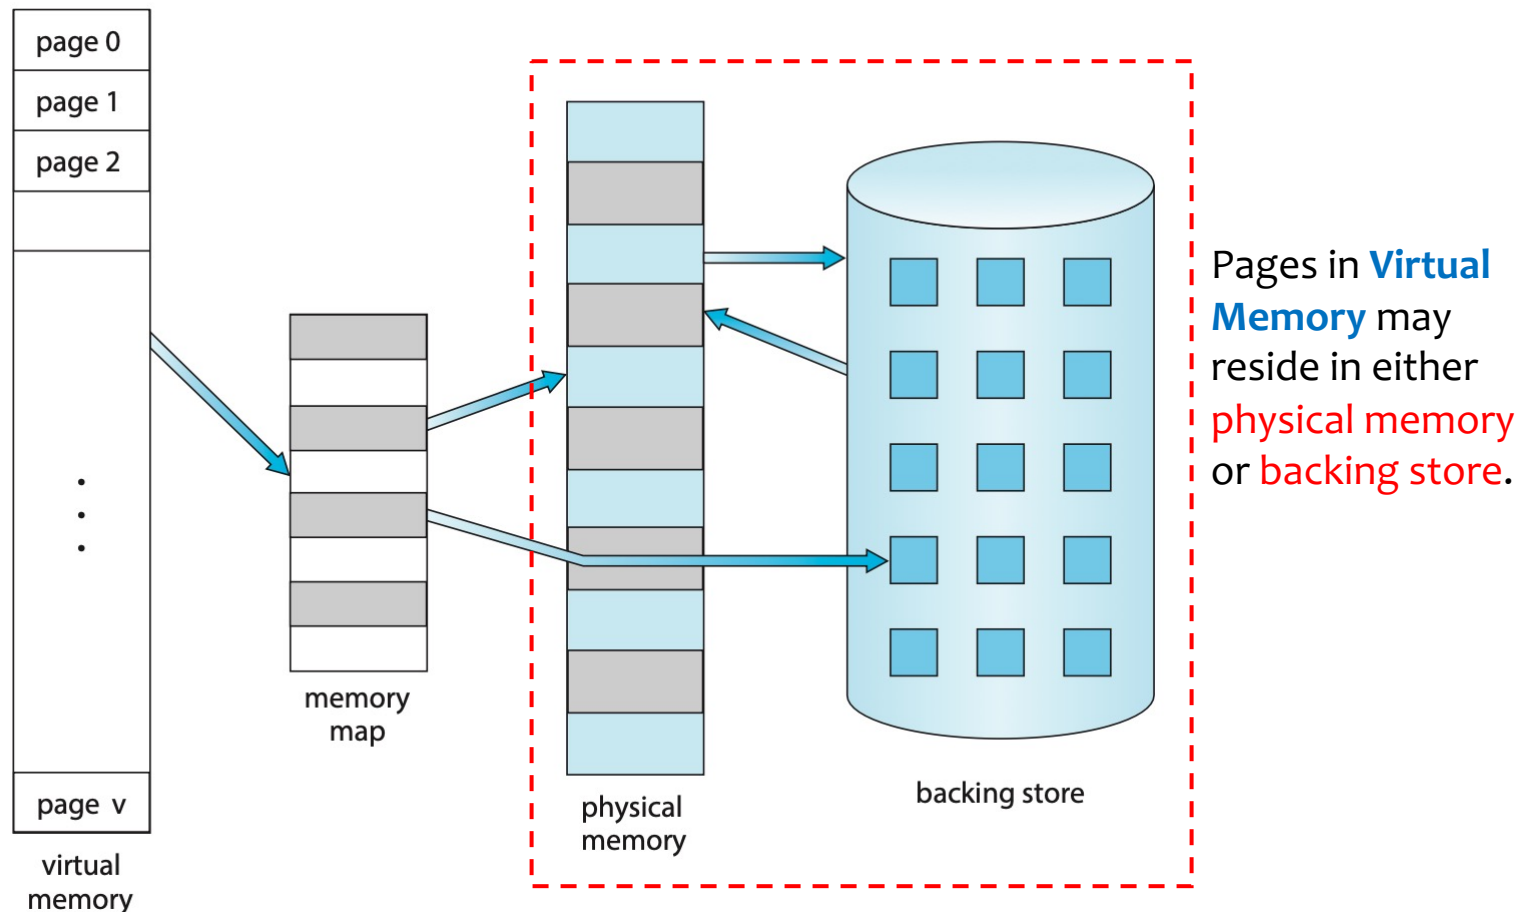
■ Background

- Even in those cases where the entire program is needed, it may not all be needed **at the same time**.
- Consider the *ability* to execute **partially-loaded** program
 - Program size no longer **constrained** by limits of physical memory
 - Each program takes **less** physical memory while running
 - \Rightarrow more programs running at the same time
 - Increased CPU utilization and throughput with no increase in response time or turnaround time
 - Less I/O needed to load or swap programs into memory
 - \Rightarrow each program runs faster



Virtual Memory

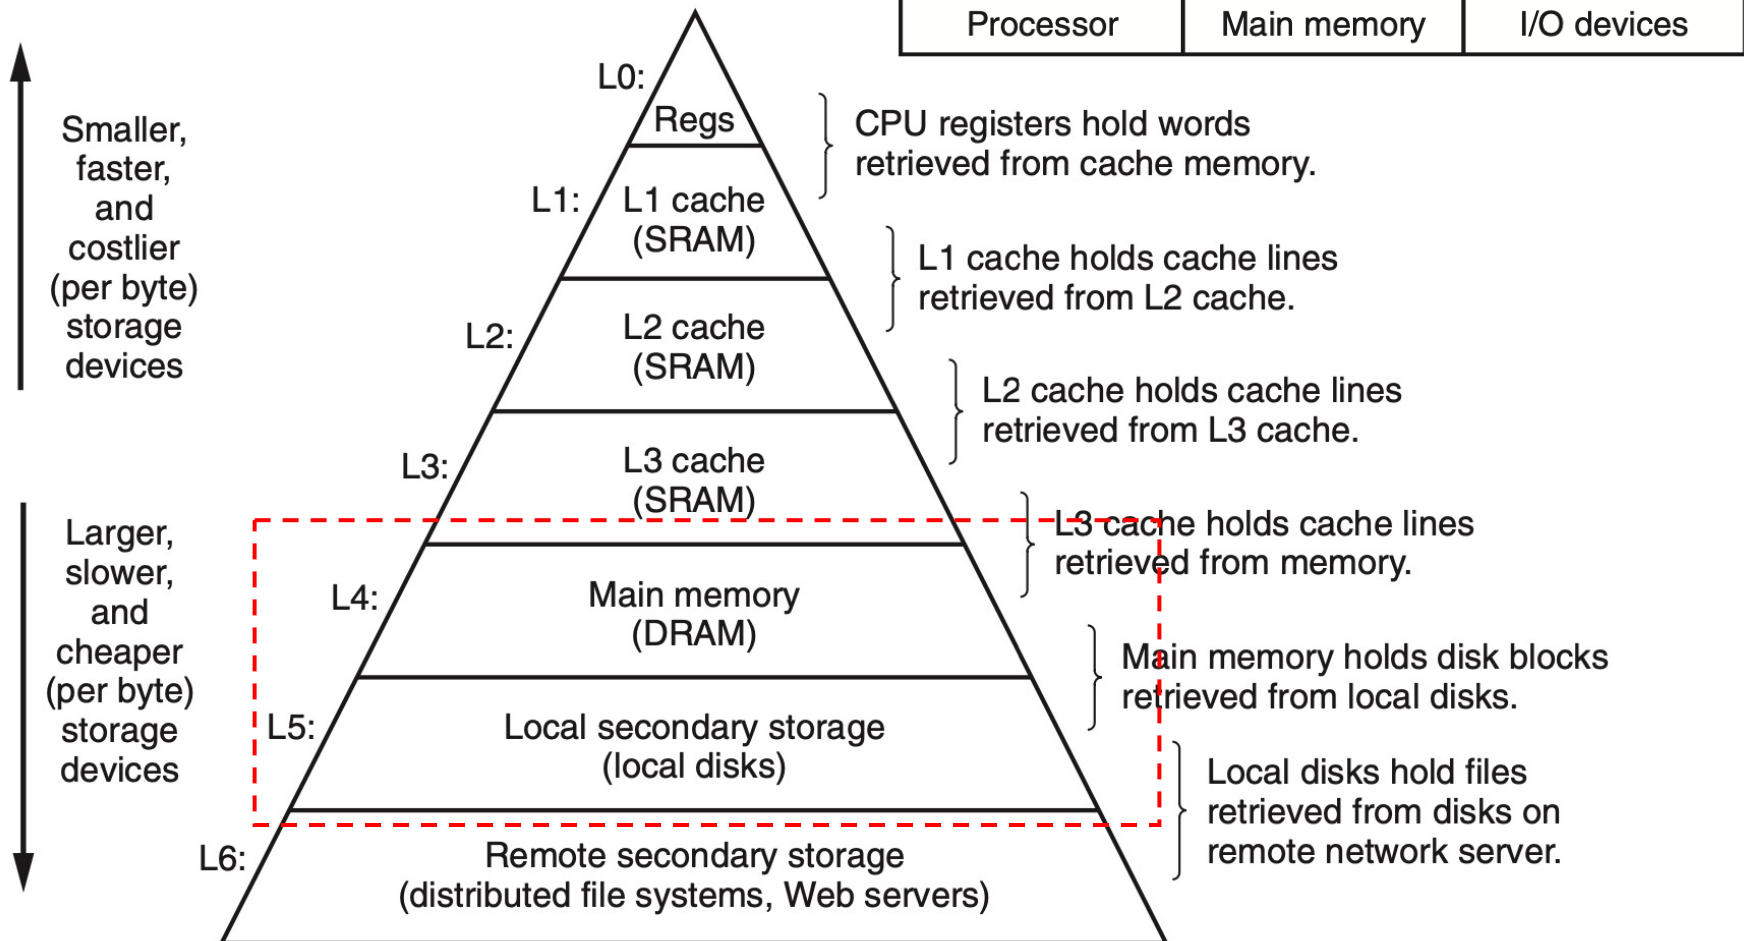
- **Virtual Memory:** separation of **logical** memory from **physical** memory
 - It allows an extremely large virtual memory to be provided for developers when only a smaller physical memory is available.





Virtual Memory

The Memory Hierarchy





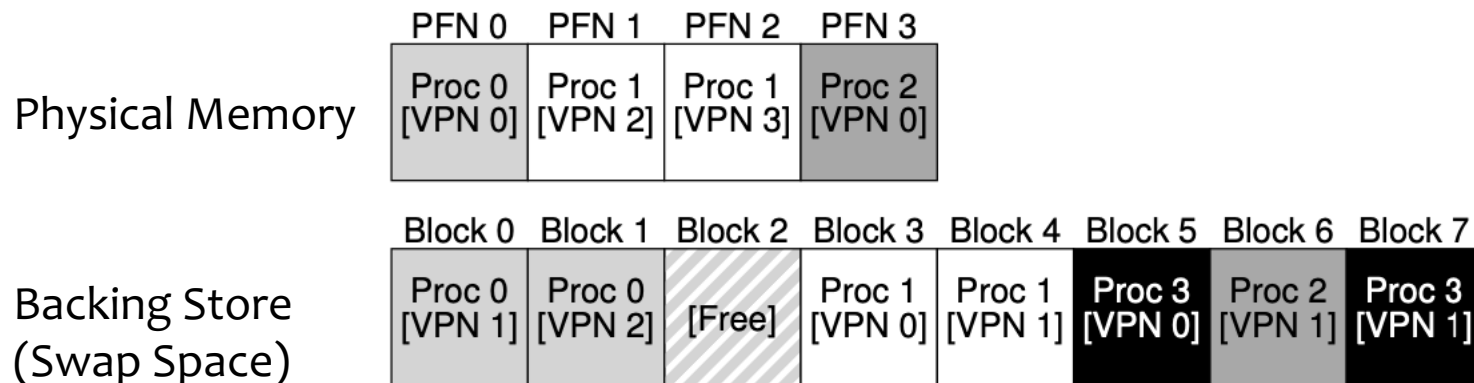
■ Virtual Memory

- **Virtual Memory:** separation of **logical** memory from **physical** memory
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation
 - More programs running concurrently
 - Less I/O needed to load or swap processes
- **Virtual Address Space:** logical view of how processes is stored in memory
 - Usually start at address 0, contiguous address until end of space
 - Meanwhile, physical memory organized in physical frames
 - MMU must map logical address into physical address.

Virtual Memory

A tiny example:


- 4-page/**frame** Physical Memory
- 8-page/**block** Backing Store (Swap Space) on disk
- 4 Processes
 - Each process has its own **Page Table**
 - i.e., mappings from VPN to PFN
 - 3 active Processes (**Proc 0**, **Proc 1**, **Proc 2**)
 - some of their valid pages in memory
 - 1 inactive Process (**Proc 3**)
 - all of its pages **swapped out** to disk.





■ Virtual Address Space

- Usually design virtual address space for stack to start at Max addr and grow downward, while heap grows upward.
 - Maximize address space use
 - Unused address space between stack and heap
 - No physical memory needed until heap or stack grows to a new page.
- Enable **sparse** address spaces with holes left for growth, dynamically linked libraries, etc.






■ Virtual Address Space

■ Shared Library Using Virtual Memory

- System libraries shared via mapping into virtual address space

■ Shared Memory by mapping pages into virtual address space


- Pages can be shared during `fork()`, speeding process creation





Basic Concepts

- Bring a page into memory only when it is needed (**on demand**).
 - Less I/O needed, no unnecessary I/O
 - Less memory needed
 - Faster response
 - More users.
- Page is needed \Rightarrow reference to it
 - invalid and illegal reference \Rightarrow abort
 - invalid not-in-memory \Rightarrow bring to memory
- Similar to paging system with swapping
 - Pager is the swapper that deals with pages
 - With swapping, pager guesses which pages will be used before swapping out again.
- Lazy swapper – Never swaps a page into memory unless page will be needed.
 - Pager brings only those pages needed into memory.





■ Basic Concepts

■ Valid/Invalid Bit

- Within each **Page Table Entry (PTE)**, a **valid** bit is needed (*also called **present** bit on some systems*)
 - **v** (or **1**): page in memory (**valid**)
 - **i** (or **0**): page not in memory (**invalid**)
- Initially set valid bits to **i (0)** for all entries

Page Table	Frame #	v-i bit
		v
		v
		v
		v
		i
	
		i
		i

- During address translation, if valid bit in PTE is **i**, then **page fault**.
 - **invalid** bit **does not** indicate **illegal** virtual address



■ Illegal Address Reference Examples

- Illegal address reference generally triggers **segmentation fault**.

```
/* illegal_addr1.c */
#include <stdio.h>

int main() {
    int *p = NULL;
    // p is a NULL pointer
    *p = 42;
    // dereferencing a NULL pointer
    //      --> segfault
    printf("*p: %d\n", *p);
    return 0;
}
```

```
/* illegal_addr2.c */
#include <stdio.h>

int main() {
    int *p = (void *)0xFFFF880000000000;
    // p points kernel address space
    *p = 42;
    // trying to access kernel address space
    //      --> segfault
    printf("*p: %d\n", *p);
    return 0;
}
```


```
/* illegal_addr3.c */
#include <stdio.h>

int main() {
    int *p = (void *)main;
    // p points to addr of main
    *p = 42;
    // trying to modify the first 4 bytes of main (read-only code segment)
    //      --> segfault
    printf("*p: %d\n", *p);
    return 0;
}
```



Basic Concepts

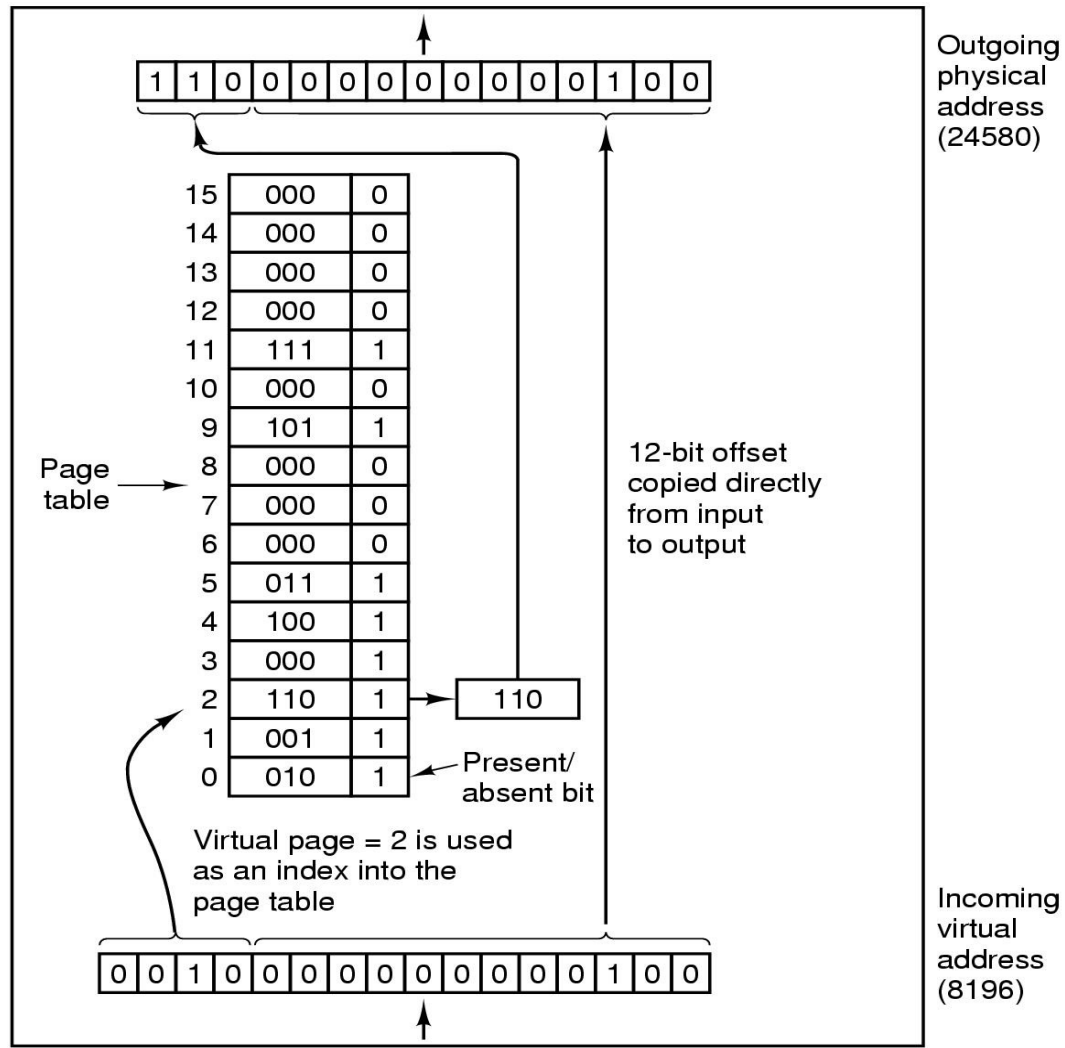
Page Table when some pages are not in main memory





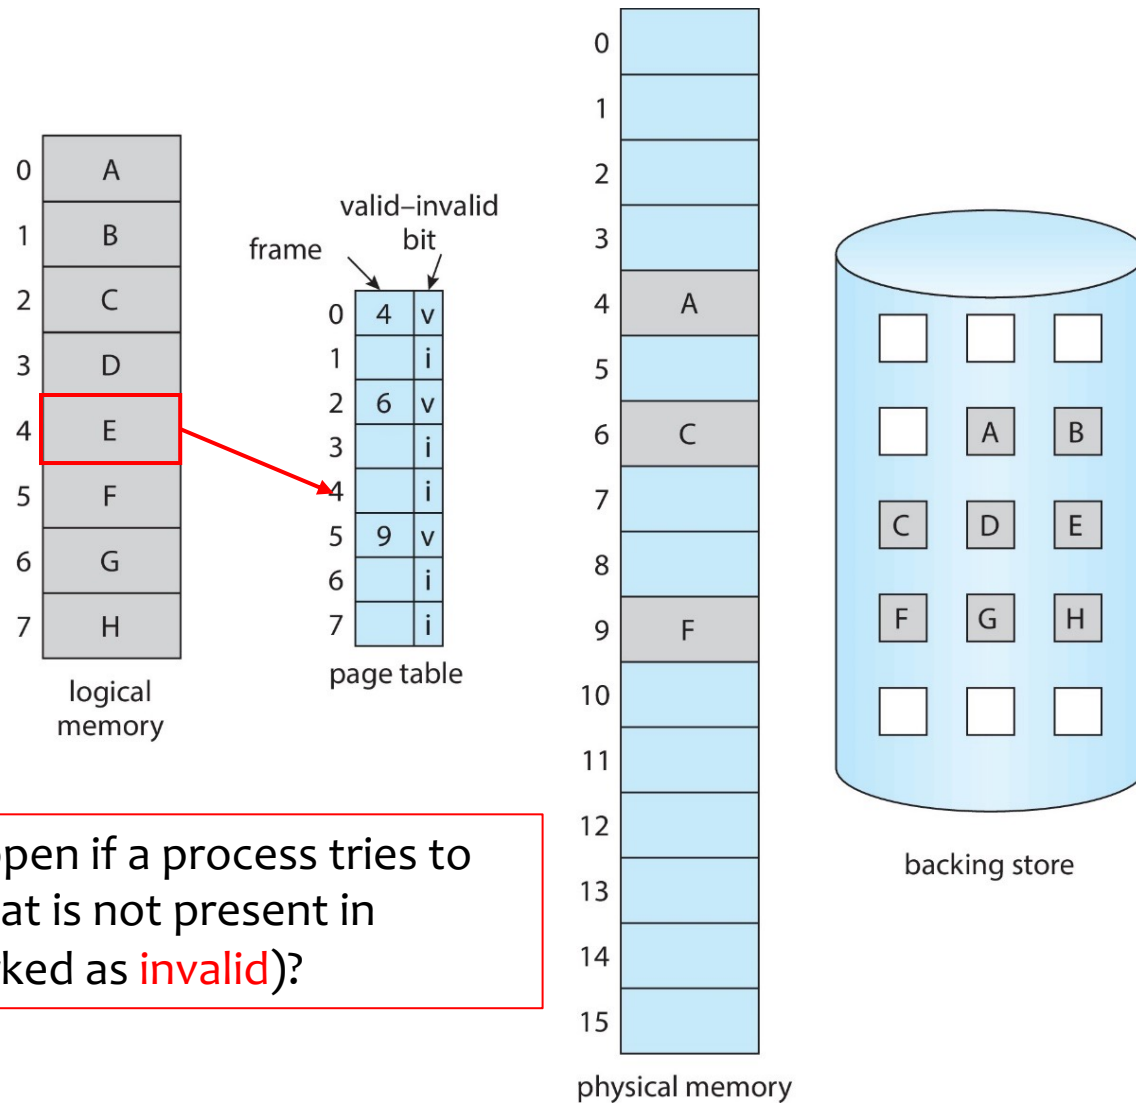
Basic Concepts

Virtual Memory Mapping Example





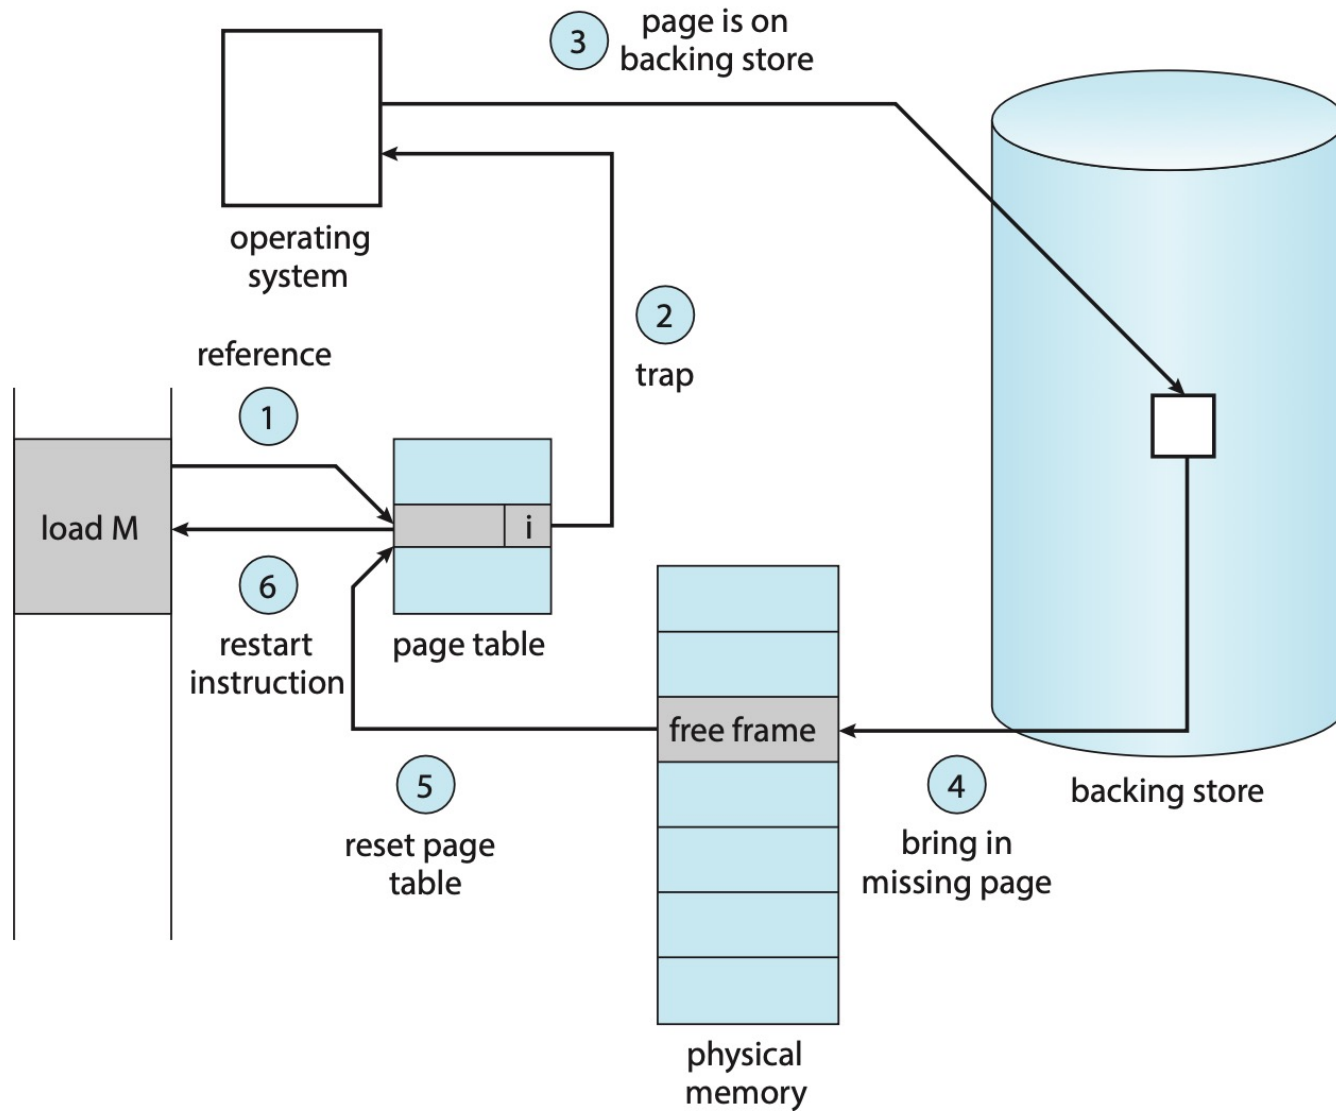
■ Page Fault



What would happen if a process tries to access a page that is not present in memory (or marked as **invalid**)?

Page Fault

- Access to a page marked as **invalid** causes a **Page Fault** (缺页错误).

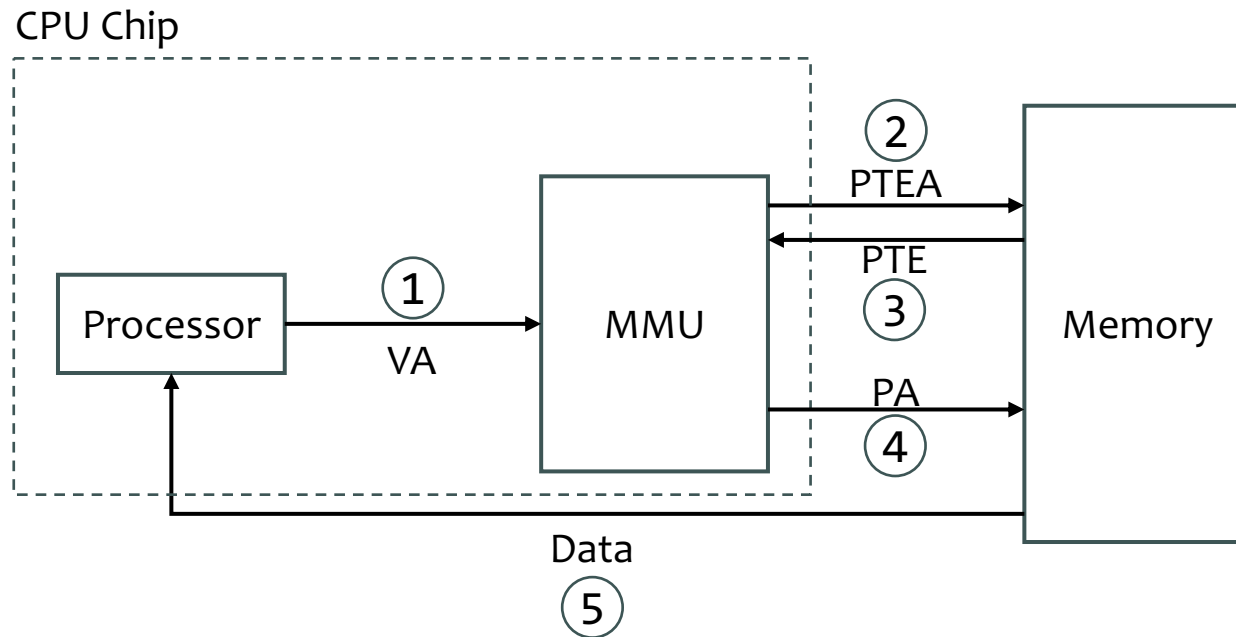




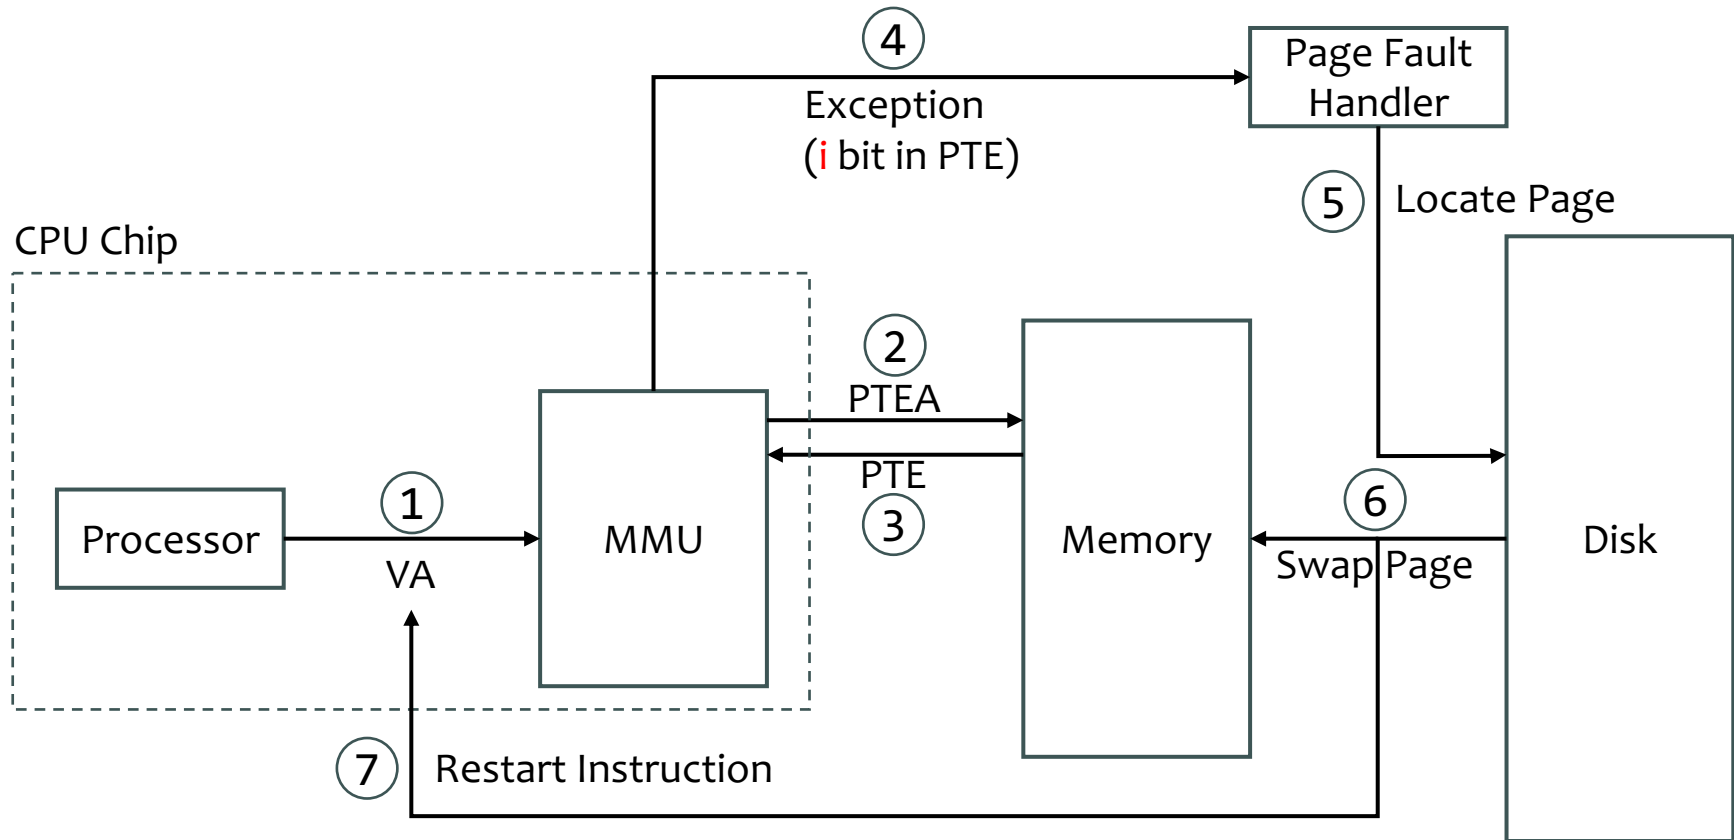
■ Page Fault

- Access to a page marked as **invalid** causes a **Page Fault**.
- Steps in handling a **Page Fault**:
 1. **Check** an internal table (usually kept within the **PCB**) for this process to determine whether the memory **reference** was **legal** or **illegal**.
 2. If the reference is **legal** and the **valid bit** in **PTE** is **invalid** (page **not** in memory, but in backing store), **trap** into the **Page Fault Handler**.
 3. **Locate** the desired page/block on **disk** (swap space).
 4. Find a **free frame** (e.g., by taking one from the **free-frame list**) and **swap** the desired page/block from disk **into** the **free frame** via **scheduled disk read operation**.
 5. When the **storage read** is complete, we modify the internal table kept within the **PCB** and the **Page Table** to indicate that the page is now in memory (i.e., by setting the valid bit to **1 (valid)** in the **PTE**).
 6. **Restart** the instruction that was interrupted by the **trap**. *The process can now access the page as though it had always been in memory.*

■ Page Hit



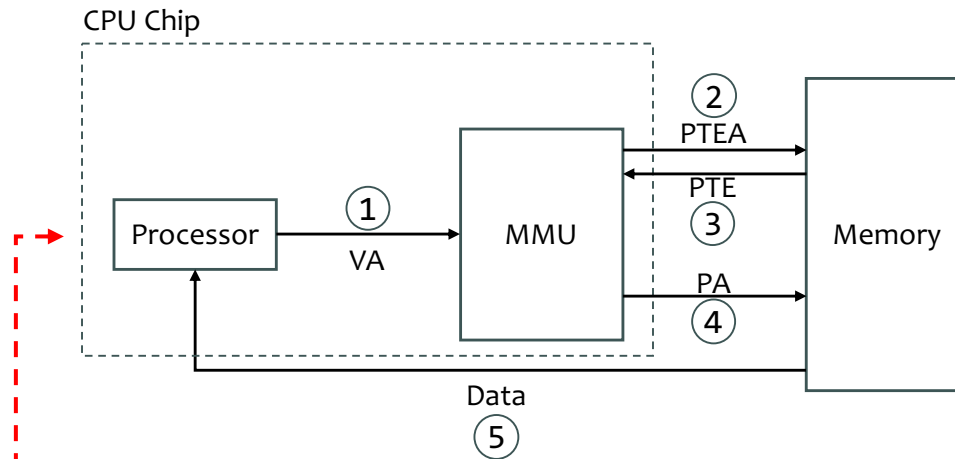
■ Page Fault



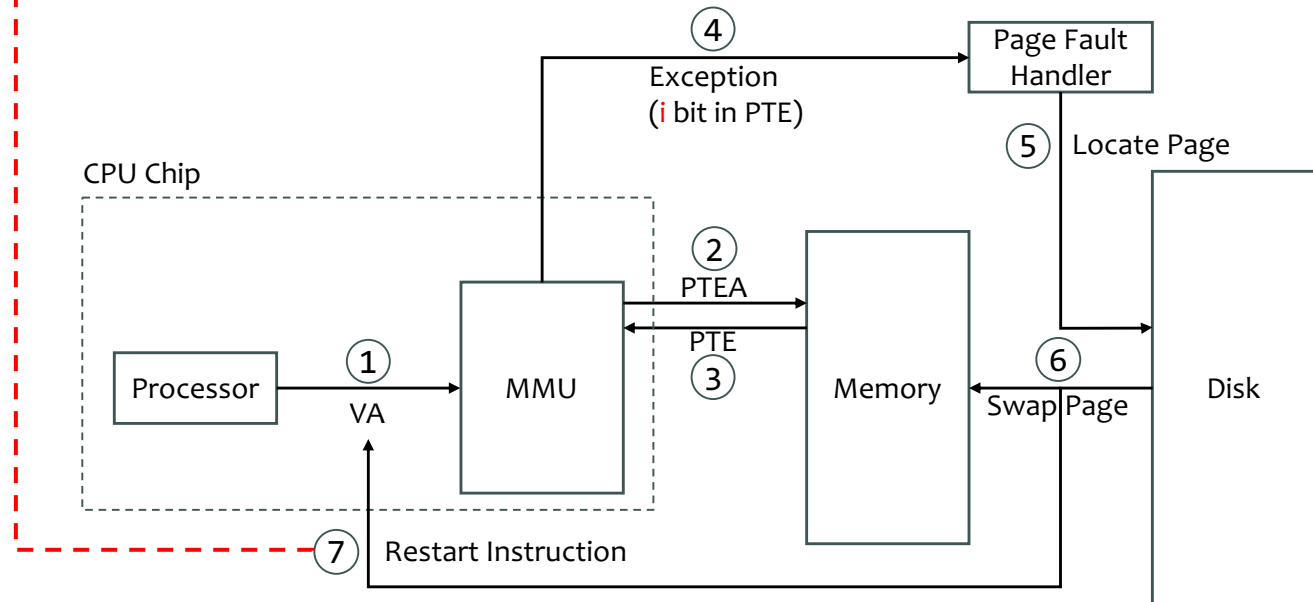


■ Page Hit vs. Page Fault

Page **Hit**:




Page **Fault**:





■ Dynamics of Demand Paging

- Typically, each process has *its own page table*.
- Each **Page Table Entry (PTE)** contains a **valid** bit to indicate whether the page is in memory or not
 - If it is in main memory, the **PTE** contains the frame number (**PFN**) of the corresponding frame in main memory
 - Otherwise, the **PTE** may contain the address of the **page block** on **disk**.
- Example: Intel x86 architecture **PTE**:
 - 2-Level Page Table (10, 10, 12-bit offset)
 - Intermediate Page Tables called "Directories"






■ Dynamics of Demand Paging

- Typically, each process has *its own page table*.
- Each **Page Table Entry (PTE)** contains a **valid** bit to indicate whether the page is in memory or not
 - If it is in main memory, the **PTE** contains the frame number (**PFN**) of the corresponding frame in main memory
 - Otherwise, the **PTE** may contain the address of the **page block on disk**.
- Example: Intel x86 architecture **PTE**:
 - 2-Level Page Table (10, 10, 12-bit offset)
 - Intermediate Page Tables called "Directories"

Some OSes store the **address** of {the corresponding block on **disk**} in the **PFN** bits

Present Bit
(valid bit)






■ Aspects of Demand Paging

- Extreme case – start process with **no** pages in main memory.
 - OS sets Instruction Pointer to the first instruction of process
 - which is on a non-memory-resident page ⇒ triggers **Page Fault**.
 - ..and for every other pages upon **first** access
 - This is called **Pure Demand Paging**.
 - **never bring a page into memory until it is required.**
- Actually, a given instruction could access multiple pages and cause multiple page faults.
 - Consider fetch and decode of instruction which adds two numbers from memory and stores result back to memory.
 - Pain decreased because of **locality of reference**.
- Hardware support needed for demand paging.
 - Page Table with valid/invalid bit
 - Secondary Memory (swap device with swap space)
 - Instruction Restart



■ Free-Frame List

- When a page fault occurs, the OS must bring the desired page from secondary storage (swap space) into main memory.
- Most OSes maintain a **free-frame list** – a pool of free frames for satisfying such requests.



- The OS typically allocates free frames using a technique known as **zero-fill-on-demand**:
 - the content of the frames zeroed-out before being allocated.
- When a system starts up, all available memory is placed on the **free-frame list**.



■ Stages in Demand Paging (the **Worst Case**)

1. **Trap** into the Operating System
2. **Save** the user registers and process state
3. **Determine** that the interrupt was a page fault
4. **Check** that the page reference was legal and **determine** that the location of the block (page) on disk.
5. **Issue** a read request from the disk to a free frame:
 1. **Wait** in a queue for this device until the read request is serviced
 2. **Wait** for the device seek and/or latency time
 3. **Begin** the transfer of the page to a free frame
6. While waiting, allocate the CPU to some other user
7. **Receive** an interrupt from the disk I/O subsystem (I/O completed)
8. **Save** the registers and process state for the other user
9. **Determine** that the interrupt was from the disk
10. **Correct** the Page Table and other tables to show page is in memory
11. **Wait** for the CPU to be allocated to this process again
12. **Restore** the user registers, process state, and new page table, and then **resume** the interrupted instruction.



■ Performance of Demand Paging

■ Three major activities:

■ Service the **page-fault** interrupt

- careful coding means just ~100 instruction cycles (~100 ns) needed

■ Read in the page (from disk)

- lots of time, e.g., ~8 milliseconds (~8,000,000 ns)

Dominant Factor

■ Restart the process

- again, just a small amount of time (~100 ns)

■ Page Fault Rate p ($0 \leq p \leq 1$)

- $p = 0 \Rightarrow$ no page faults

- $p = 1 \Rightarrow$ every reference causes a page fault.

■ Effective Access Time (EAT):

$$\text{EAT} = (1 - p) \times \text{memory_access_time} + p \times \text{page_fault_time}$$

Time spent for normal memory access

Time spent for handling page fault



■ Performance of Demand Paging

■ Example:

- $\text{memory_access_time} = 200 \text{ nanoseconds}$
- Average $\text{page_fault_time} = 8 \text{ milliseconds} = 8,000,000 \text{ ns}$
- $\text{EAT} = (1-p) \times 200 \text{ (ns)} + p \times 8 \text{ (ms)}$
 $= (1-p) \times 200 + p \times 8,000,000 \text{ (ns)}$
 $= 200 + 7,999,800 \times p \text{ (ns)}$
- **EAT** is directly proportional to the page-fault rate p .

■ If one access out of 1,000 causes a page fault ($p = 0.001$), then

- $\text{EAT} = 200 + 7,999,800 \times 0.001 = 8199.8 \text{ (ns)}$
 - \Rightarrow a **slowdown** by a factor of
 $8199.8 / 200 = 41$

■ If we want performance degradation $< 10\%$

- $\text{EAT} = 220 < 200 + 7,999,800 \times p$
 $\Rightarrow 20 > 7,999,800 \times p$
- $p < 2.5 \times 10^{-6}$.
 - i.e., less than one page fault in every **400,000** memory accesses.



■ Demand Paging Optimizations

- Swap space I/O faster than file system I/O even if on the same device.
 - Swap allocated in large chunks, less management needed than FS
- Copy entire process image to swap space at process load time
 - Then page in and out of swap space
 - Used in older BSD Unix
- Demand page in from program binary on disk, but discard rather than paging out when freeing frame
 - Used in Solaris and current BSD
 - Still need to write to swap space
 - Pages not associated with a file (like stack and heap)
 - anonymous memory
 - Pages modified in memory but not yet written back to file system
- Mobile systems
 - Typically don't support swapping
 - Instead, demand page from file system and reclaim read-only pages (such as code)




■ Copy-on-Write (CoW)

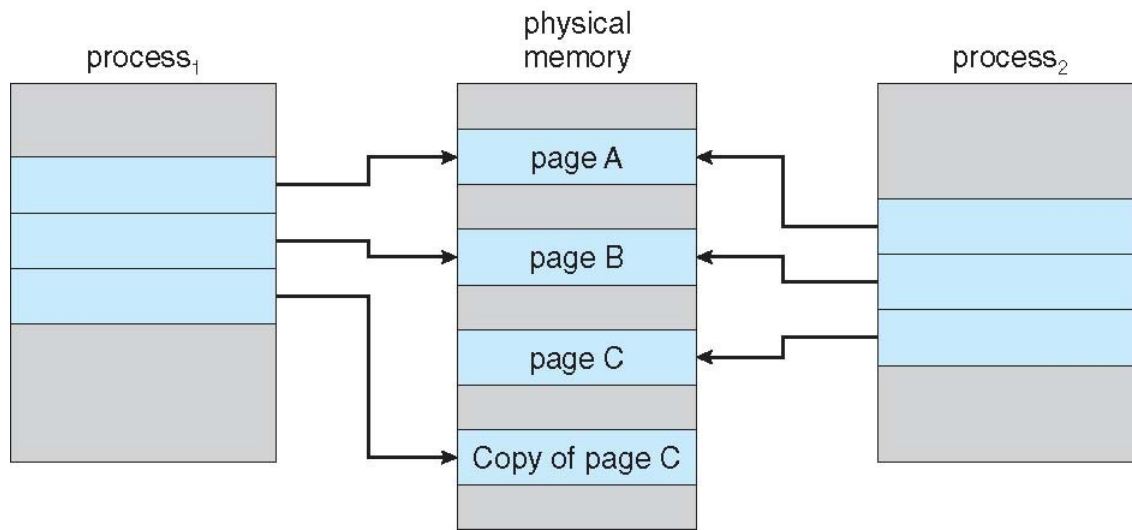
- Semantically, `fork()` syscall will create a copy of the **parent's** address space for the **child** process
 - In early UNIX systems, it will simply **duplicate** all the active pages belonging to the parent for the child process.
 - Very expensive (huge amount of memory accesses)
- Copy-on-Write (CoW, **写入时复制**) allows both parent and child processes to **initially share** the same pages in memory
 - If either process modifies a shared page, only then is the page copied.
- **CoW** allows more efficient process creation
 - only modified pages are copied.

■ Copy-on-Write (CoW)

■ Before Process 1 Modifies Page C:



■ After Process 1 Modifies Page C:





■ Copy-on-Write (CoW)

```
/* CoW.c */
#define PAGE_SHIFT 12
#define PAGE_SIZE (1 << PAGE_SHIFT)
#define PAGESIZE 4096
#define PAGESIZE_STR "4096"

uint64_t get_physical_address(uint64_t virtual_address);

int main() {
    int *addr = malloc(PAGE_SIZE); // Allocate space (PAGE_SIZE) in the heap
    *addr = 10; // Initialize addr with some value

    pid_t rc = fork();
    if (rc == 0) { /* Child process */
        printf("Child: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr before child modification: %d (%p)\n", *addr, (void
*)get_physical_address((uint64_t)addr));
        *addr = 20; /* Modify *addr */
        printf("Child: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr after child modification: %d (%p)\n", *addr, (void
*)get_physical_address((uint64_t)addr));
    } else { /* Parent process */
        printf("Parent: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr before child modification: %d (%p)\n", *addr, (void
*)get_physical_address((uint64_t)addr));
        sleep(1); /* Wait for child process to modify *addr */
        printf("Parent: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr after child modification: %d (%p)\n", *addr, (void
*)get_physical_address((uint64_t)addr));
        return 0;
    }
}
```

为addr分配一个大小为PAGE_SIZE(4096字节)的空间，以确保addr单独占有一个page



■ Copy-on-Write (CoW)

```
/* CoW.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <fcntl.h>
#include <stdint.h>
#include <sys/mman.h>

#define PAGE_SHIFT 12
#define PAGE_SIZE (1 << PAGE_SHIFT)
#define PAGEMAP_ENTRY 8

uint64_t get_physical_address(uint64_t
virtual_address) {
    int pagemap_fd;
    uint64_t offset, physical_address;
    uint64_t pagemap_entry;

    // Open the pagemap file for the current process
    pagemap_fd = open("/proc/self/pagemap",
O_RDONLY);

    // Calculate the offset in the pagemap file
    offset = (virtual_address / PAGE_SIZE) *
PAGEMAP_ENTRY;

    // Seek to the appropriate offset
    if (lseek(pagemap_fd, offset, SEEK_SET) ==
(off_t)-1) {
        perror("lseek");
        close(pagemap_fd);
        return -1;
    }

    // Read the pagemap entry
    if (read(pagemap_fd, &pagemap_entry,
PAGEMAP_ENTRY) != PAGEMAP_ENTRY) {
        perror("read");
        close(pagemap_fd);
        return -1;
    }

    close(pagemap_fd);

    // Extract the physical frame number (PFN) from
the pagemap entry
    if (pagemap_entry & (1ULL << 63)) {
        physical_address = (pagemap_entry & ((1ULL
<< 55) - 1)) * PAGE_SIZE;
        physical_address |= (virtual_address &
(PAGE_SIZE - 1));
    } else {
        return -1;
    }
    return physical_address;
}
```



Copy-on-Write (CoW)

```

/* CoW.c */
#define PAGE_SHIFT 12
#define PAGE_SIZE (1 << PAGE_SHIFT)
#define PAGES_PER_PAGE (PAGE_SIZE / (1 << PAGE_SHIFT))
#define PAGEMAP_ENTRY 8

uint64_t get_physical_address(uint64_t addr) {
    return 0;
}

int main() {
    int *addr = malloc(PAGE_SIZE); // Allocate space (PAGE_SIZE) in the heap
    *addr = 10; // Initialize addr with some value

    pid_t rc = fork();
    if (rc == 0) { /* Child process */
        printf("Child: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr before child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        *addr = 20; /* Modify *addr */
        printf("Child: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr after child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
    } else { /* Parent process */
        printf("Parent: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr before child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        sleep(1); /* Wait for child process to modify *addr */
        printf("Parent: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr after child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        return 0;
    }
}

```

```

$ gcc -g -Wall -o CoW CoW.c
$ sudo ./CoW
Parent: virtual addr before child modification: 10 (0x5555555592a0)
Parent: physical addr before child modification: 10 (0x6952822a0)
Child: virtual addr before child modification: 10 (0x5555555592a0)
Child: physical addr before child modification: 10 (0x6952822a0)

```



Copy-on-Write (CoW)

```

/* CoW.c */
#define PAGE_SHIFT 12
#define PAGE_SIZE (1 << PAGE_SHIFT)
#define PAGES_PER_PAGE (PAGE_SIZE / sizeof(uint64_t))
#define PAGEMAP_ENTRY 8

uint64_t get_physical_address(uint64_t addr) {
    return 0;
}

int main() {
    int *addr = malloc(PAGE_SIZE);
    *addr = 10; // Initialize addr with some value

    pid_t rc = fork();
    if (rc == 0) { /* Child process */
        printf("Child: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr before child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        *addr = 20; /* Modify *addr */
        printf("Child: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr after child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
    } else { /* Parent process */
        printf("Parent: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr before child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        sleep(1); /* Wait for child process to modify *addr */
        printf("Parent: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr after child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
    }
    return 0;
}

```

```

$ gcc -g -Wall -o CoW CoW.c
$ sudo ./CoW
Parent: virtual addr before child modification: 10 (0x5555555592a0)
Parent: physical addr before child modification: 10 (0x6952822a0)
Child: virtual addr before child modification: 10 (0x5555555592a0)
Child: physical addr before child modification: 10 (0x6952822a0)
Child: virtual addr after child modification: 20 (0x5555555592a0)
Child: physical addr after child modification: 20 (0xb1b4a02a0)

```



Copy-on-Write (CoW)

```

/* CoW.c */
#define PAGE_SHIFT 12
#define PAGE_SIZE (1 << PAGE_SHIFT)
#define PAGES_PER_PAGE (PAGE_SIZE / (sizeof(uint64_t) * 8))
#define PAGEMAP_ENTRY 8

uint64_t get_physical_address(uint64_t addr) {
    return (addr >> PAGE_SHIFT) * PAGE_SIZE + (addr & (PAGE_SIZE - 1));
}

int main() {
    int *addr = malloc(PAGE_SIZE);
    *addr = 10; // Initial value

    pid_t rc = fork();
    if (rc == 0) { /* Child process */
        printf("Child: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr before child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        *addr = 20; /* Modify *addr */
        printf("Child: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Child: physical addr after child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
    } else { /* Parent process */
        printf("Parent: virtual addr before child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr before child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        sleep(1); /* Wait for child process to modify *addr */
        printf("Parent: virtual addr after child modification: %d (%p)\n", *addr, addr);
        printf("Parent: physical addr after child modification: %d (%p)\n", *addr, (void *)get_physical_address((uint64_t)addr));
        return 0;
    }
}

```

```

$ gcc -g -Wall -o CoW CoW.c
$ sudo ./CoW
Parent: virtual addr before child modification: 10 (0x5555555592a0)
Parent: physical addr before child modification: 10 (0x6952822a0)
Child: virtual addr before child modification: 10 (0x5555555592a0)
Child: physical addr before child modification: 10 (0x6952822a0)
Child: virtual addr after child modification: 20 (0x5555555592a0)
Child: physical addr after child modification: 20 (0xb1b4a02a0)
Parent: virtual addr after child modification: 10 (0x5555555592a0)
Parent: physical addr after child modification: 10 (0x6952822a0)

```



■ `fork()` vs. `vfork()`

- `vfork()` is a variant of `fork()` in several versions of UNIX (including Linux, macOS, and BSD UNIX).
 - `vfork()` is **deprecated** and **not** part of the POSIX standard
 - not recommended if you wish to write portable code.
- Unlike `fork()`, `vfork()` **does not use Copy-on-Write**.
- With `vfork()`, the **parent** process is **suspended**, and the **child** process uses the **same** address space of the **parent**.
 - The child process is intended to call `exec()` **immediately** after creation.
 - Recall that `exec()` will **replace** the current child process image with a new process image specified by the arguments of `exec()`
 - \Rightarrow **new address space** that is different from the parent.
 - If the child process modifies the address space of the parent, it will lead to **undefined** behavior.
 - The parent process **resumes** after the child process either calls `exec()` or exits.




■ fork() vs. vfork()

<pre>/* ex_fork.c */ #include <stdio.h> #include <stdlib.h> #include <sys/types.h> #include <sys/wait.h> #include <unistd.h> int main() { int data = 10; pid_t rc = fork(); if (rc == 0) { /* Child */ data = 20; printf(" (Child) data: %d\n", data); _exit(0); } else { /* Parent */ wait(NULL); printf("(Parent) data: %d\n", data); } }</pre>	<pre>/* ex_vfork.c */ #define _XOPEN_SOURCE 500 #include <stdio.h> #include <stdlib.h> #include <sys/types.h> #include <sys/wait.h> #include <unistd.h> int main() { int data = 10; pid_t rc = vfork(); if (rc == 0) { /* Child */ data = 20; printf(" (Child) data: %d\n", data); _exit(0); } else { /* Parent */ wait(NULL); printf("(Parent) data: %d\n", data); } }</pre>
<pre>\$./ex_fork (Child) data: 20 (Parent) data: 10</pre>	<pre>\$./ex_vfork (Child) data: 20 (Parent) data: 20</pre>

■ Demand Paging

- What happens if there is **no free frame**?






■ What can the OS do when there is no free frame?

- Option #1: **Terminate** the user process
 - \Rightarrow Not the best choice, as it destroys the purpose of **Demand Paging**
- Option #2: The OS could **swap out** a process (the entire process), **freeing all its frames** and reducing the level of multiprogramming
 - \Rightarrow A good option under certain circumstances
- Option #3: Make use of **Page Replacement** technique.
 - \Rightarrow The most **common** solution.



■ Page Replacement


- What happens if there is **no free frame**?
 - ⇒ Page Replacement (页面置换)
 - Find some frame (**victim** frame) in memory that are *not really in use*
 - and then **swap** it out.





■ Page Replacement


- What happens if there is **no free frame**?
 - ⇒ Page Replacement (页面置换)
 - The **Question** is: How to choose the **victim** frame?
 - ⇒ Page Replacement Algorithms (more on this in the next lecture)





■ Page Replacement

- Steps in handling a **Page Replacement**.

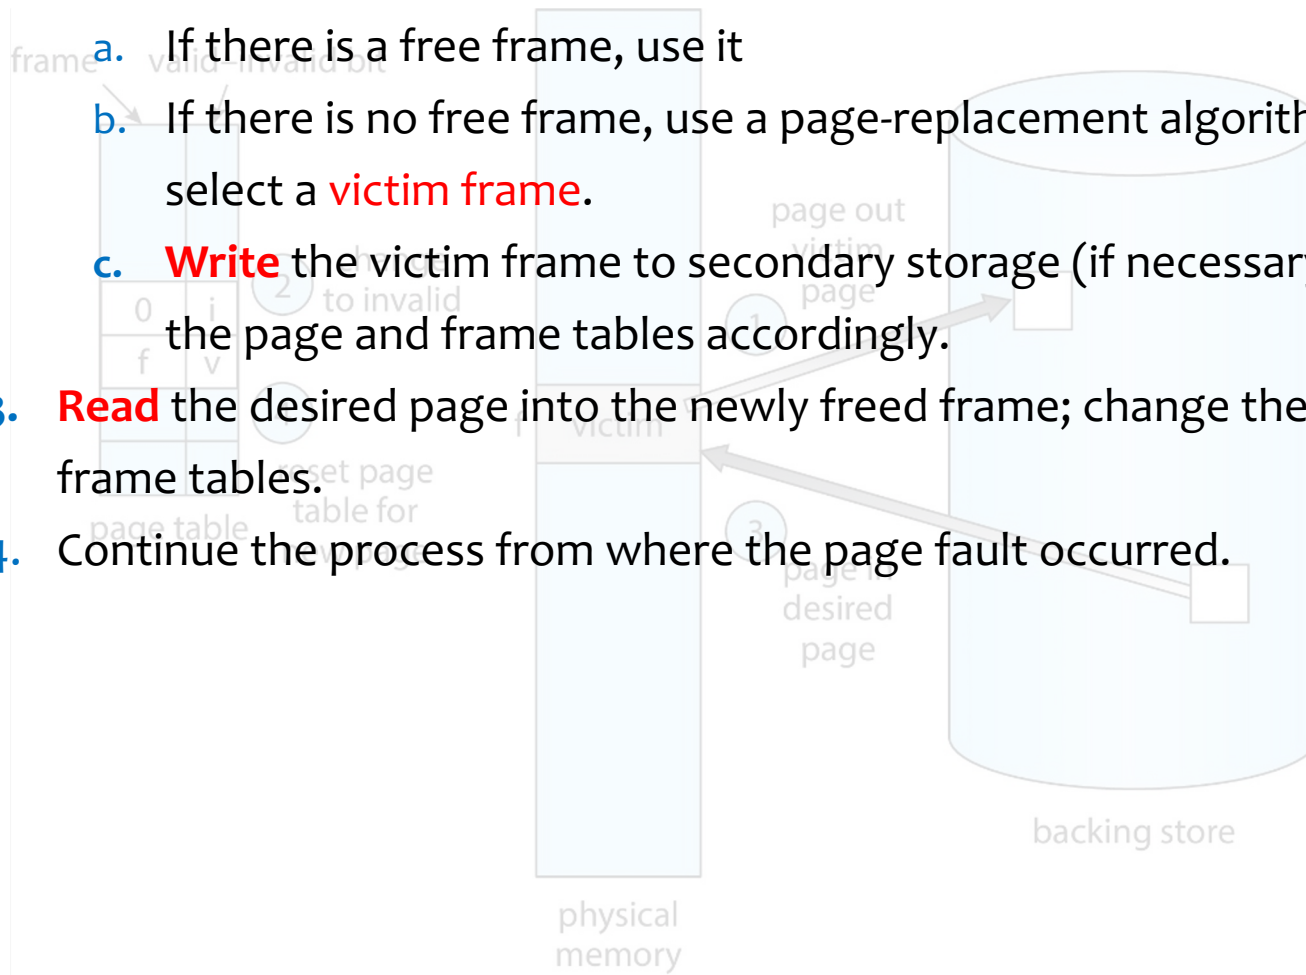




■ Page Replacement

■ Steps in handling a **Page Replacement**.

1. Find the location of the desired page/block on disk.
2. Find a free frame:
 - a. If there is a free frame, use it
 - b. If there is no free frame, use a page-replacement algorithm to select a **victim frame**.
 - c. **Write** the victim frame to secondary storage (if necessary); change the page and frame tables accordingly.
3. **Read** the desired page into the newly freed frame; change the page and frame tables.
4. Continue the process from where the page fault occurred.





■ Page Replacement

■ Steps in handling a **Page Replacement**.


1. Find the location of the desired page/block on disk.
2. Find a free frame:
 - a. If there is a free frame, use it
 - b. If there is no free frame, use a page-replacement algorithm to select a **victim frame**.
 - c. **Write** the victim frame to secondary storage (if necessary); change the page and frame tables accordingly.
3. **Read** the desired page into the newly freed frame; change the page and frame tables.
4. Continue the process from where the page fault occurred.

Notice that, if no frames are free, **two** page transfers (one for the **page-out** and one for the **page-in**) are required. This situation effectively **doubles** the page-fault service time and **increases** the effective access time accordingly.



■ The "Dirty" Bit

- We can **reduce** this overhead (of **page-out**) by using a **dirty** bit (or **modify** bit) in the **PTE**.



- The dirty bit for a page is set by the hardware whenever any byte in the page is written into, indicating that the page has been modified.
- Thus, when we select a **victim** page for replacement, we examine its **dirty** bit:
 - If the **dirty** bit is not set (0), it means this page **has not been modified**, thus we do not need to write the memory page (frame) into disk
 - If the **dirty** bit is set (1), which means that this page **has been modified** and we must write that page (frame) into disk.



Thank you!