



DCS216 Operating Systems

Lecture 22 Memory (5)

May 27th, 2024

Instructor: Xiaoxi Zhang
Sun Yat-sen University



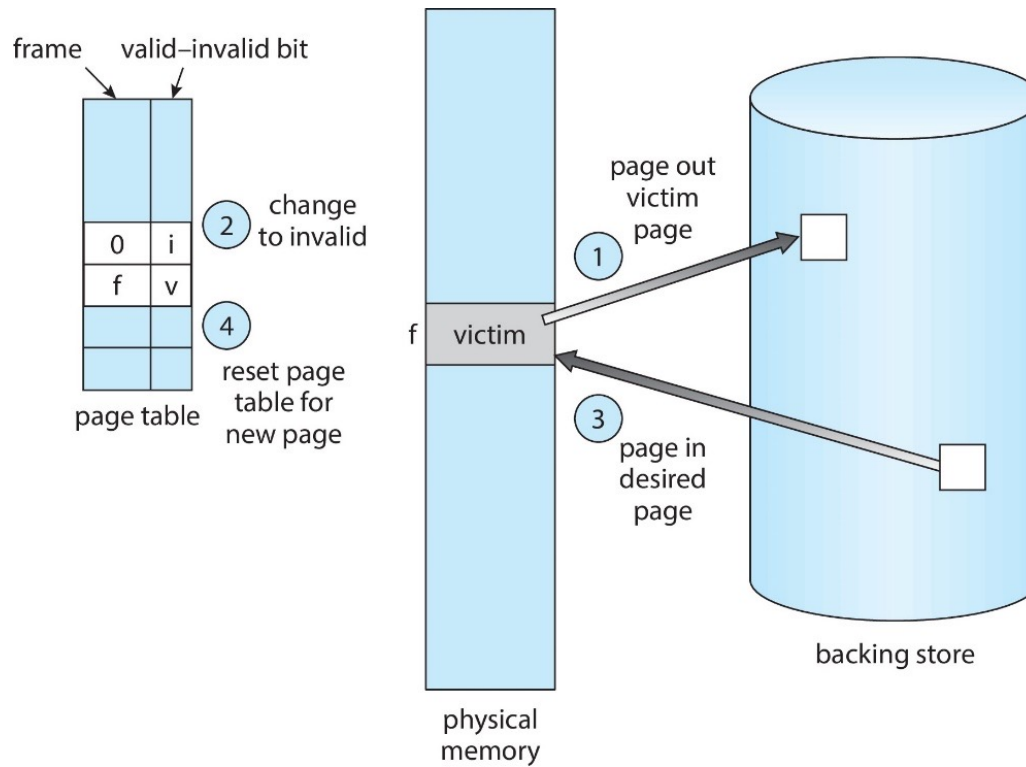
■ Content

- Background
- Demand Paging
- Copy-on-Write (**COW**)
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- OS Examples
 - Linux
 - Windows
 - Solaris



■ Page Fault

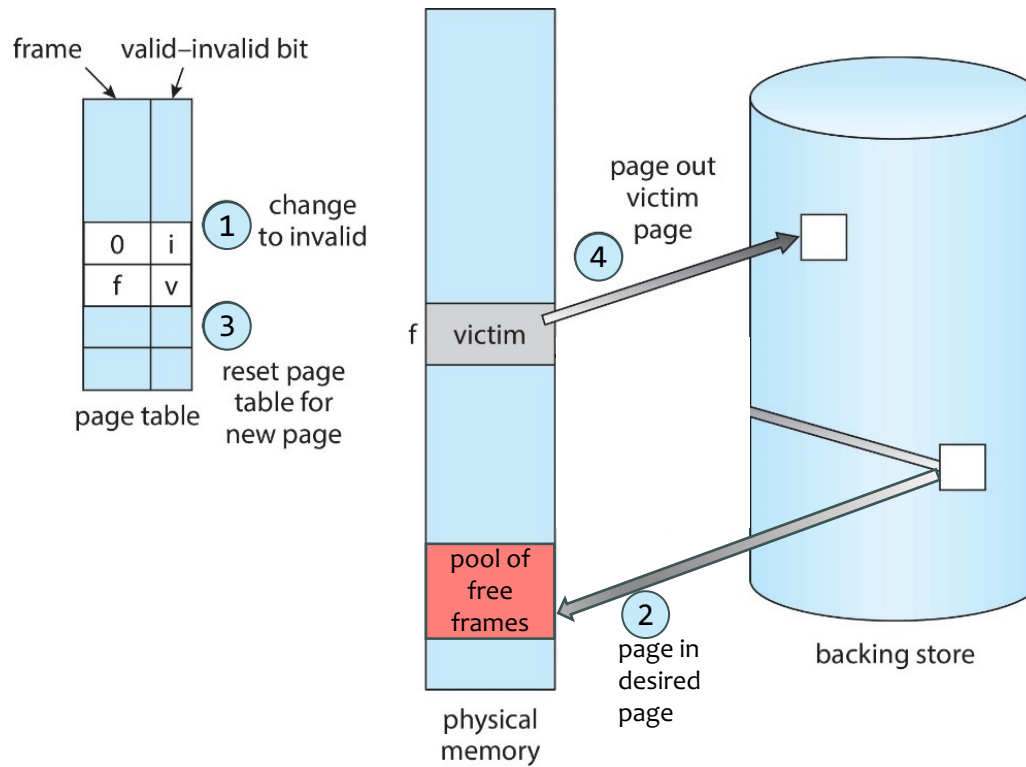
- Recall that when **Page Fault** occurs and there is no free frame, a **Page Replacement Algorithm** (e.g., LRU) normally needs to perform **two** I/O operations: **page-out** victim page and **page-in** desired page.
 - With "**Modified Bit**" set in PTE, **page-out** is mandatory.
 - How can we minimize the time for **page-out** and **page-in**?





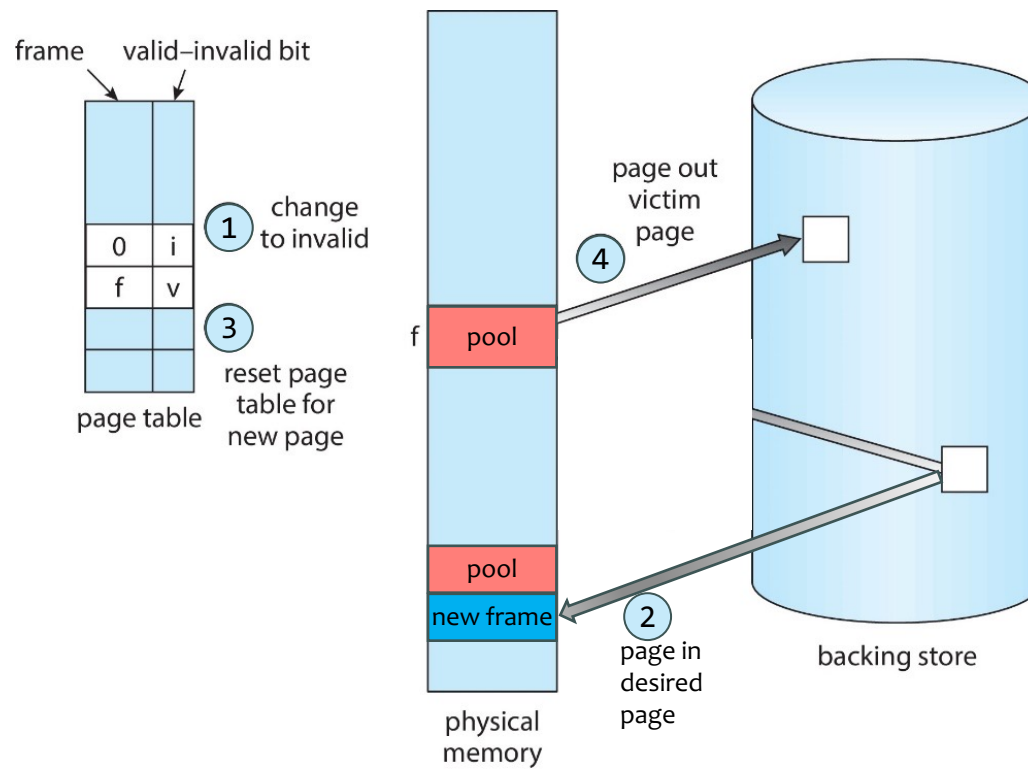
■ Page-Buffering Algorithms

- How can we minimize the time for **page-out** and **page-in**?
 - We maintain a **pool of free frames**.
 - When a **Page Fault** occurs, a victim frame is chosen. The desired page is loaded into a **frame** from the **pool of free frames** **before** the victim page is **paged out**.




■ Page-Buffering Algorithms

- How can we minimize the time for **page-out** and **page-in**?
 - Once the victim page is **paged out** (and written to disk), the **victim frame** will be **added** to the **pool of free frames**.
 - This allows the process to **restart** as soon as possible, without **waiting** for **the victim page to be written out**.



■ Page-Buffering Algorithms

- How can we minimize the time for **page-out** and **page-in**?
 - The **downside**?
 - We have to allocate an **extra** pool of free frames that **should not** be actively **used**, but rather as a **buffer** for **victim frames**.
 - Another example of **space-time** trade-off, (i.e., optimize paging **time** at the cost of extra unused memory **space**).





■ Two major problems to implement Demand Paging:

■ Frame Allocation Algorithm

- How many frames to allocate to each process?
- Which frames to replace?

■ Page Replacement Algorithm

- Which frame(s) to choose as the **victim** frame(s)?
- **Goal:** To achieve the lowest **Page Fault Rate**

■ Evaluation (by simulation)

- Running it on a particular string of memory references (**reference string**) and computing the number of **page faults** on that string
 - String is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault
 - Results depend on the number of frames available
- Our example reference string of referred page numbers is

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



■ Two major problems to implement Demand Paging:

■ Frame Allocation Algorithm

- How many frames to allocate to each process?
- Which frames to replace?

■ Page Replacement Algorithm

- Which frame(s) to choose as the **victim** frame(s)?
- **Goal:** To achieve the lowest **Page Fault Rate**

■ Evaluation (by simulation)

- Running it on a particular string of memory references (**reference string**) and computing the number of **page faults** on that string
 - String is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault
 - Results depend on the number of frames available
- Our example reference string of referred page numbers is

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



■ Allocation of Frames

- How do we allocate to **various processes** the fixed amount of free memory that is available?
- For example, consider a system with **512KB** of physical memory and the Page Size is **4KB**
 - Number of frames: **128**
 - Suppose the OS occupies 140KB \Rightarrow **35** Frames
 - Frames remaining for user processes = $128 - 35 = 93$
 - These **93** frames would be available in the **free frames list**.
- Pure Demand Paging:
 - When a user process started execution, the first **93 Page Faults** would get all free frames from the list.
 - When the free frames are exhausted, a **Page Replacement Algorithm** would be used to **select** one of the **93** in-memory pages to be **replaced** with the **94th** page, and so on.
 - When the process terminates, the **93** frames would once again be placed on the free-frames list.



■ Allocation of Frames

- Constraints for Allocation of Frames:
 - **Maximum:** We cannot allocate more than the **total number of available frames** (unless there is page sharing)
 - **Minimum:** A **minimum number of frames** should be allocated to each process.
- There should be enough frames to hold all the different pages that **any single instruction** can **reference**.
 - **Minimum** number of frames per process is **defined** by the **architecture**.
 - Suppose in a **one-level indirect** addressing, a **load** instruction (on **page 15**) refers to an address on **page 3**, which is an indirect reference to **page 22**.
 - A minimum of 3 frames are needed.
 - In IBM 370, the **SS MOVE** (Storage-to-Storage **MOVE**) instruction:
 - **MVC dest, src**



Allocation of Frames

- In IBM 370, the **SS MOVE** (Storage-to-Storage **MOVE**) instruction:
 - **MVC dest, src**
 - Instruction itself is 6 bytes, which might span 2 pages
 - E.g., the **MVC dest, src** instruction is located at **0x401FFE**
 - 2 pages to handle **dest** (**0x403100**)
 - 1 page to reference the addr of label **dest**
 - 1 page to reference the addr **dest** points to
 - 2 pages to handle **src** (**0x404500**)
 - 1 page to reference the addr of label **src**
 - 1 page to reference the addr **src** points to
 - In total: a minimum of **6 frames** needed in the **worst** case.





■ Allocation of Frames

- In **x86**, however, **direct memory-to-memory** movement is **not allowed**.
 - ``mov (%esi, %ebx, 4), (%edx)`` is **illegal**.
- The most complex instruction (in terms of # of memory accesses)
 - E.g., ``mov (%esi, %ebx, 4), %edx`` (located at **0x401FFE**).
 - requires a minimum of **4 frames**.
 - **2** pages for instruction (that **spans** two pages)
 - **2** pages for indirect memory reference of one of the operands.



■ Allocation Algorithms

- Frame Allocation Algorithms would help us in determining how many frames should be allocated to different processes in a multiprogramming environment.
- 3 most common allocation algorithms:
 - **Equal** Allocation
 - **Proportional** Allocation
 - **Priority** Allocation



■ Equal Allocation

- The frames are equally distributed among the processes.
- If we have m frames and n processes
 - allocate m/n frames to each process. (ignoring the OS for the moment)
- For example:
 - Number of frames = 93
 - Number of processes = 5
 - Each process gets $\lfloor 93/5 \rfloor = 18$ frames
 - the remaining 3 frames can be kept as **buffer pool of free frames**.
- The Problem: **memory requirement of processes are not the same.**
- For example, in a system with 64 frames with page size of 1KB. For only 2 processes: **P1(10KB)** and **P2(127KB)**, if we adopt equal allocation, then each process gets allocated 32 frames.
 - For **P1(10KB)**, 22 frames were **wasted**.
 - For **P2(127KB)**: underallocated.



■ Proportional Allocation

- The frames are distributed among processes according to their sizes.

- Let:

- m be the number of frames

- s_i be the size of process p_i

- $S = \sum s_i$

- then Number of frames allocated to p_i :

- $a_i = \frac{s_i}{S} \times m$

- For example, in a system with 64 frames with page size of 1KB. For only 2 processes: **P1(10KB)** and **P2(127KB)**, if we adopt proportional allocation, then

- Frames allocated to p_1 : $a_1 = \frac{s_1}{S} \times m = \frac{10}{137} \times 64 \approx 5$

- Frames allocated to p_2 : $a_2 = \frac{s_2}{S} \times m = \frac{127}{137} \times 64 \approx 59$



■ Priority Allocation

- The frames are distributed among processes according to their **priorities**.
 - Processes with higher priorities are allocated more frames so as to **speed up** their execution.
 - $a_i = \frac{prio_i}{\sum prio_i} \times m$
- A combination of **size and priority** can be used rather than just the process **size or priority**.




■ Global Allocation vs. Local Allocation

- Page Replacement Algorithms can also be classified into two broad categories based on the way frames are allocated to different processes:
 - **Global Allocation**
 - Global replacement allows a process to select a replacement frame from the set of all frames, **even if that frame is currently allocated to some other process**; that is, one process can take a frame from another.
 - **Local Allocation**
 - Local replacement requires that each process select from **only its own set of allocated frames**.

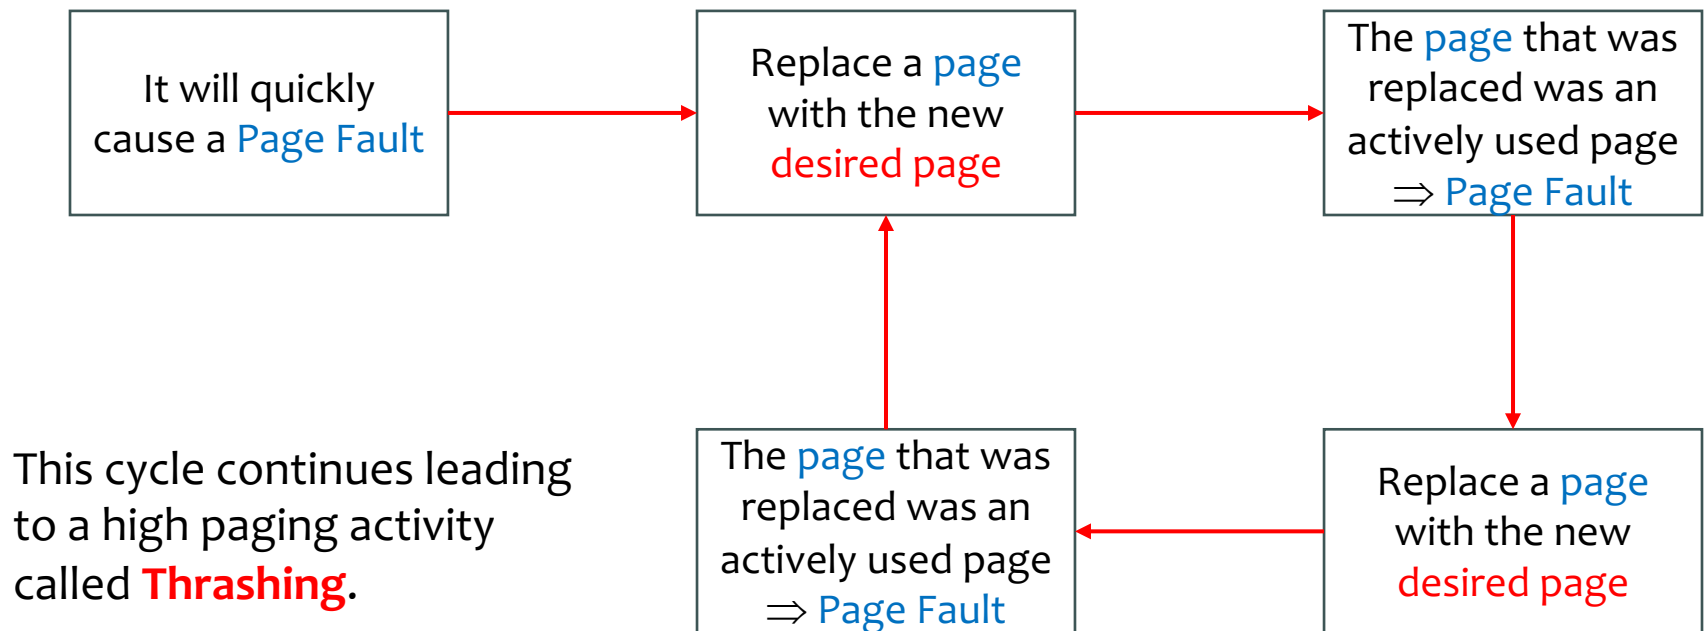


■ Global Allocation



■ Thrashing

- Consider a process that doesn't have enough frames for execution.



If a process is spending more time in paging rather than in executing, then we say that process is **Thrashing** (系统抖动).



■ Severe Performance Issue due to Thrashing

- The OS monitors CPU utilization. If CPU utilization is too low, we generally increase the degree of multiprogramming by introducing a new process into the system.
- As processes are busy swapping pages in and out, they queue up for the paging device and the ready queue (processes) becomes empty.
- Since processes are now waiting for the paging device, the CPU utilization decreases.
- The CPU scheduler observes the decreased utilization and therefore attempts to increase it by introducing new processes in the system.
- These new processes tries to take frames from the older processes, hence increasing the number of page faults.
- The CPU utilization drops further \Rightarrow the cycle repeats
- Here we see that Thrashing occurs and system throughput decreases dramatically: **No real work is being done.**



■ Severe Performance Issue due to Thrashing




Figure 10.20 Thrashing.



■ Thrashing

- If a process does not have "enough" pages, the **page-fault rate** is very high. This leads to:
 - low CPU utilization
 - OS spends most of its time swapping in and out of disk
- **Thrashing**: a process is busy swapping pages in and out with little or no actual progress.
- **Question:**
 - How do we detect **Thrashing**?
 - What is the best response to **Thrashing**?



■ Working Set Model

- How to prevent **Thrashing**?
 - We must provide a process with as many frames as it needs.
 - But how do we know how many frames it "needs"?
- We make use of the **Working Set Strategy** where it checks how many frames a process is actually using.
 - This approach defines the **locality model** of process execution.
 - As a process executes, it moves from **locality** to **locality**.
 - A **locality** is a set of pages that are actively used together.
 - A running program is generally composed of several different **localities**, which may overlap.



Working Set Model

- **Locality** (局部性) of a process changes over time.
- At time (a), the **locality** is the set of pages {18, 19, 20, 21, 22, 23, 24, 29, 30, 33}.
- At time (b), the **locality** changes to {18, 19, 20, 24, 25, 26, 27, 28, 29, 31, 32, 33}.
- Notice the overlap, as some pages are (e.g., {18, 19, 20}) are part of both localities.




Figure 10.21 Locality in a memory-reference pattern.





Working Set Model

- The **Working Set Model** is based on the assumption of **locality**.
 - We use a parameter Δ to define the **working set window**.
 - The set of pages in the most recent Δ page references is the **working set**.
 - If a page is in active use, it will be in the **working set**.
 - If it is no longer being used, it will drop from the **working set** Δ time units after its last reference.
 - Thus, the **working set** is an **approximation** of the program's **locality**.
- For example: $\Delta = 10$.
 - The working set at time t_1 is $\{1, 2, 5, 6, 7\}$, with size of **5** pages.
 - The working set at time t_2 is $\{3, 4\}$, with size of **2** pages.

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...





■ Working Set Model

- The accuracy of the **working set** depends on the selection of Δ
 - If Δ is **too large** \Rightarrow it may overlap several localities.
 - If Δ is **too small** \Rightarrow it may not cover the entire localities.
 - If Δ is **infinite** \Rightarrow the working set is the set of pages touched during the process execution.
- The most important property of the **working set** is its **size**.
 - The **working set size** of a process is denoted by WSS_i
 - so the total **demand** for frames in a system is

$$D = \sum WSS_i$$

- If the total demand is greater than the total number of available frames, i.e., $D > m$, then **Thrashing** will occur.
 - In this case, the OS will select a process to **suspend**. In other words, the process's pages are swapped out and its frames are reallocated to other processes.



■ Working Set Model

- The most important property of the **working set** is its **size**.
 - The **working set size** of a process is denoted by WSS_i
 - so the total **demand** for frames in a system is


$$D = \sum WSS_i$$

- If the total demand is greater than the total number of available frames, i.e., $D > m$, then **Thrashing** will occur.
 - In this case, the OS will select a process to **suspend**. In other words, the process's pages are swapped out and its frames are reallocated to other processes.
- The **working set strategy prevents thrashing** while keeping the degree of multiprogramming as **high** as possible
 - Thus, **optimizing CPU utilization**.



■ Working Set Model


- The working set model is successful, but it seems too clumsy.
- A strategy that uses the **Page-Fault Frequency (PFF, 缺页率)** takes a more **direct** approach: **Control** the **page-fault rate**.
 - If it is too **high**, we know the process needs more frames and we allocate the process another frame.
 - If it is too **low**, then the process may have too many frames and we remove a frame from the process.





■ Working Set Model

- There is a direct relationship between **working set** of a process and its **page-fault rate**.
 - **Working set** changes over time.
 - **Page-fault rate** peaks and valleys over time.
 - A **peak** in the page-fault rate occurs when we begin demand-paging a new locality; Once the working set of this new locality is in memory, the page-fault-rate **falls**.





■ Allocating Kernel Memory

- When a user-mode process requests additional memory (e.g., via `malloc`), pages are **allocated** from the **list of free frames maintained by the kernel**.
- However, memory allocation inside the kernel is different. When the kernel requests additional memory (e.g., via `kmalloc`), pages are often allocated from a **different free-memory pool**.
 - The kernel requests memory for data structures of **varying sizes**, some of which are **less than a page in size**. As a result, the kernel must use memory **conservatively** and attempt to **minimize waste** due to **fragmentation**.
 - Pages allocated to user-mode processes do not necessarily have to be **contiguous physical memory**. However, certain hardware devices **interact directly** with **physical memory** (*without the benefit of a virtual memory interface, since VM is part of the kernel*), and consequently may require memory residing in physically **contiguous pages**.




■ Buddy System

- Buddy system allocates memory from a **fixed-size segment** consisting of **physical contiguous pages**.
- Memory is allocated from this segment using a **power-of-2 allocator**.
 - Satisfies requests in units sizes as a **power of 2** (e.g., 4KB, 8KB, 16KB...)
 - Request **not appropriately sized** **rounded up** to next highest power of 2
 - When smaller allocation needed than is available, current chunk split into two **buddies** of next-lower power of 2
- Assume a **256KB** chunk (**A**) available, and the kernel requests **21KB**:
 - Split into **A_L** and **A_R** (each **128KB** in size)
 - **A_L** further split into **B_L** and **B_R** (each **64KB** in size)
 - **B_L** further split into **C_L** and **C_R** (each **32KB** in size)
 - **C_L** is allocated to satisfy the **21KB** request.



Buddy System


- Assume a 256KB chunk (A) available, and the kernel requests 21KB:
 - Split into A_L and A_R (each 128KB in size)
 - A_L further split into B_L and B_R (each 64KB in size)
 - B_L further split into C_L and C_R (each 32KB in size)
 - C_L is allocated to satisfy the 21KB request.





■ Buddy System

■ Example: Current state of Buddy allocation






Buddy System

Example:

- Where do we allocate a request of size **28**?






Buddy System

Example:

- Where do we allocate a request of size **28**?

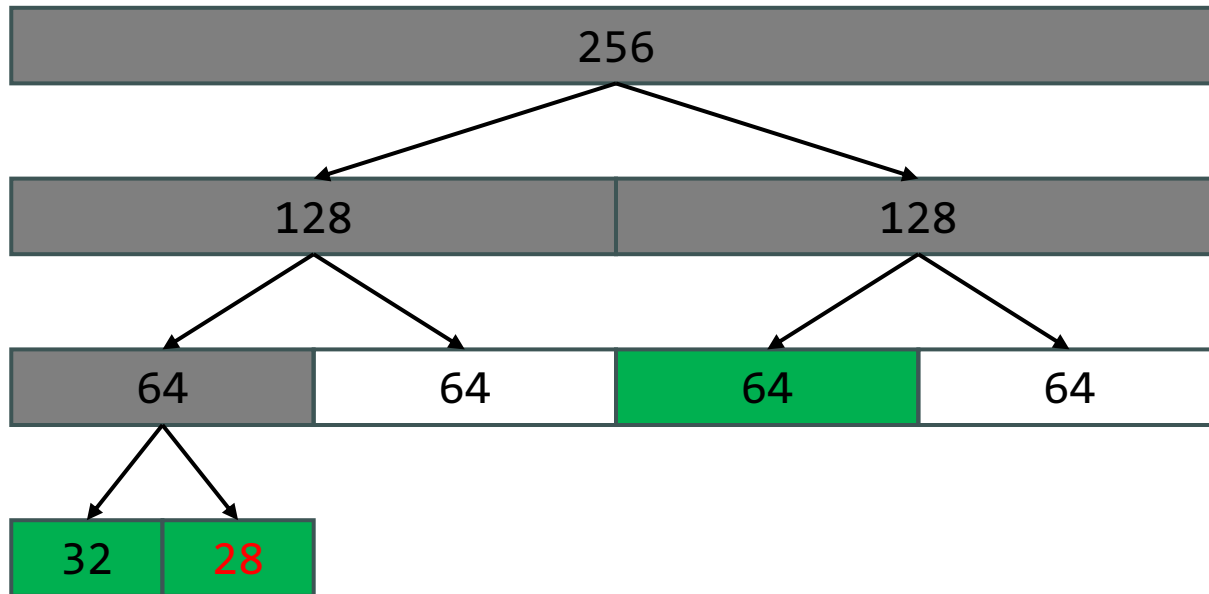




Buddy System

Example:

- Where do we allocate a request of size **28**?
- Where do we allocate a request of size **36**?






Buddy System

Example:

- Where do we allocate a request of size **28**?
- Where do we allocate a request of size **36**?



■ Buddy System


■ Example:

- Where do we allocate a request of size 28?
- Where do we allocate a request of size 36?
- What happens when we free the size 64 chunk?

 : Allocated

 : Occupied

☐ : Free

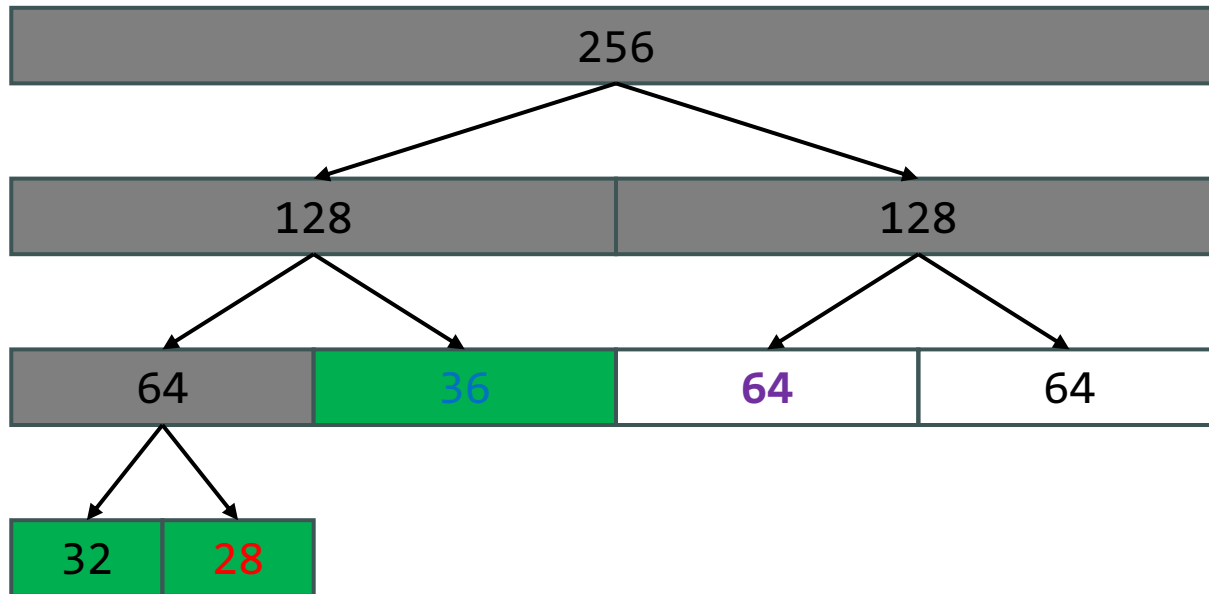





Buddy System

Example:

- Where do we allocate a request of size **28**?
- Where do we allocate a request of size **36**?
- What happens when we **free** the size **64** chunk?

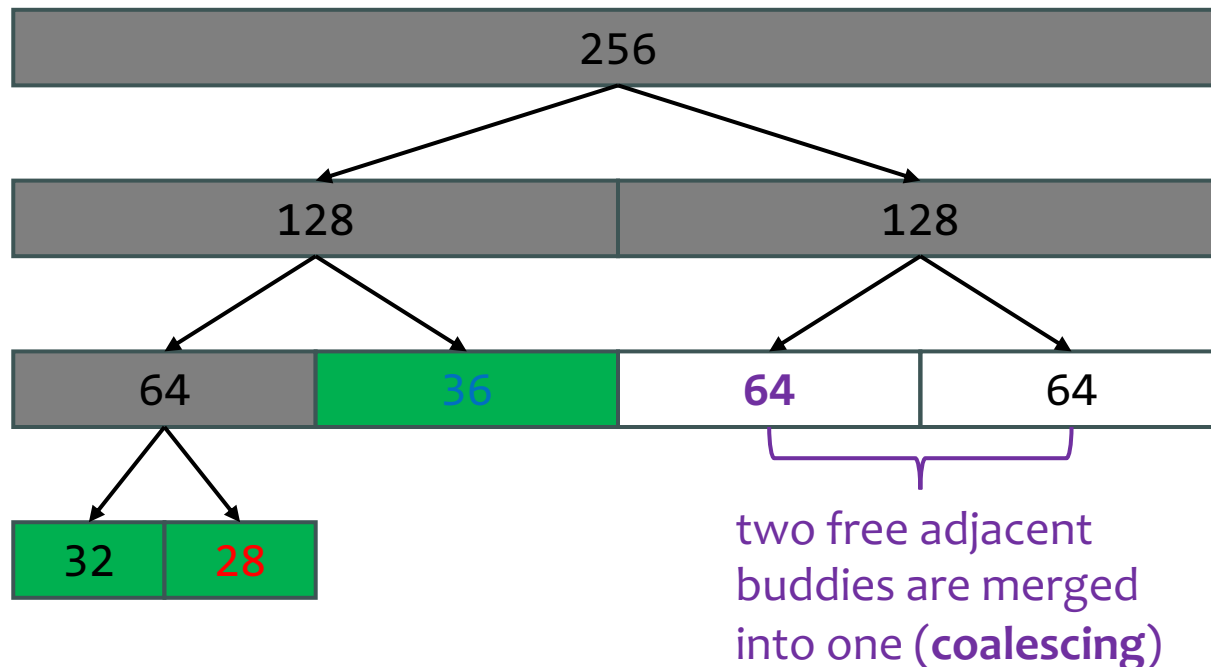





Buddy System

Example:

- Where do we allocate a request of size **28**?
- Where do we allocate a request of size **36**?
- What happens when we **free** the size **64** chunk?






Buddy System

Example:

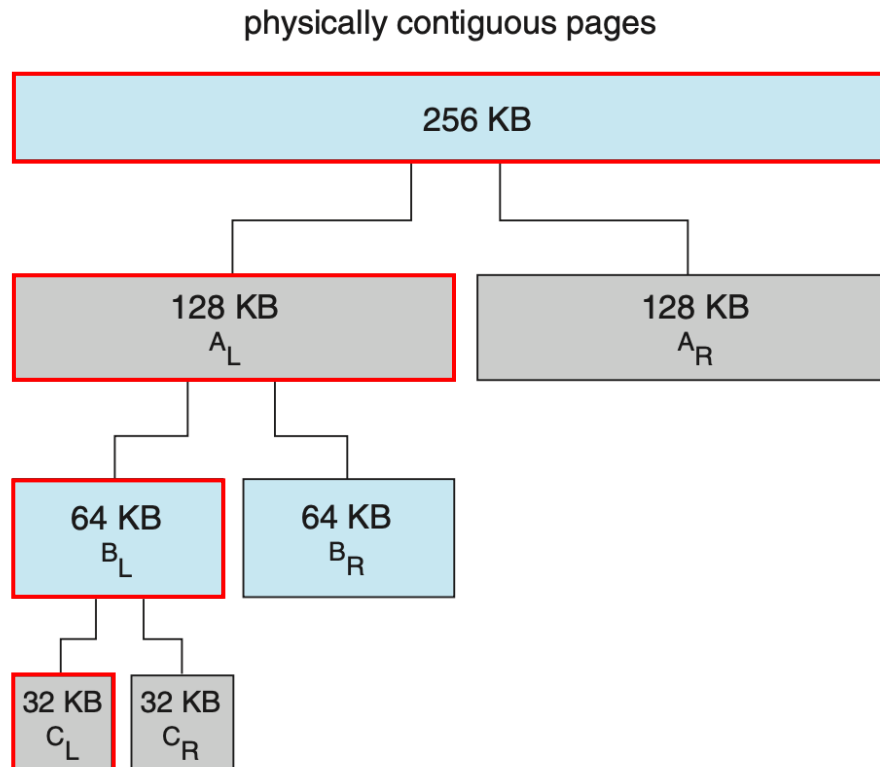
- Where do we allocate a request of size **28**?
- Where do we allocate a request of size **36**?
- What happens when we **free** the size **64** chunk?





■ Buddy System

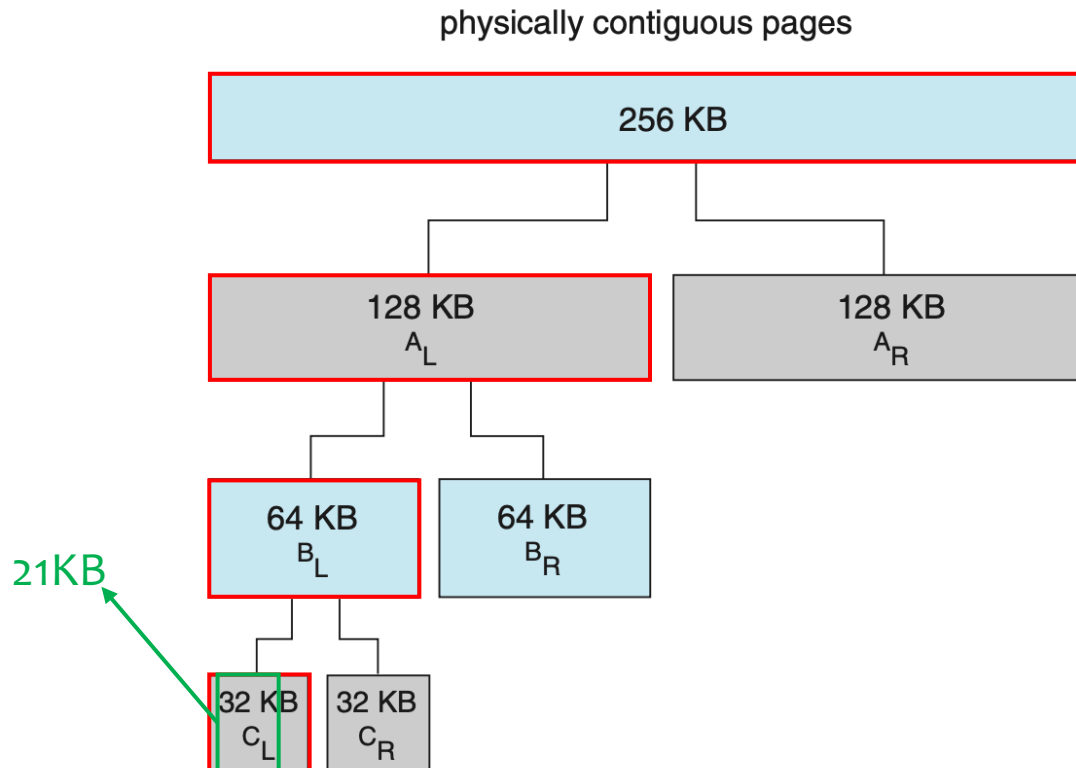
- **Advantage:** Adjacent buddies can be combined to form larger segments **quickly** using a technique known as **coalescing**.
 - For example, when the kernel releases C_L , the system can coalesce C_L and C_R into a 64KB segment B_L , which in turn can be coalesced with B_R to form a 128KB segment A_L , and so on...





■ Buddy System

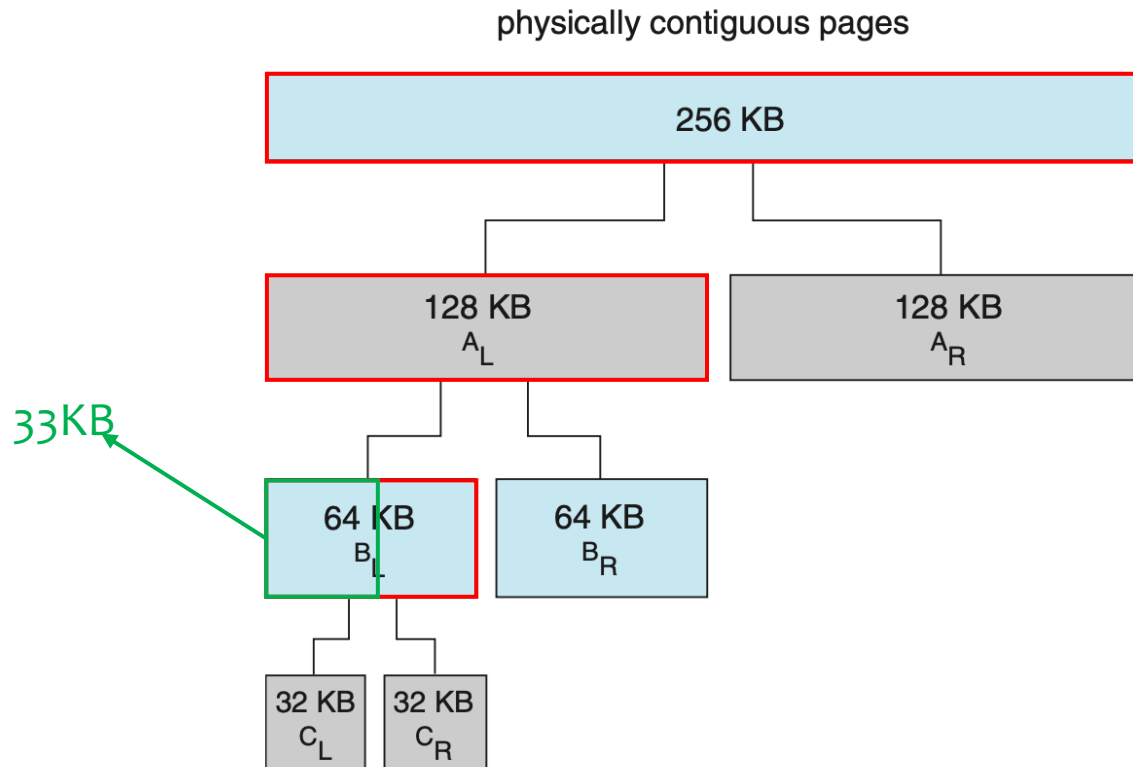
- **Disadvantage:** Rounding up to the next highest power of 2 is very likely to cause **fragmentation** within allocated segments.
 - For example, the **21KB** request is satisfied with a **32KB** segment, with **11KB** wasted.





■ Buddy System


- **Disadvantage:** Rounding up to the next highest power of 2 is very likely to cause **fragmentation** within allocated segments.
 - For example, the **21KB** request is satisfied with a **32KB** segment, with **11KB** wasted.
 - Worse, consider a **33KB** request that is satisfied with a **64KB** segment.





■ Slab Allocation


- **Slab allocation**: an alternative strategy for allocating kernel memory.
 - A **slab** is made up of one or more physically contiguous pages/frames.
 - A **cache** consists of one or more **slabs**.
 - There is a single cache for **each unique** kernel data structure





■ Slab Allocation


- **Slab allocation**: an alternative strategy for allocating kernel memory.
 - A **slab** is made up of one or more physically contiguous pages/frames.
 - A **cache** consists of one or more **slabs**.
 - There is a single cache for **each unique** kernel data structure
 - For example, all **3-KB** objects belong to the same cache.





■ Slab Allocation

- **Slab allocation**: an alternative strategy for allocating kernel memory.
 - A **slab** is made up of one or more physically contiguous pages/frames.
 - A **cache** consists of one or more **slabs**.
 - There is a single cache for **each unique** kernel data structure
 - For example, all **3-KB** objects belong to the same cache.





■ Slab Allocation

- The **slab allocation** algorithm uses **caches** to store kernel objects.
- When a **cache** is created, a number of objects (which are initially marked as free) are allocated to the **cache**.
 - The number of objects in the cache depends on the size of the associated slab.
 - For example, a 12KB slab (made up of 3 contiguous 4KB frames) could store **six** 2KB objects.
 - Initially, all objects in the cache are marked as **free**.
 - When a new object for a kernel data structure is needed, the allocator can assign any free object from the cache to satisfy the request.
 - The object assigned from the cache is marked as **used**.
- In Linux, a slab may be in one of **three** possible states:
 - **Full**. All objects in the slab are marked as **used**.
 - **Empty**. All objects in the slab are marked as **free**.
 - **Partial**. Mix of **used** and **free** objects.



■ Slab Allocation

- In Linux, a slab may be in one of **three** possible states:
 - **Full**. All objects in the slab are marked as **used**.
 - **Empty**. All objects in the slab are marked as **free**.
 - **Partial**. Mix of **used** and **free** objects.

- Upon request from the kernel:
 - The **SLAB** allocator first attempts to find a (**partial**) slab consisting **free** objects to satisfy the request.
 - If none exists, a free object is assigned from an **empty** slab
 - If no **empty** slabs are available, **a new slab is allocated** from contiguous physical frames and assigned to a cache
 - memory for the object is allocated from this new slab.



■ Slab Allocation

- The slab allocator provides two main benefits:
 - No memory is wasted due to fragmentation.
 - Memory requests can be satisfied quickly.



■ SLAB → SLOB → SLUB

- Linux originally used the buddy system. (< version 2.2)
- Beginning with version 2.2, the Linux kernel adopted the **SLAB** allocator.
- Recent distributions of Linux include the **SLOB** kernel allocators:
 - **SLOB** allocator is designed for systems with a limited amount of memory (e.g., embedded systems).
 - **SLOB** stands for **S**imple **L**ist **O**f **B**locks.
- **SLOB** maintains three lists of objects:
 - **small** (for objects < **256** bytes)
 - **medium** (for objects < **1024** bytes)
 - **large** (for all other objects < **page size**)





■ SLAB → SLOB → SLUB

- Linux originally used the buddy system. (< version 2.2)
- Beginning with version 2.2, the Linux kernel adopted the **SLAB** allocator.
- Recent distributions of Linux include the **SLOB** kernel allocators:
 - **SLOB** allocator is designed for systems with a limited amount of memory (e.g., embedded systems).
 - **SLOB** stands for **S**imple **L**ist **O**f **B**locks.
- **SLOB** maintains three lists of objects:
 - **small** (for objects < **256** bytes)
 - **medium** (for objects < **1024** bytes)
 - **large** (for all other objects < **page size**)
- Memory requests are allocated from an object on the appropriate list using a **first-fit** policy.





■ SLAB → SLOB → SLUB

- Linux originally used the buddy system. (< version 2.2)
- Beginning with version 2.2, the Linux kernel adopted the **SLAB** allocator.
- Beginning with Version 2.6.24, the **SLUB** allocator replaced **SLAB** as the default allocator for the Linux kernel.
 - **SLUB** is basically a performance-optimized **SLAB**.
 - **SLUB** reduced much of the overhead required by the **SLAB** allocator.
 - For instance, whereas **SLAB** stores certain metadata with each slab, **SLUB** stores these data in the page structure the Linux kernel uses for each page.
 - Additionally, **SLUB** does not include the per-CPU queues that the **SLAB** allocator maintains for objects in each cache. For systems with a large number of processors, the amount of memory allocated to these queues is significant. Thus, **SLUB** provides better performance as the number of processors on a system increases.



■ Other Considerations

- Prepaging
- Page Size
- TLB Reach
- Inverted Page Table
- Program Structure
- I/O Interlock and Page Locking



■ Prepaging

- To reduce the large number of page faults that occurs at process startup
- **Prepage** (预先调取页面) *all* or *some* of the pages a process will need, before they are referenced
- But if prepaged pages are unused, I/O and memory was wasted
- Assume s pages are prepaged and α of the pages is used.
 - Is cost of $s \times \alpha$ save page faults **larger** or **less** than the cost of prepaging $s \times (1 - \alpha)$ unnecessary pages?
 - $\alpha \rightarrow 0 \Rightarrow$ prepaging is not worth it.



■ Page Size

- Sometimes the OS designers have a choice in choosing the **Page Size**.
 - especially if running on custom-built CPU
- **Page size** selection must take into consideration:
 - Fragmentation
 - Page Table Size
 - Resolution
 - I/O overhead
 - Number of Page Faults
 - Locality
 - TLB Size and effectiveness
- Always a power of 2, usually in the range between 2^{12} and 2^{22} bytes.

Architecture	Huge Page Size
ARM64	4K, 2M, 1G
i386	4K, 4M
x86_64	4K, 8K, 64K, 256K, 1M, 4M, 16M, 256M
ppc64	4K, 16M



■ TLB Reach (TLB范围)

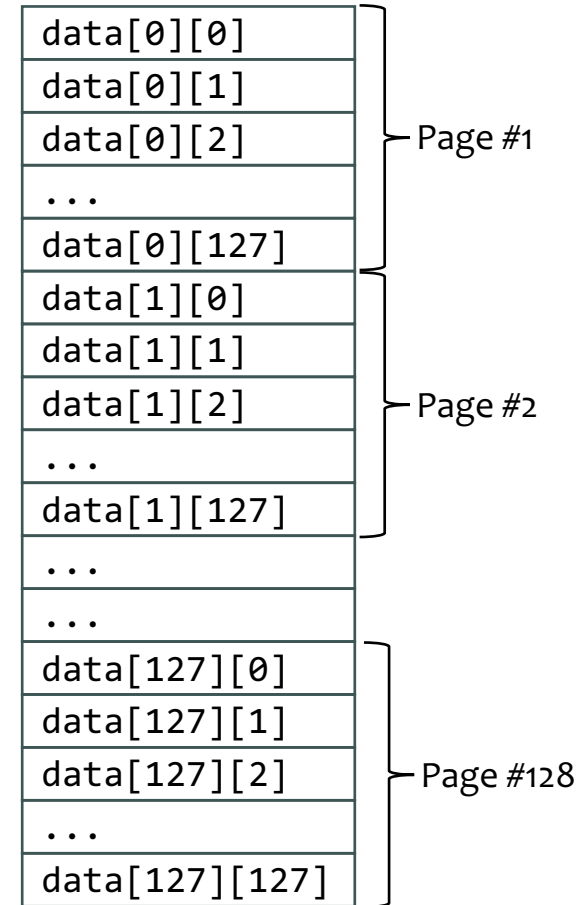
- **TLB Reach** - The amount of memory accessible from the TLB
- **TLB Reach** = (TLB Size) × (Page Size)
- Ideally, the **working set** of each process is stored in the TLB
 - Otherwise there is a high degree of page faults
- Increase the Page Size
 - \Rightarrow This may lead to an increase in fragmentation as not all applications require a large page size
- Provide Multiple Page Sizes
 - This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation



■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```



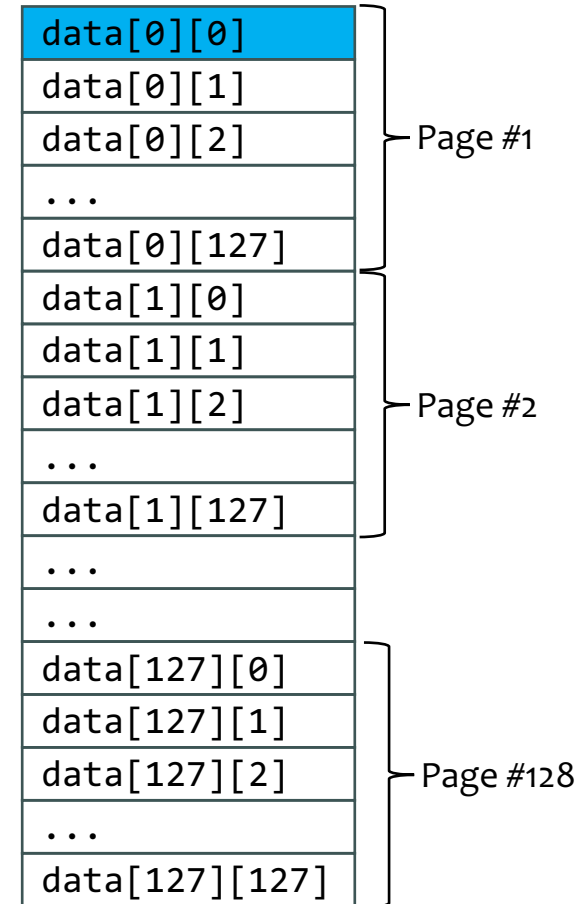


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 1



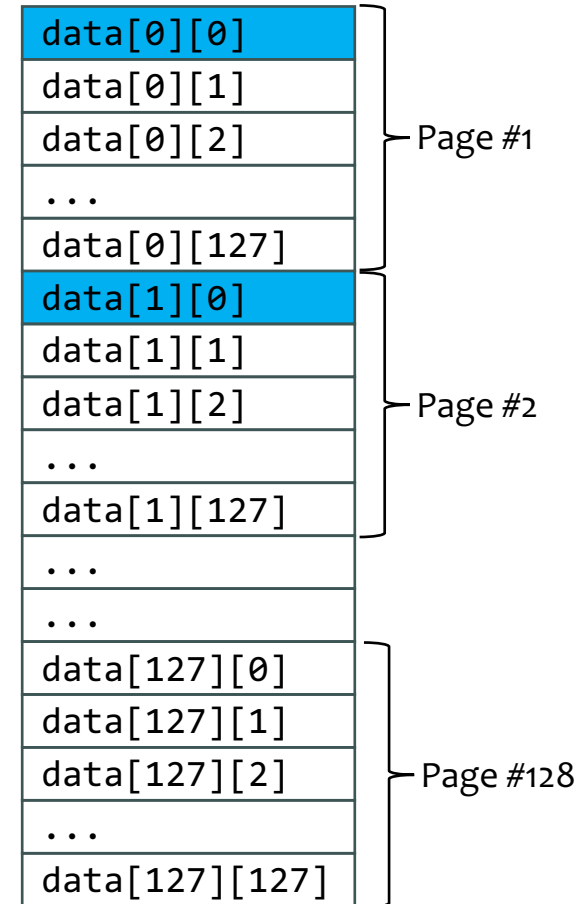


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 2



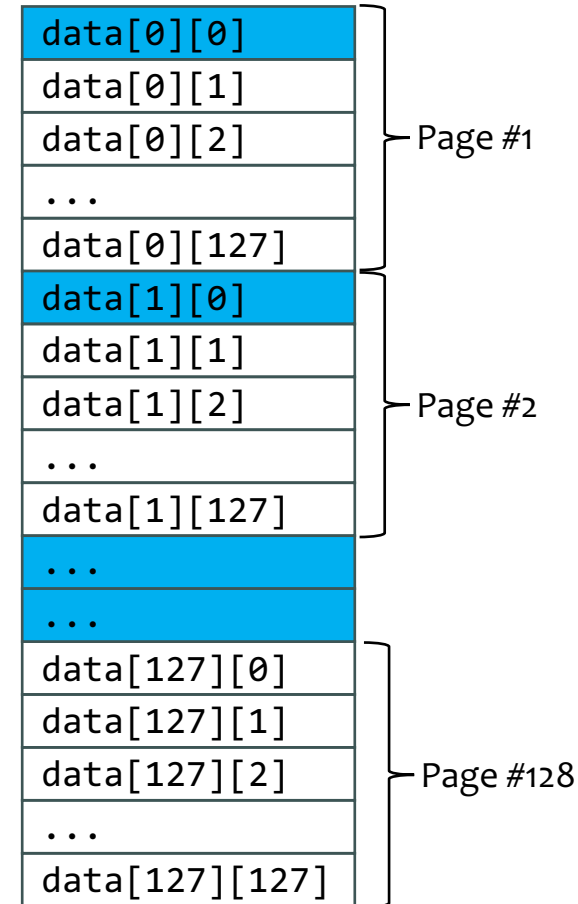


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 127



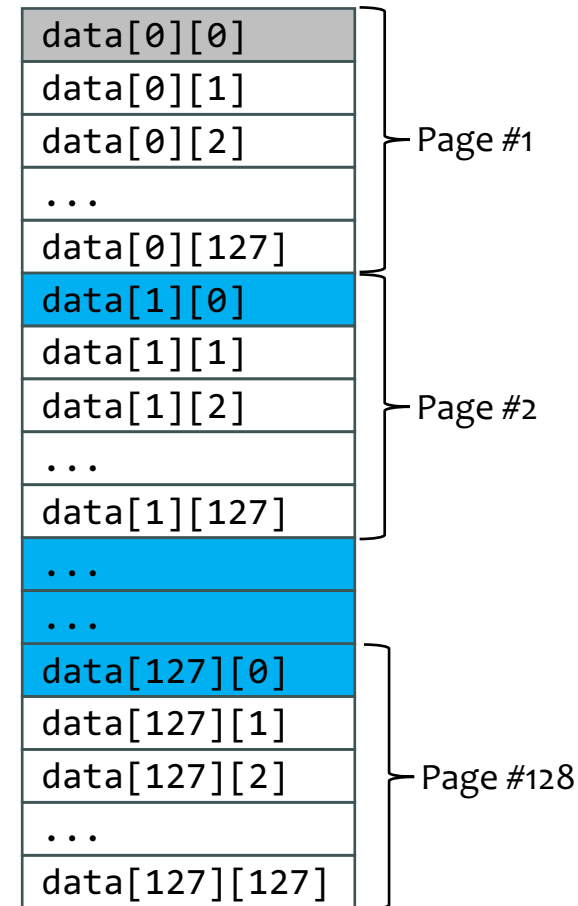


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 128
(FIFO evict
Page #1)



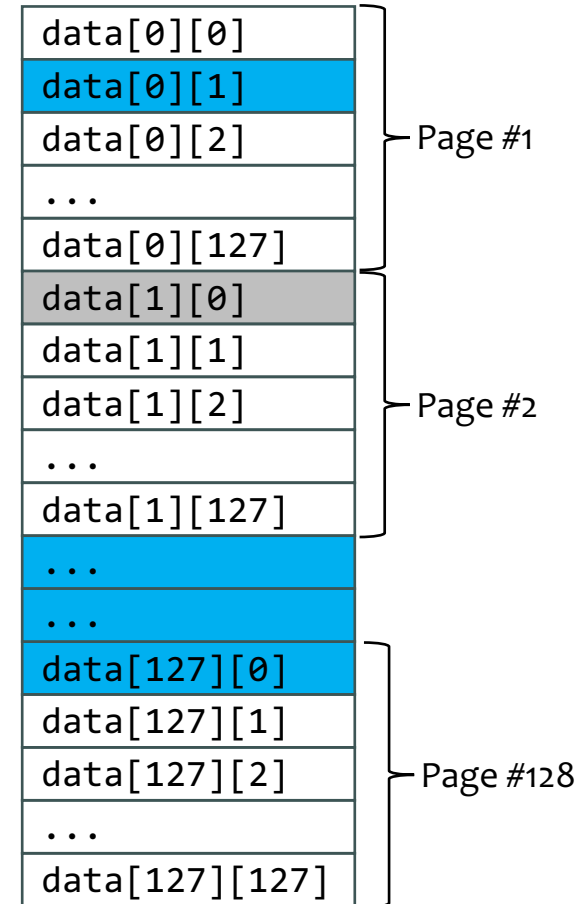


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 129



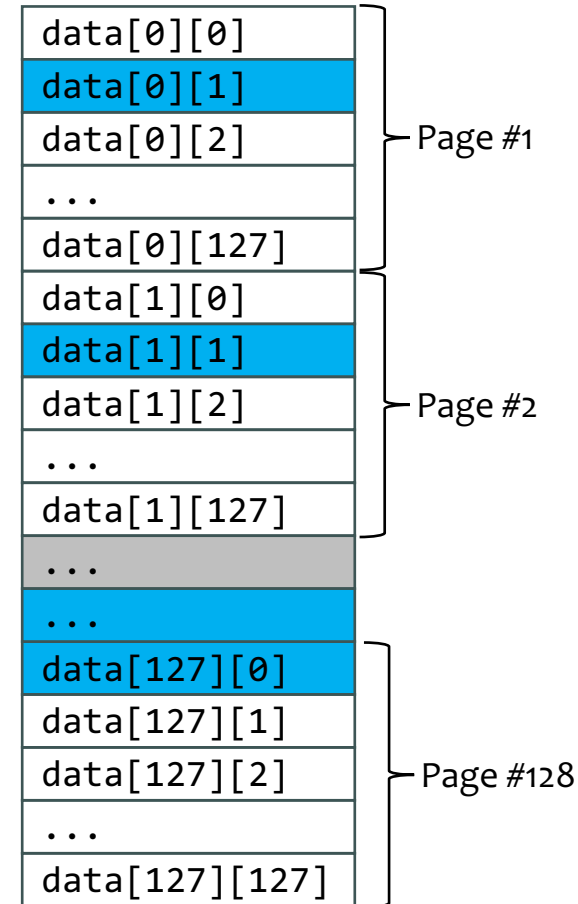


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 130






■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 254



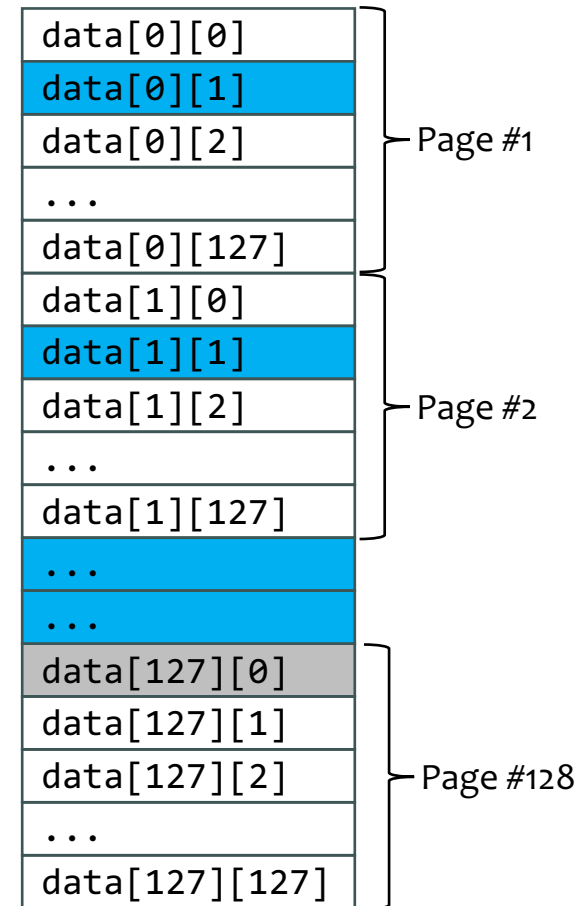


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 255



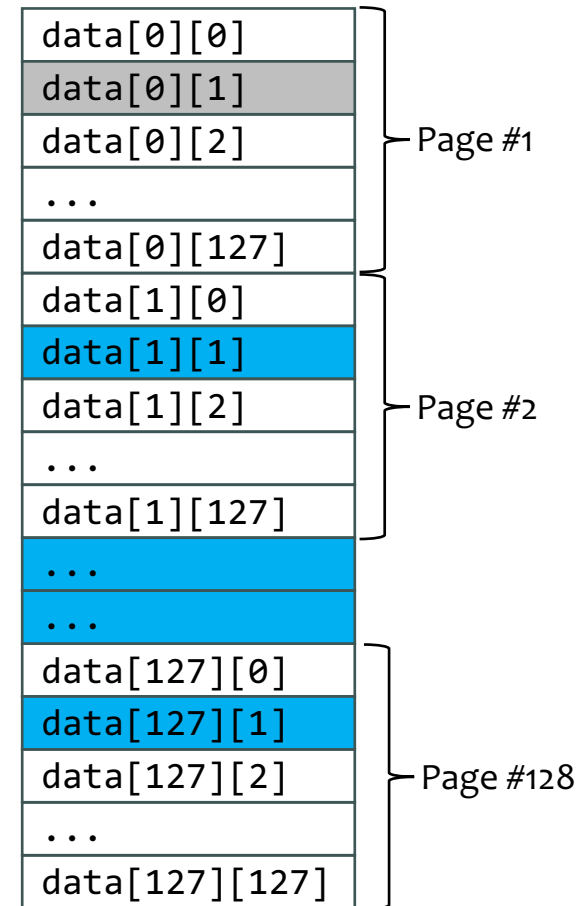


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults: 256



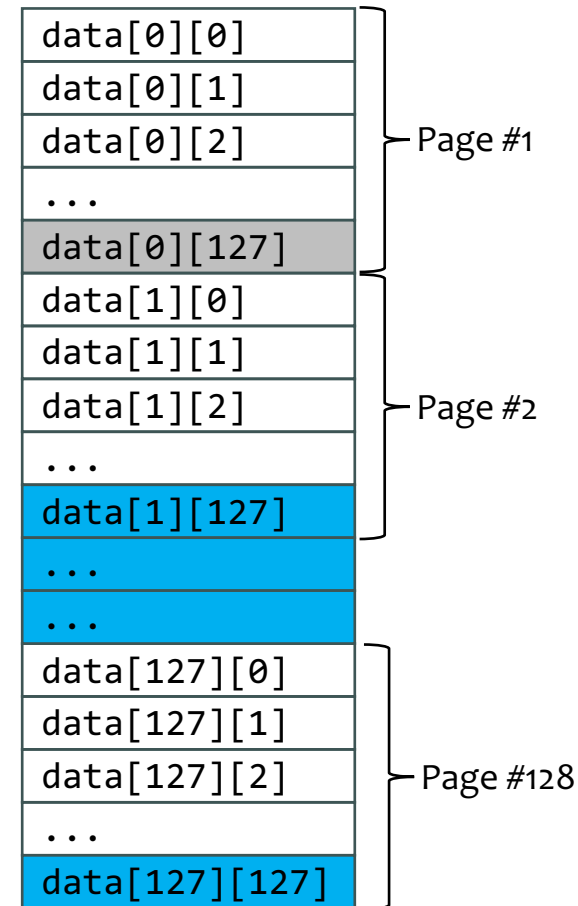


■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;  
int [128][128] data;  
  
for (j = 0; j < 128; j++) {  
    for (i = 0; i < 128; i++) {  
        data[i][j] = 0;  
    }  
}
```

of Page
Faults:
 $128 \times 128 =$
16384





■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;
int [128][128] data;


for (j = 0; j < 128; j++) {
    for (i = 0; i < 128; i++) {
        data[i][j] = 0;
    }
}
```

of Page
Faults:
 $128 \times 128 =$
16384

- An optimized (cache-friendly) version:

```
int i, j;
int [128][128] data;

for (i = 0; i < 128; i++) {
    for (j = 0; j < 128; j++) {
        data[i][j] = 0;
    }
}
```





■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;
int [128][128] data;

for (j = 0; j < 128; j++) {
    for (i = 0; i < 128; i++) {
        data[i][j] = 0;
    }
}
```

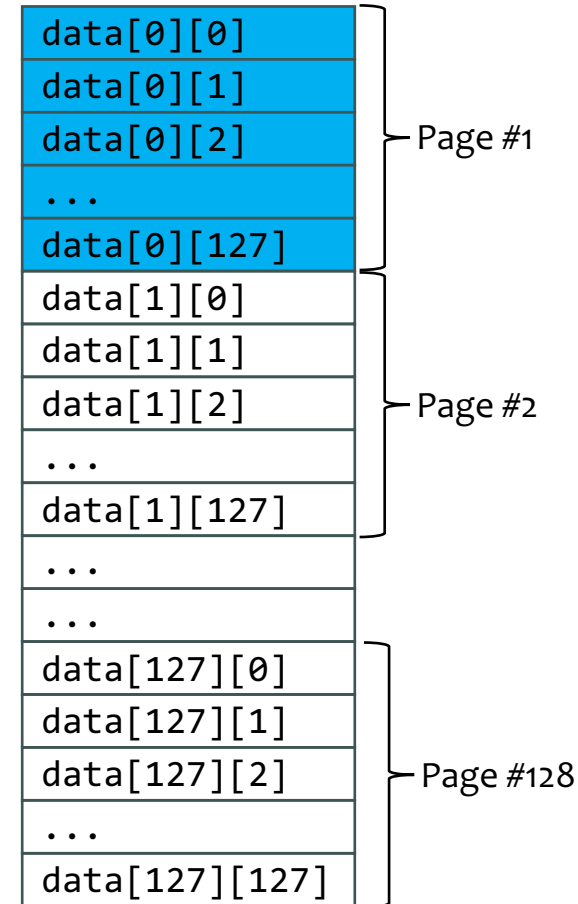
of Page
Faults:
 $128 \times 128 =$
16384

- An optimized (cache-friendly) version:

```
int i, j;
int [128][128] data;

for (i = 0; i < 128; i++) {
    for (j = 0; j < 128; j++) {
        data[i][j] = 0;
    }
}
```

of Page
Faults: **1**





Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;
int [128][128] data;

for (j = 0; j < 128; j++) {
    for (i = 0; i < 128; i++) {
        data[i][j] = 0;
    }
}
```

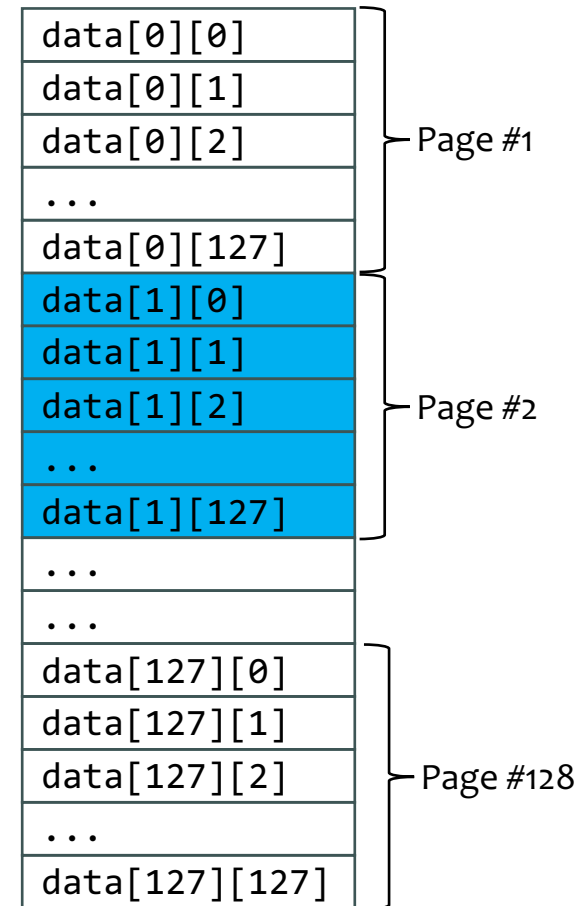
of Page
Faults:
 $128 \times 128 =$
16384

- An optimized (cache-friendly) version:

```
int i, j;
int [128][128] data;

for (i = 0; i < 128; i++) {
    for (j = 0; j < 128; j++) {
        data[i][j] = 0;
    }
}
```

of Page
Faults: **2**





■ Program Structure

- Suppose a page size of 128 words (512 Bytes), # of frames: 127.
- Consider the following program snippet:

```
int i, j;
int [128][128] data;

for (j = 0; j < 128; j++) {
    for (i = 0; i < 128; i++) {
        data[i][j] = 0;
    }
}
```

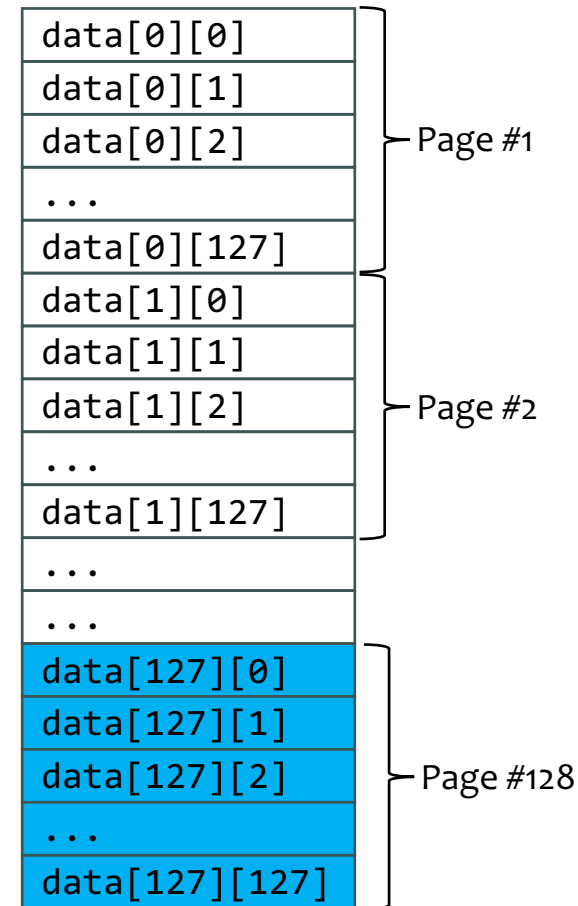
of Page
Faults:
 $128 \times 128 =$
16384

- An optimized (cache-friendly) version:

```
int i, j;
int [128][128] data;

for (i = 0; i < 128; i++) {
    for (j = 0; j < 128; j++) {
        data[i][j] = 0;
    }
}
```

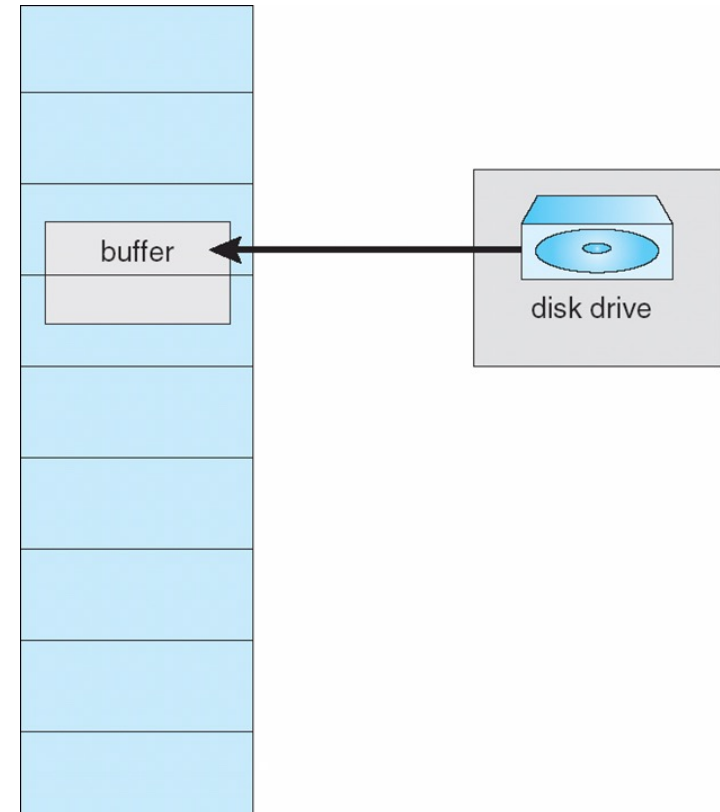
of Page
Faults: **128**





■ I/O Interlock and Page Locking

- Pages must sometimes be locked into memory.
- Consider I/O – Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm.
- **Pinning** of pages to lock into memory.





Thank you!