



DCS216 Operating Systems

Lecture 27 File Systems (3)

Jun 17th, 2024

Instructor: Xiaoxi Zhang
Sun Yat-sen University



■ Content

- File-System Structure
 - File-System Layers
 - On-Storage Structures
 - In-Memory Structures
 - Partitions and Mounting
 - Virtual File Systems
- The Very Simple File System (VSFS)
- Directory Implementation
- Allocation Methods
 - Contiguous Allocation
 - Linked Allocation
 - Indexed Allocation
- Free-Space Management
- Efficiency and Performance
- Recovery



■ Directory Implementation

- As mentioned earlier, a directory is simply a **logical** data structure that contains mappings of <filename, file_number>.
 - The directory itself is also a special type of **file**. It has an **inode** and the data **content** associated with the **inode**.
- But how is a directory actually implemented? In other words, how is the mappings of <filename, inode_num> stored in a directory?

```
$ ls -ali .
total 76
8567717 drwxrwxr-x  2 ubuntu ubuntu 4096 Jun  4 19:59 .
7236895 drwxrwxr-x 23 ubuntu ubuntu 4096 Jun  4 11:03 ..
8567720 -rw-rw-r--  1 ubuntu ubuntu   6 Jun  4 12:00 bar
8567875 -rw-rw-r--  1 ubuntu ubuntu 1781 Jun  4 19:17 file_io_syscall.c
8567901 -rw-rw-r--  1 ubuntu ubuntu   6 Jun  4 17:18 foo.txt
8567719 -rw-rw-r--  1 ubuntu ubuntu   80 Jun  4 11:33 hello.c
8567878 -rwxrwxr-x  1 ubuntu ubuntu 18600 Jun  4 14:13 lowio_creat
8567718 -rw-rw-r--  1 ubuntu ubuntu  257 Jun  4 14:11 lowio_creat.c
8567879 -rwxrwxr-x  1 ubuntu ubuntu 17632 Jun  4 16:25 lowio_lseek
8567721 -rw-rw-r--  1 ubuntu ubuntu  467 Jun  4 16:24 lowio_lseek.c
8567877 -rw-rw-r--  1 ubuntu ubuntu  153 Jun  4 14:11 Makefile
```

inode

file_name



■ Linear List

- The simplest method of implementing a directory is to use a **linear list** of **file names** with **pointers** (**inode number**) to the data blocks.
- For example

reclen	inum	strlen	name
12	5	2	.
12	2	3	..
12	19	4	foo
12	255	4	bar
36	258	28	foobar_is_a_pretty_longname



■ Linear List

- The simplest method of implementing a directory is to use a **linear list** of **file names** with **pointers** (**inode number**) to the data blocks.
- For example

reclen	inum	strlen	name
12	5	2	.
12	2	3	..
12	19	4	foo
12	255	4	bar
36	258	28	foobar_is_a_pretty_longname

12 5 0 2 . \0 ? ? ? ? ?

12 2 0 3 . . \0 ? ? ? ? ?

12 19 0 4 f o o \0 ? ? ? ?

12 FF 0 4 b a r \0 ? ? ? ?

36 2 1 28 f o o b a r _ i

s _ a _ p r e t t y _ l

o n g n a m e \0 ? ? ? ?



Linear List

- The simplest method of implementing a directory is to use a **linear list** of **file names** with **pointers** (**inode number**) to the data blocks.
- For example

reclen	inum	strlen	name
12	5	2	.
12	2	3	..
12	19	4	foo
12	255	4	bar
36	258	28	foobar_is_a_pretty_longname

12 5 0 2 . \0 ? ? ? ? ?

12 2 0 3 . . \0 ? ? ? ? ?

12 19 0 4 f o o \0 ? ? ? ? ?

12 FF 0 4 b a r \0 ? ? ? ? ?

36 2 1 28 f o o b a r _ i

s _ a _ p r e t t y _ l

o n g n a m e \0 ? ? ? ?

For simplicity, default record length is **12 bytes** for normal files (file name length less than **8**); the max **inum** is **65535** (addressable with 2 bytes); variable length filename is supported by specifying larger **reclen**.



■ Linear List

- The simplest method of implementing a directory is to use a **linear list** of **file names** with **pointers** (**inode number**) to the data blocks.
- **Advantages:** Simple to implement.



■ Linear List

- The simplest method of implementing a directory is to use a **linear list** of **file names** with **pointers** (**inode number**) to the data blocks.
- **Advantages:** Simple to implement.
- To **create** a new file in the directory:
 - **Search** the directory to be sure that no existing file has the same name.
 - Add a new entry <**filename**, **inode_num**> at the end of the directory.
- To **delete** an existing file from the directory:
 - **Search** the directory for the specific file.
 - Release the space allocated to that file.



■ Linear List

- The simplest method of implementing a directory is to use a **linear list** of **file names** with **pointers** (**inode number**) to the data blocks.
- **Advantages:** Simple to implement.
- **Disadvantages:** **Time-consuming** to operate (**linear search time**)
 - For example, in a directory with many many number of files (e.g., **60,000** files), the system will have to **iterate through** all of the files in order to locate the file that we wish to **open()**.

reclen	inum	strlen	name
12	5	2	.
12	2	3	..
12	19	4	foo
12	255	4	bar
36	258	28	foobar_is_a_pretty_longname
...			
12	65535	7	lastone



■ Linear List

- The simplest method of implementing a directory is to use a **linear list** of **file names** with **pointers** (**inode number**) to the data blocks.
- **Advantages:** Simple to implement.
- **Disadvantages:** **Time-consuming** to operate (**linear search time**)
 - **Workaround:** we could keep the **linear list ordered** (e.g., **sorted** by filename), to reduce the **search** time from $O(N)$ to $O(\log N)$ with **Binary Search**. However, the **insertion** time is increased from $O(1)$ to $O(N)$.

reclen	inum	strlen	name
12	5	2	.
12	2	3	..
12	255	4	bar
12	19	4	foo
36	258	28	foobar_is_a_pretty_longname
12	65535	7	lastone
...			



■ Hash Table

- Linear list with **hash** data structure.
 - The hash table takes a value computed from the filename and returns a pointer to the filename in the linear list.
- Decrease directory search time to $O(1)$.



■ Hash Table

- Linear list with **hash** data structure.

filename

.
..
foo
bar
foobar_is_a_pretty_longname
lastone

key

46
46 + 46 = 92
'f' + 'o' + 'o' = 324
'b' + 'a' + 'r' = 309
2859
758

Hash function: $h(\text{key}) = \text{key} \% 16$

	key	filename
00		
01		
02		
03		
04	324	foo
05	309	bar
06	758	lastone
07		
08		
09		
10		
11	2859	foobar_is_a_pretty_longname
12	92	..
13		
14	46	.
15		



■ Hash Table

- Linear list with **hash** data structure.
 - The hash table takes a value computed from the filename and returns a pointer to the filename in the linear list.
- Decrease directory search time to $O(1)$.
- **Problems:**
 - **Collisions** – situations where two filenames hash to the same location.
 - Hash tables are generally fixed size and the hash functions depends on that size.
 - Only good if entries are fixed size, or use chained-overflow method.





■ Allocation Methods

- An allocation method refers to how **disk blocks** are **allocated** for files.
 - **Contiguous** Allocation (**连续**分配)
 - **Linked** Allocation (**链接**分配)
 - **Indexed** Allocation (**索引**分配)



■ Contiguous Allocation

- Each file occupies a set of **contiguous** blocks on the disk.
 - **Simple**: only starting location (block #) and length (# of blocks) required.
 - Enables **random access**.
 - E.g., can access the 3rd block of **File E** directly via $(\&E + 2)$.
 - Best **performance** in most cases.







■ Contiguous Allocation

■ Problems with **contiguous allocation**:

- Finding **space** on the disk for a new file.
- Knowing **file size**.
- External **Fragmentation**
 - need for **compaction** off-line (**downtime**) or on-line (**file corruption**).





■ Contiguous Allocation

■ Mapping from **Logical Address** to **Physical Address**

- Assume block size == 512 bytes

- $LA / 512 = Q$


- $LA \% 512 = R$

- Block number to be accessed:

$$\text{starting_addr} + Q$$

- Displacement (**offset**) into block:

R




directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2




■ Contiguous Allocation

- The need for **compaction (碎片合并/压缩)**.
 - As files are allocated and deleted, **external fragmentation** is inevitable.
 - **Holes** (fragments) between allocated files may be unusable, leading to significant **loss of storage space**.
 - One strategy to prevent or reduce such loss is to regularly **compact** all free spaces into one contiguous space.



File allocation table		
File name	Start block	Length
File A	2	3
File B	9	5
File C	18	8
File D	30	2
File E	26	3



File allocation table		
File name	Start block	Length
File A	0	3
File B	3	5
File C	8	8
File D	19	2
File E	16	3


Figure 12.9 Contiguous File Allocation

Figure 12.10 Contiguous File Allocation (After Compaction)




Contiguous Allocation

- The need for **compaction (碎片合并/压缩)**.
 - Compaction can be done **off-line** or **on-line**.
 - off-line (停机期间)**: the file system must be **unmounted (卸载)**.
 - on-line (在线)**: special care must be taken when **compacting (defragmentation)** on-line, otherwise **data corruption** may occur.



File allocation table		
File name	Start block	Length
File A	2	3
File B	9	5
File C	18	8
File D	30	2
File E	26	3



File allocation table		
File name	Start block	Length
File A	0	3
File B	3	5
File C	8	8
File D	19	2
File E	16	3

Figure 12.9 Contiguous File Allocation

Figure 12.10 Contiguous File Allocation (After Compaction)



■ Contiguous Allocation

■ Extent-Based File Systems

- Many newer file systems (SPARC Veritas File System, Linux 2.6.19 ext4, etc.) use a modified version of contiguous allocation scheme.
- In an extent-based file system, a contiguous chunk of space is allocated **initially** for file allocation.
- Then, if that amount proves not to be large enough, another chunk of contiguous space, known as an **extent**, is added.
 - The location of a file's blocks is then recorded as a **location** and a **block count**, **plus** a **link** to the first block of the next **extent**.
 - A file may consist of **one or more extents**.



■ Linked Allocation


- **Linked Allocation** (also called **chained allocation**) solves all problems of contiguous allocation.
 - Each file is a **linked list** of storage blocks.
 - These blocks may be **scattered anywhere** on the disk.
 - Each block contains a **pointer** to the **next block**.
 - **End-of-File** could be achieved by:
 - Ending with a **NULL** pointer.



■ Linked Allocation

■ **Linked Allocation** (also called **chained allocation**) solves all problems of contiguous allocation.

- Each file is a **linked list** of storage blocks.
- These blocks may be **scattered anywhere** on the disk.
- Each block contains a **pointer** to the **next block**.
- **End-of-File** could be achieved by:
 - Ending with a **NULL** pointer.
 - Explicitly specify an **END** block.






■ Linked Allocation

- **Linked Allocation** (also called **chained allocation**) solves all problems of contiguous allocation.
 - Each file is a **linked list** of storage blocks.
 - These blocks may be **scattered anywhere** on the disk.
 - Each block contains a **pointer** to the **next block**.
 - **End-of-File** could be achieved by:
 - Ending with a **NULL** pointer.
 - Explicitly specify an **END** block.
 - or specify the **length** of the file.

File allocation table

File name	Start block	Length
• • • File B • • •	• • • 1 • • •	• • • 5 • • •





■ Linked Allocation

■ Advantages:

■ Simple:

- To **create** a new file, simply create a new entry in the directory. Each entry has a pointer to the first block of the file.
 - Initially, an **empty** file \Rightarrow **NULL** pointer in the directory entry.
- To **read** a file, simply **follow** the linked list *from block to block*.
- To **write** (**append**) to the file \Rightarrow new block allocated and append to the end of the file. (last **pointer points** to the new block).
 - The size of a file need not be declared when first created; a file can continue to **grow** as long as free blocks are available.
- No **external fragmentation** \Rightarrow No need for **compaction**.
 - Free space management: no waste of storage space.



■ Linked Allocation

■ Disadvantages:


- **Sequential access** only. No **random access** within a file.
 - To locate the i^{th} block of a file, we must start from the beginning of that file and follow the pointers until we get to the i^{th} block.



■ Linked Allocation

■ Disadvantages:

- **Sequential access** only. No **random access** within a file.
 - To locate the i^{th} block of a file, we must start from the beginning of that file and follow the pointers until we get to the i^{th} block.
- **Difficult** to make use of **Locality** (局部性).
 - The file may be **scattered around** the disk, i.e., the next block (that the pointer points to) may be located far away from the current block, making it difficult to leverage **Locality** (very **low cache hit ratio**).

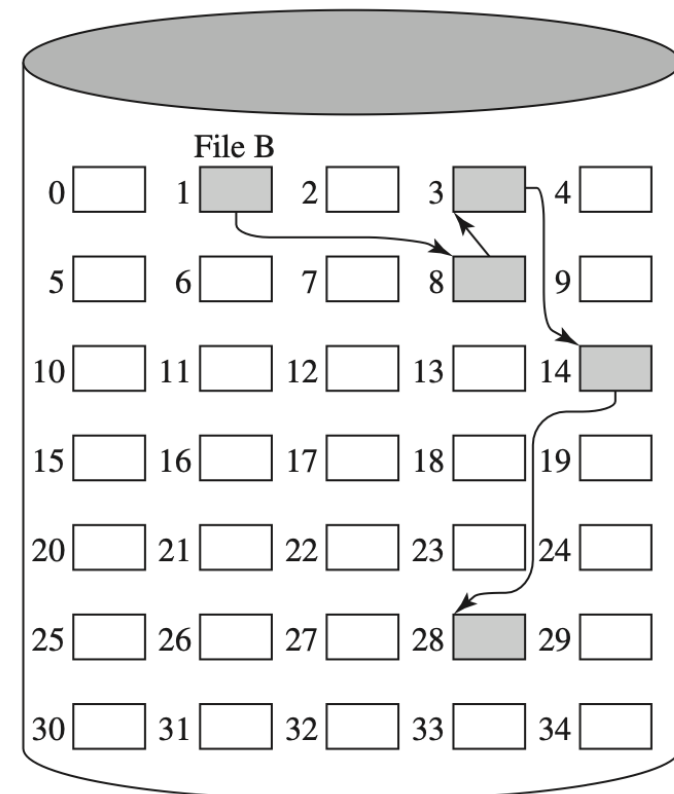




■ Linked Allocation

■ Disadvantages:

- **Sequential access** only. No **random access** within a file.
 - To locate the i^{th} block of a file, we must start from the beginning of that file and follow the pointers until we get to the i^{th} block.
- **Difficult** to make use of **Locality** (局部性).
 - The file may be **scattered around** the disk, i.e., the next block (that the pointer points to) may be located far away from the current block, making it difficult to leverage **Locality** (very **low cache hit ratio**).
- Storage Space **Utilization**.
 - extra space required for the pointers.
 - If a pointer requires **4 bytes** out of a **512-byte** block, the **0.78%** of disk is used for pointers (**overhead**).






■ File Allocation Table (FAT)

- An important variation of Linked Allocation is the use of a **File Allocation Table (FAT, 文件分配表)**.
 - a combination of **continuous** and **linked** allocation (*most of the time*).
- A section of storage at the beginning of the volume is set aside to contain the **table**.
 - The table has **one entry for each block** and is **indexed by block number**.
 - Each entry contains a pointer to the **next block** in the file.
 - a **NULL** pointer (**0**) indicates End-Of-File.
 - an **implicit linked list** via indexed table.



■ File Allocation Table (FAT)

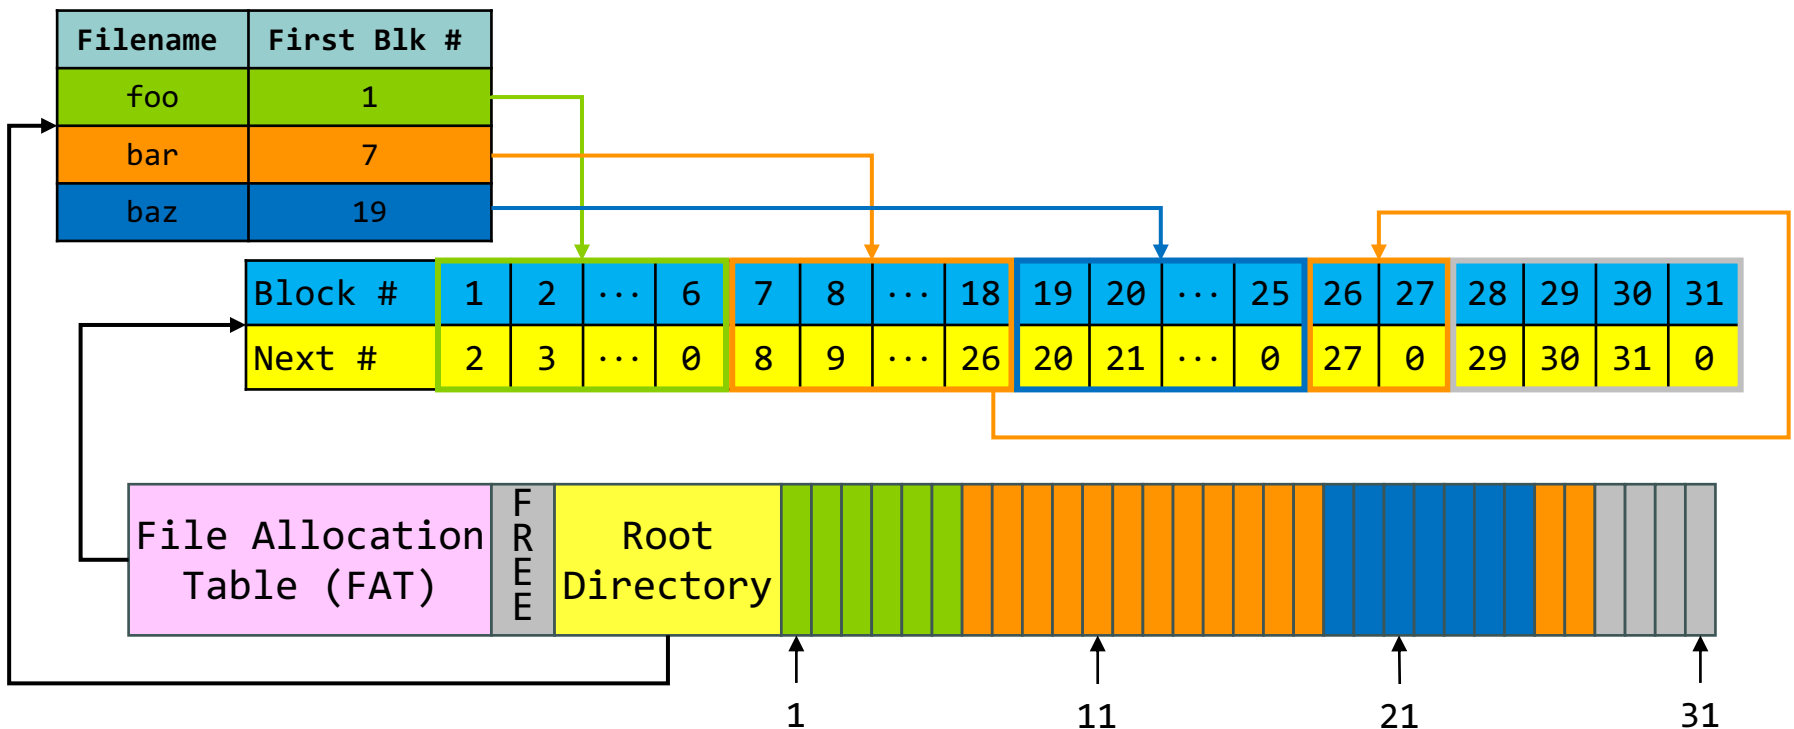
- An important variation of Linked Allocation is the use of a **File Allocation Table (FAT, 文件分配表)**.
 - a combination of **continuous** and **linked** allocation (*most of the time*).
- A section of storage at the beginning of the volume is set aside to contain the **table**.
 - The table has **one entry for each block** and is **indexed by block number**.
 - Each entry contains a pointer to the **next block** in the file.
 - a **NULL** pointer (**0**) indicates End-Of-File.
 - an **implicit linked list** via indexed table.





■ File Allocation Table (FAT)

- A section of storage at the beginning of the volume is set aside to contain the **table**.
 - The table has **one entry for each block** and is **indexed by block number**.
 - Each entry contains a pointer to the **next block** in the file.
 - a **NULL** pointer (**0**) indicates End-Of-File.
 - an **implicit linked list** via indexed table.



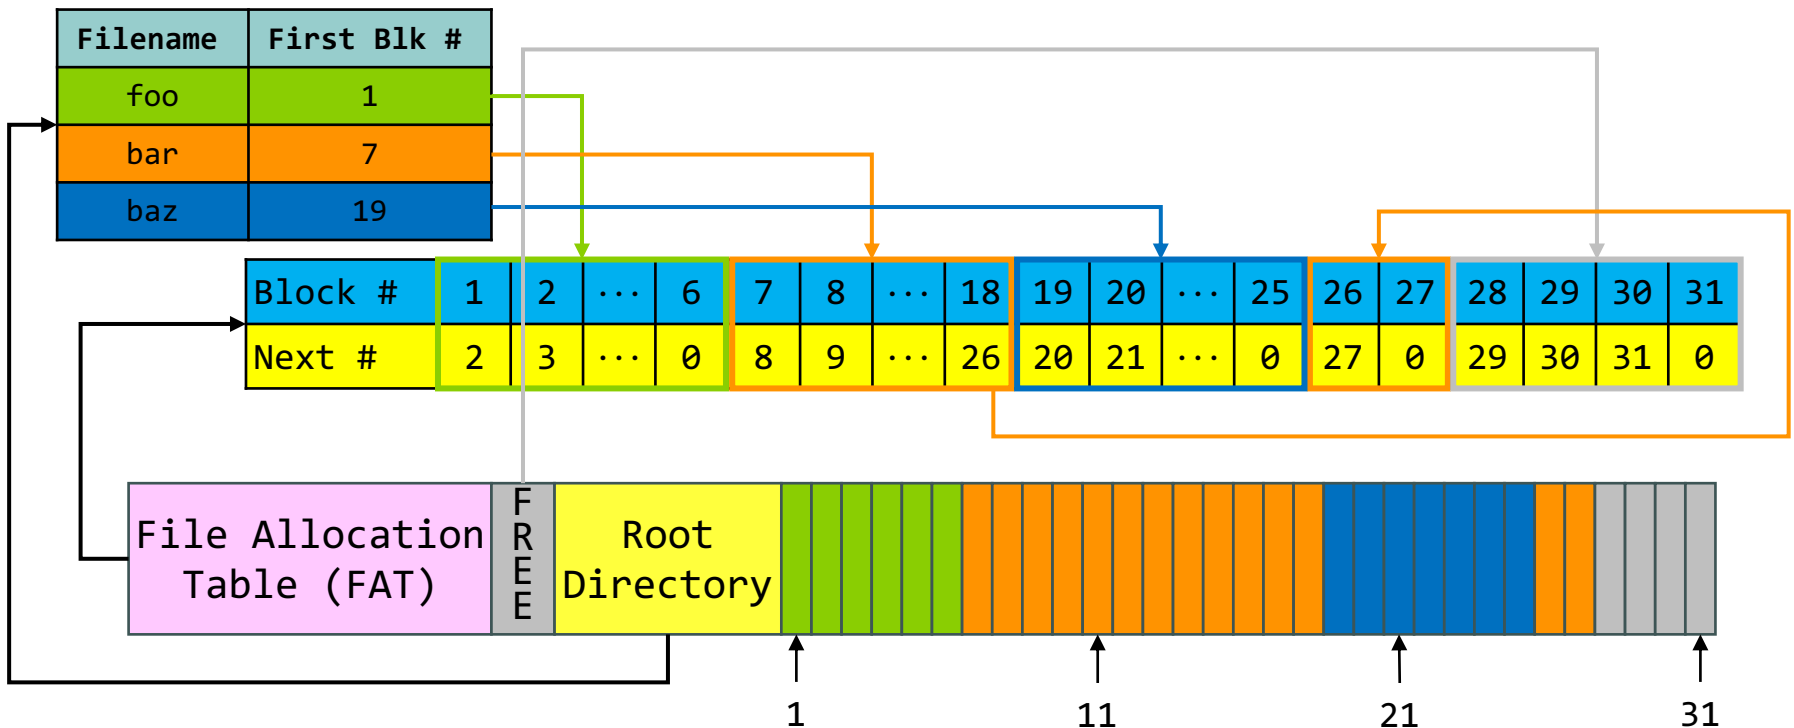


■ File Allocation Table (FAT)

■ **FAT** is basically an *implicit linked list*.

■ Instead of **allocate** (at the end of each block) a **pointer** to the next block, we delegate the linked list structure to the **FAT**.

- The **FAT** can then be **cached** in the **main memory**, which **speeds up** walking through the **linked list** structure of a file.
 - enables **random-access**.





■ FAT32 Layout

```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0 ←
mkfs.fat 4.2 (2021-01-31)
```

Format the disk, "-F32" means FAT32





■ FAT32 Layout

```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0
mkfs.fat 4.2 (2021-01-31)
$ sudo dosfsck -v /dev/ram0
```

Read the info stored in the boot sector.






■ FAT32 Layout

```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0
mkfs.fat 4.2 (2021-01-31)
$ sudo dosfsck -v /dev/ram0
...
Boot sector contents:
System ID "mkfs.fat"
Media byte 0xf8 (hard disk)
    512 bytes per logical sector
    4096 bytes per cluster
    32 reserved sectors
First FAT starts at byte 16384 (sector 32)
    2 FATs, 32 bit entries
    512000 bytes per FAT (= 1000 sectors)
Root directory start at cluster 2 (arbitrary size)
Data area starts at byte 1040384 (sector 2032)
    127738 data clusters (523214848 bytes)
...
```

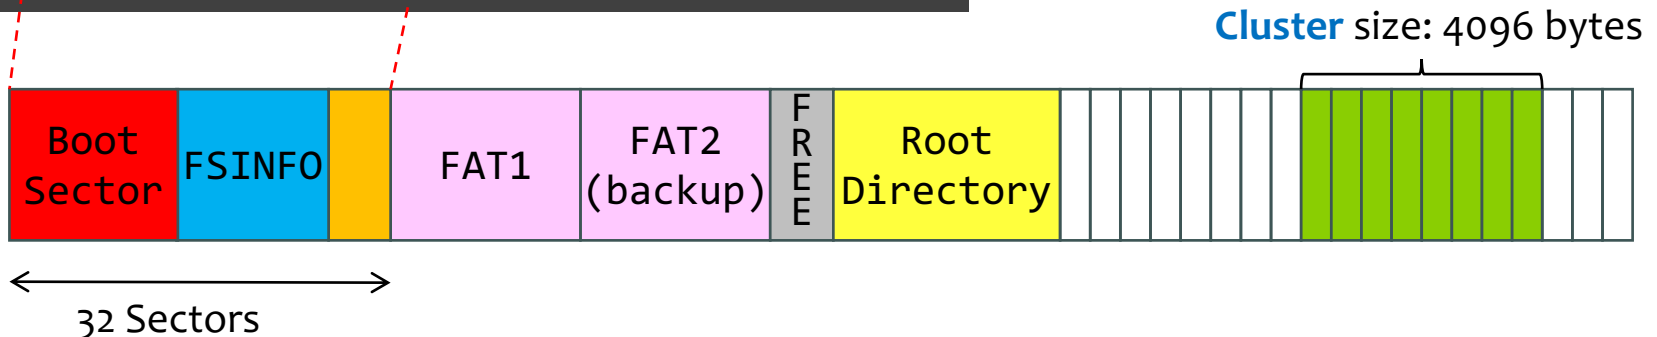
A **cluster (block)** is made of **8 sectors**.





FAT32 Layout

```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0
mkfs.fat 4.2 (2021-01-31)
$ sudo dosfsck -v /dev/ram0
...
Boot sector contents:
System ID "mkfs.fat"
Media byte 0xf8 (hard disk)
    512 bytes per logical sector
    4096 bytes per cluster
    32 reserved sectors
First FAT starts at byte 16384 (sector 32)
    2 FATs, 32 bit entries
    512000 bytes per FAT (= 1000 sectors)
Root directory start at cluster 2 (arbitrary size)
Data area starts at byte 1040384 (sector 2032)
    127738 data clusters (523214848 bytes)
...
```






FAT32 Layout

```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0
mkfs.fat 4.2 (2021-01-31)
$ sudo dosfsck -v /dev/ram0
...
Boot sector contents:
System ID "mkfs.fat"
Media byte 0xf8 (hard disk)
    512 bytes per logical sector
    4096 bytes per cluster
    32 reserved sectors
First FAT starts at byte 16384 (sector 32)
    2 FATs, 32 bit entries
    512000 bytes per FAT (= 1000 sectors)
Root directory start at cluster 2 (arbitrary size)
Data area starts at byte 1040384 (sector 2032)
    127738 data clusters (523214848 bytes)
...
```

2 FATs:

- FAT1 is the primary.
- FAT2 serves as the backup.

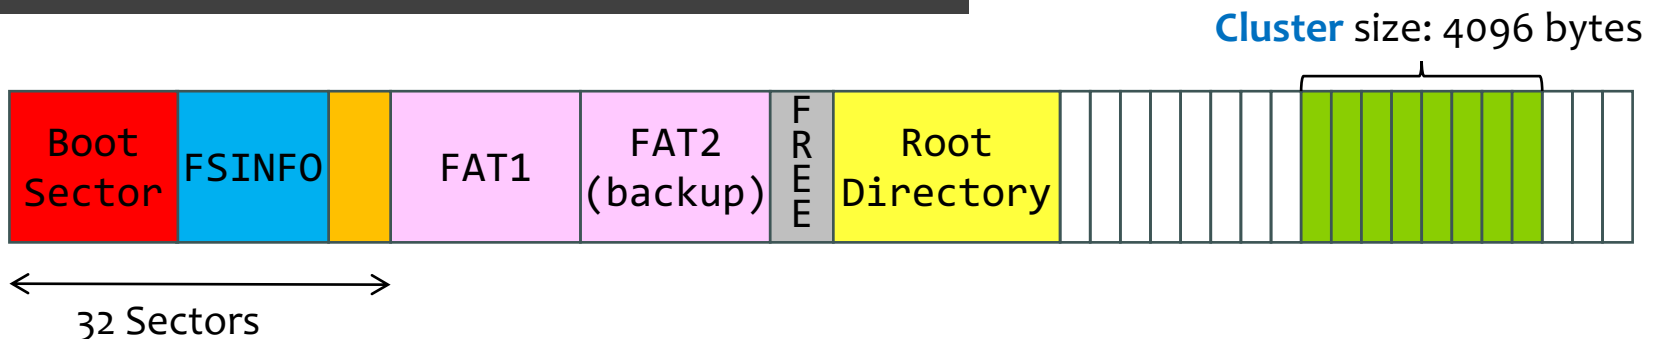




FAT32 Layout

```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0
mkfs.fat 4.2 (2021-01-31)
$ sudo dosfsck -v /dev/ram0
...
Boot sector contents:
System ID "mkfs.fat"
Media byte 0xf8 (hard disk)
    512 bytes per logical sector
    4096 bytes per cluster
    32 reserved sectors
First FAT starts at byte 16384 (sector 32)
    2 FATs, 32 bit entries
    512000 bytes per FAT (= 1000 sectors)
Root directory start at cluster 2 (arbitrary size)
Data area starts at byte 1040384 (sector 2032)
    127738 data clusters (523214848 bytes)
...
```

Each FAT is of size 512000 bytes, with a 4-byte (32-bit, hence FAT32) address, it can index up to $\frac{512000}{4} = 128000$ entries (clusters), just enough for the **127738** data clusters in the Data Area.

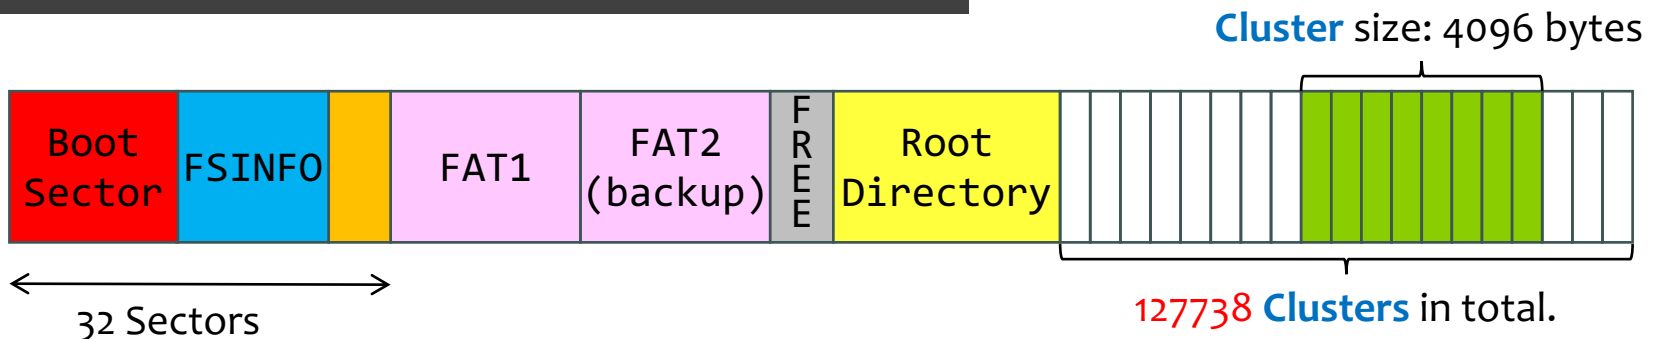




FAT32 Layout

```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0
mkfs.fat 4.2 (2021-01-31)
$ sudo dosfsck -v /dev/ram0
...
Boot sector contents:
System ID "mkfs.fat"
Media byte 0xf8 (hard disk)
    512 bytes per logical sector
    4096 bytes per cluster
    32 reserved sectors
First FAT starts at byte 16384 (sector 32)
    2 FATs, 32 bit entries
    512000 bytes per FAT (= 1000 sectors)
Root directory start at cluster 2 (arbitrary size)
Data area starts at byte 1040384 (sector 2032)
127738 data clusters (523214848 bytes)
...
```

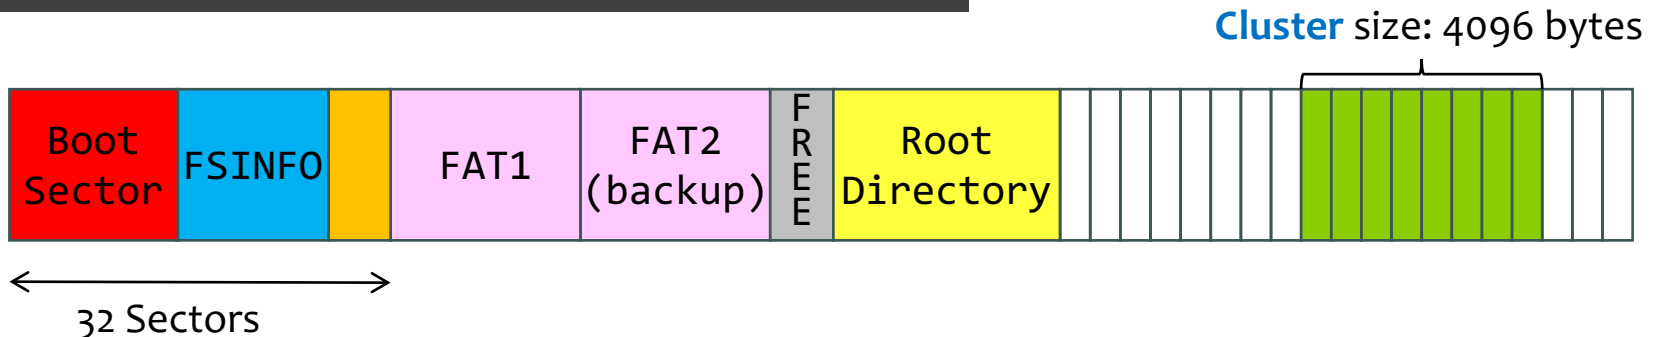
Each FAT is of size 512000 bytes, with a 4-byte (32-bit, hence FAT32) address, it can index up to $\frac{512000}{4} = 128000$ entries (clusters), just enough for the 127738 data clusters in the Data Area.





■ FAT32 Layout


```
$ sudo modprobe brd rd_nr=1 rd_size=512000
$ sudo mkfs.vfat -F32 /dev/ram0
mkfs.fat 4.2 (2021-01-31)
$ sudo dosfsck -v /dev/ram0
...
Boot sector contents:
System ID "mkfs.fat"
Media byte 0xf8 (hard disk)
    512 bytes per logical sector
    4096 bytes per cluster
    32 reserved sectors
First FAT starts at byte 16384 (sector 32)
    2 FATs, 32 bit entries
    512000 bytes per FAT (= 1000 sectors)
Root directory start at cluster 2 (arbitrary size)
Data area starts at byte 1040384 (sector 2032)
    127738 data clusters (523214848 bytes)
...
```





Indexed Allocation

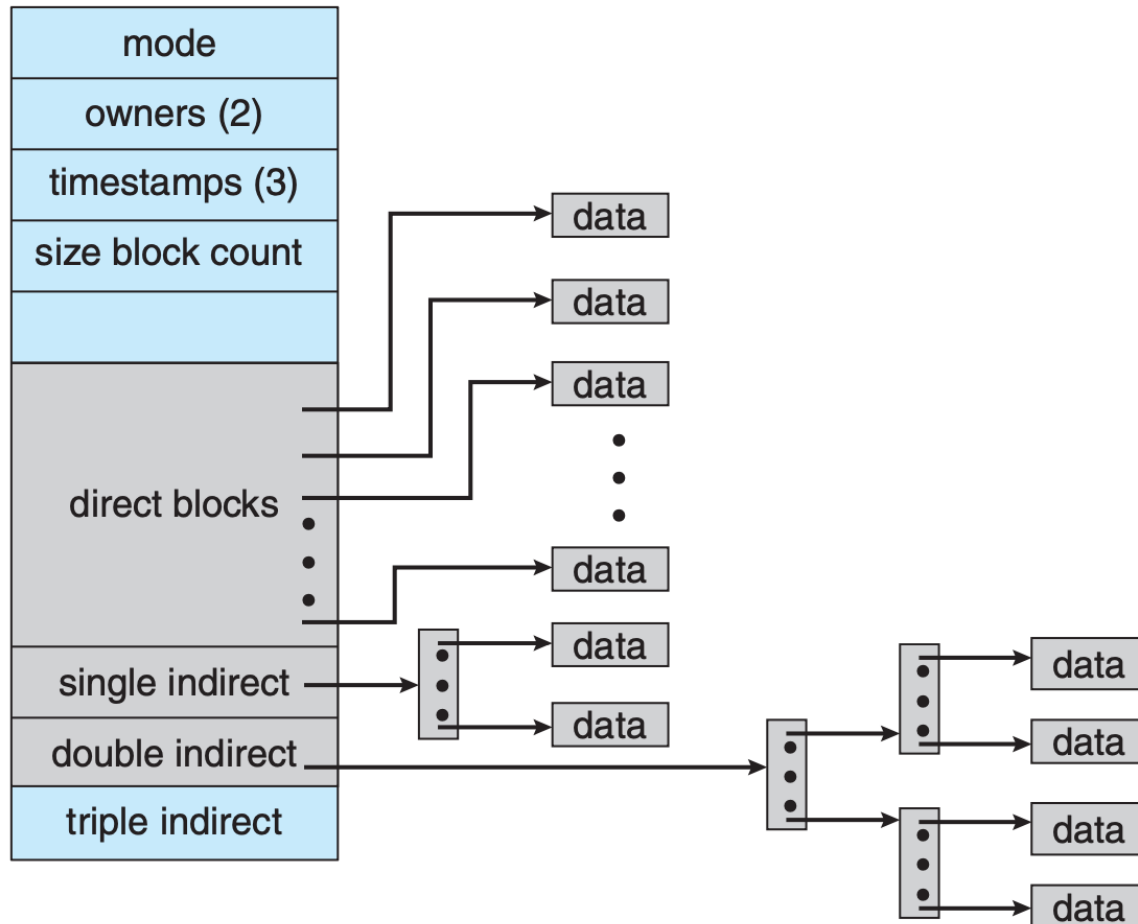
- As mentioned earlier in **VSFS (Very Simple File System)**, the UNIX File System, and ext2/ext3/ext4 in Linux, etc.





■ Indexed Allocation

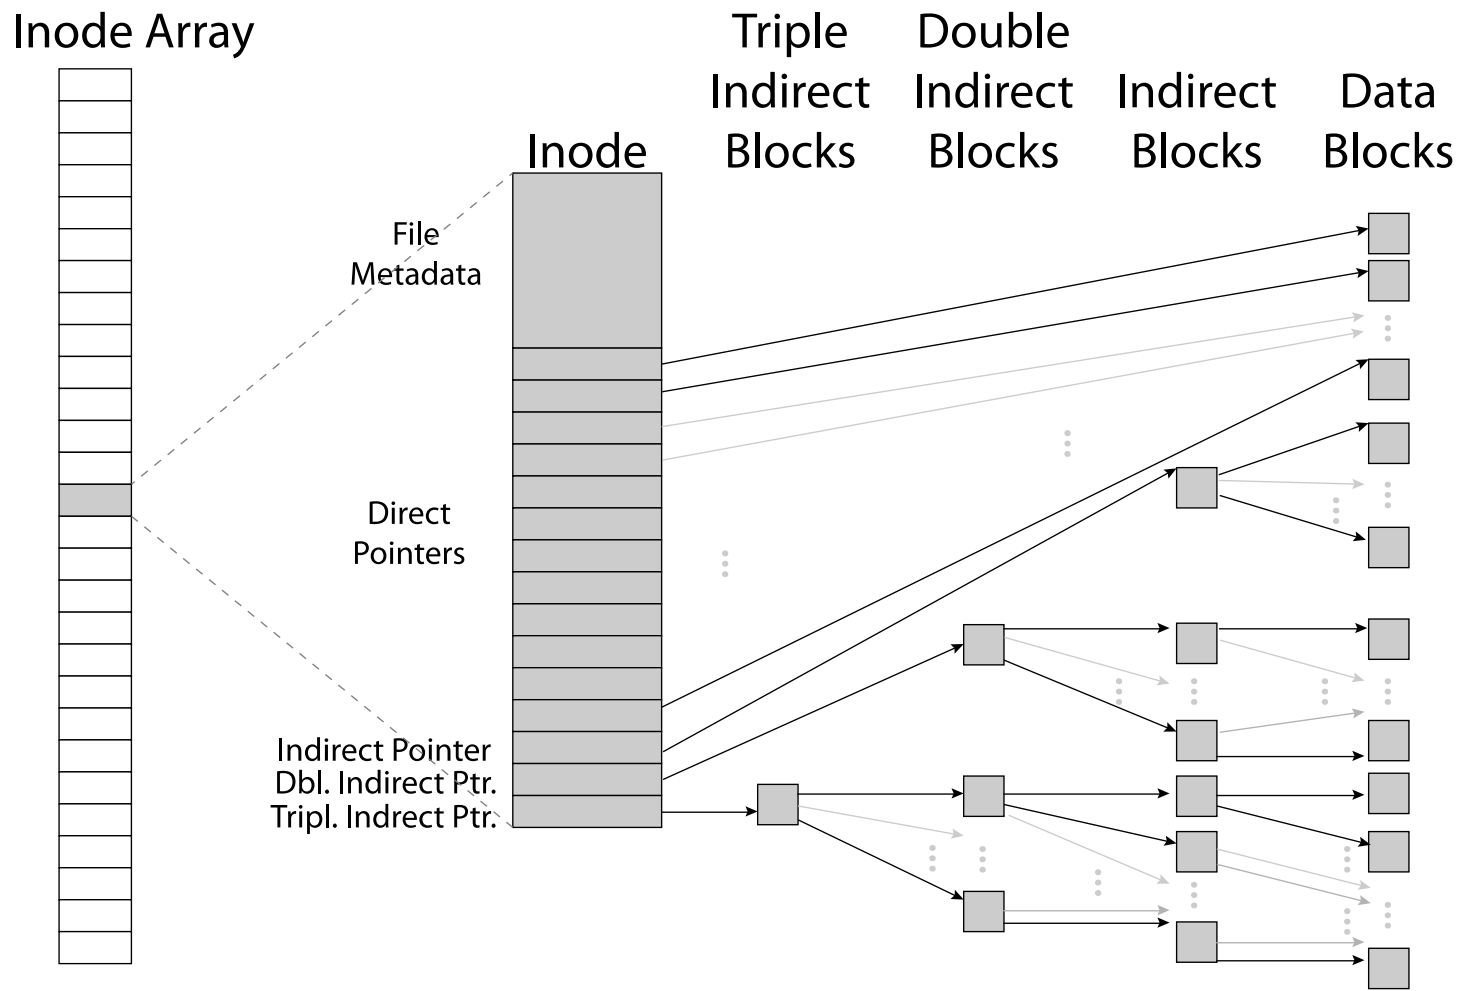
- As mentioned earlier in **VSFS (Very Simple File System)**, the UNIX File System, and ext2/ext3/ext4 in Linux, etc.





■ Indexed Allocation

- As mentioned earlier in **VSFS (Very Simple File System)**, the UNIX File System, and ext2/ext3/ext4 in Linux, etc.





■ Comparison

- Comparison of different Allocation Methods.

	Contiguous	Linked	Indexed	
Pre-allocation?	Necessary	Possible	Possible	
Fixed or variable size portions?	Variable	Fixed blocks	Fixed blocks	Variable
Portion size	Large	Small	Small	Medium
Allocation frequency	Once	Low to high	High	Low
Time to allocate	Medium	Long	Short	Medium
File allocation table size	One entry	One entry	Large	Medium



■ Free Space Management

- The system maintains a **free-space list** to keep track of free disk space, or available blocks/clusters on disk.
- There are many different schemes to implement a **free-space list**:
 - **Bit Vector** (Bitmap)
 - **Linked List**
 - **Grouping**
 - **Counting**
 - **Space Maps**



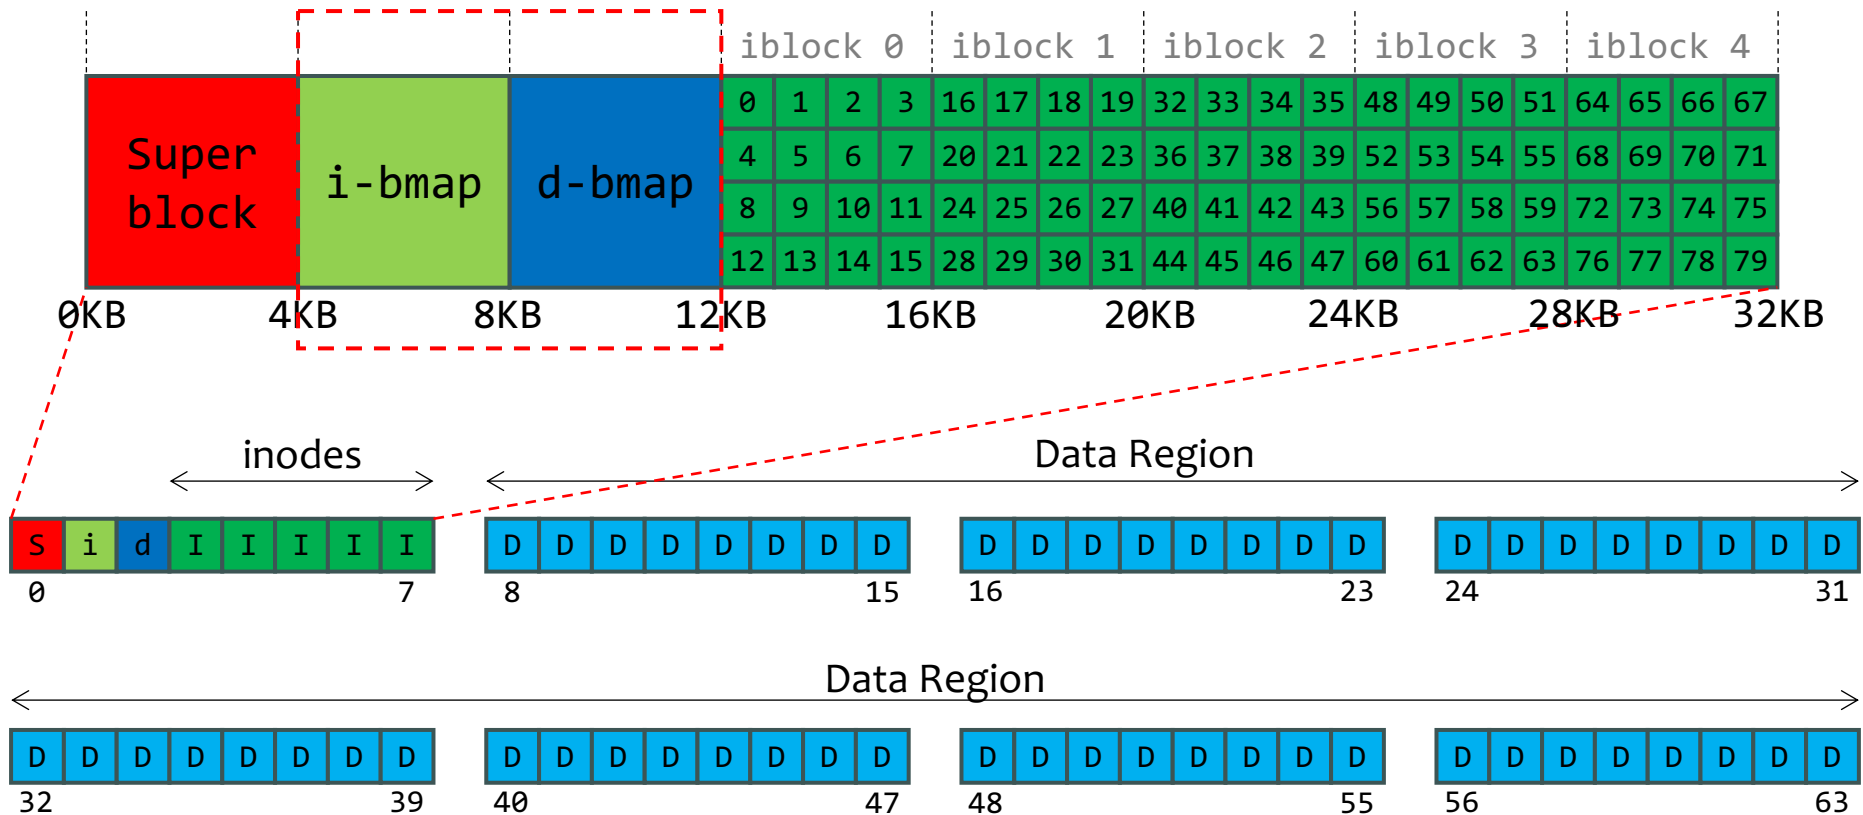
■ Bit Vector

- Most frequently, the **free-space list** is implemented as a **bitmap** or bit vector.
 - Each block is represented by one bit.
 - If the block is **free**, the bit is set to **1**.
 - If the block is **allocated**, the bit is set to **0**.
 - Bitmap requires extra space. For example, assume:
 - block size = **4KB** = 2^{12} bytes
 - disk size = **1TB** = 2^{40} bytes
 - # of blocks: $2^{40} / 2^{12} = 2^{28}$
 - Then, space requirement for bitmap: 2^{28} bits = **32MB**



■ Bit Vector

■ Example: **VSFS** or **ext2/ext3/ext4**





■ Linked List


- Another approach is to link together all the free blocks, keeping a pointer to the first free block in a special location in the file system and caching it in memory.
- For example: blocks 2, 3, 4, 5, 8, 9, 10, 11, 12, 13, 17, 18, 25, 26 and 27 are free. The number of the first free block (#2) is stored in a reserved location.

■ Disadvantages:

- Cannot get contiguous space easily (must traverse the list of free blocks).

■ Advantages:



- No waste of disk space
- No need to traverse the entire list (if # of free blocks is also recorded). **How?**





■ Grouping


- A modification of the **linked-list** approach.
- It stores the addresses of n free blocks in the first free block.
 - The first $n - 1$ of these blocks are actually free.
 - The last block contains the address of another n free blocks, and so on.
- These addresses of a large number of free blocks can now be found quickly, unlike the situation when the standard **linked-list** approach is used.
- **Example:** Suppose $n = 6$.





■ Counting

- A modification of the **Bitmap** approach.
 - space is frequently contiguously allocated and freed.
 - Each entry contains the **address** of the next free block and the **number of free blocks** that contiguously follows it.
- With simple **bitmap**, the bit vector table would be:
00111100111111000110000001110000
- With **counting**:
 $\langle 2, 4 \rangle | \langle 8, 6 \rangle | \langle 17, 2 \rangle | \langle 25, 3 \rangle$
- These entries can be stored in a **balanced tree**, rather than a **linear table** or **linked list**, for efficient **lookup, insertion and deletion**.





■ Space Maps

- ZFS (Zettabyte File System, Sun Solaris 10) was designed to encompass huge numbers of files, directories and even file systems.
 - On these scales, metadata I/O can have a huge performance impact.
 - Full data structures like bitmaps couldn't fit thousands of I/Os.
- Divide device space into **metaslab** units and manage them.
 - Given volume can contain hundreds of metaslabs.
 - Each metaslab has an associated space map.
 - Uses counting algorithm.
 - But records to log file rather than file system.
 - Log of all block activity, in time order, in counting format.
- Metaslab activity is to load space map into memory in balanced-tree structure, indexed by offset.
 - Replay log into that structure.
 - Combine contiguous free blocks into single entry.



■ Efficiency and Performance

■ Efficiency depends on

- Disk allocation and directory algorithms.
- Types of data kept in file's directory entry
- Pre-allocation or as-needed allocation of metadata structures
- Fixed-size or varying-size data structures.

■ Performance

- Disk cache
 - separate section/segment of main memory for frequently used blocks.
- Free-behind and read-ahead
 - techniques to optimize [sequential access](#).
- Improve performance by dedicating section/segment of memory as virtual disk or RAM disk.



■ Efficiency and Performance

■ Page Cache

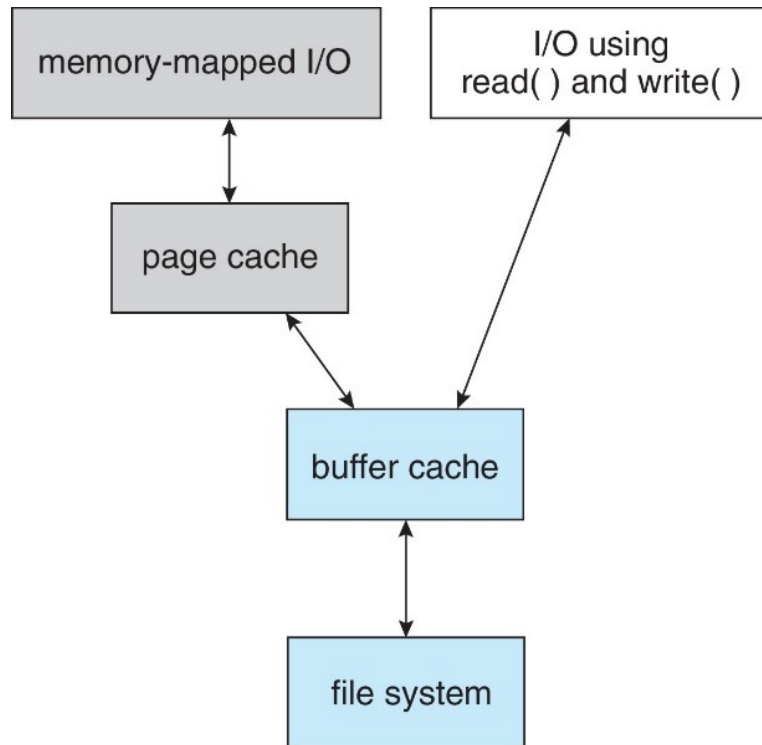
- A **page cache** caches pages rather than disk blocks using virtual memory techniques.
 - Memory-mapped I/O uses a **page cache**.
 - Routine I/O through the file system uses the buffer (disk)-cache.

■ Unified Buffer Cache (UBC)

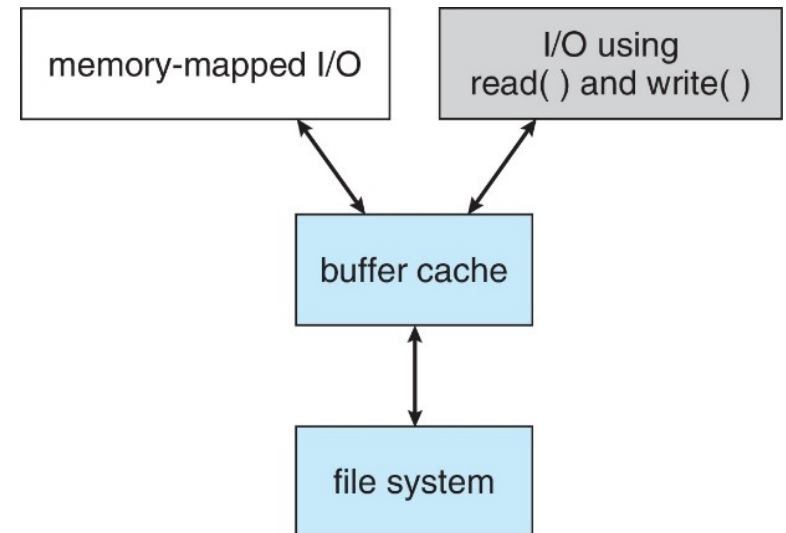
- A unified buffer cache uses the same page cache to cache both memory-mapped pages and ordinary file system I/O to avoid **double caching**.

■ Efficiency and Performance

■ Unified Buffer Cache



I/O without a unified buffer cache



I/O using a unified buffer cache



Thank you!