



DCS216 Operating Systems

Lecture 28 Virtualization

Jun 19th, 2024

Instructor: Xiaoxi Zhang
Sun Yat-sen University




■ Content

- Overview
- History
- Benefits and Features
- Building Blocks
- Types of VMs and their Implementation
- Virtualization and OS Components
- Examples


■ Overview

- The **Operating System** is all about **Abstractions**.
 - **Files** is an **abstraction** of I/O devices.
 - **Virtual Memory** is an **abstraction** of main memory and I/O devices.
 - **Processes** is an **abstraction** of CPU + Memory + I/O devices.





■ Overview

- The **Operating System** is all about **Abstractions** and **Virtualization**.
 - **Files** is an **abstraction** of I/O devices.
 - by **virtualizing** raw blocks into logical blocks (with File Systems, Device Drivers, etc.)
 - **Virtual Memory** is an **abstraction** of main memory and I/O devices.
 - by **virtualizing** physical addresses into logical addresses (via Page Tables, TLB, DMA, etc.)
 - **Processes** is an **abstraction** of CPU + Memory + I/O devices.
 - by **virtualizing** CPU states and related resources using PCBs.




■ Overview

- The **Operating System** is the collection of **software** that **manages** the underlying **hardware** and **provides services** (system call APIs) to *application programs*.



■ Overview

- The **Operating System** is the collection of **software** that **manages** the underlying **hardware** and **provides services** (system call APIs) to *application programs*.





The "**Hardware**" may not necessarily be **REAL** (in terms of *physical form*).




■ Overview

- The **Operating System** is the collection of **software** that **manages** the underlying **hardware** and **provides services** (system call APIs) to *application programs*.



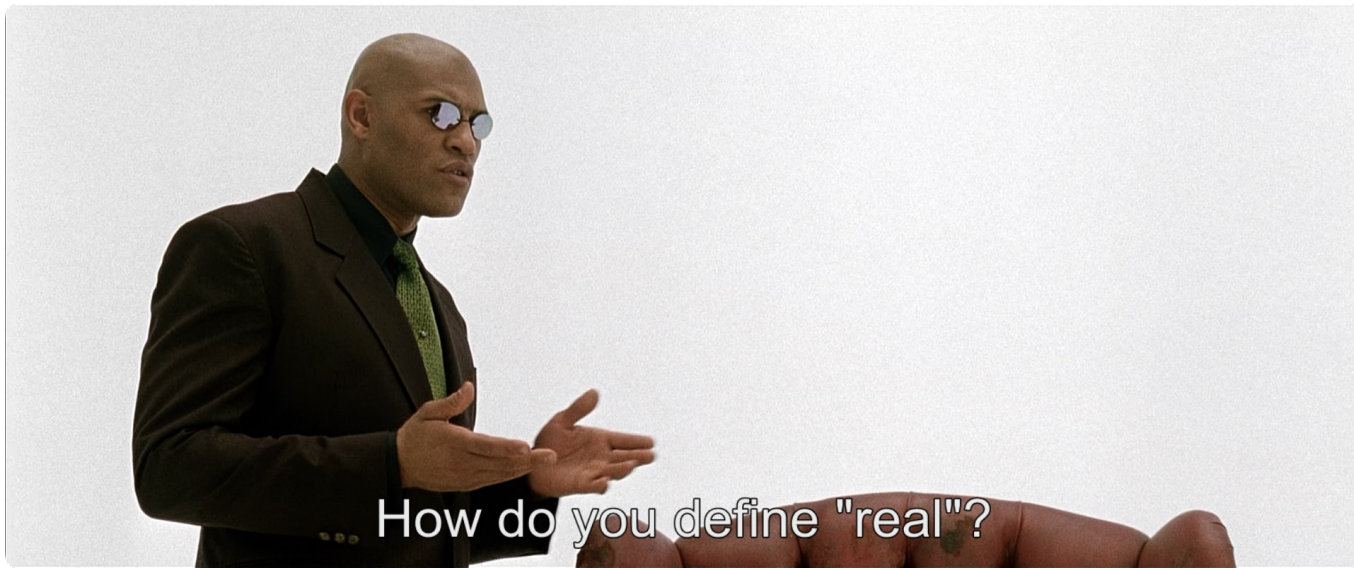

■ Overview

- The **Operating System** is the collection of **software** that **manages** the underlying **hardware** and **provides services** (system call APIs) to *application programs*.



■ Overview

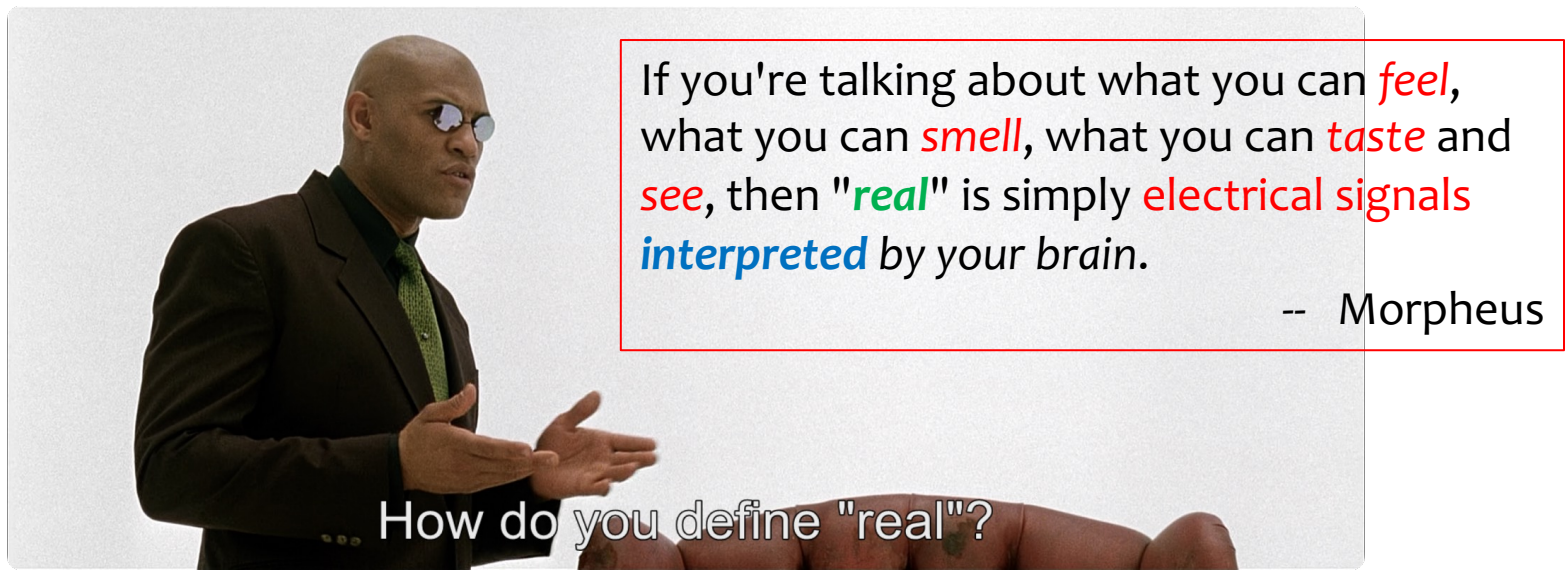

- The **Operating System** is the collection of **software** that **manages** the underlying **hardware** and **provides services** (system call APIs) to *application programs*.



How do you define "real"?

■ Overview

- The **Operating System** is the collection of **software** that **manages** the underlying **hardware** and **provides services** (system call APIs) to *application programs*.




If you're talking about what you can *feel*, what you can *smell*, what you can *taste* and *see*, then "*real*" is simply *electrical signals interpreted* by your *brain*.

-- Morpheus


How do you define "real"?

■ Overview

- The **Operating System** is the collection of **software** that **manages** the underlying **hardware** and **provides services** (system call APIs) to *application programs*.

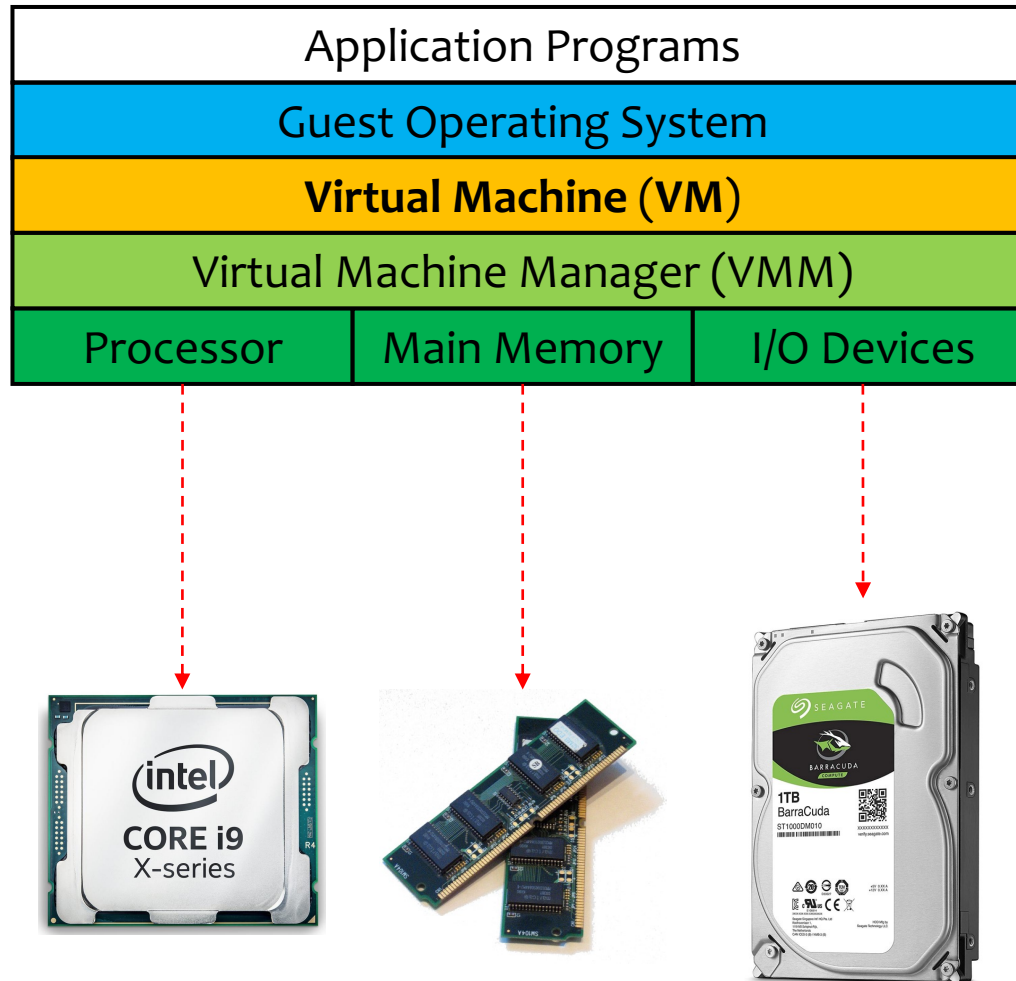


Similarly, "**hardware**" is simply **electrical signals** (0s and 1s) **interpreted** by the **OS**.



■ Overview


- A **Virtual Machine (VM)** is a software emulation of a physical computer. *It runs an OS just like a physical machine does.*



■ Overview

- A **Virtual Machine Manager (VMM)** **multiplexes** physical resources among the OSes in the VMs.

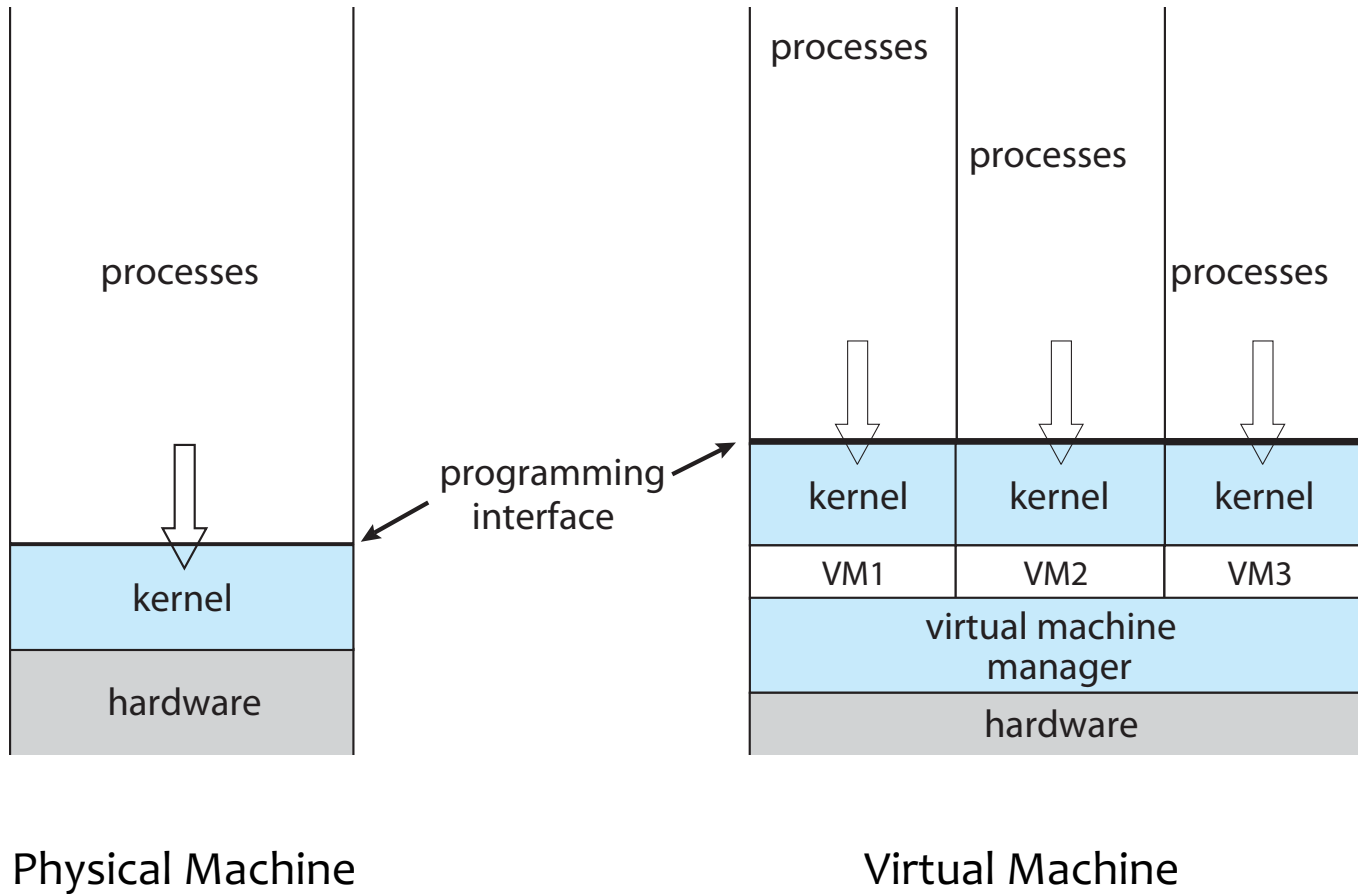
App	App	App	App
OS	OS	OS	OS
VM1	VM2	VM3	VM4
Virtual Machine Manager (VMM)			
Processor	Main Memory	I/O Devices	



■ Overview

- Fundamental idea – abstract hardware of a single computer into several different execution environments
 - Similar to layered approach in OS architecture
 - But layer creates virtual system (virtual machine) on which OSes can run
- Several components:
 - **Host**: underlying hardware(?) system.
 - **Virtual Machine Manager (VMM)** or Hypervisor: creates and runs VMs by providing an **interface** that is identical to the **host**.
 - **Guest**: Each guest process is provided with a *virtual copy* of the **host**.
Usually, the guest process is in fact an Operating System.
- Single physical machine can run multiple OSes concurrently, each in its own virtual machine.

■ System Models



■ History

- First appeared in IBM mainframes (IBM大型机) in 1972.
 - IBM **VM/370** divided (simple hardware-level multiplexing) a mainframe into multiple virtual machines, each running its own OS.
 - The hard part is **virtualizing disks** – **minidisks**.
- Formal definition of virtualization helped to establish system requirements and a target for functionality:
 - **Fidelity**. A VMM provides an environment for programs that is essentially identical to the original machine.
 - **Performance**: Programs running within that environment show only minor performance decreases.
 - **Safety**: The VMM is in complete control of system resources.
- In late 1990s, Intel CPUs fast enough for researchers to try virtualizing on general purpose PCs
 - **Xen** and **VMware** created virtualization technologies (still used today)
 - **Virtualization** has expanded into many OSes, CPUs, VMMs.



■ Benefits and Features

- Host system **protected** from VMs, VMs **protected** from each other
 - A virus that infected one of the VMs is not likely to spread to other VMs, nor the host system.
 - Sharing is provided via shared file system volume, or networking.
- **Freeze, suspend, running VMs**
 - They can move or copy somewhere else and resume.
 - Snapshot of a given state, able to restore back to that state
 - Some hypervisors allow multiple snapshots per VM
 - Clone by creating copies and running both original and the copies.
- Great for OS research, better system development efficiency
- Run multiple, different OSes on a single machine
 - Consolidation, better utilization, etc.



■ Benefits and Features

- **Templating (虚拟机模版)** – create an OS + application VM, provide it to customers
 - use it to create multiple instances of that combination quickly.
- **Live Migration (热迁移)** – move a running VM from one host to another
 - No interruption of services during migration
- All those features taken together → **Cloud Computing**
 - Use APIs to instruct cloud infrastructure (e.g., AWS, Azure) to create new VMs or web applications.



■ Implementation of VMMs

- **Type 0** hypervisors
 - hardware-based, e.g.,
- **Type 1** hypervisors
 - bare-metal, OS-level, e.g., KVM, Hyper-V, VMware ESXi, etc.
- **Type 2** hypervisors
 - application-level, process-based, e.g., VirtualBox, Parallels Desktop
- **Paravirtualization**
 - E.g., Xen
- **Programming Environment Virtualization**
 - E.g., **JVM** (Java Virtual Machine)
- **Emulators**
 - E.g., QEMU, Bochs
- **Application Containment (Containers)**
 - E.g., LXC, Docker, Kubernetes
- ...



■ Building Blocks

- Although the concept of **VM** is useful, it is difficult to implement.
 - Much work is required to provide an **exact** duplicate of the underlying (host) machine *to make the guest believe it is running on hardware*.
 - For example, on a **dual-mode** system where the CPU has only user-mode and *kernel-mode*, how can you **virtualize** the "**kernel-mode**" for the Guest OS?
- Most VMMs implement **virtual CPU (vCPU)** to represent state of CPU of the Guest.
 - The **vCPU** does not execute code.
 - It represents the state of the CPU as Guest machine *believes it to be*.
 - When the Guest is context-switched onto a CPU by the VMM, information from the **vCPU** is used to load the corresponding context.
 - much like the **PCB** in a general purpose OS.



■ Trap-and-Emulate

- Dual-mode CPU means Guest **only** executes in **physical** user-mode.
- The kernel of the Host OS runs in **physical** kernel-mode.
 - Not safe to let Guest kernel run in **physical** kernel-mode, too. **Why?**
 - No way for us to keep control of the physical machine!
 - No way to keep the Guest away from physical devices!
 - Once we let the Guest kernel run in physical kernel-mode, we can never get back control, e.g., Guest runs **I/O** instructions to **overwrite the VMM**.




■ Trap-and-Emulate

- Dual-mode CPU means Guest **only** executes in **physical** user-mode.
- The kernel of the Host OS runs in **physical** kernel-mode.
 - Not safe to let Guest kernel run in **physical** kernel-mode, too.
- Consequently, we must invent **two modes** for the **Virtual Machine**:
 - **virtual** user-mode and **virtual** kernel-mode.
 - both of which run in **physical** user-mode.
- Those actions that cause a **transfer** from user-mode to kernel-mode on a physical machine (e.g., syscalls, interrupts, or an attempt to execute a privileged instruction) must also cause a **transfer** (in the Virtual Machine) from **virtual** user-mode to **virtual** kernel-mode.
- How can such a **transfer** be accomplished?



■ Trap-and-Emulate

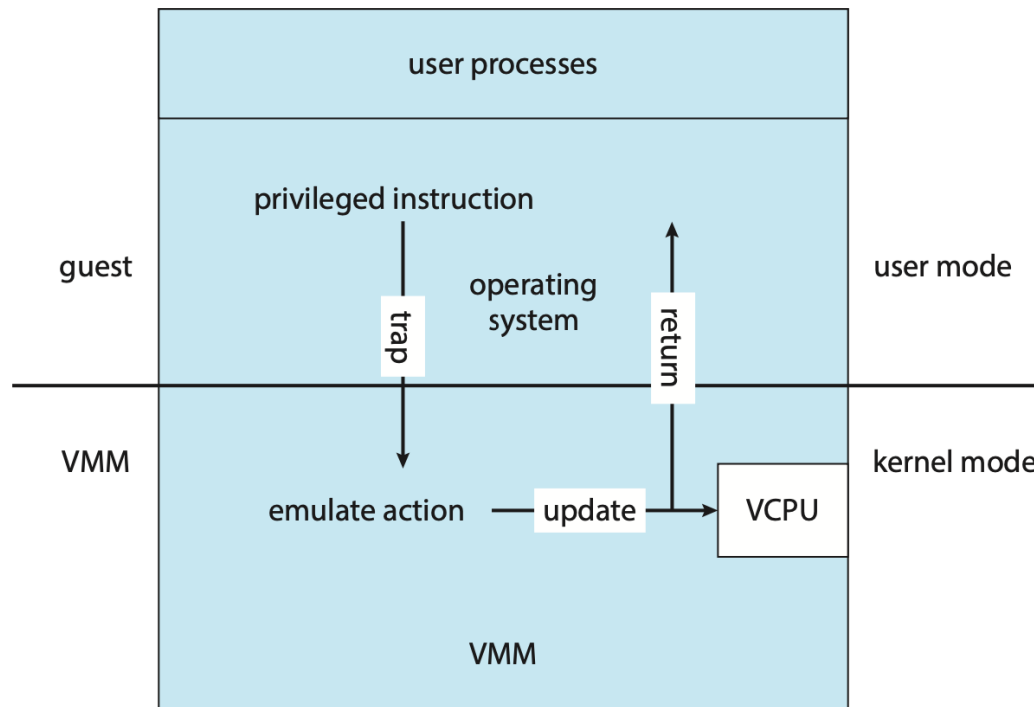
- How can such a **transfer** be accomplished? (**Trap-and-Emulate**)
 - When the **Guest** attempts a privileged instruction in user-mode, it would cause an error → **traps** into the VMM in Host machine.
 - VMM gains control, analyzes error, executes operations as attempted by the **Guest**
 - Returns control to **Guest** in user-mode.





■ Trap-and-Emulate

- How can such a **transfer** be accomplished? (**Trap-and-Emulate**)
 - For example, suppose the Guest issues a privileged instruction to modify a certain CPU flag → **traps** into the VMM
 - The VMM handles such a trap by emulating the privileged instruction on the **vCPU** that corresponds to the specific Guest, rather than executing it on *physical CPU*: **Maintain the illusion.**





■ Binary Translation

- Some CPUs don't have a clear separation between **privileged** and **non-privileged** instructions.
 - Earlier Intel x86 CPUs are among them
 - In fact, the earliest Intel CPU was designed for a calculator.
 - Intel CPUs are backward compatible → lack of separation persists
 - Consider the `popf` instruction in Intel x86:
 - Loads CPU flags register from the contents of the stack.
 - If CPU is in **privileged** mode → all flags replaced
 - If CPU is in user mode → only some flags replaced
 - No trap is generated. Therefore the aforementioned **trap-and-emulate** procedure is useless.
- Some other x86 instructions cause similar problems, we call this set of instructions ***special instructions***.



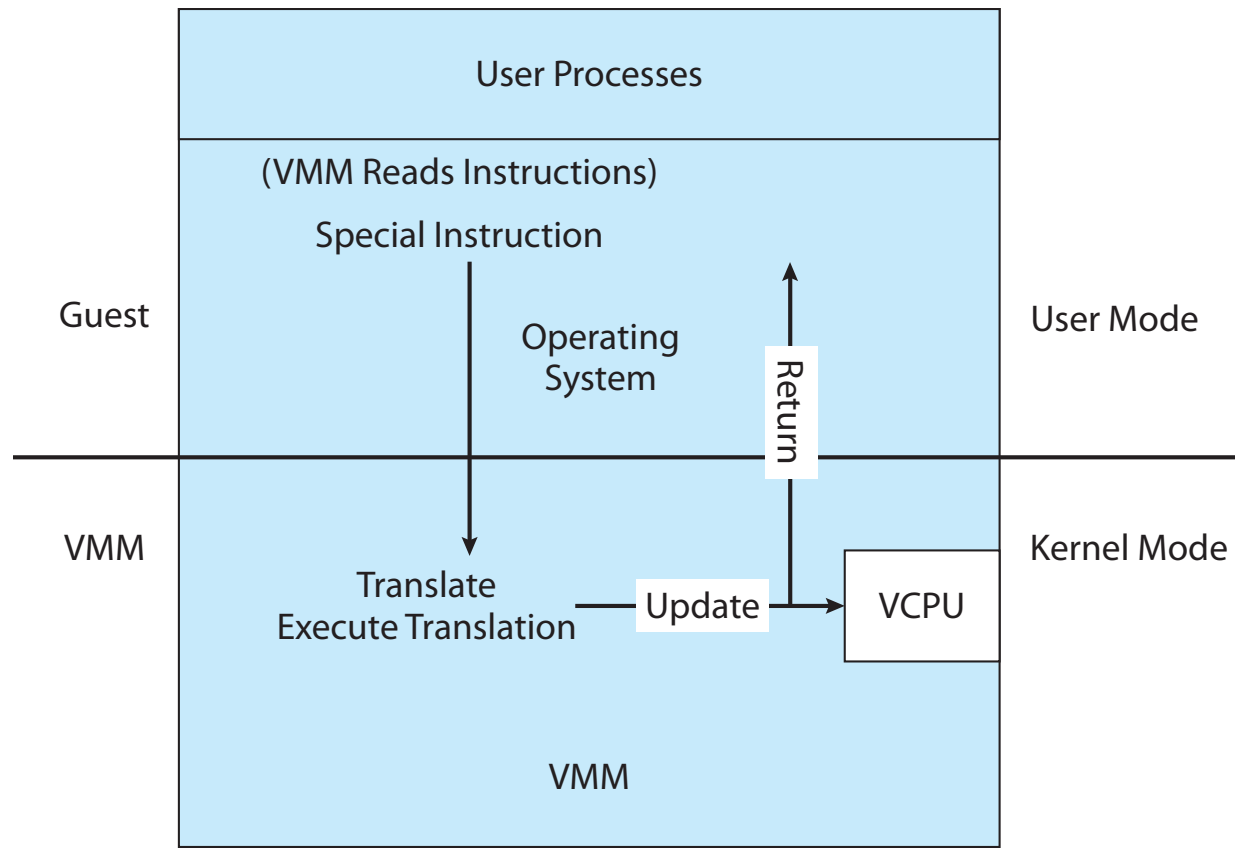
■ Binary Translation

- Some other x86 instructions cause similar problems, we call this set of instructions ***special instructions***.
 - caused **trap-and-emulate** method considered impossible until 1998.
- Binary translation solves this problem:
 - Concept is simple, but implementation very complex
 - If Guest vCPU is in user-mode, Guest can run instructions natively
 - If Guest vCPU is in kernel-mode (Guest believes it is in kernel-mode)
 - VMM examines every instruction that the Guest is about to execute by reading a few instructions ahead of program counter.
 - Non-special instructions run natively
 - ***Special instructions*** **translated** into **new set of instructions** that perform equivalent tasks (for example, changing the flags in the **vCPU**)



■ Binary Translation

- Binary translation is implemented by **translation code** within the VMM. The code reads native binary instructions dynamically from the Guest, **on demand**, and generates native binary code that executes in place of the original code.





■ Hardware Assistance

- All virtualization needs some hardware support
 - More support \Rightarrow more feature rich, stable, better performance
- Intel added new **VT-x** instructions in 2005, AMD added the **AMD-V** instructions in 2006
 - CPUs with these instructions ***remove the need for binary translation.***
 - Generally define more CPU modes: "**guest**" and "**host**"
 - from dual-mode to four-mode:
 - Host Mode; Host User Mode
 - Guest Kernel Mode; Guest User Mode
 - VMM can enable **host** mode, define characteristics of each guest VM, switch to **guest** mode, passing control to a Guest OS.
 - In **guest** mode, the Guest OS thinks it is running natively
 - Access to virtualized device, priv instructions cause trap to VMM
 - CPU maintains vCPU, context switches as needed
- Hardware support for Nested Page Tables, DMA, interrupts as well.




■ Virtual Machine Life Cycle

- Whatever the hypervisor type, a typical VM has a life cycle:
 - Created by VMM (hypervisor)
 - Resources assigned to VM
 - # of CPU or vCPU cores
 - # of memory
 - # of disk space
 - networking details (# of NICs, bandwidth requirements, etc.)
 - ...and so on
 - Resources are usually **shared** (multiplexed, or provisioned) among VMs
 - except for type 0 hypervisor, where resources are usually **dedicated**
 - When the VM is no longer needed, it can be deleted or suspended.



■ Type 0 Hypervisor

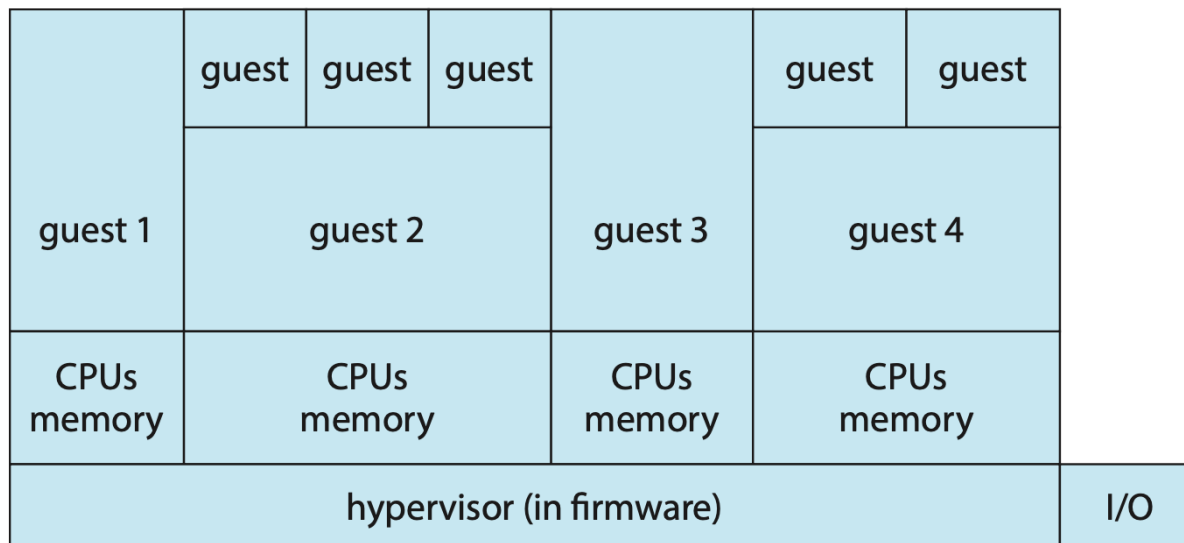
- Historically also called "*partitions*", "*domains*".
- A Hardware feature implemented by **firmware**.
 - The **VMM** itself is encoded in the **firmware** rather than in OS software.
 - At boot time, the **VMM** loads the guest images to run in each "partition"
 - Provides much less features compared with other types of hypervisors.
 - Each guest *believes* it has **dedicated** hardware (*because it really does*), simplifying many implementation details. ⇒ much **less flexible**.
- Each guest is basically running on raw hardware.





■ Type 0 Hypervisor


- Each guest is basically running on raw hardware.
 - Close or identical to the underlying hardware in performance.
 - Each guest itself can even be a VMM that supports more VMs
 - I/O is a huge **challenge**:
 - must have enough devices, controllers to dedicate to each guest.
 - For example, the underlying hardware must have at least 3 CPU cores in order to support 3 running guests, or at least 2 Ethernet ports to dedicate to two separate guests.





■ Type 0 Hypervisor

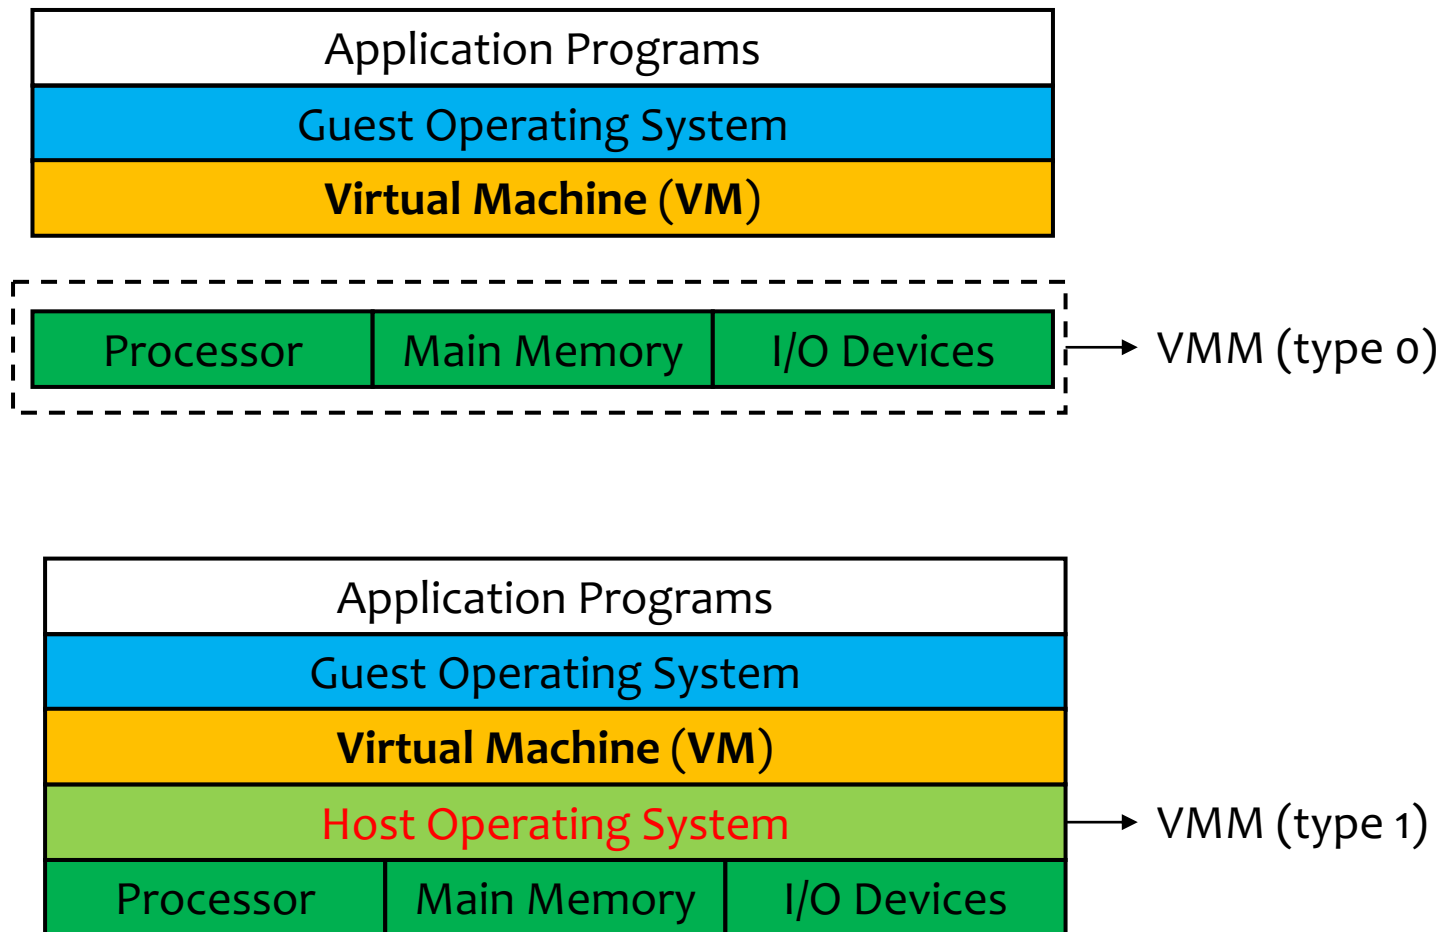
- It is rare to find a *pure* "**Type-0**" hypervisor nowadays.
 - Implemented in hardware (firmware) ⇒ Too **inflexible**.
 - Hardware resources dedicated to each guest ⇒ **Low** overall **utilization**.
 - Whatever a "**Type-0**" hypervisor can do, a "**Type-1**" can do it, too.
 - Outside the textbook, the term "**Type-0**" hypervisor is not widely recognized in the standard classification of hypervisors in the computing industry, i.e., they would instead be categorized as "**Type-1**".





■ Type 1 Hypervisor

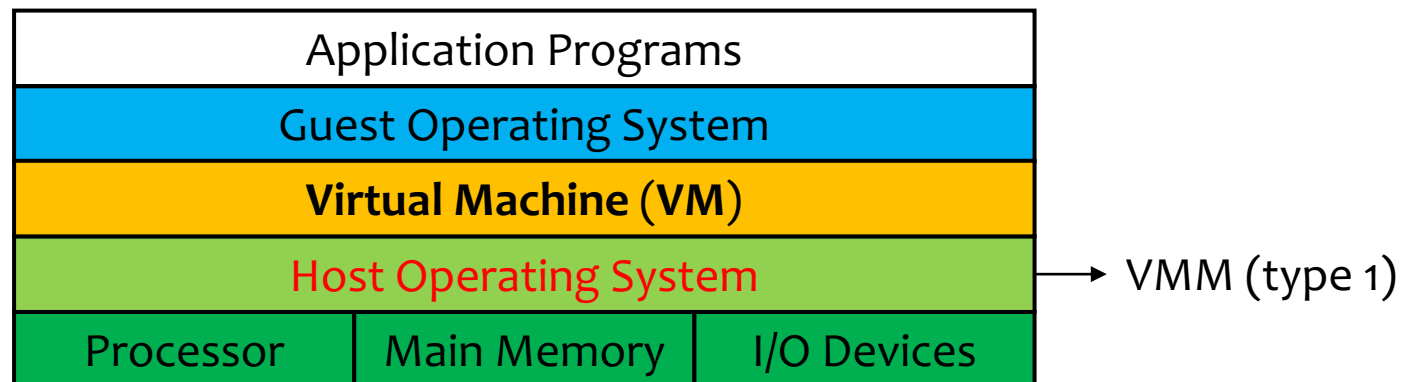
- **Type 1** Hypervisor shifts the **VMM** functionalities from *hardware* (**firmware**) to software (**OS**).





■ Type 1 Hypervisor

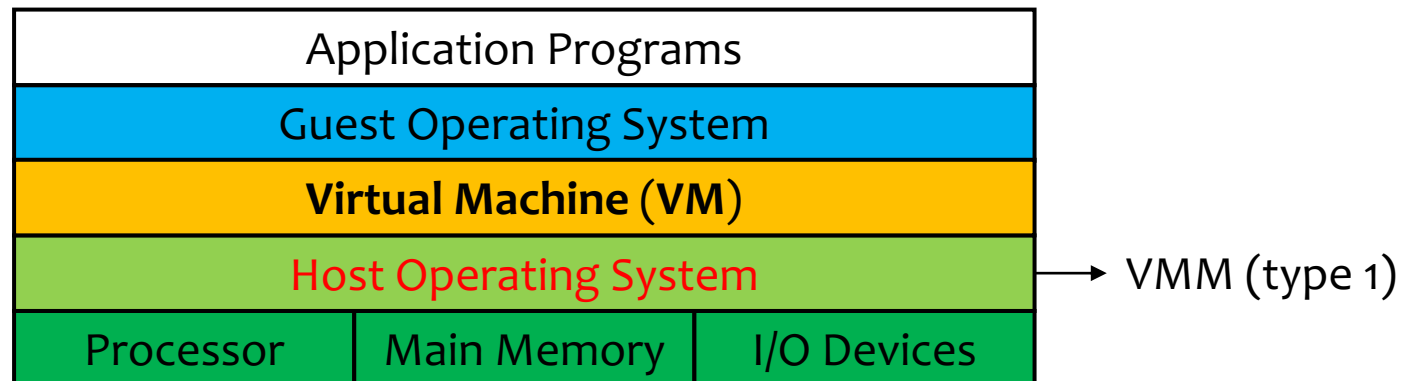
- **Type 1** Hypervisor shifts the **VMM** functionalities from *hardware* (**firmware**) to software (**OS**).
 - **Special purpose** OS that runs natively on hardware \Rightarrow in **kernel mode**.
 - Rather than providing syscall APIs for application programs, it exports **APIs** for VM management (creation, deletion, etc.)
 - Can also be **general purpose** OS that **integrates** **VMM** functionalities.
 - E.g., a typical Linux OS with **KVM** module enabled can be used as a **Type-1** hypervisor \Rightarrow the most popular Open Source hypervisor solution available today.
 - In many ways, the **guest OS** is treated as just another **process**.






■ Type 1 Hypervisor

- **Type 1** Hypervisor also called "the Data-center Operating System".
 - Data-center managers control and manage OSES in new, sophisticated ways by controlling the Type-1 hypervisor.
 - Consolidation of multiple OSES and apps onto less hardware.
 - Move guests between systems to balance performance
 - Snapshots and cloning
 - The cornerstone of **IaaS** Cloud Computing.



■ Type 1 Hypervisor

- The cornerstone of **IaaS** Cloud Computing.



Azure services


Create a
resource


Virtual
machines


Free services


All resources


Resource
groups


Azure OpenAI


Public IP
addresses


Snapshots


Network
interfaces


More services

Resources

Recent Favorite

Name	Type	Last Viewed
 winVM	Virtual machine	5 hours ago

■ Type 1 Hypervisor

- The cornerstone of **IaaS** Cloud Computing.

The screenshot shows the Microsoft Azure portal interface. The browser address bar displays 'portal.azure.com/#browse/Micr...'. The page title is 'Virtual machines'. Below the title, there are options to 'Create', 'Switch to classic', 'Reservations', 'Manage view', 'Refresh', and 'Export to CSV'. A sidebar on the left lists three options: 'Azure virtual machine' (highlighted with a red circle), 'Azure virtual machine with preset configuration', and 'More VMs and related solutions'. The main content area shows a table with columns: Subscription, Resource group, Location, and Status. The table contains two rows, both for 'Azure for Students' in the 'free' resource group, located in 'East Asia', with a status of 'Running'.

Virtual machines - Microsoft

portal.azure.com/#browse/Micr...

Microsoft Azure

Search resources, services, and docs (G+)

Home >

Virtual machines

Default Directory onmicrosoft.com)

+ Create ▾ ↺ Switch to classic ⌚ Reservations ▾ ⚙ Manage view ▾ ↻ Refresh ⬇ Export to CSV ⋮

- Azure virtual machine**
Create a virtual machine hosted by Azure
- Azure virtual machine with preset configuration**
Create a virtual machine with presets based on your workloads
- More VMs and related solutions**
Discover and deploy full workloads and Azure products for your business needs

Type equals **all** + Add filter More (2)

No grouping ▾ List view ▾


Subscription ↑↓	Resource group ↑↓	Location ↑↓	Status ↑↓
Azure for Students	free	East Asia	Running
Azure for Students	free	East Asia	Running

< Page 1 ▾ of 1 >

Give feedback

■ Type 1 Hypervisor

- The cornerstone of **IaaS** Cloud Computing.



Home > Virtual machines >

Create a virtual machine

✓ Validation passed

Basics Disks Networking Management Monitoring Advanced Tags Review + create

i Cost given below is an estimate and not the final price. For all your pricing needs, please use the pricing calculator. [↗](#)

Price

1 X Standard B1ls
by Microsoft
[Terms of use](#) | [Privacy policy](#)

Subscription credits apply ^①

0.0073 USD/hr
[Pricing for other VM sizes](#)

■ Type 1 Hypervisor

- The cornerstone of **IaaS** Cloud Computing.

The screenshot shows the Microsoft Azure portal interface. The browser address bar displays 'virtualMachines/freeVM/overview'. The page title is 'freeVM - Microsoft Azure'. The navigation bar includes 'Microsoft Azure' and a search bar. The main content area is divided into two columns: 'Virtual machine' and 'Networking'. The 'Virtual machine' column lists properties such as Computer name, Operating system (Linux), VM generation (V1), VM architecture (x64), Agent status (Not Ready), Agent version (Unknown), Hibernation (Disabled), Host group, Host, Proximity placement group, Colocation status (N/A), and Capacity reservation group. The 'Networking' column lists properties such as Public IP address, Public IP address (IPv6) (2603:1040:200:3::129), Private IP address (10.1.1.4), Private IP address (IPv6) (fd3b:7311:5041::4), Virtual network/subnet (freeVM-vnet/default), and DNS name. A 'Size' section is also visible, showing the VM size as Standard B1s, with 1 vCPU and 1 GiB RAM.

Virtual machine	
Computer name	-
Operating system	Linux
VM generation	V1
VM architecture	x64
Agent status	Not Ready
Agent version	Unknown
Hibernation	Disabled
Host group	-
Host	-
Proximity placement group	-
Colocation status	N/A
Capacity reservation group	-

Networking	
Public IP address	-
Public IP address (IPv6)	2603:1040:200:3::129 (Network interface freevm242)
Private IP address	10.1.1.4
Private IP address (IPv6)	fd3b:7311:5041::4
Virtual network/subnet	freeVM-vnet/default
DNS name	-

Size	
Size	Standard B1s
vCPUs	1
RAM	1 GiB



■ Type 1 Hypervisor

- **Recommended:** Create an Azure Student account

- <https://azure.microsoft.com/en-us/free/students>



Azure

Explore ▾

Products ▾

Solutions ▾

More ▾



Learn

Contact Sales

Support

Sign in

Build in the cloud free with Azure for Students

Use your university or school email to sign up and renew each year you're a student

Start free

Learn about eligibility

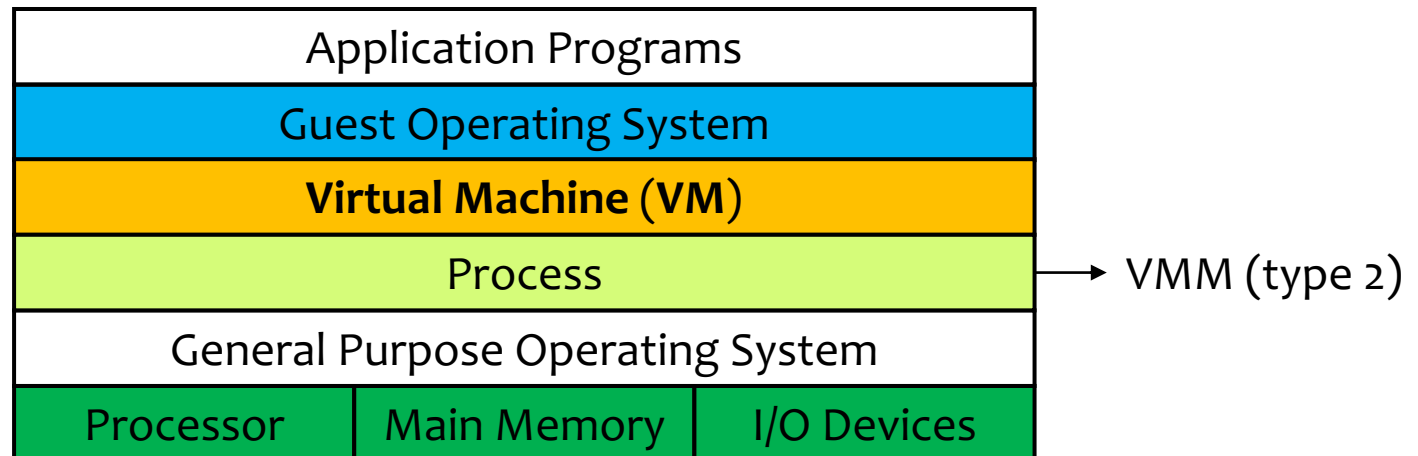
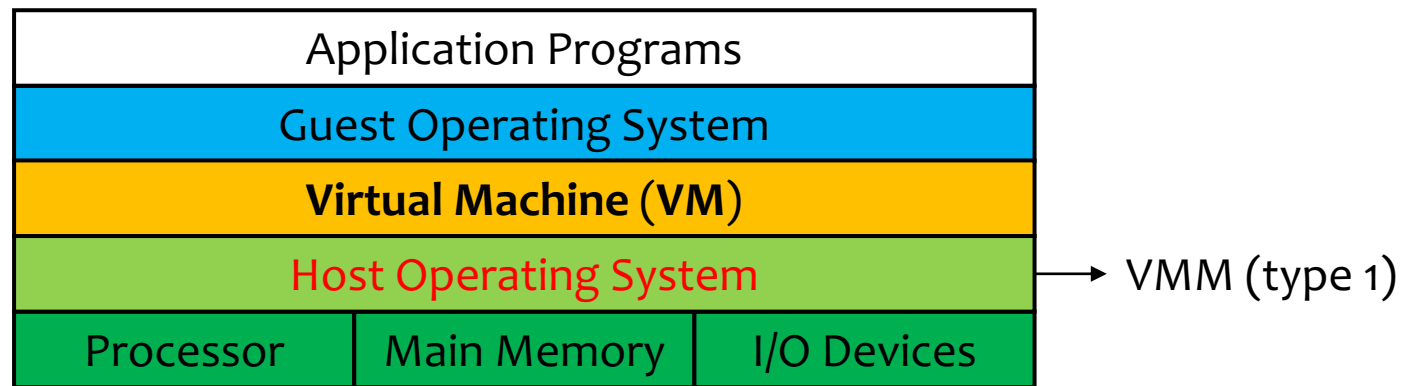
Start with
\$100 Azure
credit

No credit
card
required



■ Type 2 Hypervisor

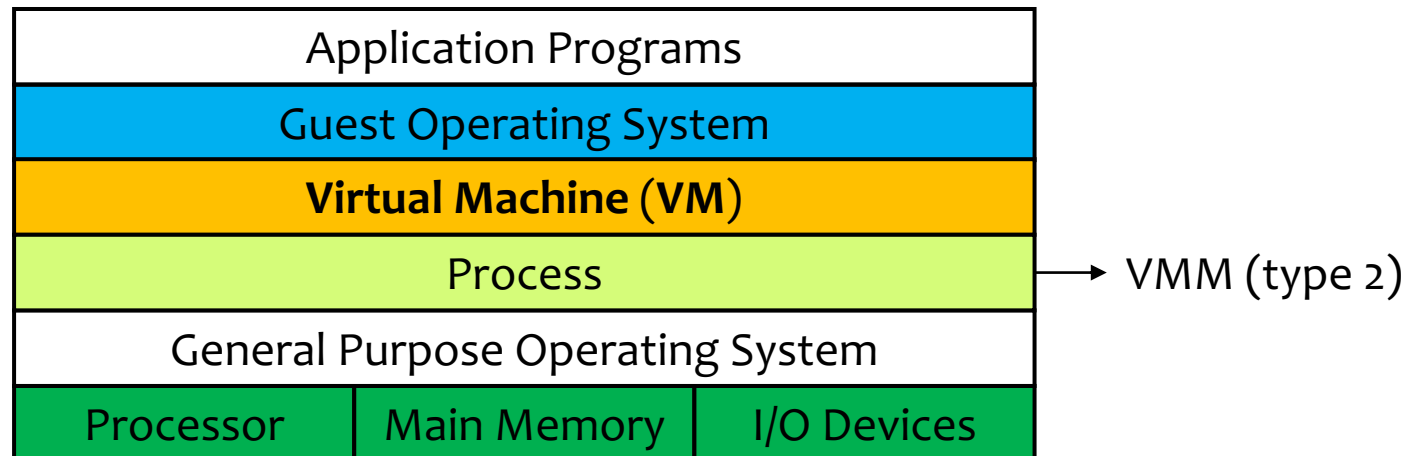
- **Type 2** Hypervisor shifts the **VMM** functionalities from *kernel-mode OS* to *user-mode Processes*.





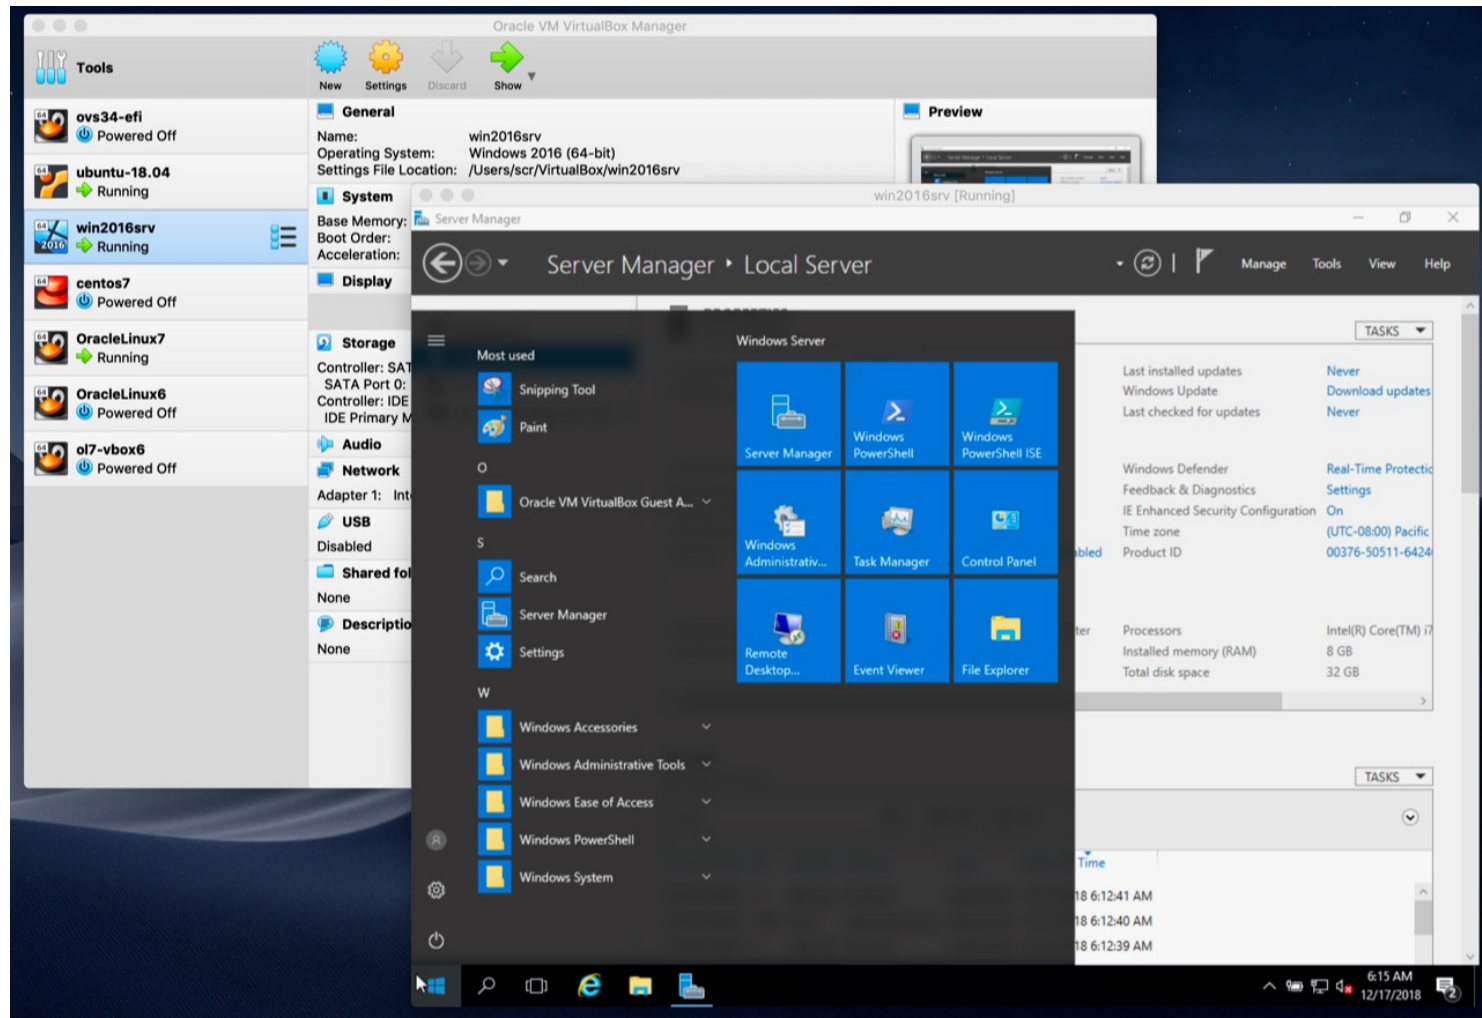
■ Type 2 Hypervisor

- **Type 2** Hypervisor shifts the **VMM** functionalities from *kernel-mode OS* to user-mode **Processes**.
 - Very little OS involvement in virtualization.
 - **VMM** is simply another process, run and managed by host OS.
 - Even the host doesn't know it is a hypervisor that manages guests.
 - **Disadvantage:** **poor performance** (can't take advantage of HW features)
 - **Advantages:** **Flexible** (a normal user can install a **Type-2** hypervisor software, e.g., VirtualBox, Parallels Desktop, on any supported OS).



■ Type 2 Hypervisor

- A normal user can install a Type-2 hypervisor software, e.g., **VirtualBox**, Parallels Desktop, on any supported OS.



■ Type 2 Hypervisor

- A normal user can install a Type-2 hypervisor software, e.g., VirtualBox, **Parallels Desktop**, on any supported OS.





■ Type 3 Hypervisor

- **Type 3** Hypervisor shifts the **VMM** functionalities from *user-mode Processes* to ???



■ Type 3 Hypervisor

- ~~Type 3~~ Hypervisor shifts the ~~VMM~~ functionalities from ~~user-mode~~ **Processes** to ???
- No, there is no such thing as a **Type-3** Hypervisor, since there is no more higher-level abstractions than **Processes** that can *fully* sustain a Virtual Machine.



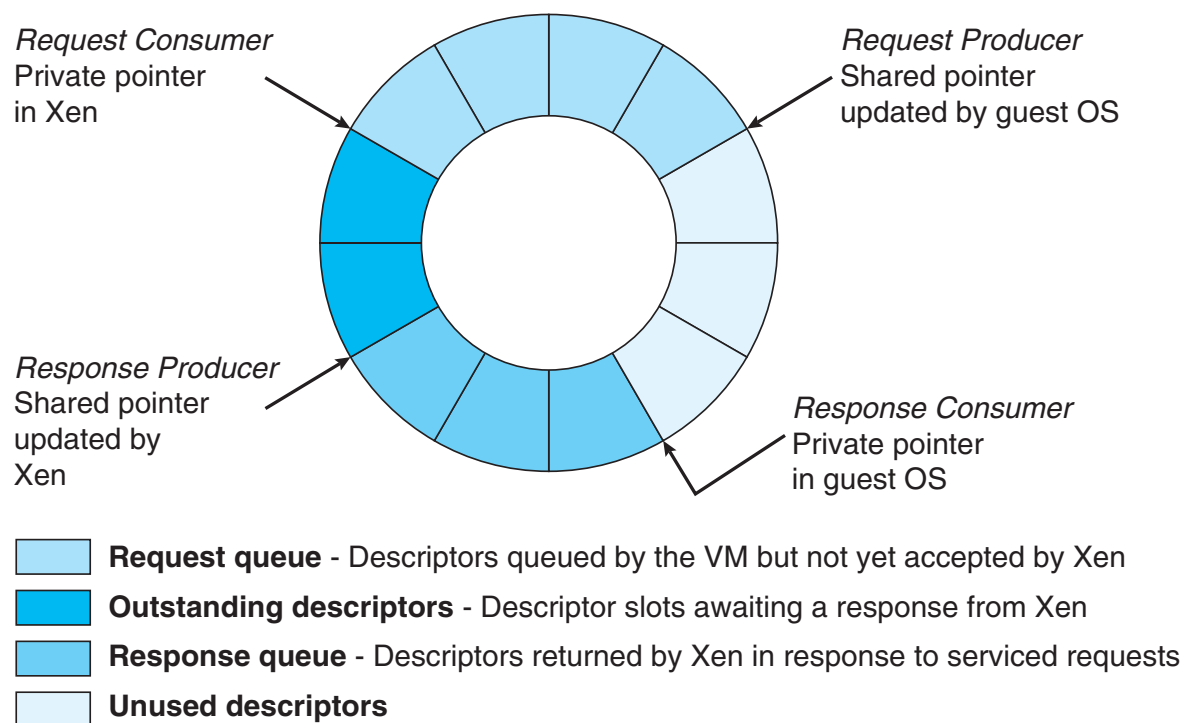
■ Paravirtualization

- A **hypervisor** uses a technique called "**Full Virtualization**" (**全虚拟化**) when it comes to managing Virtual Machines, which means that the Virtual Machine OSes need not be altered or modified.
 - In other words, in "Full Virtualization", the Guest OS does not know whether it is running on real hardware or a VMM.
- Apparently, this is not the only way. **Paravirtualization** (**半虚拟化**) works differently by providing the Guest OS with specialized APIs.
 - The Guest **knows** it is not running on hardware.
 - The Guest OS must be **modified** to accommodate the specialized APIs.
 - The specialized APIs would normally enable more **efficient use of resources** and **a smaller virtualization layer**.



■ Paravirtualization

- The Xen VMM was the leader in paravirtualization by implementing several techniques to optimize the performance of **guests** and **host**. For example, clean and simple device abstractions:
 - More **Efficient** I/O
 - Better **communication** between guest and VMM about device I/O
 - Each device has circular buffer shared by guest and VMM (*shared memory*).





■ Paravirtualization

- The Xen VMM was the leader in paravirtualization by implementing several techniques to optimize the performance of guests and host. For example, clean and simple device abstractions:
 - **Memory Management** does not include **Nested Page Tables**.
 - Each guest has its own read-only tables
 - Guest uses hypercall (call to hypervisor) when page-table changes needed
- Paravirtualization allowed virtualization of older x86 CPUs (and others) without binary translation.
- But on modern CPUs, Xen no longer required guest modification
 - Hence, no longer paravirtualization.



■ Programming Environment Virtualization

- Not really virtualization, but using same techniques, providing similar features.
- Programming language is designed to run within custom-built virtualized environment
 - For example, Java has many features that depend on running in **Java Virtual Machine (JVM)**.
- In this case, virtualization is defined as providing APIs that define a set of features made available to a language and programs written in that language to provide an improved execution environment.
- JVM compiled to run on many systems.
- Programs written in Java run in the JVM
 - compiled into Java Bytecode
 - can run on anywhere with JVM (no matter the underlying system).
 - "Write Once, Run Anywhere".



■ Programming Environment Virtualization

Java Source Code

```
outer:
for (int i = 2; i < 1000; i++) {
    for (int j = 2; j < i; j++) {
        if (i % j == 0)
            continue outer;
    }
    System.out.println (i);
}
```

Java Bytecode

```
0:  iconst_2
1:  istore_1
2:  iload_1
3:  sipush 1000
6:  if_icmpge 44
9:  iconst_2
10: istore_2
11: iload_2
12: iload_1
13: if_icmpge 31
16: iload_1
17: iload_2
18: irem
19: ifne 25
22: goto 38
25: iinc 2, 1
28: goto 11
31: getstatic #84
34: iload_1
35: invokevirtual #85
38: iinc 1, 1
41: goto 2
44: return
```



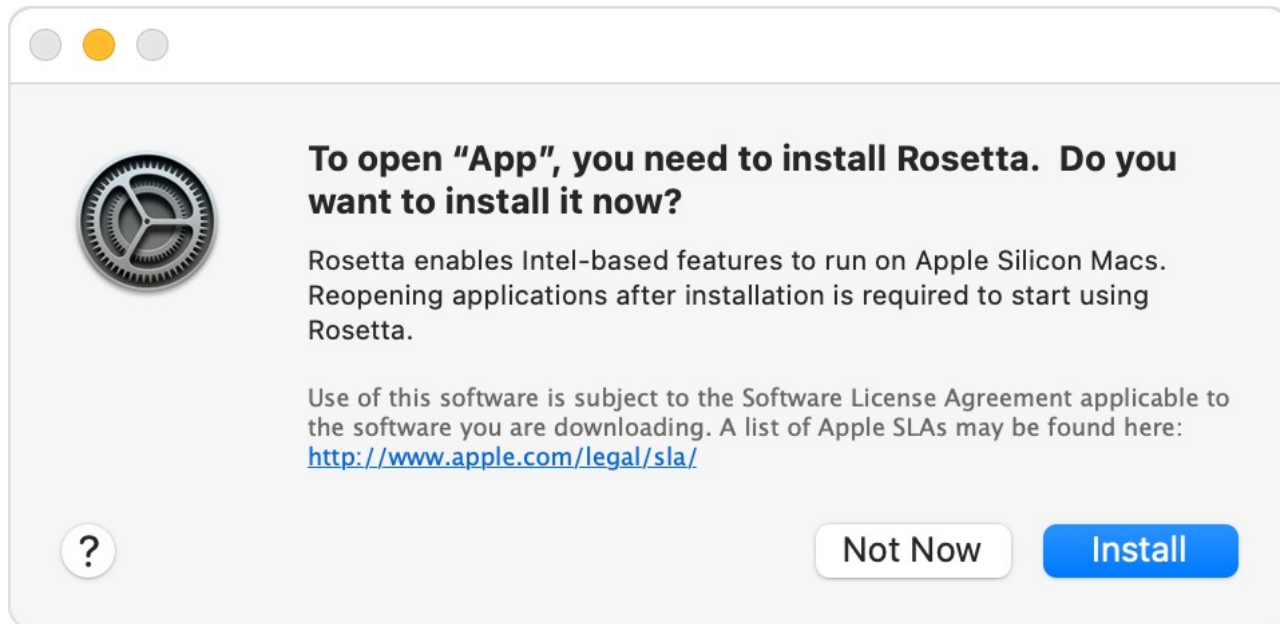
■ Emulation

- Another (older) way for running one OS on a different OS
 - Virtualization requires underlying CPU (or more precisely, the Instruction Set Architecture, ISA) to be the same as guest was compiled for
 - Emulation allows guest to run on different CPU
- Need to translate all guest instructions from Guest CPU to native CPU
 - **Emulation**, not **Virtualization**
- Useful when host system has one architecture, guest compiled for other architecture
 - Company replacing outdated servers with new servers containing different CPU architecture, but still want to run old applications
- Performance challenge: order of magnitude slower than native code
 - New machines faster than older machines so can reduce slowdown
 - Bochs (pure software emulator) is way slower than QEMU (hardware-assisted emulator)



■ Emulation

- Not all emulators are poor in performance. For example, **Rosetta 2**, developed by Apple Inc for the transition from **Intel processors** to **Apple Silicon** chips, can have even better performance when running x86 programs on M1 macOS than on native Intel CPUs!



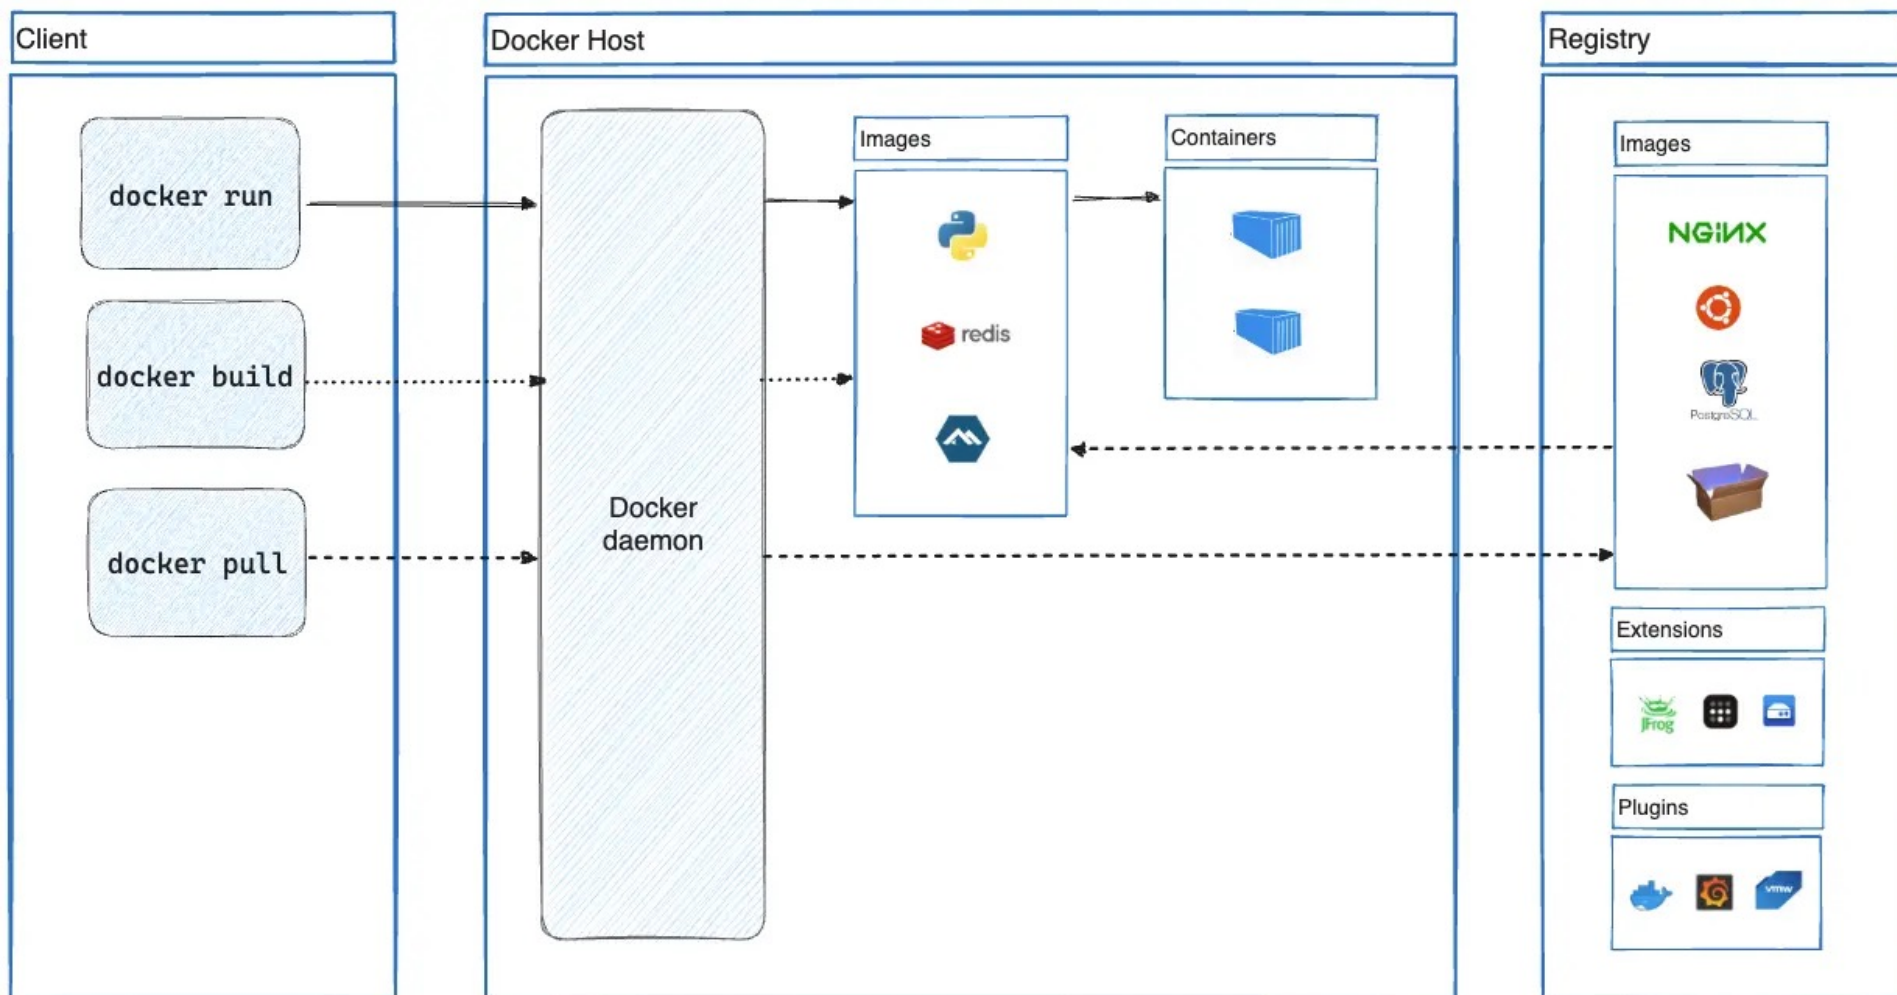


■ Application Containment

- Some goals of **virtualization** are segregation of apps, performance and resource management, easy start, stop, move, and management.
- Can do those things without full-fledged virtualization
 - If applications compiled for the host operating system, don't need full virtualization to meet these goals.
- Containers create virtual layer between OS and apps.
 - Only one kernel running – host OS
 - OS and devices are virtualized, providing resources within zone with impression that they are only processes on system.
 - Each zone has its own applications; networking stack, addresses and ports; user accounts, etc.
 - CPU and memory resources divided between zones
 - Zone can have its own scheduler to use those resources.

■ Docker

- A successful example of Application Containment.
- much more **lightweight** than traditional Virtual Machines.





Thank you!